

# GNU/Linux Application Programming

Lecture Handouts

WANG Xiaolin

`wx672ster+linux@gmail.com`

August 12, 2019

## Contents





<b>1</b>	<b>Getting Started</b>	<b>2</b>
<b>2</b>	<b>Shell Basics</b>	<b>3</b>
<b>3</b>	<b>Shell Programming</b>	<b>7</b>
<b>4</b>	<b>C Programming Basics</b>	<b>11</b>
4.1	Programming Environment . . . . .	11
4.2	OS Basics . . . . .	19
<b>5</b>	<b>The Linux Environment</b>	<b>26</b>
<b>6</b>	<b>Working With Files</b>	<b>28</b>
6.1	File . . . . .	28
6.2	Directory . . . . .	35
<b>7</b>	<b>Processes and Threads</b>	<b>38</b>
7.1	Process . . . . .	38
7.2	Thread . . . . .	40
7.3	Signals . . . . .	44
<b>8</b>	<b>IPC</b>	<b>47</b>
8.1	Pipes and FIFOs . . . . .	47
<b>9</b>	<b>User Interface</b>	<b>52</b>
9.1	Dialog, Zenity . . . . .	52
9.2	Ncurses . . . . .	52
9.3	GTK+ . . . . .	52
9.4	Qt . . . . .	52
<b>10</b>	<b>Terminal</b>	<b>52</b>
<b>11</b>	<b>IDE</b>	<b>52</b>

## References

- [1] MATTHEW N, STONES R. Beginning linux programming. John Wiley & Sons, 2008.
- [2] COOPER M. Advanced Bash Scripting Guide 5.3 Volume 1. Lulu.com, 2010.
- [3] RAYMOND E S. The art of Unix programming. Addison-Wesley, 2003.

- [4] STEVENS W R, RAGO S A. Advanced programming in the UNIX environment. Addison-Wesley, 2013.
- [5] LOVE R. Linux System Programming: Talking Directly to the Kernel and C Library. O'Reilly Media, Inc., 2007.
- [6] KERRISK M. The Linux Programming Interface: A Linux and UNIX System Programming Handbook. No Starch Press, 2010.
- [7] BRYANT R E, O'HALLARON D R. Computer Systems: A Programmer's Perspective. 2nd ed. Addison-Wesley, 2010.
- [8] TANENBAUM A S. Modern Operating Systems. 4th ed. Prentice Hall Press, 2015.
- [9] Wikipedia. File Allocation Table — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=File\\_Allocation\\_Table&oldid=661104239](http://en.wikipedia.org/w/index.php?title=File_Allocation_Table&oldid=661104239).
- [10] Wikipedia. Inode — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Inode&oldid=647736522>.

## Course Web Links

 <https://cs6.swfu.edu.cn/moodle>  
 [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/linux-app/slides/](https://cs2.swfu.edu.cn/~wx672/lecture_notes/linux-app/slides/)  
 [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/linux-app/src/](https://cs2.swfu.edu.cn/~wx672/lecture_notes/linux-app/src/)  
 <https://cs3.swfu.edu.cn/tech>

## /etc/hosts

```
202.203.132.241  cs6.swfu.edu.cn
202.203.132.242  cs2.swfu.edu.cn
202.203.132.245  cs3.swfu.edu.cn
```

## Homework

### Weekly tech question

1. What was I trying to do?
2. How did I do it? (steps)
3. The expected output? The real output?
4. How did I try to solve it? (steps, books, web links)
5. How many hours did I struggle?

✉ [wx672ster+linux@gmail.com](mailto:wx672ster+linux@gmail.com)

🌐 Preferably in English

📖 in stackoverflow style

Or simply show me the tech questions you asked on any website

# 1 Getting Started

## Linux Commands

**Where to find them?** /bin, /usr/bin, /usr/local/bin,  
~/bin, ...  
\$ echo \$PATH

**How to find them?** which, whereis, type

## Command not found?



**First** double check your spelling


**Then** try:

```

🔍 aptitude search xxx
🔍 apt-cache search xxx
🔍 apt-file search xxx
🔍 sudo apt install packagename
🔍 Google "linux command xxx"
```

## Text Editors

 vs. 





## Help Your Editor

### Suffix matters

```
$ vim ✗  
$ vim hello ✗  
$ vim hello.c ✓  
$ vim hello.py ✓  
$ emacs ✗  
$ emacs hello ✗  
$ emacsclient hello.c ✓  
$ emacsclient hello.py ✓
```



## Keyboard




 vimtutor  
 `Ctrl` + `h` `t`

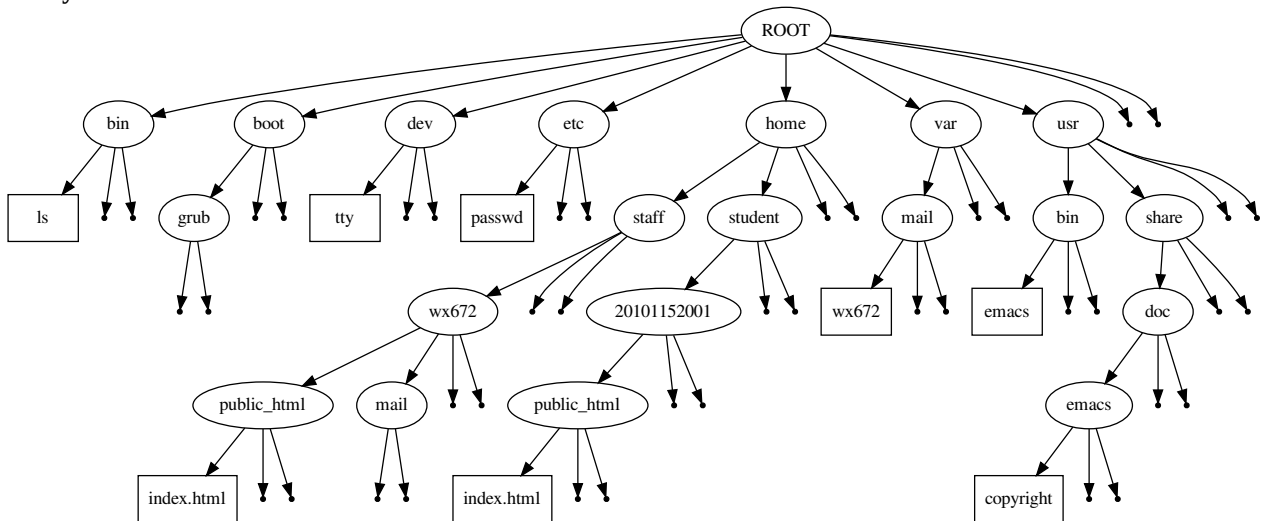
## 2 Shell Basics

### Shell

-  A command line interpreter
-  A programming language



## Directory Structure



Todo	How
Where am I?	pwd
What's in it?	ls
Move around?	cd
Disk usage?	du, df
USB drive?	lsblk, mount
New folder?	mkdir

## File Operations

### Ways to create a file

- ✎ Using an editor (vim, emacs, nano...), or
- \$ cat > filename
- \$ echo "hello, world" > filename
- \$ touch filename

More file operations:

Todo	How	Todo	How
Copy?	cp	Move/Rename?	mv
Delete?	rm	What's it?	file
Link?	ln	Permission?	chmod, chown
Count?	wc	Archive?	tar, gzip, 7z, ...
Sort?	sort, uniq	Search?	find, grep

## Redirection

### Redirecting output

```
$ ls -l > output.txt
$ ps aux >> output.txt
```

### Redirecting input

```
$ more < output.txt
```

## Process Operations

Todo	How	Todo	How
Kill?	kill, Ctrl-c	suspend?	Ctrl-z
background?	bg, &	foreground?	fg, jobs
status?	ps, top		

## System Info

Todo	How	Todo	How
who?	w, who, whoami	how long?	uptime
software?	apt, aptitude, dpkg	kernel?	uname, lsmod
hardware?	lspci, lsusb, lscpu	memory?	free, lsmem

## APT — package management

Todo	How
upgrading?	apt update && apt upgrade
install?	apt install xxx
remove?	apt purge xxx
search?	apt search xxx
details?	apt show xxx
friendly UI?	aptitude

## CLI Shortcuts

<b>Ctrl</b> + <b>a</b> : beginning of line	<b>Ctrl</b> + <b>e</b> : end of line
<b>Ctrl</b> + <b>f</b> : forward	<b>Ctrl</b> + <b>b</b> : backward
<b>Ctrl</b> + <b>n</b> : next	<b>Ctrl</b> + <b>p</b> : previous
<b>Ctrl</b> + <b>r</b> : reverse search	<b>Ctrl</b> + <b>u</b> : cut to beginning
<b>Ctrl</b> + <b>k</b> : kill (cut to end)	<b>Ctrl</b> + <b>y</b> : yank (paste)
<b>Ctrl</b> + <b>d</b> : delete a character	<b>Tab</b> : completion

## Tmux

<b>Ctrl</b> + <b>a c</b> : create window	<b>Ctrl</b> + <b>a</b> <b>Ctrl</b> + <b>a</b> : switch window
<b>Ctrl</b> + <b>a n</b> : next window	<b>Ctrl</b> + <b>a p</b> : previous window
<b>Ctrl</b> + <b>a -</b> : split window	<b>Ctrl</b> + <b>a l</b> : split window
<b>Ctrl</b> + <b>a j</b> : go down	<b>Ctrl</b> + <b>a k</b> : go up
<b>Ctrl</b> + <b>a l</b> : go right	<b>Ctrl</b> + <b>a h</b> : go left

## Understanding “ls -l”

```
-rw----- 1 sam sam 57 Apr 17 1998 weather.txt
drwxr-xr-x 6 sam sam 102 Oct 9 1999 web_page
-rw-rw-r-- 1 sam sam 7648 Feb 11 20:41 web_site.tar
-rw----- 1 sam sam 574 Dec 16 1998 file.txt
```

File Name  
Modification Time  
Size (in bytes)  
Group  
Owner  
Number of hard links  
File Permissions  
File types

d - directory  
- - regular file  
l - soft link  
c - character device  
b - block device  
s - socket  
p - named pipe (FIFO)

## 9-bit permission

```
7 5 5
111 101 101
rwx r-x r-x
```

Other  
Group  
User

```
$ chmod 755 foo      $ chmod 644 foo
$ chmod 000 foo      $ chmod 777 foo
$ chmod a-r foo      $ chmod u+x foo
$ chmod g+w foo      $ chmod go=rx foo
```

## Wildcard Expansion

Character	Meaning	Example
?	any one	\$ ls ????.txt
*	zero or more	\$ ls *.c
[]	or	\$ ls *. [ch]
{ }	and	\$ ls *. {c,h,cpp}

## Example

```
$ touch {2,3,4,234}. {jpg,png} && ls
```

output:

2.jpg	234.jpg	3.jpg	4.jpg
2.png	234.png	3.png	4.png

```
$ rm [234].jpg      $ rm ?.jpg
$ rm {2,3,4,234}.jpg $ rm ?.*
$ rm 2*             $ rm *
```

## Everything Is A File

```
$ cat /dev/null > /var/log/messages # empty a file
$ : > /var/log/messages # no new process
$ ls > /dev/null
$ dd if=/dev/zero of=/tmp/clean bs=1k count=1k
$ dd if=/dev/urandom of=/tmp/random bs=1k count=1k
```

## /proc

Allow higher-level access to driver and kernel information

```
$ cat /proc/cpuinfo
$ cat /proc/meminfo
$ cat /proc/version
$ cat /proc/1/status
# echo 100000 > /proc/sys/kernel/pid_max
```

## Pipe

*Chain processes together*

### Unnamed pipe

```
$ ps aux | sort | less
```

### Named pipe

1. `$ mkfifo mypipe`
2. `$ gzip -9 -c < mypipe > out.gz`
3. `$ cat file > mypipe`

- [https://en.wikipedia.org/wiki/Named\\_pipe](https://en.wikipedia.org/wiki/Named_pipe)

## 3 Shell Programming

### \$ — Give Me The Value Of ...

`$var` Give me the value of variable “var”

`$(echo hello)` Give me the value (output) of command “echo hello”

`$((1+1))` Give me the value (result) of “1+1”

`$$` Give me the value of special variable “\$”

`$?` Give me the value of special variable “?”

`$0` Give me the value of special variable “0”

`@` Give me the value of special variable “@”

### Variables

```
$ a=8; b=2
$ a=a+5; a=$a+5 🙄
$ let a=a+5; let a+=5 😊
$ let b=b+a; let b+=a 😊
$ echo a; echo $a
$ (( a=5, b=6, a+=b )) 😊
$ (( b=a<5?8:9 )) 😊
$ r=$(( RANDOM%100 )) 😊
$ echo "$a" # partial quoting
$ echo '$a' # full quoting
$ a=$(ls -l); echo $a; echo "$a"
$ a=hello; b=world; let a+=b 🙄
```

### Positional Parameters

`$0`, `$1`, `$2`, ..., `$@`,  `$#`

```
1  #!/bin/bash
2
3  echo "You said:"
4
5  echo -e "\t$@"
6  echo
7  echo -e "\targc = $# "
8  echo -e "\targv[0] = $0 "
9
10 i=1
11 for arg in $@; do
12     echo -e "\targv[$i] = $arg"
13     let i++
14 done
```

```
1  #include <stdio.h>
2
3  int main(int argc, char *argv[])
4  {
5      int i;
6      printf("You said:\n\t");
7
8      for(i=1; i<argc; i++)
9          printf("%s ",argv[i]);
10
11     printf("\n\n\targc = %d\n", argc);
12
13     for(i=0; i<argc; i++)
14         printf("\targv[%d] = %s\n",i,argv[i]);
15
16     return 0;
17 }
```

### Parameter Substitution

#### Default value

```
$ echo ${s:=abc}
$ echo ${s:=xyz}

$ echo ${v:-8}
$ echo ${v:-10}
```

## Example

```
1  #!/bin/bash
2
3  echo -n Hello, ${1:-world}
4  echo !
```

## Parameter Substitution

### Substring removal

```
$ for f in *.pbm; do ppm2tiff $f ${f%pbm}.tif; done
```

### Substring replacement

```
$ for f in *.jpg; do mv $f ${f/jpg/JPG}; done
```

## Environmental Variables

### Each process has an environment

\$PATH	\$PWD	\$HOME	\$UID	\$USER
\$GROUPS	\$SHELL	\$TERM	\$DISPLAY	\$TEMP
\$HOSTNAME	\$HOSTTYPE	\$IFS	\$EDITOR	\$BROWSER
\$HISTSIZE	\$FUNCNAME	\$TMOUT	...	

```
export HISTSIZE=2000
export BROWSER='/usr/bin/x-www-browser'
export EDITOR='vim'
export ALTERNATE_EDITOR="vim"
export PDFVIEWER='/usr/bin/zathura'

$ env
$ declare
```

## Tests

```
$ (( 5 < 6 )) && echo should be
$ [[ 1 < 2 ]] && echo of course
$ [[ $a -lt $b ]] && echo yes || echo no
$ if [[ $a -lt $b ]]; then echo yes; else echo no; fi
$ if test $a -lt $b; then echo of course; fi
$ if a = 5; then echo a=$a; fi # whitespace matters ✗
$ if a=5; then echo a=$a; fi ☹
$ if test a=5; then echo a=$a; fi ☹
$ if test a = 5; then echo a=$a; fi ☹
$ if test $a = 5; then echo a=$a; fi ✓
$ test $a = 5 && echo a=$a ✓
$ [[ $a = 5 ]] && echo a=$a ✓
$ [[ cmp a b ]] && echo same file ✗
$ if test cmp a b; then echo same file; fi ✗
$ if cmp a b; then echo same file; fi ✓
$ [[ -f ~/.bash_aliases ]] && . ~/.bash_aliases
$ [[ -x /usr/bin/xterm ]] && /usr/bin/xterm -e tmux &
$ [[ "$pass" != "$MYPASS" ]] && echo 'Wrong password!' && exit 1
$ help test
```



```

1  #!/bin/bash
2
3  words=$@
4  string=linux
5  if echo "$words" | grep -q "$string"
6  then
7      echo "$string found in $words"
8  else
9      echo "$string not found in $words"
10 fi

```

## Loops

```

for ARG in LIST; do COMMAND(s); done
$ for i in 1 2 3; do echo -n i="$i "; done
$ for i in {1..10}; do echo $i; done
$ for i in $(seq 10); do echo $i; done
$ for ((i=1; i<=10; i++)); do echo $i; done
$ for ((i=1, j=1; i<=10; i++, j++)); do
    echo $i-$j 😞
    echo $((i-$j)) 😊
done
$ for ((i=1; i<=10; i++)) { echo $i; } # C style
$ for i in hello world; do echo -n "$i "; done

```

## Loops

```

while CONDITION; do COMMAND(s); done
$ a=0; while [[ a < 10 ]]; do echo $a; ((a++)); done 😞
$ while [[ $a < 10 ]]; do echo $a; ((a++)); done 😞
$ while [[ $a -lt 10 ]]; do echo $a; ((a++)); done ✓
$ while [ $a -lt 10 ]; do echo $a; ((a++)); done ✓
$ while (( a < 10 )); do echo $a; ((a++)); done ✓
$ until (( a = 10 )); do echo $a; ((a++)); done 😞
$ until (( a == 10 )); do echo $a; ((a++)); done ✓
$ while read n; do n2 $n; done
$ while read n; do n2 $n; done < datafile
$ until (( n == 0 )); do read n; n2 $n; done

```

## case

```

1  #!/bin/bash
2
3  [ -z "$1" ] && echo "Usage: `basename $0` [dhb]<number>" && exit 0;
4
5  case "$1" in
6      [dD]*)
7          NUM=$(echo $1 | cut -b 2-)
8          printf "\tDec\tHex\tBin\n"
9          printf "\t%d\t0x%02X\t%s\n" $NUM $NUM $(bc <<< "obase=2;$NUM")
10         ;;
11     [hH]*)
12         NUM=$(echo $1 | cut -b 2-)
13         NUM=$(echo $NUM | tr [:lower:] [:upper:])
14         printf "\tHex\t\tDec\t\tBin\n"
15         printf "\t0x%s\t\t%s\t\t%s\n" $NUM $(bc <<< "ibase=16;obase=A;$NUM") \
16             $(bc <<< "ibase=16;obase=2;$NUM")
17         ;;
18     0[xX]*)
19         NUM=$(echo $1 | cut -b 3-)
20         NUM=$(echo $NUM | tr [:lower:] [:upper:])
21         printf "\tHex\t\tDec\t\tBin\n"
22         printf "\t0x%s\t\t%s\t\t%s\n" $NUM $(bc <<< "ibase=16;obase=A;$NUM") \
23             $(bc <<< "ibase=16;obase=2;$NUM")
24         ;;
25     [bB]*)
26         NUM=$(echo $1 | cut -b 2-)
27         printf "\tBin\t\tHex\t\tDec\n"
28         printf "\t%s\t\t0x%s\t\t%s\n" $NUM $(bc <<< "ibase=2;obase=10000;$NUM") \
29             $(bc <<< "ibase=2;obase=1010;$NUM")
30         ;;
31     *)
32         printf "Dec\tHex\tBin\n"
33         printf "%d\t0x%08X\t%08d\n" $1 $1 $(bc <<< "obase=2;$1")
34         ;;
35 esac

```

## select

```

1  #!/bin/bash
2
3  PS3='Your favorite OS? '
4
5  select OS in "Linux" "Mac OSX" "Windows"
6  do
7      [[ "$OS" = "Linux" ]] && echo wise guy.
8      [[ "$OS" = "Mac OSX" ]] && echo rich guy.
9      [[ "$OS" = "Windows" ]] && echo patient guy.
10     break
11 done

```

## Functions

```

1  #!/bin/bash
2
3  function screencapture(){
4      ffmpeg -f x11grab -s 1920x1080 -r 30 -i :0.0 \
5          -c:v libx264 -crf 0 -preset ultrafast screen.mkv
6  }
7
8  w2pdf(){
9      libreoffice --convert-to pdf:writer_pdf_Export "$1"
10 }
11
12 rfc(){
13     [[ -n "$1" ]] || {
14         cat <<EOF
15         rfc - Command line RFC viewer
16         Usage: rfc <index>
17     EOF
18         return 1
19     }
20     find /usr/share/doc/RFC/ -type f -iname "rfc$1.*" | xargs less
21 }

```

## Array

<https://www.tutorialspoint.com/unix/unix-using-arrays.htm>

# 4 C Programming Basics

## 4.1 Programming Environment

### Program Languages

#### Machine code

The *binary numbers* that the CPUs can understand.

100111000011101111001111 ... and so on ...

#### Assembly language — friendly to humans

People don't think in numbers.

```


1  MOV A,47 ;1010 1111
2  ADD A,B  ;0011 0111
3  HALT     ;0111 0110

```

The ASM programs are translated to machine code by *assemblers*.

#### High level languages

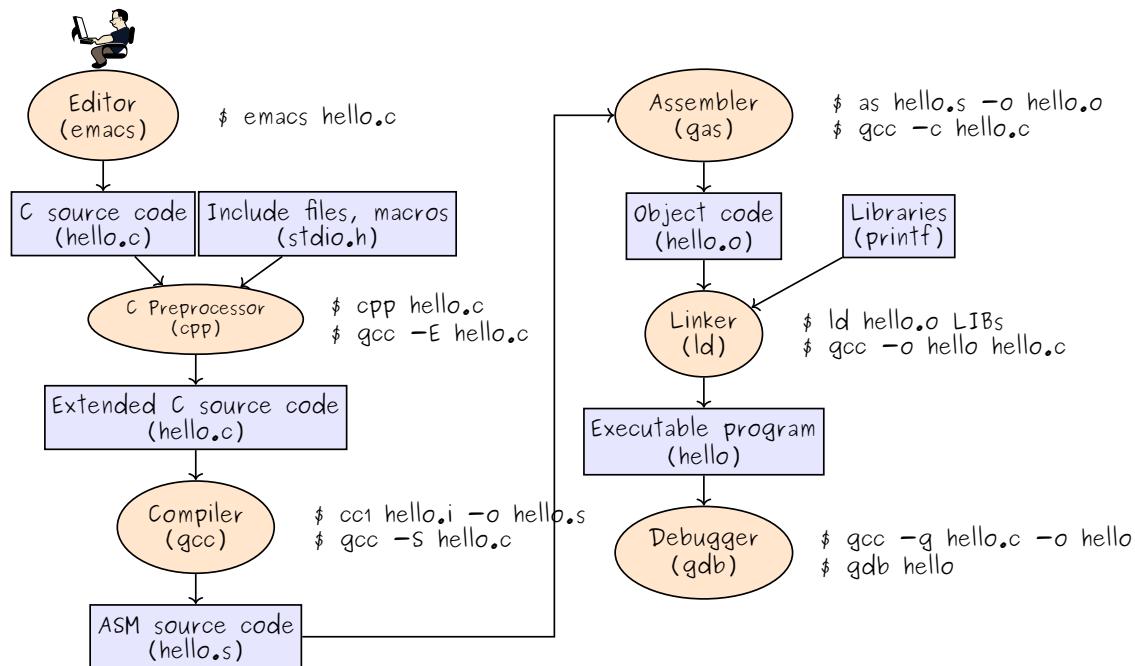
Even easier to understand for humans. Examples:



*Compilers* do the translation work.

### 4.1.1 The Tool Chain

#### Compilation



**Source code** written by programmer in high-level language, in our case in C. We write c source code with a *text editor*, such as emacs, vim, etc.

**Preprocessing** is the first pass of any C compilation. It processes include-files, conditional compilation instructions and macros.

**cpp** The GNU C preprocessor  
\$ gcc -E hello.c

**Compilation** is the second pass. It takes the output of the preprocessor, and the source code, and generates assembly source code.

**gcc/g++** GNU C/C++ compiler  
\$ gcc -S hello.c

**Assembly** is the third stage of compilation. It takes the assembly source code and produces an assembly listing with offsets. The assembler output is stored in an object file.

**as** the portable GNU assembler  
\$ gcc -c hello.c

**Linking** is the final stage of compilation. It combines object code with predefined routines from libraries and produces the executable program.

**ld** The GNU linker  
\$ gcc hello.c -lm

**Wrapper** The whole compilation process is usually not done 'by hand', but using a wrapper program that combines the functions of preprocessor(cpp), compiler(gcc/g++), assembler(as) and linker(ld).

\$ gcc -Wall hello.c -lm -o hello

#### Compiler vs. Interpreter

```
hello.c
#include <stdio.h>
int main()
{
    printf("Hello, world!\n");
    return 0;
}
```

hello.sh

```
$ gcc -o hello hello.c
$ ./hello
```

```
#!/bin/bash
echo 'Hello, world!'
```

```
$ chmod +x hello.sh
$ ./hello.sh
```

```
hello.py
#!/usr/bin/python
print "Hello, world!"
```

```
$ chmod +x hello.py
$ ./hello.py
```

#### 4.1.2 Header Files

##### Header Files

###### Why?

```
#include "add.h"
```

```
int triple(int x)
{
    return add(x, add(x,x));
}
```

- Ensure everyone use the same code
- Easy to share, upgrade, reuse

###### Why not?

```
int add(int, int);
```

```
int triple(int x)
{
    return add(x, add(x, x));
}
```

##### In the header files...

- function declarations
- macro definitions
- constants
- system wide global variables

```
$ ls /usr/include/
```

#### 4.1.3 Library Files

##### Library Files

**Static libraries** .a files. Very old ones, but still alive.

```
$ find /usr/lib -name "*.a"
```

**Shared libraries** .so files. The preferred ones.

```
$ find /usr/lib -name "*.so.*"
```

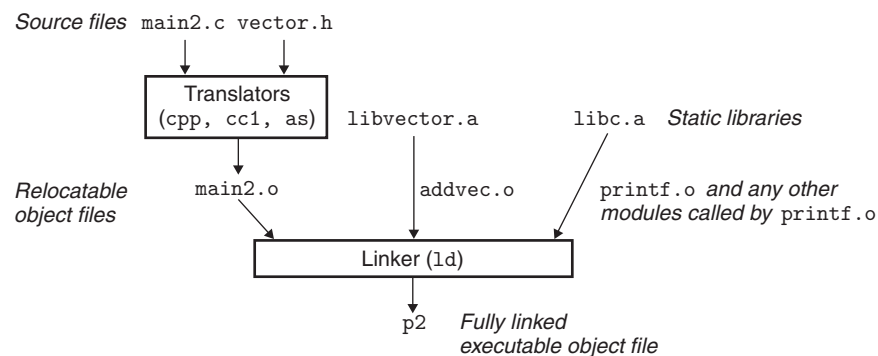
Examples:

```
$ gcc -o hello hello.c /usr/lib/libm.a
```

```
$ gcc -o hello hello.c -lm
```

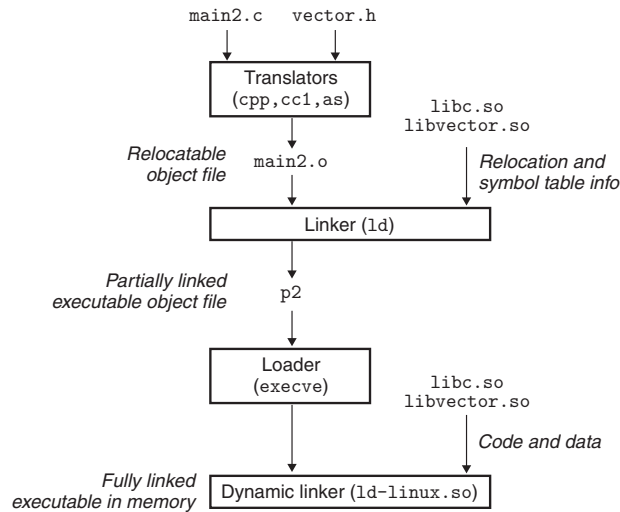
##### Static Linking

- The entire program and all data of a process must be in physical memory for the process to execute
- The size of a process is thus limited to the size of physical memory



##### Dynamic Linking

- Only one copy in memory
- Don't have to re-link after a library update



## Build A Static Library

### Source codes

#### main.c

```

1  #include "lib.h"
2
3  int main(int argc, char* argv[])
4  {
5      int i=1;
6
7      for (; i<argc; i++)
8      {
9          hello(argv[i]);
10         hi(argv[i]);
11     }
12     return 0;
13 }

```

#### lib.h

```

1  #include <stdio.h>
2
3  void hello(char *);
4  void hi(char *);

```

#### hello.c

```

1  #include <stdio.h>
2
3  void hello(char *arg)
4  {
5      printf("Hello, %s!\n", arg);
6  }

```

#### hi.c

```

1  #include <stdio.h>
2
3  void hi(char *arg)
4  {
5      printf("Hi, %s!\n", arg);
6  }

```

## Build A Static Library

### Step by step

1. Get *hello.o* and *hi.o*  
`$ gcc -c hello.c hi.c`
2. Put \*.o into *libhi.a*  
`$ ar crv libhi.a hello.o hi.o`
3. Use *libhi.a*  
`$ gcc main.c libhi.a`

## Build A Static Library

### Makefile

```

1  main: main.c lib.h libhi.a
2      gcc -Wall -o main main.c libhi.a
3
4  libhi.a: hello.o hi.o
5      ar crv libhi.a hello.o hi.o
6
7  hello.o: hello.c
8      gcc -Wall -c hello.c
9
10 hi.o: hi.c
11     gcc -Wall -c hi.c
12
13 clean:
14     rm -f *.o *.a main

```

## Build A Shared Library

Source codes

hello.c

```

1  #include "hello.h"
2
3  int main(int argc, char *argv[])
4  {
5      if (argc != 2)
6          printf ("Usage: %s needs an argument.\n", argv[0]);
7      else
8          hi(argv[1]);
9      return 0;
10 }

```

hello.h

```

1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int hi(char*);

```

hi.c

```

1  #include "hello.h"
2
3  int hi(char* s)
4  {
5      printf ("Hi, %s\n",s);
6      return 0;
7  }

```

## Build A Shared Library

Step by step

1. Get *hi.o*  
\$ gcc -fPIC -c hi.c
2. Get *libhi.so*  
\$ gcc -shared -o libhi.so hi.o
3. Use *libhi.so*  
\$ gcc -L. -Wl,-rpath=. hello.c -lhi
4. Check it  
\$ ldd a.out

## Build A Shared Library

Makefile

```

1  # http://www.cprogramming.com/tutorial/shared-libraries-linux-gcc.html
2  # http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html
3  #
4  # gcc -fPIC -c hi.c
5  # gcc -shared -o libhi.so hi.o
6  # gcc -L/current/dir -Wl,option -Wall -o hello hello.c -lhi
7  #
8  # -L          - tells ld where to search libraries
9  # -Wl,option - pass option as an option to the linker (ld)
10 # -rpath=dir - Add a directory to the runtime library search path
11
12 hello: hello.c hello.h libhi.so
13         gcc -L. -Wl,-rpath=. -Wall -o hello hello.c -lhi
14 libhi.so: hi.o hello.h
15         gcc -shared -o libhi.so hi.o
16 hi.o: hi.c hello.h
17         gcc -fPIC -c hi.c
18 clean:
19         rm *.o *.so hello

```


## GNU C Library

Linux API > POSIX API

```

$ man 7 libc
$ man 3 intro
$ man gcc
$ info gcc
🔴 sudo apt install gcc-doc

```



### 4.1.4 Error Handling

errno.h

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/stat.h>
4  #include <fcntl.h>
5
6  int main(int argc, char *argv[])
7  {
8      if ( open(argv[1], O_RDONLY) == -1 ){
9          perror("open");
10         exit(EXIT_FAILURE);
11     }
12     return 0;
13 }

```

```

$ man errno
$ man errno.h
$ man perror
• [Advanced programming in the UNIX environment, Sec. 1.7]
• https://stackoverflow.com/questions/30078281/raise-error-in-a-bash-script

```

### 4.1.5 The Make Utility

#### The Make Utility


To compile a single C program:

```
$ gcc hello.c -o hello
```

✓ OK. But...

What if you have a large project with 1000+ files?





**Linux 4.9 source tree:** 3799 directories, 55877 files

**make:** help you maintain your programs.

## Makefile

```
1 target: dependencies
2 |-----TAB--> command
```

## Example

```
1 hello: hello.c
2 |-----TAB--> gcc -o hello hello.c
```

\$ info make makefiles

## Makefile

```
1 edit: main.o kbd.o command.o display.o \
2       insert.o search.o files.o utils.o
3       gcc -Wall -o edit main.o kbd.o command.o display.o \
4       insert.o search.o files.o utils.o
5
6 main.o: main.c defs.h
7       gcc -c -Wall main.c
8 kbd.o : kbd.c defs.h command.h
9       gcc -c -Wall kbd.c
10 command.o: command.c defs.h command.h
11       gcc -c -Wall command.c
12 display.o : display.c defs.h buffer.h
13       gcc -c -Wall display.c
14 insert.o: insert.c defs.h buffer.h
15       gcc -c -Wall insert.c
16 search.o: search.c defs.h buffer.h
17       gcc -c -Wall search.c
18 files.o: files.c defs.h buffer.h command.h
19       gcc -c -Wall files.c
20 utils.o: utils.c defs.h
21       gcc -c -Wall utils.c
22
23 clean:
24       rm edit main.o kbd.o command.o display.o \
25       insert.o search.o files.o utils.o
```

```
./
├── command.c
├── display.c
├── files.c
├── insert.c
├── kbd.c
├── main.c
├── search.c
├── utils.c
├── buffer.h
├── command.h
├── defs.h
└── Makefile
```

## 4.1.6 Version Control

### git

#### To create a new local git repo

In your source code directory, do:

```
$ git init
$ git add .
$ git commit -m "something to say..."
```

#### To clone a remote repo

Example:



```
$ git clone https://github.com/wx672/lecture-notes.git
$ git clone https://github.com/wx672/dotfile.git
```

#### Most commonly used git Commands

```
$ git add filename[s]
$ git rm filename[s]
$ git commit
```

```
$ git status
$ git log
$ git diff
$ git push
$ git pull
$ git help {add,rm,commit,...}
```

```
$ man gittutorial
$ man gittutorial-2
```

```
 sudo apt install git
 https://github.com
```

#### 4.1.7 Manual Pages

##### Man page

##### Layout

```
1 NAME
2     A one-line description of the command.
3 SYNOPSIS
4     A formal description of how to run it and what
5     command line options it takes.
6 DESCRIPTION
7     A description of the functioning of the command.
8 EXAMPLES
9     Some examples of common usage.
10 SEE ALSO
11     A list of related commands or functions.
12 BUGS
13     List known bugs.
14 AUTHOR
15     Specify your contact information.
16 COPYRIGHT
17     Specify your copyright information.
```

##### Man Page

##### Groff source code

```
1 .\" Text automatically generated by txt2man
2 .TH untitled "06 August 2019" "" ""
3 .SH NAME
4 \fBA one-line description of the command.
5 .SH SYNOPSIS
6 .nf
7 .fam C
8 \fBA formal description of how to run it and what command line options it takes.
9 .fam T
10 .fi
11 .fam T
12 .fi
13 .SH DESCRIPTION
14 \fBA description of the functioning of the command.
15 .SH EXAMPLES
16 Some examples of common usage.
17 .SH SEE ALSO
18 \fBA list of related commands or functions.
19 .SH BUGS
20 List known bugs.
21 .SH AUTHOR
22 Specify your contact information.
23 .SH COPYRIGHT
24 Specify your copyright information.
```

```
$ man 7 groff
$ man txt2man
$ man a2x
$ ls /usr/share/man
```

#### 4.1.8 A Sample GNU Package


##### How to “Do one thing, and do it well”?

```
$ apt source hello
```

## 4.2 OS Basics

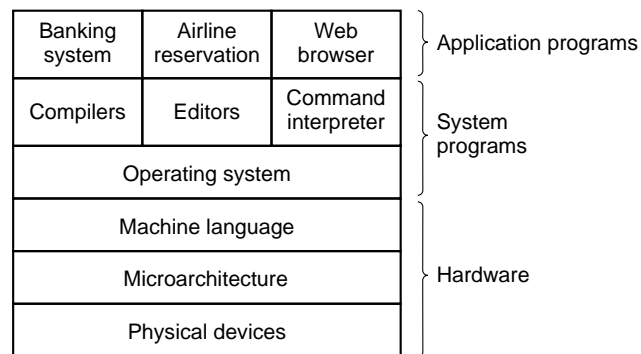
### Abstractions

To hide the complexity of the actual implementations



See also: [Computer Systems: A Programmer's Perspective, Sec. 1.9.2, The Importance of Abstractions in Computer Systems].

### A Computer System



#### 4.2.1 Hardware

##### CPU Working Cycle



1. Fetch the first instruction from memory
2. Decode it to determine its type and operands
3. execute it

##### Special CPU Registers

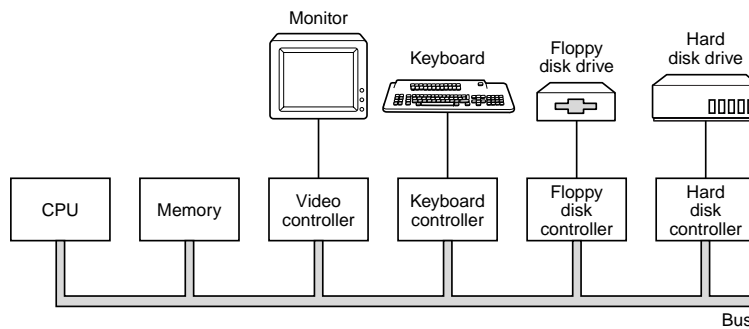
**Program counter (PC):** keeps the memory address of the next instruction to be fetched

**Stack pointer (SP):** points to the top of the current stack in memory

**Program status (PS):** holds

- condition code bits
- processor state

##### System Bus



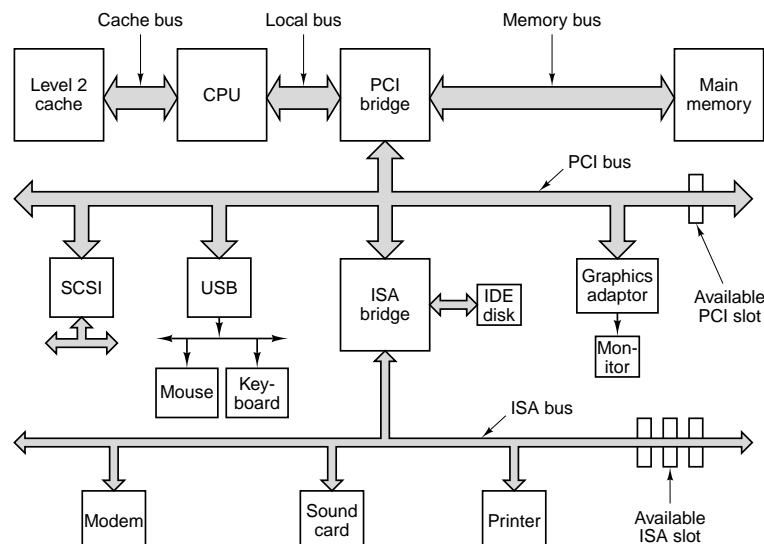
**Address Bus:** specifies the memory locations (addresses) for the data transfers

**Data Bus:** holds the data transferred. Bidirectional

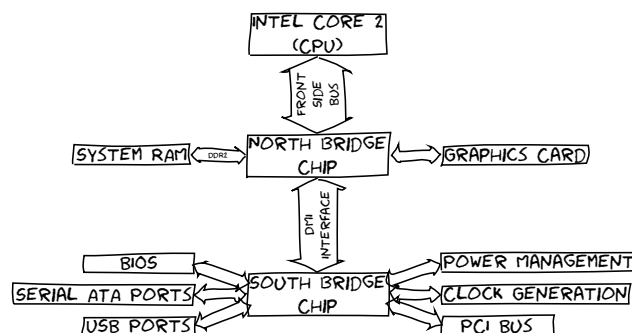
**Control Bus:** contains various lines used to route timing and control signals throughout the system

### Controllers and Peripherals

- Peripherals are real devices controlled by controller chips
- Controllers are processors like the CPU itself, have control registers
- Device driver writes to the registers, thus control it
- Controllers are connected to the CPU and to each other by a variety of buses



### Motherboard Chipsets



See also: *Motherboard Chipsets And The Memory Map* <sup>1</sup>.

- The CPU doesn't know what it's connected to

<sup>1</sup><http://duartes.org/gustavo/blog/post/motherboard-chipsets-memory-map>

- CPU test bench? network router? toaster? brain implant?
- The CPU talks to the outside world through its pins
  - some pins to transmit the physical memory address
  - other pins to transmit the values
- The CPU's gateway to the world is the *front-side bus*

### Intel Core 2 QX6600

- 33 pins to transmit the physical memory address
  - so there are  $2^{33}$  choices of memory locations
- 64 pins to send or receive data
  - so data path is 64-bit wide, or 8-byte chunks

This allows the CPU to physically address 64GB of memory ( $2^{33} \times 8B$ )

See also: *Datasheet for Intel Core 2 Quad-Core Q6000 Sequence* <sup>2</sup>.

### Some physical memory addresses are mapped away!

- only the addresses, not the spaces
- Memory holes
  - 640 KiB ~ 1 MiB
  - /proc/iomem

### Memory-mapped I/O

- BIOS ROM
- video cards
- PCI cards
- ...

This is why 32-bit OSes have problems using 4 GiB of RAM.

0xFFFFFFFF	Reset vector	JUMP to 0xF0000	4GB
0xFFFFFFF0		Unaddressable memory, real mode is limited to 1MB.	4GB-16B
0x100000		System BIOS	1MB
0xF0000		Ext. System BIOS	960KB
0xE0000		Expansion Area (maps ROMs for old peripheral cards)	896KB
0xC0000		Legacy Video Card Memory Access	768KB
0xA0000		Accessible RAM (640KB is enough for anyone – old DOS area)	640KB
0			0

### the northbridge

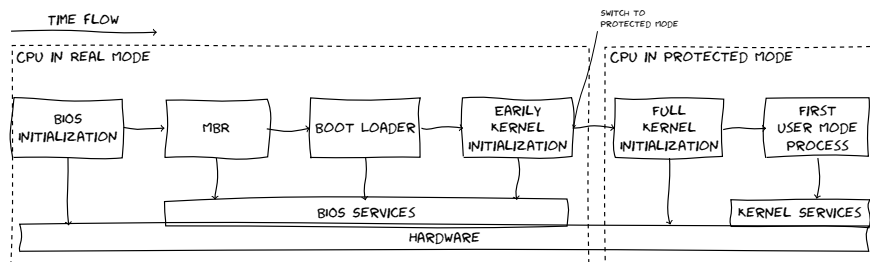
1. receives a physical memory request
  2. decides where to route it
    - to RAM? to video card? to ...?
    - decision made via the *memory address map*
- When is the memory address map built? `setup()`.

## 4.2.2 Bootstrapping

### Bootstrapping

#### Can you pull yourself up by your own bootstraps?


A computer cannot run without first loading software but must be running before any software can be loaded.




### Intel x86 Bootstrapping

1. BIOS (0xfffffff0)
  - ➡ POST ➡ HW init ➡ Find a boot device (FD,CD,HD...) ➡ Copy sector zero (MBR) to RAM (0x00007c00)
2. MBR – the first 512 bytes, contains
  - Small code ( $< 446B$ ), e.g. GRUB stage 1, for loading GRUB stage 2
  - the primary partition table ( $16 \times 4 = 64B$ )

<sup>2</sup><http://download.intel.com/design/processor/datashts/31559205.pdf>

- its job is to load the second-stage boot loader.
3. GRUB stage 2 — load the OS kernel into RAM
  4.  startup
  5. init — the first user-space process



```
$ sudo hd -n512 /dev/sda
```

### 4.2.3 Interrupt

#### Why Interrupt?

While a process is reading a disk file, can we do...

```
1 while(!done_reading_a_file())
2 {
3     let_CPU_wait();
4     // or...
5     lend_CPU_to_others();
6 }
7 operate_on_the_file();
```

#### Modern OS are Interrupt Driven

**HW INT** by sending a signal to CPU

**SW INT** by executing a *system call*

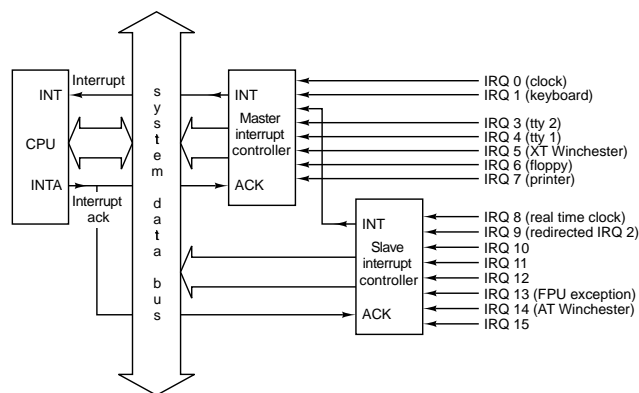
**Trap (exception)** is a software-generated INT caused by an error or by a specific request from an user program

**Interrupt vector** is an array of pointers to the memory addresses of *interrupt handlers*. This array is indexed by a unique device number


```
$ less /proc/devices
```

```
$ less /proc/interrupts
```

#### Programmable Interrupt Controllers




#### Interrupt Processing



Detailed explanation: in [Modern Operating Systems, Sec. 1.3.5, I/O Devices].

## Interrupt Timeline



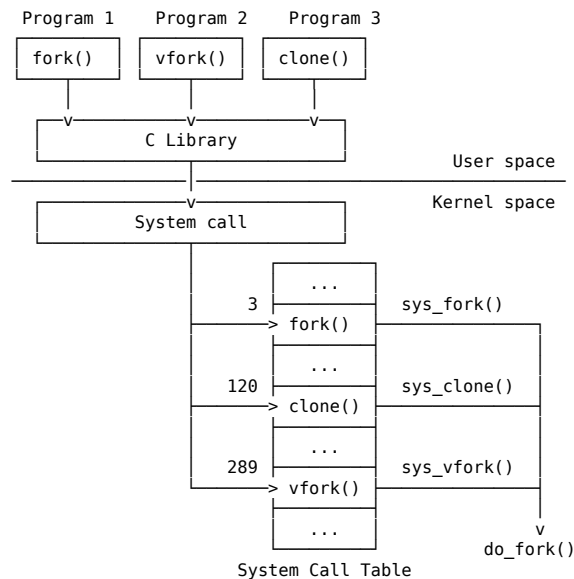
## 4.2.4 System Calls

### System Calls

#### A System Call


- is how a program requests a service from an OS kernel
- provides the interface between a process and the OS

```
$ man 2 intro
$ man 2 syscalls
```



## The 11 steps in making a system call

*read(fd, buffer, nbytes)*



## Example

Linux `INT 80h`

**Interrupt Vector Table:** The very first 1KiB of x86 memory.

- 256 entries  $\times$  4B = 1KiB
- Each entry is a complete memory address (segment:offset)
- It's populated by Linux and BIOS
- Slot 80h: address of the kernel services dispatcher (☛ sys-call table)

## Example

```

1  Msg: db 'Hello, world'
2  MsgLen: equ $-Msg
3  mov eax,4      ; sys_write syscall = 4
4  mov ebx,1      ; 1 = STDOUT
5  mov ecx,Msg    ; offset of the message
6  mov edx,MsgLen ; length of string
7  int 80h        ; call the kernel

```

```

$ nasm -f elf64 hello.asm -o hello.o
$ ld hello.o -o hello
$ ./hello

```



Process management	
Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &amp;statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

File management	
Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &amp;buf)</code>	Get a file's status information

Directory and file system management	
Call	Description
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

Miscellaneous	
Call	Description
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&amp;seconds)</code>	Get the elapsed time since Jan. 1, 1970

Fig. 1-18. Some of the major POSIX system calls. The return code *s* is  $-1$  if an error has occurred. The return codes are as follows: *pid* is a process id, *fd* is a file descriptor, *n* is a byte count, *position* is an offset within the file, and *seconds* is the elapsed time. The parameters are explained in the text.

## System Call Examples

`fork()`

```

1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main ()
5  {
6      printf("Hello World!\n");
7      fork();
8      printf("Goodbye Cruel World!\n");
9      return 0;
10 }
```

`$ man 2 fork`

`execve()`

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("Hello World!\n");
7     if( fork() != 0 )
8         printf("I am the parent process.\n");
9     else {
10         printf("A child is listing the directory contents...\n");
11         execl("/bin/ls", "ls", "-al", NULL);
12     }
13     return 0;
14 }

```

\$ man 2 execve

\$ man 3 exec

Quoted from [stackoverflow](#): What is the difference between the functions of the *exec* family of system calls:

There is no *exec* system call — this is usually used to refer to all the *execXX* calls as a group. They all do essentially the same thing: loading a new program into the current process, and provide it with arguments and environment variables. The differences are in how the program is found, how the arguments are specified, and where the environment comes from.

- The calls with *v* in the name take an array parameter to specify the `argv[]` array (*vector*) of the new program.
- The calls with *l* in the name take the arguments of the new program as a variable-length argument *list* to the function itself.
- The calls with *e* in the name take an extra argument to provide the *environment* of the new program; otherwise, the program inherits the current process's environment.
- The calls with *p* in the name search the *PATH* environment variable to find the program if it doesn't have a directory in it (i.e. it doesn't contain a / character). Otherwise, the program name is always treated as a path to the executable.

## 5 The Linux Environment

### Command Line Options

*getopt.c*

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main(int argc, char* argv[]) {
5     int opt;
6
7     while ((opt = getopt(argc, argv, "hf:l")) != -1) {
8         switch (opt) {
9             case 'h':
10                 printf("Usage: %s [-h] [-f file] [-l]\n", argv[0]);
11                 break;
12             case 'l':
13                 printf("option: %c\n", opt);
14                 break;
15             case 'f':
16                 printf("filename: %s\n", optarg);
17                 break;
18             }
19     }
20     return 0;
21 }

```

```
$ man 3 getopt
```

## Command Line Options

*getopt.sh*

```
1  #!/bin/bash
2
3  while getopts hf:l OPT; do
4      case $OPT in
5          h) echo "usage: `basename $0` [-h] [-f file] [-l]"
6              exit 1 ;;
7          l) echo "option: l" ;;
8          f) echo "filename: $OPTARG" ;;
9      esac
10 done
```

```
$ ./getopt.sh -h
$ ./getopt.sh -lf filename
$ ./getopt.sh -l -f filename
$ ./getopt.sh -f filename -l
```

## Environment Variable

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
4  extern char** environ;
5
6  int main() {
7      char** env = environ;
8
9      while (*env) {
10         printf("%s\n", *env);
11         env++;
12     }
13
14     return 0;
15 }
```

\$ env  
\$ man 3 getenv  
\$ man 3 putenv

[Beginning linux programming, p.147, Sec. The environ Variable]

## Time and Date

```
1  #include <time.h>
2  #include <stdio.h>
3
4  int main(void)
5  {
6      time_t t = time(NULL); /* long int */
7
8      printf("epoch time:\t%ld\n",t);
9      printf("calendar time:\t%s", ctime(&t));
10
11     return 0;
12 }
```

- January 1 1970 — start of the Unix epoch

```
$ man 3 time
```

```
$ man 3 ctime
```

## Temporary Files

### mkstemp.c

```
1  #include <stdlib.h>
2  #include <unistd.h>
3  #define _GNU_SOURCE
4  #include <stdio.h>
5
6  int main(int argc, char *argv[])
7  {
8      char c, *f;
9
10     asprintf(&f, "%sXXXXXX", argv[1]);
11     int tmp = mkstemp(f);
12
13     while ( read(0, &c, 1) == 1)
14         write(tmp, &c, 1);
15
16     unlink(f);
17     free(f);
18     return 0;
19 }
```

### mktemp.sh

```
1  #!/bin/bash
2
3  tmp=$(mktemp)
4
5  while read LINE; do
6      echo $LINE >> $tmp
7  done
8
9
$ man 3 mkstemp
$ man 3 tmpfile
$ man 3 asprintf
```

## Logging

### syslog.c

```
1  #include <syslog.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4
5  int main(int argc, char *argv[])
6  {
7      if ( open(argv[1], O_RDONLY) < 0 )
8          syslog(LOG_ERR | LOG_USER, "%s - %m\n", argv[1]);
9      else
10         syslog(LOG_INFO | LOG_USER, "%s - %m\n", argv[1]);
11     return 0;
12 }
```

### logger.sh

```
1  #!/bin/bash
2
3  [[ -f "$1" ]] && logger "$1 exists." || logger "$1 not found."
```

## 6 Working With Files

### 6.1 File

#### File

*A logical view of information storage*

#### User's view

A file is the smallest storage unit on disk.

- Data cannot be written to disk unless they are within a file

## UNIX view

Each file is a sequence of 8-bit bytes

- It's up to the application program to interpret this byte stream.

## File

What is stored in a file?

Source code, object files, executable files, shell scripts, PostScript...

## Different type of files have different structure

- UNIX looks at contents to determine type
  - Shell scripts** start with “#!”
  - PDF** start with “%PDF...”
  - Executables** start with *magic number*
- Windows uses file naming conventions
  - executables** end with “.exe” and “.com”
  - MS-Word** end with “.doc”
  - MS-Excel** end with “.xls”

## File Types

**Regular files:** ASCII, binary

**Directories:** Maintaining the structure of the FS

**In UNIX, everything is a file**

**Character special files:** I/O related, such as terminals, printers ...

**Block special files:** Devices that can contain file systems, i.e. disks

Disks — logically, linear collections of blocks; disk driver translates them into physical block addresses


## File Operations

POSIX file system calls

<code>creat(name, mode)</code>	<code>read(fd, buffer, byte_count)</code>
<code>open(name, flags)</code>	<code>write(fd, buffer, byte_count)</code>
<code>close(fd)</code>	<code>lseek(fd, offset, whence)</code>
<code>link(oldname, newname)</code>	<code>chown(name, owner, group)</code>
<code>unlink(name)</code>	<code>fchown(fd, owner, group)</code>
<code>truncate(name, size)</code>	<code>chmod(name, mode)</code>
<code>ftruncate(fd, size)</code>	<code>fchmod(fd, mode)</code>
<code>stat(name, buffer)</code>	<code>utimes(name, times)</code>
<code>fstat(fd, buffer)</code>	

## File System Implementation

### A typical file system layout



On-Disk Information Structure

Boot block a MBR copy

Superblock Contains volume details

number of blocks	size of blocks
free-block count	free-block pointers
free FCB count	free FCB pointers

I-node Organizes the files FCB (File Control Block), contains file details (metadata).


Superblock

Keeps information about the file system

- Type — ext2, ext3, ext4...
- Size
- Status — how it's mounted, free blocks, free inodes, ...
- Information about other metadata structures

```
$ sudo dumpe2fs /dev/sda1 | less
```


Implementing Files



Contiguous Allocation

- 😊 simple
- 😊 good for read only

☹ fragmentation



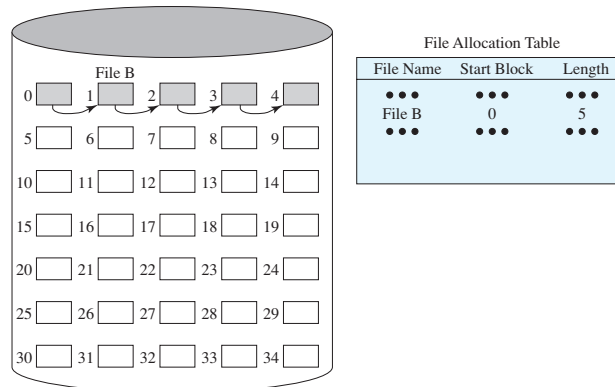
Linked List (Chained) Allocation

A pointer in each disk block

- 😊 no waste block
- 😞 slow random access

⊗ not  $2^n$


**Linked List (Chained) Allocation** Though there is no external fragmentation, consolidation is still preferred.



FAT: Linked list allocation with a table in RAM

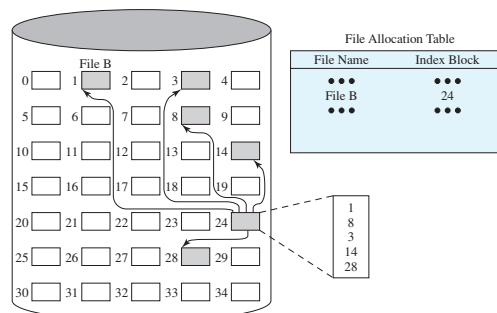
- Taking the pointer out of each disk block, and putting it into a table in memory
- fast random access (chain is in RAM)
- is  $2^n$
- the entire table must be in RAM

$$disk \nearrow \Rightarrow FAT \nearrow \Rightarrow RAM_{used} \nearrow$$



See also: [File Allocation Table — Wikipedia, The Free Encyclopedia, Wikipedia:FAT].

## Indexed Allocation

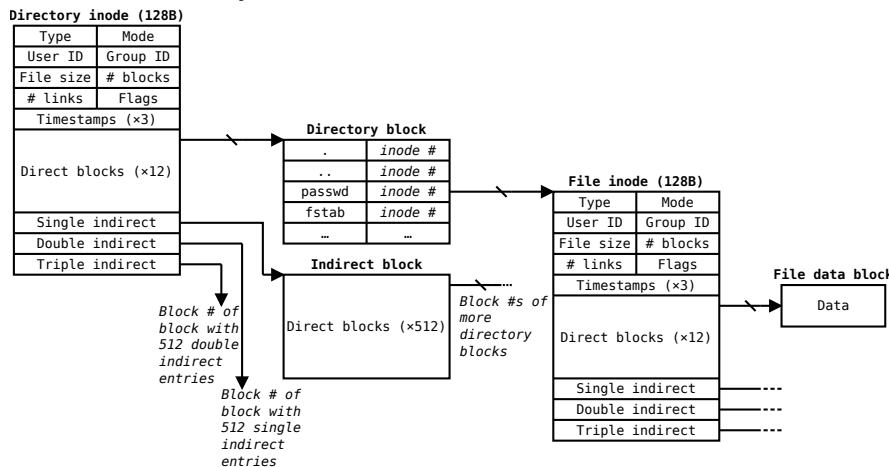


**I-node** A data structure for each file. An i-node is in memory *only if* the file is open

$$files_{opened} \nearrow \Rightarrow RAM_{used} \nearrow$$

See also: [Inode — Wikipedia, The Free Encyclopedia, Wikipedia:inode].

## UNIX Treats a Directory as a File

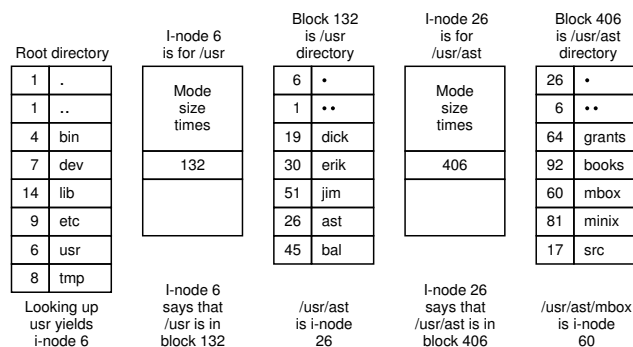


`open()`

**Why?** To avoid constant searching

- Without `open()`, every file operation involves searching the directory for the file.

The steps in looking up `/usr/ast/mbox`



`fd open(pathname, flags)`

A per-process *open-file table* is kept in the OS

- upon a successful `open()` syscall, a new entry is added into this table
- indexed by *file descriptor (fd)*
- `close()` to remove an entry from the table

To see files opened by a process, e.g. `init`

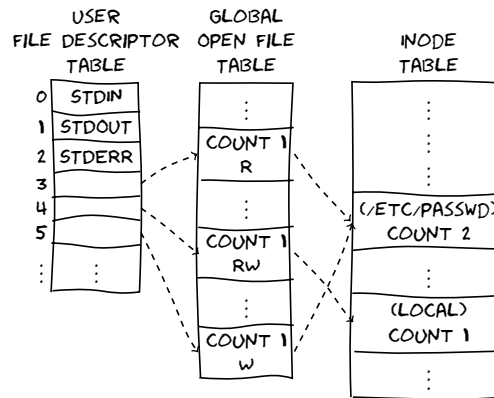
`$ lsof -p 1`

`$ man 2 open`

**A process executes the following code:**


```
fd1 = open("/etc/passwd", O_RDONLY);
fd2 = open("local", O_RDWR);
fd3 = open("/etc/passwd", O_WRONLY);
```





### One more process B:

```
fd1 = open("/etc/passwd", O_RDONLY);
fd2 = open("private", O_RDONLY);
```



**write()**

```
1 #include <unistd.h>
2
3 int main(void)
4 {
5     write(1, "Hello, world!\n", 14);
6
7     return 0;
8 }
```

```
$ man 2 write
$ man 3 write
```

**read()**

```
1 #include <unistd.h>
2
3 int main(void)
4 {
5     char buffer[10];
6
7     read(0, buffer, 10);
8
9     write(1, buffer, 10);
10
11     return 0;
12 }
```

```
$ man 2 read
$ man 3 read
```

cp

```

1  #include <sys/types.h>   /* include necessary header files */
2  #include <fcntl.h>
3  #include <stdlib.h>
4  #include <unistd.h>
5
6  #define BUF_SIZE 4096    /* use a buffer size of 4096 bytes */
7  #define OUTPUT_MODE 0700 /* protection bits for output file */
8
9  int main(int argc, char *argv[])
10 {
11     int in, out, rbytes, wbytes;
12     char buf[BUF_SIZE];
13
14     if (argc != 3) exit(1);
15
16     if ( (in = open(argv[1], O_RDONLY)) < 0 ) exit(2); /* open source file */
17
18     if ( (out = creat(argv[2], OUTPUT_MODE)) < 0 ) exit(3); /* create destination file */
19
20     while (1) { /* Copy loop */
21         if ( (rbytes = read(in, buf, BUF_SIZE)) <= 0 ) break; /* read a block of data */
22         if ( (wbytes = write(out, buf, rbytes)) <= 0 ) exit(4); /* write data */
23     }
24
25     close(in);
26     close(out);
27     if (rbytes == 0) exit(0); /* no error on last read */
28     else exit(5);           /* error on last read */
29 }

```


## stdio — The Standard I/O Library

**System calls:** open(), read(), write(), close()...

**Library functions:** fopen(), fread(), fwrite, fclose()...

**Avoid calling syscalls directly as much as you can**

- Portability
- Buffered I/O



**open() vs. fopen()**

open()

```
1  #include <unistd.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4  #include <stdio.h>
5
6  int main()
7  {
8      char c;
9      int in;
10     in = open("/tmp/1m.test", O_RDONLY);
11
12     while (read(in, &c, 1) == 1);
13
14     return 0;
15 }
```

\$ strace -c ./open

\$ dd if=/dev/zero of=/tmp/1m.test bs=1k count=1024

<https://stackoverflow.com/questions/1658476/c-fopen-vs-open>

fopen() — Buffered I/O

```
1  #include <stdio.h>
2
3  int main(void)
4  {
5      FILE *stream;
6
7      stream = fopen("/tmp/1m.test", "r");
8
9      while ( fgetc(stream) != EOF );
10
11     fclose(stream);
12
13     return 0;
14 }
```

\$ strace -c ./fopen

cp — With stdio


```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char *argv[])
5  {
6      FILE *in, *out;
7      int c=0;
8
9      if (argc != 3) exit(1);
10
11     in = fopen(argv[1], "r");
12     out = fopen(argv[2], "w");
13
14     while ( (c = fgetc(in)) != EOF )
15         fputc(c, out);
16
17     return 0;
18 }
```

**Homework:** Try fread()/fwrite() instead.

<https://stackoverflow.com/questions/32742430/is-getc-a-macro-or-a-function> | <https://stackoverflow.com/questions/9104568/macro-vs-function-in-c>

## 6.2 Directory

### Implementing Directories



- (a) A simple directory (Windows)
  - fixed size entries
  - disk addresses and attributes in directory entry
- (b) Directory in which each entry just refers to an i-node (UNIX)

### Directory entry in glibc

```

1 struct dirent {
2     ino_t      d_ino;      /* Inode number */
3     off_t      d_off;      /* Not an offset; see below */
4     unsigned short d_reclen; /* Length of this record */
5     unsigned char d_type;   /* Type of file; not supported
6                             by all filesystem types */
7     char        d_name[256]; /* Null-terminated filename */
8 };

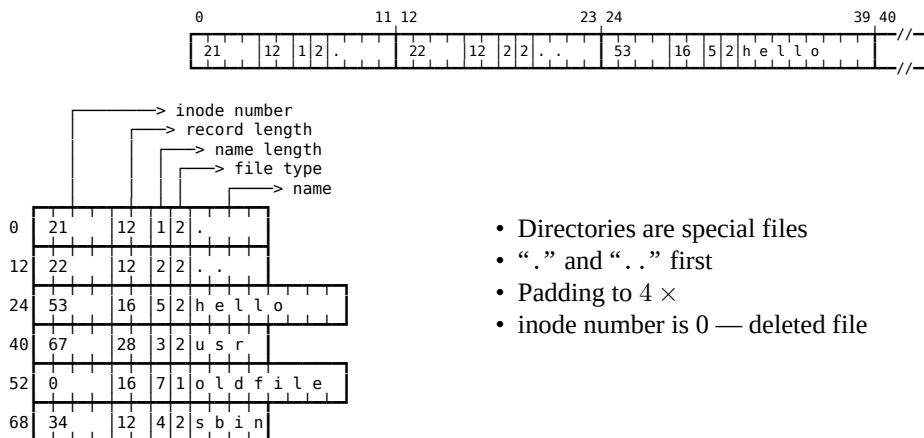
```

```

$ man readdir
$ view /usr/include/x86_64-linux-gnu/bits/dirent.h

```

### Ext2 Directories



- Directories are special files
- “.” and “..” first
- Padding to 4 ×
- inode number is 0 — deleted file

### ls

```

1 #include <sys/types.h>
2 #include <dirent.h>
3 #include <stddef.h>
4 #include <stdio.h>
5
6 int main(int argc, char *argv[])
7 {
8     DIR *dp;
9     struct dirent *entry;
10
11     dp = opendir(argv[1]);
12
13     while ( (entry = readdir(dp)) != NULL ){
14         printf("%s\n", entry->d_name);
15     }
16
17     closedir(dp);
18
19     return 0;
20 }

```

### The real ls.c?

- 116 A4 pages
- 5308 lines

Do one thing, and do it really well.

\$ apt source coreutils

mkdir(), chdir(), rmdir(), getcwd()

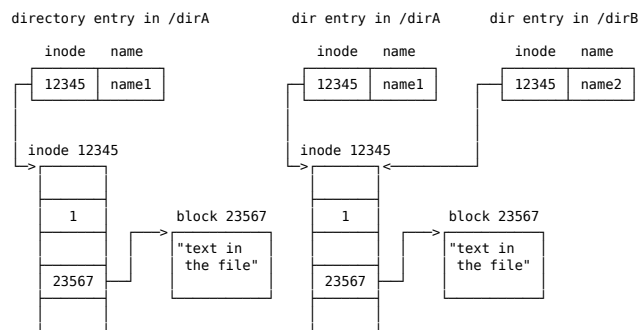
```

1  #include <sys/stat.h>
2  #include <sys/types.h>
3  #include <unistd.h>
4  #include <stdio.h>
5
6  int main(int argc, char *argv[])
7  {
8      char s[100];
9      if ( mkdir(argv[1], S_IRUSR|S_IXUSR) == 0 )
10         chdir(argv[1]);
11         printf("PWD = %s\n", getcwd(s,100));
12         rmdir(argv[1]);
13         return 0;
14     }


```

## Hard Links

Hard links → the same inode




## Drawback



## Symbolic Links

A symbolic link has its own inode → a directory entry



**Fast symbolic link:** Short path name (< 60 chars) needs no data block. Can be stored in the 15 pointer fields

link(), unlink(), symlink()

```
1  #include <unistd.h>
2  #include <stdio.h>
3
4  int main(int argc, char *argv[])
5  {
6      /* link(argv[1], argv[2]); */
7      /* symlink(argv[1], argv[2]); */
8      unlink(argv[1]);
9      perror(argv[0]);
10     return 0;
11 }
```

## 7 Processes and Threads

### 7.1 Process

#### Process


A **process** is an instance of a program in execution

**Processes are like human beings:**

- they are generated
- they have a life
- they optionally generate one or more child processes, and
- eventually they die

A small difference:


- sex is not really common among processes
- each process has just one parent



#### Process Control Block (PCB)


##### Implementation

A process is *the collection of data structures* that fully describes how far the execution of the program has progressed.

- Each process is represented by a *PCB*
- `task_struct` in 


process state
PID
program counter
registers
memory limits
list of open files
...

#### Process Creation



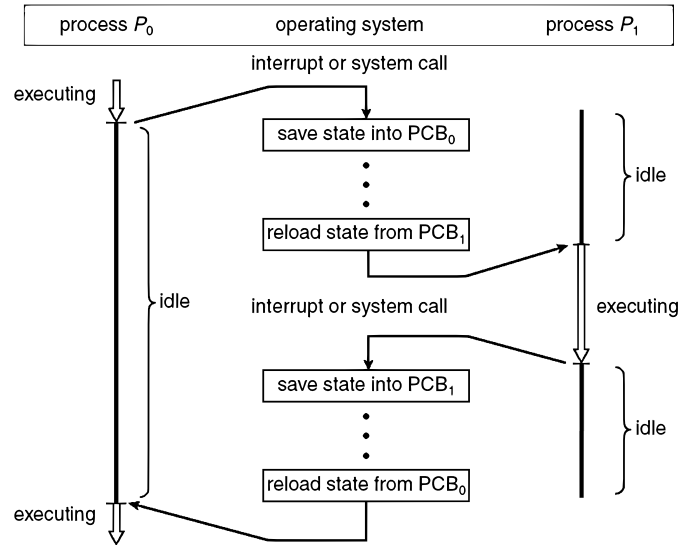
- When a process is created, it is almost identical to its parent
  - It receives a (logical) copy of the parent's address space, and
  - executes the same code as the parent
- The parent and child have separate copies of the data (stack and heap)

## Process State Transition



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

## CPU Switch From Process To Process



## Forking in C

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main ()
5 {
6     printf("Hello World!\n");
7     fork();
8     printf("Goodbye Cruel World!\n");
9     return 0;
10 }
```

`exec()`

```

1  #include <unistd.h>
2
3  int main(void) {
4      char *argv[] = {NULL};
5      char *envp[] = {NULL};
6      execve("callee.out", argv, envp);
7  }

```

Fig. 1: caller.c

```

1  #include <stdio.h>
2
3  int main(int argc, char **argv) {
4      if (argc == 0 && argv[0] == NULL)
5          puts("yup");
6  }

```

Fig. 2: callee.c

```

$ gcc -Wall caller.c -o caller.out
$ gcc -Wall callee.c -o callee.out
$ ./caller.out ./callee.out

```

Fig. 3: argv[0] can be NULL

```

1  int main()
2  {
3      pid_t pid;
4      /* fork another process */
5      pid = fork();
6      if (pid < 0) { /* error occurred */
7          fprintf(stderr, "Fork Failed");
8          exit(-1);
9      }
10     else if (pid == 0) { /* child process */
11         execlp("/bin/ls", "ls", NULL);
12     }
13     else { /* parent process */
14         /* wait for the child to complete */
15         wait(NULL);
16         printf ("Child Complete");
17         exit(0);
18     }
19     return 0;
20 }

```

**More about argv[0]** `int execve(const char *pathname, char *const argv[], char *const envp[]);`

- pathname should be the binary image of a program. Or it can be a script (man 2 execve);
- argv[0] is the new process name, usually the same as the basename of pathname, though it can be any other string. It can even be NULL (see Figure 3 for example).

The fact that argv[0] contains the name used to invoke the program can be employed to perform a useful trick. We can create multiple links to (i.e., names for) the same program, and then have the program look at argv[0] and take different actions depending on the name used to invoke it. An example of this technique is provided by the *gzip(1)*, *gunzip(1)*, and *zcat(1)* commands, all of which are links to the same executable file. [The Linux Programming Interface: A Linux and UNIX System Programming Handbook, Sec. 6.6]

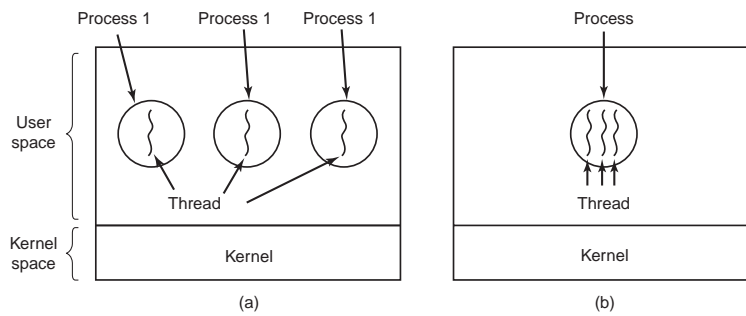
- <https://stackoverflow.com/questions/2794150/when-can-argv0-have-null>
- <https://stackoverflow.com/questions/36673765/why-can-the-execve-system-call-run-bin-sh-without-any-argv-arguments-but-not>

## 7.2 Thread

### Process vs. Thread

a single-threaded process = resource + execution  
a multi-threaded process = resource + executions





**A process** = a unit of resource ownership, used to group resources together;

**A thread** = a unit of scheduling, scheduled for execution on the CPU.

## Threads

code, data, open files, signals...		
thread ID	thread ID	thread ID
program counter	program counter	program counter
register set	register set	register set
stack	stack	stack

## POSIX Threads

**IEEE 1003.1c** The standard for writing portable threaded programs. The threads package it defines is called *Pthreads*, including over 60 function calls, supported by most UNIX systems.

### Some of the Pthreads function calls

Thread call	Description
<code>pthread_create</code>	Create a new thread
<code>pthread_exit</code>	Terminate the calling thread
<code>pthread_join</code>	Wait for a specific thread to exit
<code>pthread_yield</code>	Release the CPU to let another thread run
<code>pthread_attr_init</code>	Create and initialize a thread's attribute structure
<code>pthread_attr_destroy</code>	Remove a thread's attribute structure

## Pthreads

### Example 1

```
#include <pthread.h>
#include <stdlib.h>
#include <unistd.h>
#include <stdio.h>

void *thread_function(void *arg){
    int i;
    for( i=0; i<20; i++ ){
        printf("Thread says hi!\n");
        sleep(1);
    }
    return NULL;
}

int main(void){
    pthread_t mythread;
    if(pthread_create(&mythread, NULL, thread_function, NULL)){
        printf("error creating thread.");
        abort();
    }

    if(pthread_join(mythread, NULL)){
        printf("error joining thread.");
        abort();
    }

    exit(0);
}
```

## Pthreads

`pthread_t` defined in `pthread.h`, is often called a "thread id" (`tid`);  
`pthread_create()` returns zero on success and a non-zero value on failure;  
`pthread_join()` returns zero on success and a non-zero value on failure;

### How to use pthread?

- `#include <pthread.h>`
- ```
$ gcc thread1.c -o thread1 -pthread
$ ./thread1
```

## Pthreads

### Example 2

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4
5  #define NUMBER_OF_THREADS 10
6
7  void *print_hello_world(void *tid)
8  {
9      /* prints the thread's identifier, then exits.*/
10     printf ("Thread %d: Hello World!\n", tid);
11     pthread_exit(NULL);
12 }
13
14 int main(int argc, char *argv[])
15 {
16     pthread_t threads[NUMBER_OF_THREADS];
17     int status, i;
18     for (i=0; i<NUMBER_OF_THREADS; i++)
19     {
20         printf ("Main: creating thread %d\n",i);
21         status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);
22
23         if(status != 0){
24             printf ("Oops. pthread_create returned error code %d\n",status);
25             exit(-1);
26         }
27     }
28     exit(NULL);
29 }
```

## Linux Threads


### To the Linux kernel, there is no concept of a thread

- Linux implements all threads as standard processes
- To Linux, a thread is merely a process that shares certain resources with other processes
- Some OS (MS Windows, Sun Solaris) have cheap threads and expensive processes.
- Linux processes are already quite lightweight

On a 75MHz Pentium      thread:  $1.7\mu s$   
                             fork:  $1.8\mu s$

## Linux Threads

`clone()` creates a separate process that shares the address space of the calling process. The cloned task behaves *much like* a separate thread.



`clone()`

```

1 #include <sched.h>
2 int clone(int (*fn) (void *), void *child_stack,
3           int flags, void *arg, ...);

```

**arg 1** the function to be executed, i.e. `fn(arg)`, which returns an `int`;

**arg 2** a pointer  $\blacktriangleright$  a (usually malloced) memory space to be used as the stack for the new thread;

**arg 3** a set of flags used to indicate how much the calling process is to be shared. In fact,

`clone(0) == fork()`

**arg 4** the arguments passed to the function.

It returns the PID of the child process or -1 on failure.

\$ man clone

## The `clone()` System Call

Some flags:

| flag          | Shared                |
|---------------|-----------------------|
| CLONE_FS      | File-system info      |
| CLONE_VM      | Same memory space     |
| CLONE_SIGHAND | Signal handlers       |
| CLONE_FILES   | The set of open files |

## In practice, one should try to avoid calling `clone()` directly

Instead, use a threading library (such as `pthread`) which use `clone()` when starting a thread (such as during a call to `pthread_create()`)

## `clone()` Example

```

1 #include <unistd.h>
2 #include <sched.h>
3 #include <sys/types.h>
4 #include <stdlib.h>
5 #include <string.h>
6 #include <stdio.h>
7 #include <fcntl.h>
8
9 int variable;
10
11 int do_something()
12 {
13     variable = 42;
14     _exit(0);
15 }
16
17 int main(void)
18 {
19     void *child_stack;
20     variable = 9;
21
22     child_stack = (void *) malloc(16384);
23     printf("The variable was %d\n", variable);
24
25     clone(do_something, child_stack,
26          CLONE_FS | CLONE_VM | CLONE_FILES, NULL);
27     sleep(1);
28
29     printf("The variable is now %d\n", variable);
30     return 0;
31 }

```

### Stack Grows Downwards

```
child_stack = (void**)malloc(8192) + 8192/sizeof(*child_stack);
```

## 7.3 Signals

- Signals are software interrupts. Every signal has a name (SIGXXXX). Signals are classic examples of asynchronous events. They occur at what appear to be random times to the process. The process can't simply test a variable (such as `errno`) to see whether a signal has occurred; instead, the process has to tell the kernel "if and when this signal occurs, do the following." [*Advanced programming in the UNIX environment*, chap. 10]
  - Signals are software interrupts sent to a program to indicate that an important event has occurred. The events can vary from user requests to illegal memory access errors. Some signals, such as the interrupt signal, indicate that a user has asked the program to do something that is not in the usual flow of control. (<https://www.tutorialspoint.com/unix/unix-signals-traps>)
  - Signals are similar to interrupts, the difference being that interrupts are mediated by the processor and handled by the kernel while signals are mediated by the kernel (possibly via system calls) and handled by processes. The kernel may pass an interrupt as a signal to the process that caused it (typical examples are SIGSEGV, SIGBUS, SIGILL and SIGFPE). ([https://en.wikipedia.org/wiki/Signal\\_\(IPC\)](https://en.wikipedia.org/wiki/Signal_(IPC)))
  - `signal(7)`
- \$ trap -1

### Signals

- Signals are software interrupts
- Every signal has a name (SIGXXXX)
- One process can send a signal to another process

### Sending signals

```
$ [Ctrl]+[c], [Ctrl]+[z], ...
$ kill -signal <pid>
```

### Trapping signals

```
#!/ trap <command> <signals>
```

## Trap

```
1  #!/bin/bash
2
3  sigint(){
4      echo -e "Why Ctrl-c?\n-> "
5  }
6
7  trap sigint SIGINT
8
9  echo -n "-> "
10
11 while read CMD; do
12     $CMD
13     echo -n "-> "
14 done
```

```
#! trap "rm -rf $tmpfiles" EXIT
```

## Example

*SIGINT*

```

1  #include <stdio.h>
2  #include <string.h>           /* for strlen() */
3  #include <stdlib.h>
4  #include <unistd.h>           /* for fork() */
5  #include <sys/wait.h>         /* for waitpid() */
6  #include <signal.h>
7
8  #define MAXLINE 4096
9
10 void sig_int(int signo)
11 {
12     printf("Why Ctrl-c?\n ");
13 }
14
15 int main(void)
16 {
17     char buf[MAXLINE];
18     pid_t pid;
19     int status;
20
21     if (signal(SIGINT, sig_int) == SIG_ERR)
22         perror("signal");
23
24     printf(" ");
25     while (fgets(buf, MAXLINE, stdin) != NULL) {
26         buf[strlen(buf) - 1] = '\0'; /* null */
27
28         if ( (pid = fork()) == 0 ) { /* child */
29             execlp(buf, buf, (char *)0);
30             perror("execlp");
31             exit(127);
32         }
33
34         if ((pid = waitpid(pid, &status, 0)) < 0)
35             perror("waitpid");
36         printf(" ");
37     }
38     exit(0);
39 }

```

- <https://stackoverflow.com/questions/840501/how-do-function-pointers-in-c-work>

**Example**  
*SIGUSR1*

```

1  #include <signal.h>
2  #include <stdio.h>
3  #include <unistd.h>
4
5  void sig_usr(int);
6
7  int main(void)
8  {
9      printf("PID = %d\n", getpid());
10
11     if (signal(SIGUSR1, sig_usr) == SIG_ERR)
12         perror("signal<SIGUSR1>");
13
14     for ( ; ; )
15         pause();
16 }
17
18 void sig_usr(int signo)
19 {
20     if (signo == SIGUSR1)
21         printf("received SIGUSR1\n");
22     else
23         perror("sig_usr");
24 }

```

\$ kill -USR1 <PID>

### Example SIGALRM

|                                                                                                                                                                                                                                                                                                                        |                                                                                                                                                                                                                            |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre> 1  #include &lt;signal.h&gt; 2  #include &lt;stdio.h&gt; 3  #include &lt;unistd.h&gt; 4  #include &lt;stdlib.h&gt; 5 6  void cry(int sig) 7  { 8      puts("C: I'm crying..."); 9      kill(getppid(),sig); 10 } 11 12 void complain(int sig) 13 { 14     puts("P: You're noisy."); 15 } 16 17 int main() </pre> | <pre> 16 17 int main() 18 { 19     if ( fork() == 0 ){ 20         signal(SIGALRM, cry); 21         alarm(2); 22         pause(); 23     } 24 25     signal(SIGALRM, complain); 26     pause(); 27     exit(0); 28 } </pre> |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|


## 8 IPC

### 8.1 Pipes and FIFOs

- [The Linux Programming Interface: A Linux and UNIX System Programming Handbook, chap. 44]

#### Pipe


```
$ ls | wc -l
```




- A pipe is a byte stream
- Unidirectional
- `read()` would be blocked if nothing written at the other end

`tee`

```
$ ls | tee ls.out
```



- When we say that a pipe is a byte stream, we mean that there is no concept of messages or message boundaries when using a pipe. The process reading from a pipe can read blocks of data of any size, regardless of the size of blocks written by the writing process. Furthermore, the data passes through the pipe sequentially — bytes are read from a pipe in exactly the order they were written. It is not possible to randomly access the data in a pipe using `lseek()`. [The Linux Programming Interface: A Linux and UNIX System Programming Handbook, chap. 44]




- No direct link between A and B
- A pipe is simply a buffer maintained in kernel memory
 

```
$ cat /proc/sys/fs/pipe-max-size
```
- Need system calls

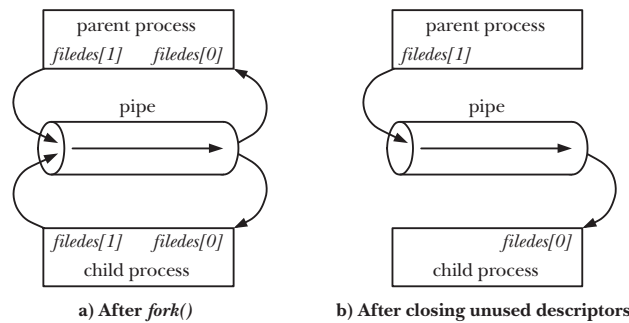
`pipe()`

```
1 #include <unistd.h>
2
3 int pipe(int fd[2]);
```



`pipe() + fork()`





- Pipes must have a reader and a writer. If a process tries to write to a pipe that has no reader, it will be sent the SIGPIPE signal from the kernel. This is imperative when more than two processes are involved in a pipeline. (<http://www.tldp.org/LDP/lpg/node20.html>)
- While it is possible for the parent and child to both read from and write to the pipe, this is not usual. Therefore, immediately after the `fork()`, one process closes its descriptor for the write end of the pipe, and the other closes its descriptor for the read end. For example, if the parent is to send data to the child, then it would close its read descriptor for the pipe, `filedes[0]`, while the child would close its write descriptor for the pipe, `filedes[1]`. One reason that it is not usual to have both the parent and child reading from a single pipe is that if two processes try to simultaneously read from a pipe, we can't be sure which process will be the first to succeed—the two processes race for data. Preventing such races would require the use of some synchronization mechanism. However, if we require bidirectional communication, there is a simpler way: just create two pipes, one for sending data in each direction between the two processes. (If employing this technique, then we need to be wary of deadlocks that may occur if both processes block while trying to read from empty pipes or while trying to write to pipes that are already full.) [*The Linux Programming Interface: A Linux and UNIX System Programming Handbook*, chap. 44, p. 893]
- Pipes can be used for communication between any two (or more) related processes, as long as the pipe was created by a common ancestor before the series of `fork()` calls that led to the existence of the processes. For example, a pipe could be used for communication between a process and its grandchild. The first process creates the pipe, and then forks a child that in turn forks to yield the grandchild. A common scenario is that a pipe is used for communication between two siblings — their parent creates the pipe, and then creates the two children. This is what the shell does when building a pipeline.
- **Closing unused pipe file descriptors.** The process reading from the pipe closes its write descriptor for the pipe, so that, when the other process completes its output and closes its write descriptor, the reader sees end-of-file (once it has read any outstanding data in the pipe). If the reading process doesn't close the write end of the pipe, then, after the other process closes its write descriptor, the reader won't see end-of-file, even after it has read all data from the pipe. Instead, a `read()` would block waiting for data, because the kernel knows that there is still at least one write descriptor open for the pipe. That this descriptor is held open by the reading process itself is irrelevant; in theory, that process could still write to the pipe, even if it is blocked trying to read. For example, the `read()` might be interrupted by a signal handler that writes data to the pipe. The writing process closes its read descriptor for the pipe for a different reason. When a process tries to write to a pipe for which no process has an open read descriptor, the kernel sends the SIGPIPE signal to the writing process. By default, this signal kills a process. A process can instead arrange to catch or ignore this signal, in which case the `write()` on the pipe fails with the error EPIPE (broken pipe). Receiving the SIGPIPE signal or getting the EPIPE error is a useful indication about the status of the pipe, and this is why unused read descriptors for the pipe should be closed. If the writing process doesn't close the read end of the pipe, then, even after the other process closes the read end of the pipe, the writing process will still be able to write to the pipe. Eventually, the writing process will fill the pipe, and a further attempt to write will block indefinitely. One final reason for closing unused file descriptors is that it is only after all file descriptors in all processes that refer to a pipe are closed that the pipe is destroyed and its resources released for reuse by other processes. At this point, any unread data in the pipe is lost.

```

1  #include <sys/wait.h>
2  #include <unistd.h>
3  #include <string.h>
4  #include <stdlib.h>
5  #include <stdio.h>
6
7  #define BUF_SIZE 10
8
9  int main(int argc, char *argv[]) /* Over-simplified! */
10 {
11     int pfd[2]; /* Pipe file descriptors */
12     char buf[BUF_SIZE];
13     ssize_t numRead;
14
15     pipe(pfd); /* Create the pipe */
16
17     switch (fork()) {
18     case 0: /* Child - reads from pipe */
19         close(pfd[1]); /* Write end is unused */
20
21         for (;;) { /* Read data from pipe, echo on stdout */
22             if ( (numRead = read(pfd[0], buf, BUF_SIZE)) == 0)
23                 break; /* End-of-file */
24             if (write(1, buf, numRead) != numRead)
25                 perror("child - partial/failed write");
26         }
27         puts(""); /* newline */
28
29         close(pfd[0]);
30
31         _exit(EXIT_SUCCESS);
32
33     default: /* Parent - writes to pipe */
34         close(pfd[0]); /* Read end is unused */
35
36         if (write(pfd[1], argv[1], strlen(argv[1])) != strlen(argv[1]))
37             perror("parent - partial/failed write");
38
39         close(pfd[1]); /* Child will see EOF */
40
41         wait(NULL); /* Wait for child to finish */
42         exit(EXIT_SUCCESS);
43     }
44 }

```

- <https://stackoverflow.com/questions/5422831/what-is-the-difference-between-using-exit-exit-in-a-conventional-linux-fo>
- `_exit(2)`

`popen()`

### Named Pipe (FIFO)

**PIPEs** pass data between related processes.


**FIFOs** pass data between any processes.

`$ mkfifo myfifo`

```

$ echo hello > myfifo
$ cat myfifo

```




`tee`

```

$ echo hello | tee myfifo
$ wc myfifo

```



- [https://en.wikipedia.org/wiki/Named\\_pipe](https://en.wikipedia.org/wiki/Named_pipe)

## IPC With FIFO

```
1  #include <unistd.h>
2  #include <stdlib.h>
3  #include <stdio.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/types.h>
7  #include <sys/stat.h>
8
9  #define FIFO_NAME "/tmp/myfifo"
10
11 int main(int argc, char *argv[]) /* Oversimplified */
12 {
13     int fd, i, mode = 0;
14     char c;
15
16     if (argc < 2) {
17         fprintf(stderr, "Usage: %s <O_RDONLY | O_WRONLY | O_NONBLOCK>\n", argv[0]);
18         exit(EXIT_FAILURE);
19     }
20
21     for(i = 1; i < argc; i++) {
22         if (strncmp(++argv, "O_RDONLY", 8) == 0)
23             mode |= O_RDONLY;
24         if (strncmp(*argv, "O_WRONLY", 8) == 0)
25             mode |= O_WRONLY;
26         if (strncmp(*argv, "O_NONBLOCK", 10) == 0)
27             mode |= O_NONBLOCK;
28     }
29
30     if (access(FIFO_NAME, F_OK) == -1) mkfifo(FIFO_NAME, 0777);
31
32     printf("Process %d: FIFO(fd %d, mode %d) opened.\n",
33           getpid(), fd = open(FIFO_NAME, mode), mode);
34
35     if( (mode == 0) | (mode == 2048) )
36         while( read(fd,&c,1) == 1 ) putchar(c);
37
38     if( (mode == 1) | (mode == 2049) )
39         while( (c = getchar()) != EOF ) write(fd,&c,1);
40
41     exit(EXIT_SUCCESS);
42 }
```

```
$ watch 'ls -l /tmp/myfifo'
$ ./a.out O_RDONLY
$ ./a.out O_WRONLY
$ ./a.out O_RDONLY O_NONBLOCK
$ ./a.out O_WRONLY O_NONBLOCK
```

## O\_NONBLOCK

- A read()/write() will wait on an empty blocking FIFO

- A `read()` on an empty nonblocking FIFO will return 0 bytes
- `open(const char *path, O_WRONLY | O_NONBLOCK);`
  - Returns an error (-1) if FIFO not open
  - Okay if someone's reading the FIFO

### Opening a FIFO with `open()`

- If opened with `O_RDWR`, the result is undefined. If you do want to pass data in both directions, it's much better to use a pair of FIFOs or pipes, one for each direction.
- There are four legal combinations of `O_RDONLY`, `O_WRONLY`, and the `O_NONBLOCK` flag.

```

1  open(const char *path, O_RDONLY);
2  /* In this case, the open call will block; it will not return until a process opens the
3     same FIFO for writing. */
4
5  open(const char *path, O_RDONLY | O_NONBLOCK);
6  /* The open call will now succeed and return immediately, even if the FIFO has not been
7     opened for writing by any process. */
8
9  open(const char *path, O_WRONLY);
10 /* In this case, the open call will block until a process opens the same FIFO for
11    reading. */
12
13 open(const char *path, O_WRONLY | O_NONBLOCK);
14 /* This will always return immediately, but if no process has the FIFO open for reading,
15    open will return an error, -1, and the FIFO won't be opened. If a process does have the
16    FIFO open for reading, the file descriptor returned can be used for writing to the
17    FIFO. */

```

## 9 User Interface

### 9.1 Dialog, Zenity

### 9.2 Ncurses

### 9.3 GTK+

### 9.4 Qt

## 10 Terminal

- chap 62 of [*The Linux Programming Interface: A Linux and UNIX System Programming Handbook*]
  - [<http://tldp.org/HOWTO/Text-Terminal-HOWTO-16.html>]
  - <http://pubs.opengroup.org/onlinepubs/9699919799/>
- \$ `man infocmp`

## 11 IDE

Makefile, git