

GNU/Linux Application Programming

WANG Xiaolin

`wx672ster+linux@gmail.com`

July 31, 2019

Textbooks



MATTHEW N, STONES R. Beginning linux programming. John Wiley & Sons, 2008.



COOPER M. Advanced Bash Scripting Guide 5.3 Volume 1. Lulu.com, 2010.



RAYMOND E S. The art of Unix programming. Addison-Wesley, 2003.



STEVENS W R, RAGO S A. Advanced programming in the UNIX environment. Addison-Wesley, 2013.



LOVE R. Linux System Programming: Talking Directly to the Kernel and C Library. O'Reilly Media, Inc., 2007.



KERRISK M. The Linux Programming Interface: A Linux and UNIX System Programming Handbook. No Starch Press, 2010.



BRYANT R E, O'HALLARON D R. Computer Systems: A Programmer's Perspective. 2nd ed. Addison-Wesley, 2010.

Course Web Links

 <https://cs6.swfu.edu.cn/moodle>

 https://cs2.swfu.edu.cn/~wx672/lecture_notes/linux-app/slides/

 https://cs2.swfu.edu.cn/~wx672/lecture_notes/linux-app/src/

 <https://cs3.swfu.edu.cn/tech>

/etc/hosts

```
202.203.132.241  cs6.swfu.edu.cn
```

```
202.203.132.242  cs2.swfu.edu.cn
```

```
202.203.132.245  cs3.swfu.edu.cn
```

Homework

Weekly tech question

1. What was I trying to do?
2. How did I do it? (steps)
3. The expected output? The real output?
4. How did I try to solve it? (steps, books, web links)
5. How many hours did I struggle?

✉ wx672ster+linux@gmail.com

🇬🇧 Preferably in English

📖 in [stackoverflow](#) style

Or simply show me the tech questions you asked on any website

Getting Started

Linux Programs

Where to find them? /bin, /usr/bin, /usr/local/bin,
~/bin, ...

```
$ echo $PATH
```

How to find them? which, whereis, type

Command not found?

First double check your spelling

Then try:

```
④ aptitude search xxx
```

```
④ apt-cache search xxx
```

```
④ apt-file search xxx
```

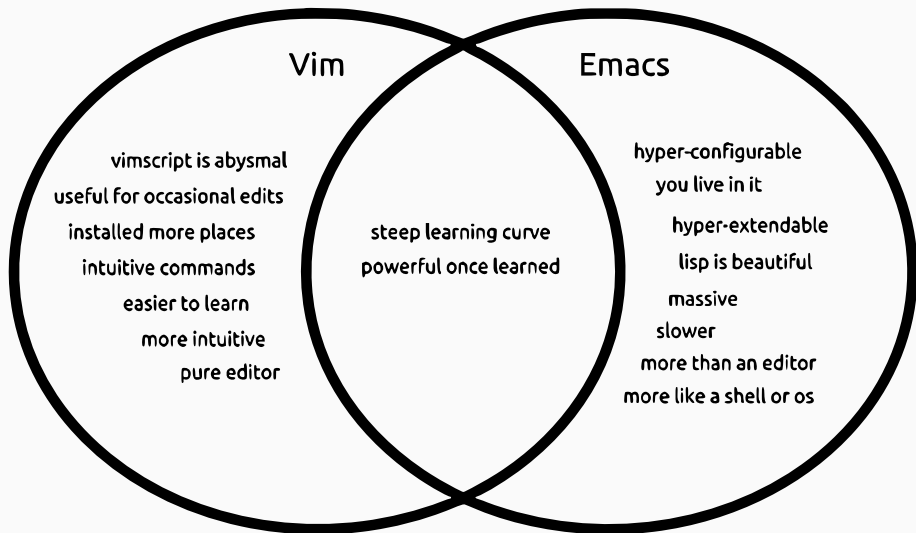
```
④ sudo apt install packagename
```

```
G Google "linux command xxx"
```

Text Editors



vs.



Help Your Editor

Suffix matters

\$ vim ✗

\$ vim hello.c ✓


\$ vim hello.py ✓

\$ emacs ✗

\$ emacsclient hello.c ✓

\$ emacsclient hello.py ✓

Keyboard



vimtutor



C-h t

Compiler vs. Interpreter

hello.c

```
#include <stdio.h>
```

```
int main()
```

```
{  
    printf("Hello, world!\n");  
    return 0;  
}
```

```
$ gcc -o hello hello.c
```

```
$ ./hello
```

hello.sh

```
#!/bin/bash
```

```
echo 'Hello, world!'
```

```
$ chmod +x hello.sh
```

```
$ ./hello.sh
```

hello.py


```
#!/usr/bin/python
```

```
print "Hello, world!"
```

```
$ chmod +x hello.py
```

```
$ ./hello.py
```

Compilation



Header Files

Why?

```
#include "add.h"
```

```
int triple(int x)
{
    return add(x, add(x,x));
}
```

- ▶ Ensure everyone use the same code
- ▶ Easy to share, upgrade, reuse

Why not?

```
int add(int, int);
```

```
int triple(int x)
{
    return add(x, add(x, x));
}
```

In the header files...

- ▶ function declarations
- ▶ macro definitions
- ▶ constants
- ▶ system wide global variables

```
$ ls /usr/include/
```

Library Files

Static libraries `.a` files. Very old ones, but still alive.

```
$ find /usr/lib -name "*.a"
```

Shared libraries `.so` files. The preferred ones.

```
$ find /usr/lib -name "*.so.*"
```

Examples:


```
$ gcc -o hello hello.c /usr/lib/libm.a
```

```
$ gcc -o hello hello.c -lm
```

Static Linking


- ▶ The entire program and all data of a process must be in physical memory for the process to execute
- ▶ The size of a process is thus limited to the size of physical memory

Source files `main2.c` `vector.h`



Dynamic Linking

- ▶ Only one copy in memory
- ▶ Don't have to re-link after a library update



Build A Static Library

Source codes

main.c

```
1  #include "lib.h"
2
3  int main(int argc, char* argv[])
4  {
5      int i=1;
6
7      for (; i<argc; i++)
8      {
9          hello(argv[i]);
10         hi(argv[i]);
11     }
12     return 0;
13 }
```

lib.h

```
1  #include <stdio.h>
2
3  void hello(char *);
4  void hi(char *);
```

hello.c

```
1  #include <stdio.h>
2
3  void hello(char *arg)
4  {
5      printf("Hello, %s!\n", arg);
6  }
```

hi.c

```
1  #include <stdio.h>
2
3  void hi(char *arg)
4  {
5      printf("Hi, %s!\n", arg);
6  }
```


Build A Static Library

Step by step

1. Get `hello.o` and `hi.o`

```
$ gcc -c hello.c hi.c
```

2. Put `*.o` into `libhi.a`

```
$ ar crv libhi.a hello.o hi.o
```

3. Use `libhi.a`

```
$ gcc main.c libhi.a
```

Build A Static Library

Makefile

```
1  main: main.c lib.h libhi.a
2      gcc -Wall -o main main.c libhi.a
3
4  libhi.a: hello.o hi.o
5      ar crv libhi.a hello.o hi.o
6
7  hello.o: hello.c
8      gcc -Wall -c hello.c
9
10 hi.o: hi.c
11     gcc -Wall -c hi.c
12
13 clean:
14     rm -f *.o *.a main
```

Build A Shared Library

Source codes

hello.c

```
1  #include "hello.h"
2
3  int main(int argc, char *argv[])
4  {
5      if (argc != 2)
6          printf ("Usage: %s needs an argument.\n", argv[0]);
7      else
8          hi(argv[1]);
9      return 0;
10 }
```

hello.h

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int hi(char*);
```

hi.c

```
1  #include "hello.h"
2
3  int hi(char* s)
4  {
5      printf ("Hi, %s\n",s);
6      return 0;
7  }
```

Build A Shared Library

Step by step

1. Get `hi.o`

```
$ gcc -fPIC -c hi.c
```

2. Get `libhi.so`

```
$ gcc -shared -o libhi.so hi.o
```

3. Use `libhi.so`

```
$ gcc -L. -Wl,-rpath=. hello.c -lhi
```

4. Check it

```
$ ldd a.out
```

Build A Shared Library

Makefile

```
1  # http://www.cprogramming.com/tutorial/shared-libraries-linux-gcc.html
2  # http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html
3  #
4  # gcc -fPIC -c hi.c
5  # gcc -shared -o libhi.so hi.o
6  # gcc -L/current/dir -Wl,option -Wall -o hello hello.c -lhi
7  #
8  # -L          - tells ld where to search libraries
9  # -Wl,option - pass option as an option to the linker (ld)
10 # -rpath=dir - Add a directory to the runtime library search path
11
12 hello: hello.c hello.h libhi.so
13         gcc -L. -Wl,-rpath=. -Wall -o hello hello.c -lhi
14 libhi.so: hi.o hello.h
15         gcc -shared -o libhi.so hi.o
16 hi.o: hi.c hello.h
17         gcc -fPIC -c hi.c
18 clean:
19         rm *.o *.so hello
```

Getting Help

```
⌘ sudo apt install gcc-doc
```


```
$ man gcc
```

```
$ info gcc
```

Shell Basics

Shell

- ▶ A command line interpreter
- ▶ A programming language



Redirection

Redirecting output

```
$ ls -l > output.txt
```

```
$ ps aux >> output.txt
```

Redirecting input

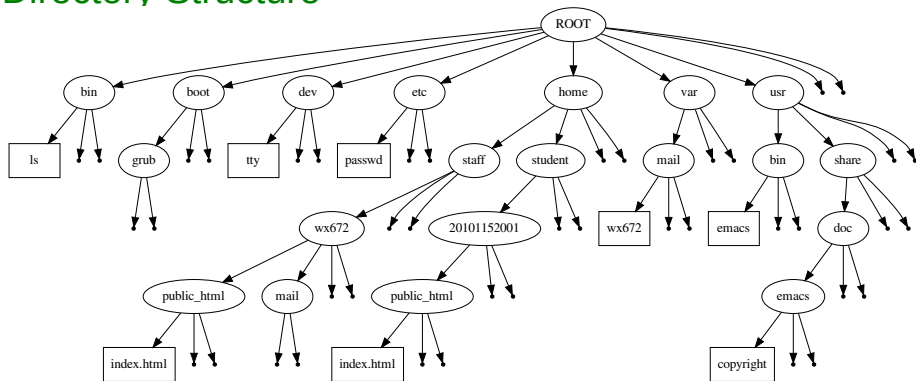
```
$ more < output.txt
```

Pipe

Chain processes together

```
$ ps aux | sort | less
```

Directory Structure



Todo	How
Where am I?	<code>pwd</code>
What's in it?	<code>ls</code>
Move around?	<code>cd</code>
Disk usage?	<code>du</code> , <code>df</code>
USB drive?	<code>lsblk</code> , <code>mount</code>
New folder?	<code>mkdir</code>

File Operations

Ways to create a file

 Using an editor (vim, emacs, nano...), or

```
$ cat > filename
```

```
$ echo "hello, world" > filename
```

```
$ touch filename
```

More file operations:

Todo	How	Todo	How
Copy?	cp	Move/Rename?	mv
Delete?	rm	What's it?	file
Link?	ln	Permission?	chmod, chown
Count?	wc	Archive?	tar, gzip, 7z, ...
Sort?	sort, uniq	Search?	find, grep

Process Operations

Todo	How	Todo	How
Kill?	kill, Ctrl-c	suspend?	Ctrl-z
background?	bg, &	foreground?	fg, jobs
status?	ps, top		

System Info

Todo	How	Todo	How
who?	w, who, whoami	how long?	uptime
software?	apt, aptitude, dpkg	kernel?	uname, lsmod
hardware?	lspci, lsusb, lscpu	memory?	free, lsmem

APT — package management

Todo	How
upgrading?	apt update && apt upgrade
install?	apt install xxx
remove?	apt purge xxx
search?	apt search xxx
details?	apt show xxx
friendly UI?	aptitude

CLI Shortcuts


Ctrl-a:	beginning of line	Ctrl-e:	end of line
Ctrl-f:	forward	Ctrl-b:	backward
Ctrl-n:	next	Ctrl-p:	previous
Ctrl-r:	reverse search	Ctrl-u:	cut to beginning
Ctrl-k:	kill (cut to end)	Ctrl-y:	yank (paste)
Ctrl-d:	delete a character	TAB:	completion

Tmux

C-a c:	create window	C-a C-a:	switch window
C-a n:	next window	C-a p:	previous window
C-a -:	split window	C-a :	split window
C-a j:	go down	C-a k:	go up
C-a l:	go right	C-a h:	go left

Understanding “ls -l”


```
-rw----- 1 sam sam 57 Apr 17 1998 weather.txt
drwxr-xr-x 6 sam sam 102 Oct 9 1999 web_page
-rw-rw-r-- 1 sam sam 7648 Feb 11 20:41 web_site.tar
-rw----- 1 sam sam 574 Dec 16 1998 file.txt
```



d - directory
- - regular file
l - soft link
c - character device
b - block device
s - socket
p - named pipe (FIFO)

9-bit permission

```
7 5 5
111 101 101
rwx r-x r-x
```



```
$ chmod 755 foo
$ chmod 000 foo
$ chmod a-r foo
$ chmod g+w foo
$ chmod 644 foo
$ chmod 777 foo
$ chmod u+x foo
$ chmod go=rx foo
```


Wildcard Expansion

Character	Meaning	Example
?	any one	\$ ls ????.txt
*	zero or more	\$ ls *.c
[]	or	\$ ls *. [ch]
{}	and	\$ ls *.{c,h,cpp}

Example

```
$ touch {2,3,4,234}.{jpg,png} && ls
```

```
output:  2.jpg  234.jpg  3.jpg  4.jpg
         2.png  234.png  3.png  4.png
```

```
$ rm [234].jpg
```

```
$ rm ?.jpg
```

```
$ rm {2,3,4,234}.jpg
```

```
$ rm ?.*
```

```
$ rm 2*
```

```
$ rm *
```

Everything Is A File

```
$ cat /dev/null > /var/log/messages # empty a file
```

```
  $ : > /var/log/messages # no new process
```

```
$ ls > /dev/null
```

```
$ dd if=/dev/zero of=/tmp/clean bs=1k count=1k
```

```
$ dd if=/dev/urandom of=/tmp/random bs=1k count=1k
```

Shell Programming

Shell Programming

why? handy, powerful

which? bash

Variables

```
$ a=8; b=2
```

```
$ a=a+5; a=$a+5 ☹️
```

```
$ let a=a+5; let a+=5 😊
```

```
$ let b=b+a; let b+=a 😊
```

```
$ echo a; echo $a
```

```
$ (( a=5, b=6, a+=b )) 😊
```

```
$ (( b=a<5?8:9 )) 😊
```

```
$ r=$(( RANDOM%100 )) 😊
```

```
$ echo "$a" # partial quoting
```

```
$ echo '$a' # full quoting
```

```
$ a=$(ls -l); echo $a; echo "$a"
```

```
$ a=hello; b=world; let a+=b ☹️
```

Positional Parameters

\$0, \$1, \$2, ..., \$@, \$#

```
1  #!/bin/bash
2
3  echo "You said:"
4
5  echo -e "\t$@"
6  echo
7  echo -e "\targc = $#"
```

```
8  echo -e "\targv[0] = $0"
```

```
9
10 i=1
11 for arg in $@; do
12     echo -e "\targv[$i] = $arg"
13     let i++
14 done
```

```
1  #include <stdio.h>
2
3  int main(int argc, char *argv[])
4  {
5      int i;
6      printf("You said:\n\t");
7
8      for(i=1; i<argc; i++)
9          printf("%s ",argv[i]);
10
11     printf("\n\n\targc = %d\n", argc);
12
13     for(i=0; i<argc; i++)
14         printf("\targv[%d] = %s\n",i,argv[i]);
15
16     return 0;
17 }
```

Parameter Substitution

Default value

```
$ echo ${s:=abc}
```

```
$ echo ${v:-8}
```

```
$ echo ${s:=xyz}
```

```
$ echo ${v:-10}
```

Example

```
1  #!/bin/bash
2
3  echo -n Hello, ${1:-world}
4  echo !
```

Parameter Substitution

Substring removal

```
$ for f in *.pbm; do ppm2tiff $f ${f%.pbm}.tif; done
```

Substring replacement

```
$ for f in *.jpg; do mv $f ${f/.jpg/JPG}; done
```


Environmental Variables

Each process has an environment

\$PATH	\$PWD	\$HOME	\$UID	\$USER
\$GROUPS	\$SHELL	\$TERM	\$DISPLAY	\$TEMP
\$HOSTNAME	\$HOSTTYPE	\$IFS	\$EDITOR	\$BROWSER
\$HISTSIZE	\$FUNCNAME	\$TMOUT	...	

```
export HISTSIZE=2000
```

```
export BROWSER='/usr/bin/x-www-browser'
```

```
export EDITOR='vim'
```

```
export ALTERNATE_EDITOR="vim"
```

```
export PDFVIEWER='/usr/bin/zathura'
```

\$ env

\$ declare

Tests I

```
$ (( 5 < 6 )) && echo should be
```

```
$ [[ 1 < 2 ]] && echo of course
```

```
$ [[ $a -lt $b ]] && echo yes || echo no
```

```
$ if [[ $a -lt $b ]]; then echo yes; else echo no; fi
```

```
$ if test $a -lt $b; then echo of course; fi
```

```
$ if a = 5; then echo a=$a; fi # whitespace matters ❌
```

```
$ if a=5; then echo a=$a; fi 😞
```

```
$ if test a=5; then echo a=$a; fi 😞
```

```
$ if test a = 5; then echo a=$a; fi 😞
```

```
$ if test $a = 5; then echo a=$a; fi ✓
```

```
$ test $a = 5 && echo a=$a ✓
```

```
$ [[ $a = 5 ]] && echo a=$a ✓
```

Tests II

```
$ [[ cmp a b ]] && echo same file ✗
```

```
$ if test cmp a b; then echo same file; fi ✗
```

```
$ if cmp a b; then echo same file; fi ✓
```

```
$ [[ -f ~/.bash_aliases ]] && . ~/.bash_aliases
```

```
$ [[ -x /usr/bin/xterm ]] && /usr/bin/xterm -e tmux &
```

```
$ [[ "$pass" != "$MYPASS" ]] && echo 'Wrong password!' && exit  
1
```

```
$ help test
```

Tests III

```
1  #!/bin/bash
2
3  words=$@
4  string=linux
5  if echo "$words" | grep -q "$string"
6  then
7      echo "$string found in $words"
8  else
9      echo "$string not found in $words"
10 fi
```

Loops

```
for ARG in LIST; do COMMAND(s); done
```

```
$ for i in 1 2 3; do echo -n i="$i "; done
```

```
$ for i in {1..10}; do echo $i; done
```

```
$ for i in $(seq 10); do echo $i; done
```

```
$ for ((i=1; i<=10; i++)); do echo $i; done
```

```
$ for ((i=1, j=1; i<=10; i++, j++)); do  
    echo $i-$j  
    echo $((i-$j))  
done
```

```
$ for ((i=1; i<=10; i++)) { echo $i; }
```

```
$ for i in hello world; do echo -n "$i "; done
```

Loops

`while` `CONDITION`; `do` `COMMAND(s)`; `done`

`$ a=0; while [[a < 10]]; do echo $a; ((a++)); done` ☹️

`$ while [[$a < 10]]; do echo $a; ((a++)); done` ☹️

`$ while [[$a -lt 10]]; do echo $a; ((a++)); done` ✓

`$ while [$a -lt 10]; do echo $a; ((a++)); done` ✓

`$ while ((a < 10)); do echo $a; ((a++)); done` ✓

`$ until ((a = 10)); do echo $a; ((a++)); done` ☹️

`$ until ((a == 10)); do echo $a; ((a++)); done` ✓

`$ while read n; do n2 $n; done`

`$ while read n; do n2 $n; done < datafile`

`$ until ((n == 0)); do read n; n2 $n; done`

case

```
1  #!/bin/bash
2
3  [ -z "$1" ] && echo "Usage: `basename $0` [d|h] <number>" && exit 0;
4
5  case "$1" in
6      [dD]*)
7          NUM=$(echo $1 | cut -b 2-)
8          printf "\tDec\tHex\tBin\n"
9          printf "\t%d\t0x%02X\t%s\n" $NUM $NUM $(bc <<< "obase=2;$NUM")
10         ;;
11     [hH]*)
12         NUM=$(echo $1 | cut -b 2-)
13         NUM=$(echo $NUM | tr [:lower:] [:upper:])
14         printf "\tHex\t\tDec\t\tBin\n"
15         printf "\t0x%s\t\t%s\t\t%s\n" $NUM $(bc <<< "ibase=16;obase=A;$NUM") \
16             $(bc <<< "ibase=16;obase=2;$NUM")
17         ;;
18     0[xX]*)
19         NUM=$(echo $1 | cut -b 3-)
20         NUM=$(echo $NUM | tr [:lower:] [:upper:])
21         printf "\tHex\t\tDec\t\tBin\n"
22         printf "\t0x%s\t\t%s\t\t%s\n" $NUM $(bc <<< "ibase=16;obase=A;$NUM") \
23             $(bc <<< "ibase=16;obase=2;$NUM")
24         ;;
25     [bB]*)
26         NUM=$(echo $1 | cut -b 2-)
27         printf "\tBin\t\tHex\t\tDec\n"
28         printf "\t%s\t\t0x%s\t\t%s\n" $NUM $(bc <<< "ibase=2;obase=10000;$NUM") \
29             $(bc <<< "ibase=2;obase=1010;$NUM")
30         ;;
31     *)
32         printf "Dec\tHex\tBin\n"
33         printf "%d\t0x%08X\t%08d\n" $1 $1 $(bc <<< "obase=2;$1")
34         ;;
35  esac
```

select

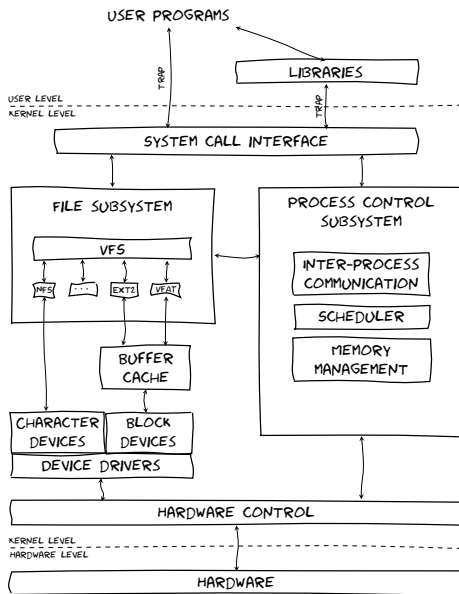
```
1  #!/bin/bash
2
3  PS3='Your favorite OS? '
4
5  select OS in "Linux" "Mac OSX" "Windows"
6  do
7      [[ "$OS" = "Linux" ]] && echo wise guy.
8      [[ "$OS" = "Mac OSX" ]] && echo rich guy.
9      [[ "$OS" = "Windows" ]] && echo patient guy.
10     break
11 done
```


Functions

```
1  #!/bin/bash
2
3  function screencapture(){
4      ffmpeg -f x11grab -s 1920x1080 -r 30 -i :0.0 \
5          -c:v libx264 -crf 0 -preset ultrafast screen.mkv
6  }
7
8  w2pdf(){
9      libreoffice --convert-to pdf:writer_pdf_Export "$1"
10 }
11
12 rfc(){
13     [[ -n "$1" ]] || {
14         cat <<EOF
15         rfc - Command line RFC viewer
16         Usage: rfc <index>
17     EOF
18         return 1
19     }
20     find /usr/share/doc/RFC/ -type f -iname "rfc$1.*" | xargs less
21 }
```

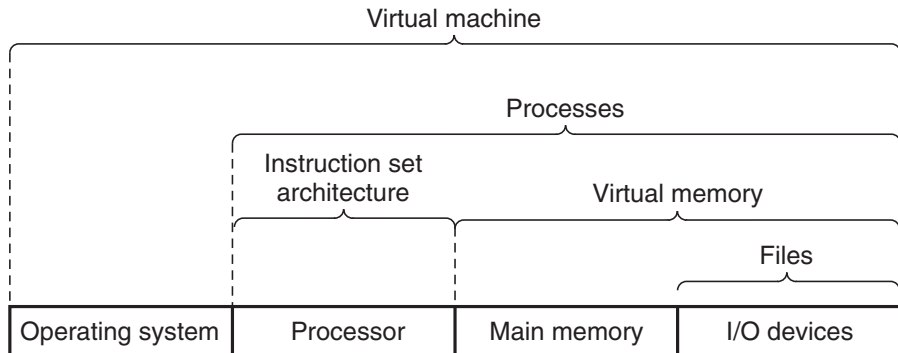
OS Basics

What's in the OS?

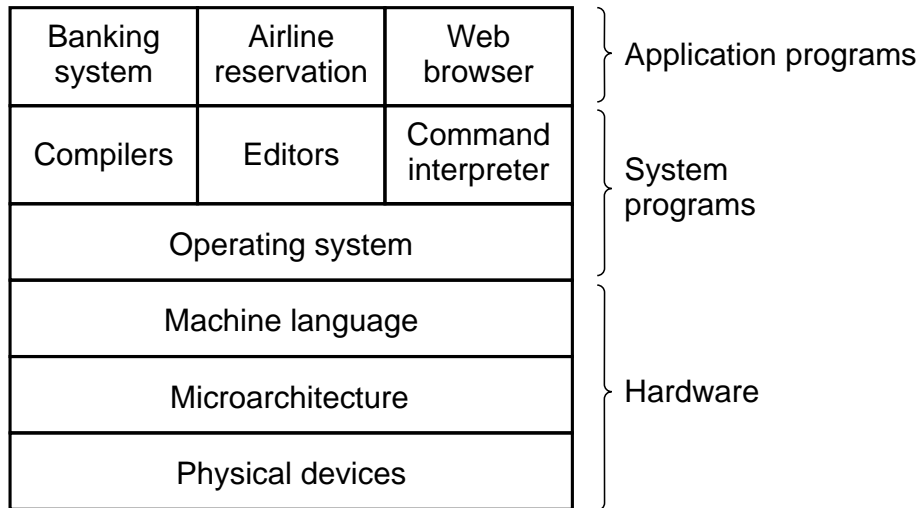


Abstractions

To hide the complexity of the actual implementations



A Computer System



CPU Working Cycle



1. Fetch the first instruction from memory
2. Decode it to determine its type and operands
3. execute it

Special CPU Registers

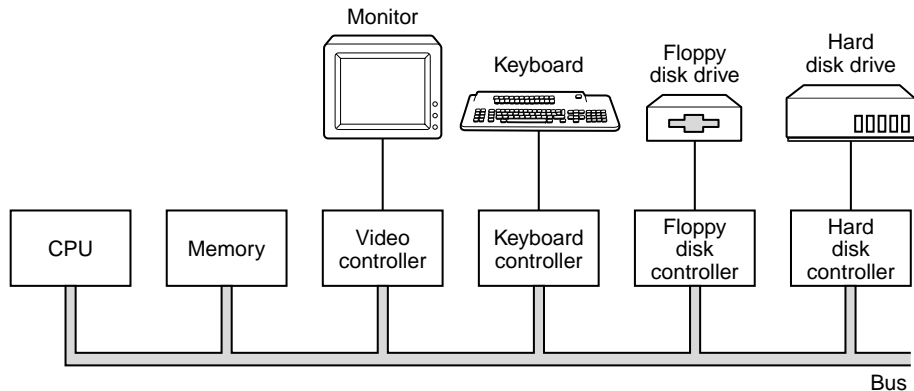
Program counter (PC): keeps the memory address of the next instruction to be fetched

Stack pointer (SP): ➡ the top of the current stack in memory

Program status (PS): holds

- condition code bits
- processor state

System Bus



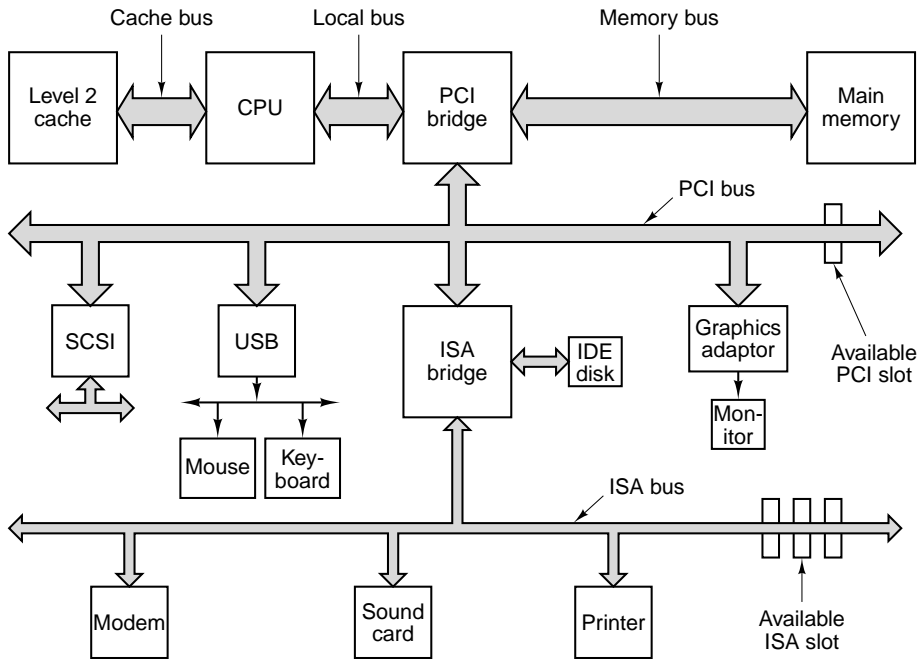
Address Bus: specifies the memory locations (addresses) for the data transfers

Data Bus: holds the data transferred. Bidirectional

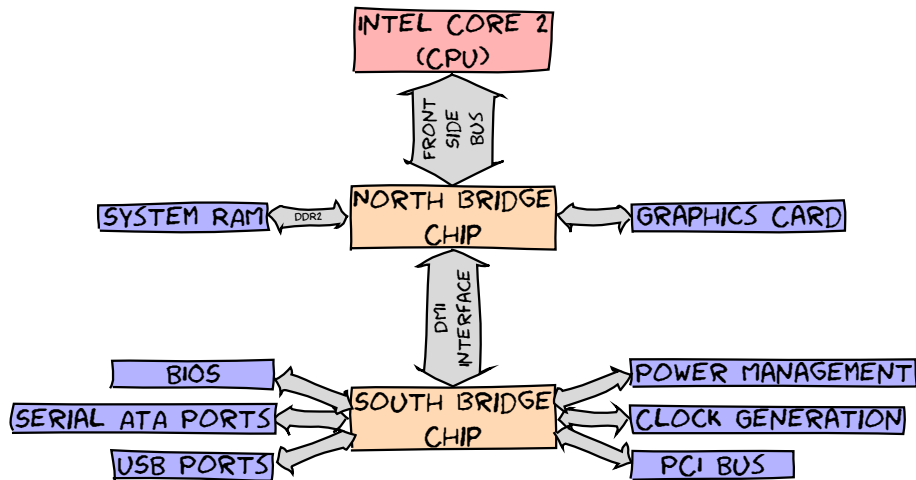
Control Bus: contains various lines used to route timing and control signals throughout the system

Controllers and Peripherals

- ▶ Peripherals are real devices controlled by controller chips
- ▶ Controllers are processors like the CPU itself, have control registers
- ▶ Device driver writes to the registers, thus control it
- ▶ Controllers are connected to the CPU and to each other by a variety of buses



Motherboard Chipsets



- ▶ The CPU doesn't know what it's connected to
 - CPU test bench? network router? toaster? brain implant?
- ▶ The CPU talks to the outside world through its pins
 - some pins to transmit the physical memory address
 - other pins to transmit the values
- ▶ The CPU's gateway to the world is the **front-side bus**

Intel Core 2 QX6600

- ▶ 33 pins to transmit the physical memory address
 - so there are 2^{33} choices of memory locations
- ▶ 64 pins to send or receive data
 - so data path is 64-bit wide, or 8-byte chunks

This allows the CPU to physically address 64GB of memory ($2^{33} \times 8B$)

Some physical memory addresses are mapped away!

- ▶ only the addresses, not the spaces
- ▶ Memory holes
 - 640 KiB ~ 1 MiB
 - /proc/iomem

Memory-mapped I/O

- ▶ BIOS ROM
- ▶ video cards
- ▶ PCI cards
- ▶ ...

This is why 32-bit OSes have problems using 4 GiB of RAM.

0xFFFFFFFF	JUMP to 0xF0000	4GB
Reset vector		
0xFFFFFFFF	Unaddressable memory, real mode is limited to 1MB.	4GB-16B
0x100000	System BIOS	1MB
0xF0000	Ext. System BIOS	960KB
0xE0000	Expansion Area (maps ROMs for old peripheral cards)	896KB
0xC0000	Legacy Video Card Memory Access	768KB
0xA0000	Accessible RAM (640KB is enough for anyone – old DOS area)	640KB
0		0

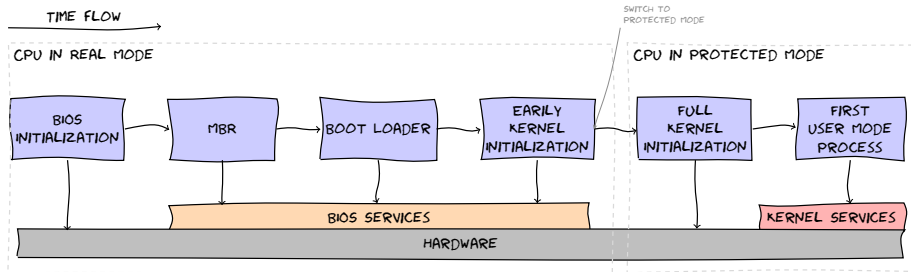
the northbridge

1. receives a physical memory request
2. decides where to route it
 - to RAM? to video card? to ...?
 - decision made via the [memory address map](#)


Bootstrapping


Can you pull yourself up by your own bootstraps?

A computer cannot run without first loading software but must be running before any software can be loaded.



Intel x86 Bootstrapping

1. BIOS (0xfffffff0)
 - ➡ POST ➡ HW init ➡ Find a boot device (FD,CD,HD...) ➡ Copy **sector zero (MBR)** to RAM (0x00007c00)
2. MBR – the first 512 bytes, contains
 - ▶ Small code (< 446 B), e.g. GRUB stage 1, for loading GRUB stage 2
 - ▶ the primary partition table ($16 \times 4 = 64$ B)
 - ▶ its job is to load the second-stage boot loader.
3. GRUB stage 2 — load the OS kernel into RAM
4.  startup
5. init — the first user-space process



```
$ sudo hd -n512 /dev/sda
```

Why Interrupt?

While a process is reading a disk file, can we do...

```
1 while(!done_reading_a_file())
2 {
3     let_CPU_wait();
4     // or...
5     lend_CPU_to_others();
6 }
7 operate_on_the_file();
```


Modern OS are Interrupt Driven

HW INT by sending a signal to CPU

SW INT by executing a **system call**

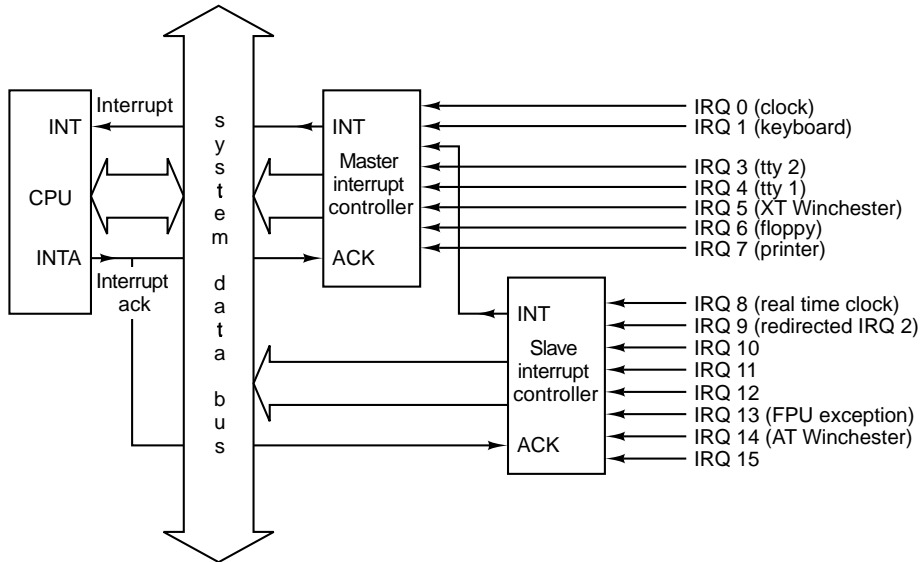
Trap (exception) is a software-generated INT caused by an error or by a specific request from an user program

Interrupt vector is an array of pointers ➡ the memory addresses of **interrupt handlers**. This array is indexed by a unique device number

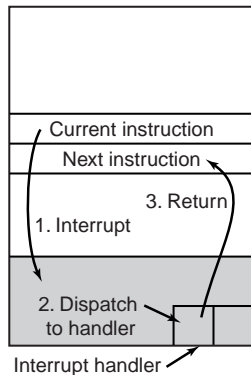
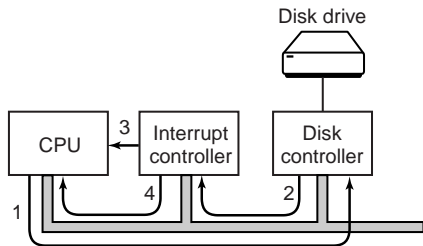
```
$ less /proc/devices
```

```
$ less /proc/interrupts
```

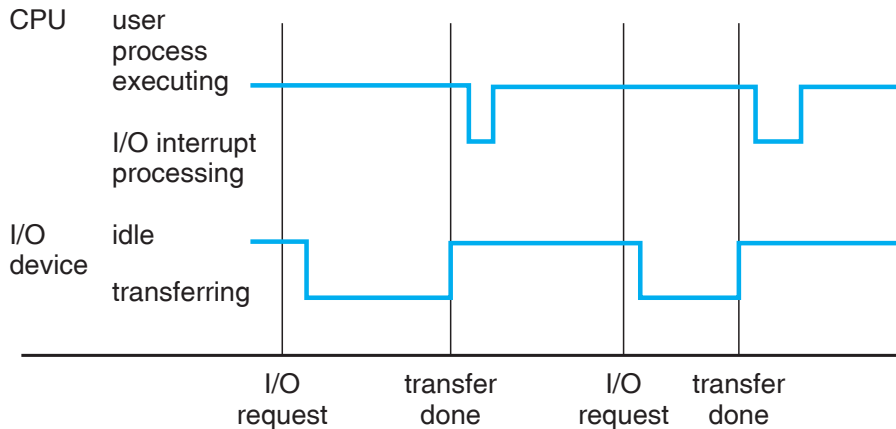
Programmable Interrupt Controllers



Interrupt Processing



Interrupt Timeline



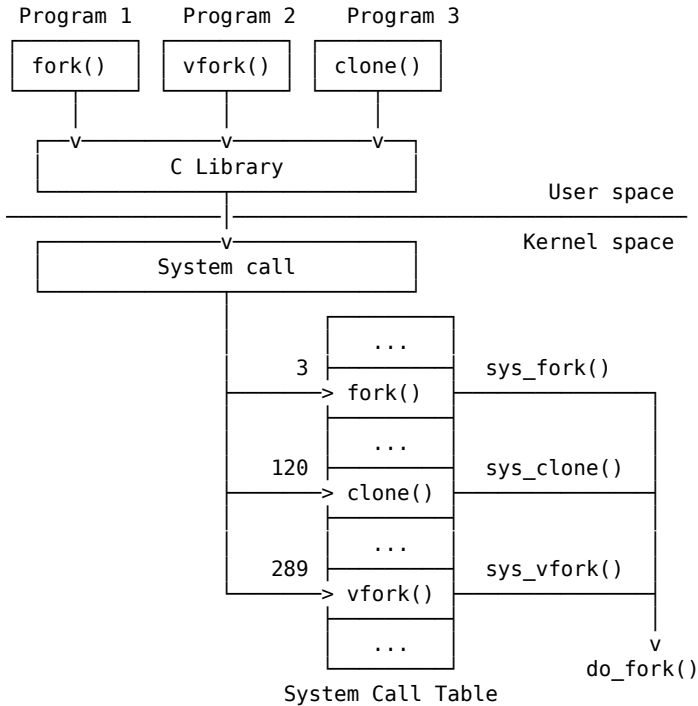
System Calls

A System Call

- ▶ is how a program requests a service from an OS kernel
- ▶ provides the interface between a process and the OS

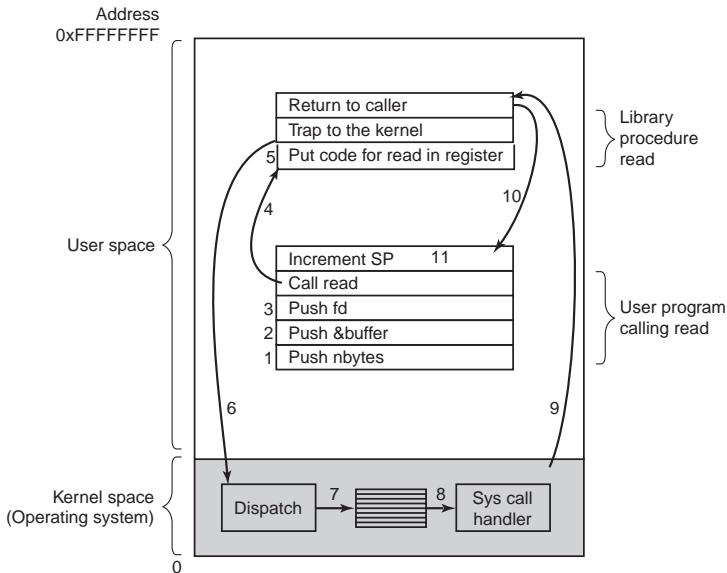
```
$ man 2 intro
```

```
$ man 2 syscalls
```



The 11 steps in making a system call

`read(fd,buffer,nbytes)`



Example

Linux INT 80h

Interrupt Vector Table: The very first 1KiB of x86 memory.

- ▶ $256 \text{ entries} \times 4\text{B} = 1\text{KiB}$
- ▶ Each entry is a complete memory address (segment:offset)
- ▶ It's populated by Linux and BIOS
- ▶ Slot 80h: address of the kernel services dispatcher (☛ sys-call table)

Example

```
1  Msg: db 'Hello, world'
2  MsgLen: equ $-Msg
3  mov eax,4          ; sys_write syscall = 4
4  mov ebx,1          ; 1 = STDOUT
5  mov ecx,Msg        ; offset of the message
6  mov edx,MsgLen     ; length of string
7  int 80h            ; call the kernel
```

```
$ nasm -f elf64 hello.asm -o hello.o
$ ld hello.o -o hello
$ ./hello
```

Process management

Call	Description
pid = fork()	Create a child process identical to the parent
pid = waitpid(pid, &statloc, options)	Wait for a child to terminate
s = execve(name, argv, environp)	Replace a process' core image
exit(status)	Terminate process execution and return status

File management

Call	Description
fd = open(file, how, ...)	Open a file for reading, writing or both
s = close(fd)	Close an open file
n = read(fd, buffer, nbytes)	Read data from a file into a buffer
n = write(fd, buffer, nbytes)	Write data from a buffer into a file
position = lseek(fd, offset, whence)	Move the file pointer
s = stat(name, &buf)	Get a file's status information

Directory and file system management

Call	Description
s = mkdir(name, mode)	Create a new directory
s = rmdir(name)	Remove an empty directory
s = link(name1, name2)	Create a new entry, name2, pointing to name1
s = unlink(name)	Remove a directory entry
s = mount(special, name, flag)	Mount a file system
s = umount(special)	Unmount a file system

Miscellaneous

Call	Description
s = chdir(dirname)	Change the working directory
s = chmod(name, mode)	Change a file's protection bits
s = kill(pid, signal)	Send a signal to a process
seconds = time(&seconds)	Get the elapsed time since Jan. 1, 1970

System Call Examples

fork()

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      printf("Hello World!\n");
7      fork();
8      printf("Goodbye Cruel World!\n");
9      return 0;
10 }
```

\$ man 2 fork

exec()

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main ()
5  {
6      printf("Hello World!\n");
7      if(fork() != 0 )
8          printf("I am the parent process.\n");
9      else {
10         printf("A child is listing the directory contents...\n");
11         execl("/bin/ls", "ls", "-al", NULL);
12     }
13     return 0;
14 }
```

\$ man 3 exec

GNU C Library

Linux API > POSIX API

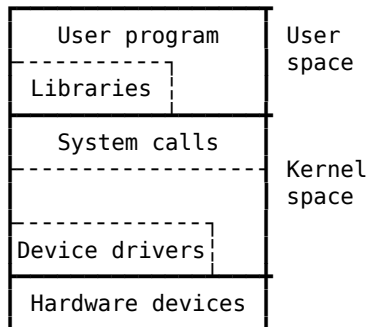
\$ man 7 libc

\$ man 3 intro

\$

\$

\$



Working With Files

File

A logical view of information storage

User's view

A file is the smallest storage unit on disk.

- ▶ Data cannot be written to disk unless they are within a file

UNIX view

Each file is a sequence of 8-bit bytes

- ▶ It's up to the application program to interpret this byte stream.

File

What is stored in a file?

Source code, object files, executable files, shell scripts, PostScript...

Different type of files have different structure

- ▶ UNIX looks at contents to determine type

Shell scripts start with “#!”

PDF start with “%PDF . . .”

Executables start with magic number

- ▶ Windows uses file naming conventions

executables end with “.exe” and “.com”

MS-Word end with “.doc”

MS-Excel end with “.xls”

File Types

Regular files: ASCII, binary

Directories: Maintaining the structure of the FS

In UNIX, everything is a file

Character special files: I/O related, such as terminals, printers ...

Block special files: Devices that can contain file systems, i.e. disks

Disks — logically, linear collections of blocks; disk driver translates them into physical block addresses

File Operations


POSIX file system calls

```
creat(name, mode)
open(name, flags)
close(fd)
link(oldname, newname)
unlink(name)
truncate(name, size)
ftruncate(fd, size)
stat(name, buffer)
fstat(fd, buffer)
```

```
read(fd, buffer, byte_count)
write(fd, buffer, byte_count)
lseek(fd, offset, whence)
chown(name, owner, group)
fchown(fd, owner, group)
chmod(name, mode)
fchmod(fd, mode)
utimes(name, times)
```

File System Implementation

A typical file system layout



On-Disk Information Structure

Boot block a MBR copy

Superblock Contains volume details

number of blocks	size of blocks
free-block count	free-block pointers
free FCB count	free FCB pointers

I-node Organizes the files **FCB (File Control Block)**, contains file details (metadata).

Superblock


Keeps information about the file system

- ▶ Type — ext2, ext3, ext4...
- ▶ Size
- ▶ Status — how it's mounted, free blocks, free inodes, ...
- ▶ Information about other metadata structures

```
$ sudo dumpe2fs /dev/sda1 | less
```

Implementing Files

Contiguous Allocation




File Allocation Table

File Name	Start Block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3

- simple;
- good for read only;

- fragmentation

Linked List (Chained) Allocation A pointer in each disk block




File Allocation Table

File Name	Start Block	Length
• • • File B • • •	• • • 1 • • •	• • • 5 • • •

- no waste block; access;
- slow random - not 2^n

Linked List (Chained) Allocation Though there is no external fragmentation, consolidation is still preferred.



File Allocation Table

File Name	Start Block	Length
• • • File B • • •	• • • 0 • • •	• • • 5 • • •

FAT: Linked list allocation with a table in RAM

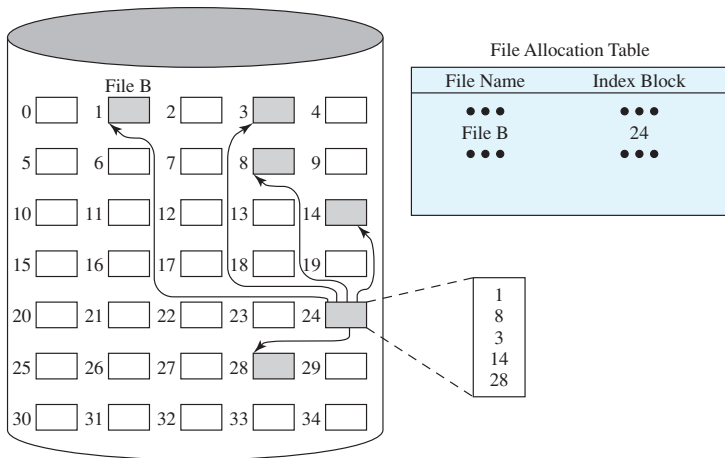
- ▶ Taking the pointer out of each disk block, and putting it into a table in memory
- ▶ fast random access (chain is in RAM)
- ▶ is 2^n
- ▶ the entire table must be in RAM

$disk \nearrow \Rightarrow FAT \nearrow \Rightarrow RAM_{used} \nearrow$

Physical
block

0		
1		
2	10	
3	11	
4	7	← File A starts here
5		
6	3	← File B starts here
7	2	
8		
9		
10	12	
11	14	
12	-1	
13		
14	-1	
15		← Unused block

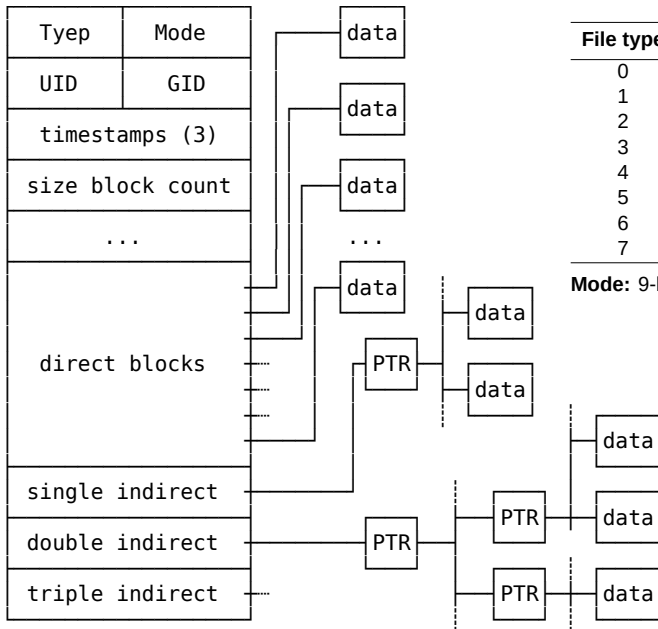
Indexed Allocation



I-node A data structure for each file. An i-node is in memory *only* if the file is open

$$files_{opened} \nearrow \Rightarrow RAM_{used} \nearrow$$

I-node



File type	Description
0	Unknown
1	Regular file
2	Directory
3	Character device
4	Block device
5	Named pipe
6	Socket
7	Symbolic link

Mode: 9-bit pattern

UNIX Treats a Directory as a File

.	2
..	2
bin	11116545
boot	2
cdrom	12
dev	3
:	:

Directory inode (128B)

Type	Mode
User ID	Group ID
File size	# blocks
# links	Flags
Timestamps (x3)	
Direct blocks (x12)	
Single indirect	
Double indirect	
Triple indirect	

Directory block

.	inode #
..	inode #
passwd	inode #
fstab	inode #
...	...

Indirect block

Direct blocks (x512)

File inode (128B)

Type	Mode
User ID	Group ID
File size	# blocks
# links	Flags
Timestamps (x3)	
Direct blocks (x12)	
Single indirect	
Double indirect	
Triple indirect	

File data block

Data

Block # of
block with
512 double
indirect
entries

Block # of
block with
512 single
indirect
entries

Block #s of
more
directory
blocks

open()

Why? To avoid constant searching

- ▶ Without open(), every file operation involves searching the directory for the file.

The steps in looking up /usr/ast/mbox

Root directory

1	.
1	..
4	bin
7	dev
14	lib
9	etc
6	usr
8	tmp

Looking up
usr yields
i-node 6

I-node 6
is for /usr

Mode size times
132

I-node 6
says that
/usr is in
block 132

Block 132
is /usr
directory

6	.
1	..
19	dick
30	erik
51	jim
26	ast
45	bal

/usr/ast
is i-node
26

I-node 26
is for
/usr/ast

Mode size times
406

I-node 26
says that
/usr/ast is in
block 406

Block 406
is /usr/ast
directory

26	.
6	..
64	grants
92	books
60	mbox
81	minix
17	src

/usr/ast/mbox
is i-node
60

```
fd open(pathname, flags)
```

A per-process **open-file table** is kept in the OS

- ▶ upon a successful `open()` syscall, a new entry is added into this table
- ▶ indexed by **file descriptor (fd)**
- ▶ `close()` to remove an entry from the table


To see files opened by a process, e.g. `init`

```
$ lsof -p 1
```

```
$ man 2 open
```

A process executes the following code:

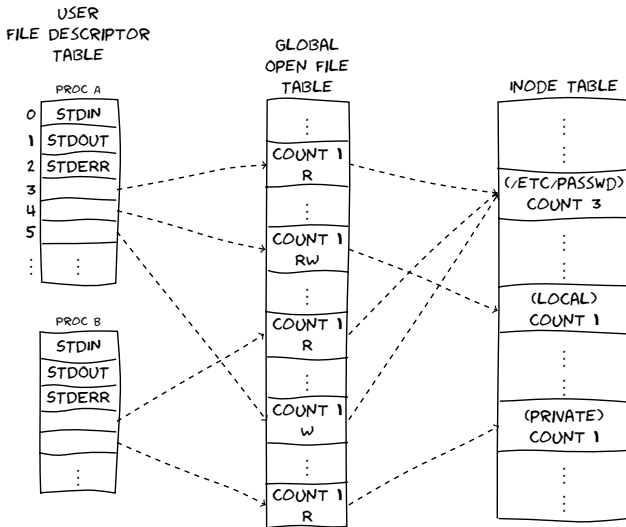
```
fd1 = open("/etc/passwd", O_RDONLY);  
fd2 = open("local", O_RDWR);  
fd3 = open("/etc/passwd", O_WRONLY);
```



One more process B:

```
fd1 = open("/etc/passwd", O_RDONLY);
```

```
fd2 = open("private", O_RDONLY);
```



write()

```
1  #include <unistd.h>
2
3  int main(void)
4  {
5      write(1, "Hello, world!\n", 14);
6
7      return 0;
8  }
```

\$ man 2 write

\$ man 3 write

read()

```
1  #include <unistd.h>
2
3  int main(void)
4  {
5      char buffer[10];
6
7      read(0, buffer, 10);
8
9      write(1, buffer, 10);
10
11     return 0;
12 }
```

\$ man 2 read

\$ man 3 read

cp

```
1  #include <sys/types.h>    /* include necessary header files */
2  #include <fcntl.h>
3  #include <stdlib.h>
4  #include <unistd.h>
5
6  #define BUF_SIZE 4096      /* use a buffer size of 4096 bytes */
7  #define OUTPUT_MODE 0700  /* protection bits for output file */
8
9  int main(int argc, char *argv[])
10 {
11     int in, out, rbytes, wbytes;
12     char buf[BUF_SIZE];
13
14     if (argc != 3) exit(1);
15
16     if ( (in = open(argv[1], O_RDONLY)) < 0 ) exit(2); /* open source file */
17
18     if ( (out = creat(argv[2], OUTPUT_MODE)) < 0 ) exit(3); /* create destination file */
19
20     while (1) { /* Copy loop */
21         if ( (rbytes = read(in, buf, BUF_SIZE)) <= 0 ) break; /* read a block of data */
22         if ( (wbytes = write(out, buf, rbytes)) <= 0 ) exit(4); /* write data */
23     }
24
25     close(in);
26     close(out);
27     if (rbytes == 0) exit(0); /* no error on last read */
28     else exit(5);           /* error on last read */
29 }
```


stdio — The Standard I/O Library

System calls: `open()`, `read()`, `write()`, `close()`...

Library functions: `fopen()`, `fread()`, `fwrite`, `fclose()`...

Avoid calling syscalls directly as much as you can

- ▶ Portability
- ▶ Buffered I/O



open() vs. fopen()

open()

```
1  #include <unistd.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4
5  int main()
6  {
7      char c;
8      int in;
9      in = open("/tmp/1m.test", O_RDONLY);
10
11     while (read(in, &c, 1) == 1);
12
13     return 0;
14 }
```

```
$ strace -c ./open
```

fopen() — Buffered I/O

```
1  #include <stdio.h>
2
3  int main(void)
4  {
5      FILE *stream;
6
7      stream = fopen("/tmp/1m.test", "r");
8
9      while ( fgetc(stream) != EOF );
10
11     fclose(stream);
12
13     return 0;
14 }
```

```
$ strace -c ./fopen
```

```
$ dd if=/dev/zero of=/tmp/1m.test bs=1k count=1024
```

cp — With stdio

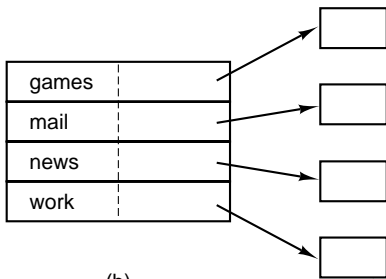
```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(int argc, char *argv[])
5  {
6      FILE *in, *out;
7      int c=0;
8
9      if (argc != 3) exit(1);
10
11     in = fopen(argv[1], "r");
12     out = fopen(argv[2], "w");
13
14     while ( (c = fgetc(in)) != EOF )
15         fputc(c, out);
16
17     return 0;
18 }
```

Homework: Try fread()/fwrite() instead.

Implementing Directories

games	attributes
mail	attributes
news	attributes
work	attributes

(a)



(b)

Data structure
containing the
attributes

(a) A simple directory (Windows)

- ▶ fixed size entries
- ▶ disk addresses and attributes in directory entry

(b) Directory in which each entry just refers to an i-node (UNIX)


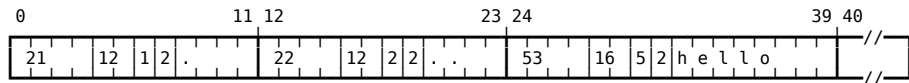
Directory entry in `glibc`

```
1  struct dirent {
2      ino_t      d_ino;          /* Inode number */
3      off_t      d_off;          /* Not an offset; see below */
4      unsigned short d_reclen;    /* Length of this record */
5      unsigned char d_type;       /* Type of file; not supported
6                                   by all filesystem types */
7      char       d_name[256];    /* Null-terminated filename */
8  };
```

\$ `man readdir`

\$ `view /usr/include/x86_64-linux-gnu/bits/dirent.h`

Ext2 Directories



- ▶ Directories are special files
- ▶ “.” and “..” first
- ▶ Padding to $4 \times$
- ▶ inode number is 0 — deleted file

ls

```
1  #include <sys/types.h>
2  #include <dirent.h>
3  #include <stddef.h>
4  #include <stdio.h>
5
6  int main(int argc, char *argv[])
7  {
8      DIR *dp;
9      struct dirent *entry;
10
11     dp = opendir(argv[1]);
12
13     while ( (entry = readdir(dp)) != NULL ){
14         printf("%s\n", entry->d_name);
15     }
16
17     closedir(dp);
18
19     return 0;
20 }
```

The real ls.c?

116 A4 pages
5308 lines

Do one thing, and do
it really well.

\$ apt source coreutils

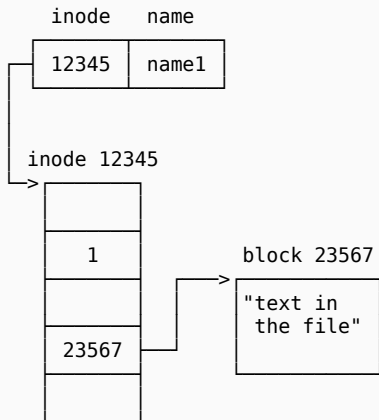
mkdir(), chdir(), rmdir(), getcwd()

```
1  #include <sys/stat.h>
2  #include <sys/types.h>
3  #include <unistd.h>
4  #include <stdio.h>
5
6  int main(int argc, char *argv[])
7  {
8      char s[100];
9      if ( mkdir(argv[1], S_IRUSR|S_IXUSR) == 0 )
10         chdir(argv[1]);
11     printf("PWD = %s\n", getcwd(s,100));
12     rmdir(argv[1]);
13     return 0;
14 }
```


Hard Links

Hard links → the same inode

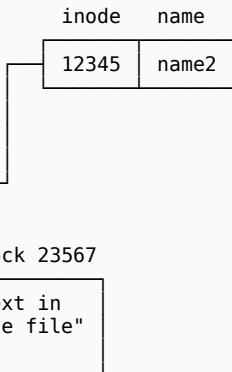
directory entry in /dirA



dir entry in /dirA




dir entry in /dirB




Drawback

C's directory



(a)


B's directory



(b)

C's directory

B's directory



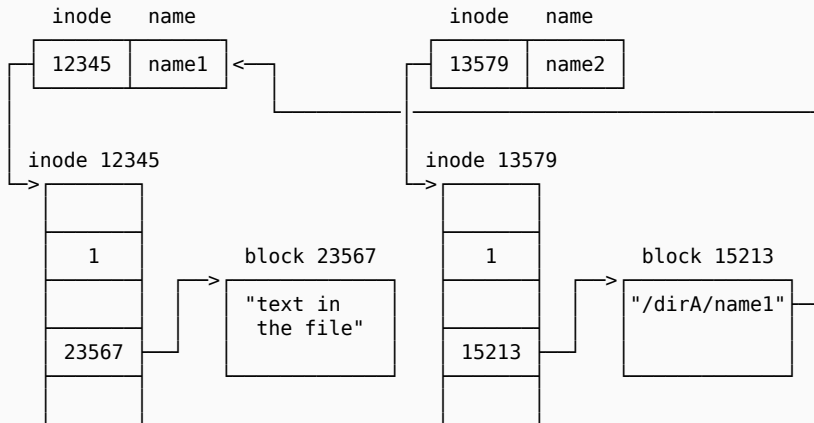
(c)

Symbolic Links

A symbolic link has its own inode → a directory entry

directory entry in /dirA

directory entry in /dirB



Fast symbolic link: Short path name (< 60 chars) needs no data block.
Can be stored in the 15 pointer fields

link(), unlink(), symlink()

```
1  #include <unistd.h>
2  #include <stdio.h>
3
4  int main(int argc, char *argv[])
5  {
6      link(argv[1], argv[2]);
7      perror(argv[0]);
8      return 0;
9  }
10
11  /* symlink(argv[1], argv[2]); */
12  /* unlink(argv[1]); */
```

The Linux Environment

Command Line Options

getopt.c

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(int argc, char* argv[]) {
5      int opt;
6
7      while ((opt = getopt(argc, argv, "hf:l")) != -1) {
8          switch (opt) {
9              case 'h':
10                 printf("Usage: %s [-h] [-f file] [-l]\n", argv[0]);
11                 break;
12              case 'l':
13                 printf("option: %c\n", opt);
14                 break;
15              case 'f':
16                 printf("filename: %s\n", optarg);
17                 break;
18             }
19         }
20         return 0;
21     }
```

\$ man 3 getopt

Command Line Options

getopt.sh

```
1  #!/bin/bash
2
3  while getopts hf:l OPT; do
4      case $OPT in
5          h) echo "usage: `basename $0` [-h] [-f file] [-l]"
6              exit 1 ;;
7          l) echo "option: l" ;;
8          f) echo "filename: $OPTARG" ;;
9      esac
10  done
```

```
$ ./getopt.sh -h
```

```
$ ./getopt.sh -lf filename
```

```
$ ./getopt.sh -l -f filename
```

```
$ ./getopt.sh -f filename -l
```

Environment Variable

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
4  extern char** environ;
5
6  int main() {
7      char** env = environ;          $ env
8                                      $ man 3 getenv
9      while (*env) {                $ man 3 putenv
10         printf("%s\n", *env);
11         env++;
12     }
13
14     return 0;
15 }
```

Time and Date

```
1  #include <time.h>
2  #include <stdio.h>
3
4  int main(void)
5  {
6      time_t t = time(NULL); /* long int */
7
8      printf("epoch time:\t%ld\n",t);
9      printf("calendar time:\t%s", ctime(&t));
10
11     return 0;
12 }
```

► January 1 1970 — start of the Unix epoch

\$ man 3 time

\$ man 3 ctime

Temporary Files

mkstemp.c

```
1  #include <stdlib.h>
2  #include <unistd.h>
3  #define _GNU_SOURCE
4  #include <stdio.h>
5
6  int main(int argc, char *argv[])
7  {
8      char c, *f;
9
10     asprintf(&f, "%sXXXXXX", argv[1]);
11     int tmp = mkstemp(f);
12
13     while ( read(0, &c, 1) == 1)
14         write(tmp, &c, 1);
15
16     unlink(f);
17     free(f);
18     return 0;
19 }
```

mktemp.sh

```
1  #!/bin/bash
2
3  tmp=$(mktemp)
4
5  while read LINE; do
6      echo $LINE >> $tmp
7  done
8
9
```

\$ man 3 mkstemp

\$ man 3 tmpfile

\$ man 3 asprintf

Logging

syslog.c

```
1  #include <syslog.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4
5  int main(int argc, char *argv[])
6  {
7      if ( open(argv[1], O_RDONLY) < 0 )
8          syslog(LOG_ERR | LOG_USER, "%s - %m\n", argv[1]);
9      else
10         syslog(LOG_INFO | LOG_USER, "%s - %m\n", argv[1]);
11     return 0;
12 }
```

logger.sh

```
1  #!/bin/bash
2
3  [[ -f "$1" ]] && logger "$1 exists." || logger "$1 not found."
```

Terminal

IDE

Processes and signals

IPC

User Interface