

# Processes

Wang Xiaolin  
wx672ster@gmail.com

June 13, 2013

## Contents

<b>Todo list</b>	<b>1</b>
<b>1 Processes, Lightweight Processes, and Threads</b>	<b>2</b>
<b>2 Process Descriptor</b>	<b>2</b>
2.1 Process State . . . . .	3
2.2 Identifying a Process . . . . .	4
2.3 Relationships Among Processes . . . . .	14
2.4 How Processes Are Organized . . . . .	17
2.5 Process Resource Limits . . . . .	18
<b>3 Process Switch</b>	<b>19</b>
<b>4 Creating Processes</b>	<b>21</b>
<b>5 Destroying Processes</b>	<b>25</b>

### Textbook:

- Chapter 3, *Processes*, [BC05]
- Chapter 3, *Process Management*, [Lov10]
- Chapter 2, *Process Management and Scheduling*, [Mau08]

## References

- [BC05] D.P. Bovet and M. Cesatí. *Understanding The Linux Kernel*. 3rd ed. O'Reilly, 2005.
- [Lov10] R. Love. *Linux Kernel Development*. Developer's Library. Addison-Wesley, 2010.
- [Mau08] W. Mauerer. *Professional Linux Kernel Architecture*. John Wiley & Sons, 2008.
- [RFS05] C.S. Rodriguez, G. Fischer, and S. Smolski. *The Linux kernel primer: a top-down approach for x86 and PowerPC architectures*. Prentice Hall Professional Technical Reference, 2005.

## Todo list

Draw a process relationship graph . . . . .	4
---	---

# 1 Processes, Lightweight Processes, and Threads

---

## Processes

A **process** is

- an instance of a program in execution
- a dynamic entity (has lifetime)
- a collection of data structures describing the execution progress
- the unit of system resources allocation

The Linux kernel internally refers to processes as *tasks*.



---

## When A Process Is created

The child

- is almost identical to the parent
  - has a logical copy of the parent's address space
  - executes the same code
- has its own data (stack and heap)



---

## Multithreaded Applications

### Threads

- are execution flows of a process
- share a large portion of the application data structures

### Lightweight processes (LWP) — Linux way of multithreaded applications

- each LWP is scheduled individually by the kernel
  - no nonblocking syscall is needed
- LWPs may share some resources, like the address space, the open files, and so on.



---

## 2 Process Descriptor

---

## Process Descriptor

To manage processes, the kernel must have a clear picture of what each process is doing.

- the process's priority
- running or blocked
- its address space
- files it opened
- ...

**Process descriptor:** a `task_struct` type structure containing all the information related to a single process.

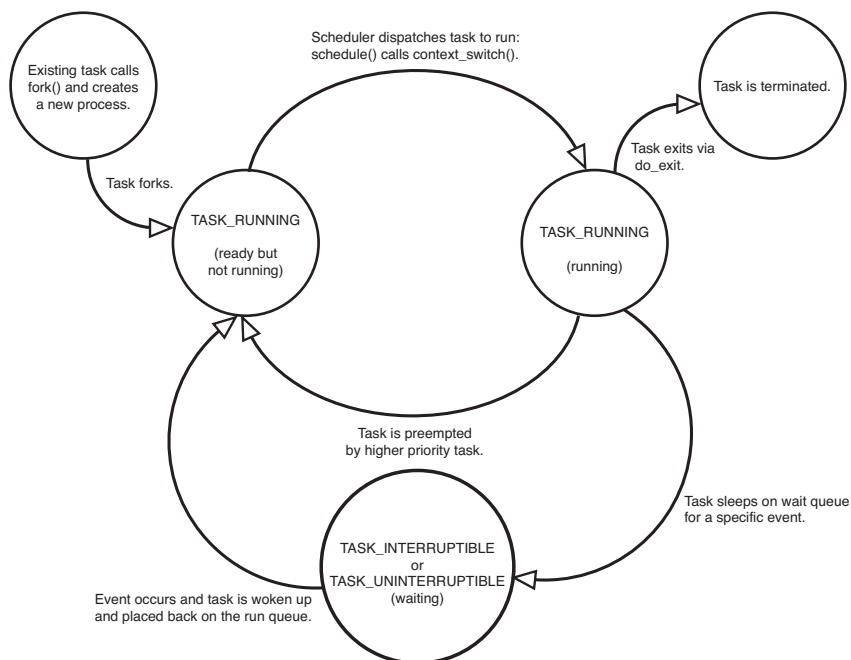
```
struct task_struct {  
    /* 160 lines of code in 2.6.11 */  
};
```



## 2.1 Process State

---

### Process State



(sec 3.2.1 in [BC05]) The `state` field of `task_struct` is an array of flags, each of which describes a possible process state. In the current Linux version, these states are mutually exclusive, and hence exactly one flag of state always is set; the remaining flags are cleared. The following are the possible process states:

**TASK\_RUNNING** The process is either executing on a CPU or waiting to be executed.

**TASK\_INTERRUPTIBLE** The process is suspended (sleeping) until some condition becomes true. Raising a hardware interrupt, releasing a system resource the process is waiting for, or delivering a signal are examples of conditions that might wake up the process (put its state back to **TASK\_RUNNING**).

**TASK\_UNINTERRUPTIBLE** Like **TASK\_INTERRUPTIBLE**, except that delivering a signal to the sleeping process leaves its state unchanged. This process state is seldom used. It is valuable, however, under certain specific conditions in which a process must wait until a given event occurs without being interrupted. For instance, this state may be used when a process opens a device file and the corresponding device driver starts probing for a corresponding hardware device. The device driver must not be interrupted until the probing is complete, or the hardware device could be left in an unpredictable state.

**TASK\_STOPPED** Process execution has been stopped; the process enters this state after receiving a **SIGSTOP**, **SIGTSTP**, **SIGTTIN**, or **SIGTTOU** signal.

**TASK\_TRACED** Process execution has been stopped by a debugger. When a process is being monitored by another (such as when a debugger executes a **ptrace()** system call to monitor a test program), each signal may put the process in the **TASK\_TRACED** state.

**EXIT\_ZOMBIE** Process execution is terminated, but the parent process has not yet issued a **wait4()** or **waitpid()** system call to return information about the dead process. Before the **wait()**-like call is issued, the kernel cannot discard the data contained in the dead process descriptor because the parent might need it.

**EXIT\_DEAD** The final state: the process is being removed by the system because the parent process has just issued a **wait4()** or **waitpid()** system call for it. Changing its state from **EXIT\_ZOMBIE** to **EXIT\_DEAD** avoids race conditions due to other threads of execution that execute **wait()**-like calls on the same process.

## 2.2 Identifying a Process

---

### PID AND TGID

- kernel finds a process by its *process descriptor pointer* pointing to a **task\_struct**
- users find a process by its **PID**
- all the threads of a multithreaded application share the same identifier

**tgid**: the PID of the thread group leader

```
struct task_struct {  
    ...  
    pid_t pid;  
    pid_t tgid;  
    ...  
};
```

```
~$ ps -eo pgid,ppid,pid,tgid,tid,nlwp,comm -sort pid
```



TODO: Draw a process relationship graph

---

**How many PIDs can there be?**

- `#define PID_MAX_DEFAULT 0x8000`
- Max PID number = `PID_MAX_DEFAULT - 1 = 32767`
- `$ cat /proc/sys/kernel/pid_max`

### Which are the free PIDs?

```
static pidmap_t pidmap_array[PIDMAP_ENTRIES] =
{
    [ 0 ... PIDMAP_ENTRIES-1 ] =
    { ATOMIC_INIT(BITS_PER_PAGE), NULL }
};
```

`pidmap_array` consumes a single page.




---

### Process Descriptor Handling

**thread\_union:** 2 consecutive page frames (8K) containing


- a process kernel stack
- a `thread_info` structure

```
union thread_union {
    struct thread_info thread_info;
    unsigned long stack[2048]; /* 1024 for 4KB stacks */
};
```



- The kernel uses the `alloc_thread_info` and `free_thread_info` macros to allocate and release the memory area storing a `thread_info` structure and a kernel stack.
  - Benefits of storing `thread_info` and kernel mode stack together (Sec 3.2.2.2, *Identifying the current process*, [BC05]):
    - Efficiency: the kernel can easily obtain the address of the `thread_info` structure of the process currently running on a CPU from the value of the `esp` register.
    - For multi-processor systems: the kernel can easily obtain the address of the `thread_info` structure of the process currently running on a CPU from the value of the `esp` register. Earlier versions of Linux did not store the kernel stack and the process descriptor together. Instead, they were forced to introduce a global static variable called `current` to identify the process descriptor of the running process. On multi-processor systems, it was necessary to define `current` as an array one element for each available CPU.
-

## Kernel Stack



### More about kernel stack:

- stackoverflow: **Does there exist Kernel stack for each process?** In Linux, each task (userspace or kernel thread) has a kernel stack of either 8kb or 4kb, depending on kernel configuration. There are indeed separate stack pointers, however, only one is present in the CPU at any given time; if userspace code is running, the kernel stack pointer to be used on exceptions or interrupts is specified by the task-state segment, and if kernel code is running, the user stack pointer is saved in the context structure located on the kernel stack.
- stackoverflow: **kernel stack for Linux process?**
  - There is just one common kernel memory. In it each process has it's own **task\_struct** + kernel stack (by default 8K).
  - In a context switch
    1. the old stack pointer value is stored in the **task\_struct** of the old process;
    2. the stack pointer for the new process is read from the **task\_struct** of this new process.
  - The kernel stack is a non-shared space where system calls can put their data. If you'd share them between processes, several kernel routines could use the same stack at the same time ⇒ data corruption.

- stackoverflow: **kernel stack vs user-mode application stack?** Conceptually, both are the same data structure: a stack.

*security concern*

The reason why there are two different stack per thread is because in user mode, code must not be allowed to mess up kernel memory. When switching to kernel mode, a different stack in memory only accessible in kernel mode is used for return addresses and so on.

If the user mode had access to the kernel stack, it could modify a jump address (for instance), then do a system call; when the kernel jumps to the previously modified address, your code is executed in kernel mode!

Also, security-related information/information about other processes (for synchronization) might be on the kernel stack, so the user mode should not have read access to it either.

- Kernel stack vs. User stack
- Kernel Stack Initialization

---

```

struct thread_info {
    struct task_struct    *task;           /* main task structure */
    struct exec_domain    *exec_domain;    /* execution domain */
    unsigned long         flags;           /* low level flags */
    unsigned long         status;          /* thread-synchronous flags */
    __u32                 cpu;             /* current CPU */
    __s32                 preempt_count;   /* 0 => preemptable, <0 => BUG */

    mm_segment_t          addr_limit;      /* thread address space:
                                           0-0xBFFFFFFF for user-thread
                                           0-0xFFFFFFFF for kernel-thread
                                           */
    struct restart_block   restart_block;
    unsigned long         previous_esp;    /* ESP of the previous stack in case
                                           of nested (IRQ) stacks
                                           */
    __u8                  supervisor_stack[0];
};

```

## Why both `task_struct` and `thread_info`?

- There wasn't a `thread_info` in pre-2.6 kernel
- Size matters



More about `thread_info`:

- (www.linuxjournal.com: The Linux Process Model) Since 2.2.x, the `task_struct` is allocated at the bottom of the kernel stack. We can overlap the `task_struct` on the kernel stack because the `task_struct` is a per-task structure exactly as the kernel stack.
- Why are `task_struct` and `thread_info` separate?

Quote from <http://www.spinics.net/lists/newbies/msg22259.html>:

Q: There is one `task_struct` and `thread_info` for each process, and they link to each other, right? So why are they separate structs?

A: `thread_info` is embedded in kernel stack, so we can easily get the `task_struct` using kernel stack.

Quote from <http://www.spinics.net/lists/newbies/msg22263.html>

Q: That is a fine workaround for the structs being separate. But why not keep all the data in one struct on the kernel stack, like it was done in 2.4?

A: `task_struct` is huge. it's around 1.7KB on a 32 bit machine. on the other hand, you can easily see that `thread_info` is much slimmer.

*size matters*

kernel stack is either 4 or 8KB, and either way a 1.7KB is pretty much, so storing a slimmer struct, that points to `task_struct`, immediately saves a lot of stack space and is a scalable solution.

- `supervisor_stack[0]` in `struct thread_info` is a *flexible array member*.
  - Sec 6.17, *Arrays of Length Zero*, GCC Manual
  - What is the advantage of using zero-length arrays in C?
  - MEMxx-C. Understand how flexible array members are to be used
  - (Kerneltrap.org: Regarding `thread_info`) In the case of `tread_info`: `supervisor_stack` seems to be the kernel stack, in this case an integer number of pages (i.e. 4096 or 8192 bytes) is allocated for the rest of `struct thread_info` + the stack, so

- a) you don't need the real size because you allocate the fixed amount anyway, and
- b) you would have to declare an array of `(4096 - offsetof(struct thread_info, supervisor_stack))` bytes inside the struct itself and that is just not possible.

Or not declare the stack inside the struct but do offset calculations, but the zero sized array is more elegant and maintainable. There is nothing to fear about that, this is perfectly normal C.

- process descriptors (`task_structs`) are stored in dynamic memory (`ZONE_HIGHMEM`) rather than in the memory area permanently assigned to the kernel
- Kernel stack size, 8K vs. 4K
  - (Sec 3.2.2.1, *Process descriptors handling*, [BC05]) In Sec 2.3, *Segmentation in Linux*, [BC05], we learned that a process in Kernel Mode accesses a stack contained in the kernel data segment, which is different from the stack used by the process in User Mode. Because kernel control paths make little use of the stack, only a few thousand bytes of kernel stack are required. Therefore, 8 KB is ample space for the stack and the `thread_info` structure. However, when stack and `thread_info` structure are contained in a single page frame, the kernel uses a few additional stacks to avoid the overflows caused by deeply nested interrupts and exceptions (see Chapter 4, *Interrupts and Exceptions*, [BC05]).
  - (lwn.net: 4K stacks in 2.6) Each process on the system has its own kernel stack, which is used whenever the system goes into kernel mode while that process is running. Since each process requires a kernel stack, the creation of a new process requires an order 1 allocation. So the two-page kernel stacks can limit the creation of new processes, even though the system as a whole is not particularly short of resources. Shrinking kernel stacks to a single page eliminates this problem and makes it easy for Linux systems to handle far more processes at any given time.
  - in `linux-2.6.11/include/asm-i386/thread_info.h`

*problem in finding 2 consecutive page frames*

```
#ifndef CONFIG_4KSTACKS
#define THREAD_SIZE (4096)
#else
#define THREAD_SIZE (8192)
#endif
```

---

## `thread_info` and `task_struct` are mutually linked

```
struct thread_info {
    struct task_struct *task; /* main task structure */
    ...
};

struct task_struct {
    ...
    struct thread_info *thread_info;
    ...
};
```





## Identifying The Current Process

Efficiency benefit from `thread_union`

- Easy get the base address of `thread_info` from `esp` register by masking out the 13 least significant bits of `esp`

### `current_thread_info()`

```
/* how to get the thread information struct from C */
static inline struct thread_info *current_thread_info(void)
{
    struct thread_info *ti;
    __asm__ ("andl_%%esp,_%0;" : "=r" (ti) : "0" (~(THREAD_SIZE - 1)));
    return ti;
}
```

Can be seen as:

```
movl $0xffffe000,%ecx /* or 0xfffff000 for 4KB stacks */
andl %esp,%ecx
movl %ecx,p
```



(Stackoverflow: Understanding the getting of `task_struct` pointer from process kernel stack) Each process only gets 8192 bytes of kernel stack, aligned to a 8192-byte boundary, so whenever the stack pointer is altered by a `push` or a `pop`, the low 13 bits are the only part that changes.  $2^{13} == 8192$ .

$$\begin{aligned}\sim (\text{THREAD\_SIZE} - 1) &= \sim (8k - 1) \\ &= \sim (0x00002000 - 1) \\ &= \sim 0x00001fff \\ &= 11111111111111111110000000000000b\end{aligned}$$

---

## To get the process descriptor pointer

`current_thread_info()->task`

```
movl $0xffffe000,%ecx /* or 0xfffff000 for 4KB stacks */
andl %esp,%ecx
movl (%ecx),p
```

Because the `task` field is at offset 0 in `thread_info`, after executing these 3 instructions `p` contains the process descriptor pointer.

### `current` — a marco pointing to the current running task

```
static inline struct task_struct * get_current(void)
{
    return current_thread_info()->task;
}

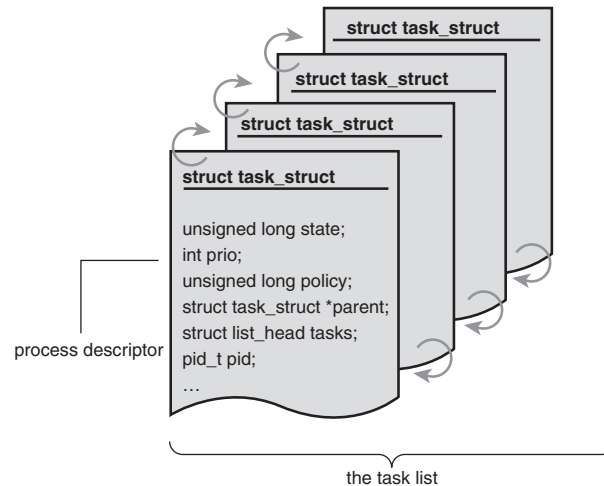
#define current get_current()
```



## Task List

The kernel stores the list of processes in a [circular doubly linked list](#) called the [task list](#).

**Swapper** The head of this list, `init_task`, process 0.



**Q:** Stackoverflow: Why do we need a swapper task in linux?

*swapper*

**A:** The reason is historical and programatic. The idle task is the task running, if no other task is runnable, like you said it. It has the lowest possible priority, so that's why it's running of no other task is runnable.

**Programatic reason:** This simplifies process scheduling a lot, because you don't have to care about the special case: "What happens if no task is runnable?", because there always is at least one task runnable, the idle task. Also you can count the amount of cpu time used per task. Without the idle task, which task gets the cpu-time accounted no one needs?

**Historical reason:** Before we had cpus which are able to step-down or go into power saving modes, it HAD to run on full speed at any time. It ran a series of `NOP`-instructions, if no tasks were runnable. Today the scheduling of the idle task usually steps down the cpu by using `HLT`-instructions (halt), so power is saved. So there is a functionality somehow in the idle task in our days.

In Windows you can see the idle task in the process list, it's the idle process.

**Q:** superuser.com: What is the main purpose of the swapper process in Unix?

**A:** It hasn't been a swapper process since the 1990s, and swapping hasn't really been used since the 1970s.

Unices stopped using swapping a long time ago. They've been demand-paged operating systems for a few decades — since System V R2V5 and 4.0BSD. The swapper process, as was, used to perform process swap operations. It used to swap entire processes — including all of the kernel-space data structures for the process — out to disc and swap them back in again. It would be woken up, by the kernel, on a regular basis, and would scan the process table to determine what swapped-out-and-ready-to-run processes could be swapped in and what swapped-in-but-asleep processes could be swapped out. Any textbook on Unix from the 1980s will go into this in more detail, including the swap algorithm. But it's largely irrelevant to demand-paged Unices,

even though they retained the old swap mechanism for several years. (The BSDs tried quite hard to avoid swapping, in favour of paging, for example.)

Process #0 is the first process in the system, hand-crafted by the kernel. It `fork()`s process 1, the first user process. What it does other than that is dependent from what Unix the operating system actually is. As mentioned, the BSDs before FreeBSD 5.0 retained the old swapper mechanism, and process #0 simply dropped into the swapper code in the kernel, the `scheduler()` function, after it had finished system initialization. System V was much the same, except that process #0 was conventionally named `sched` rather than `swapper`. (The names are pretty much arbitrary choices.) In fact, most — possibly all — Unices had a (largely unused) old swapper mechanism that hung around as process #0.

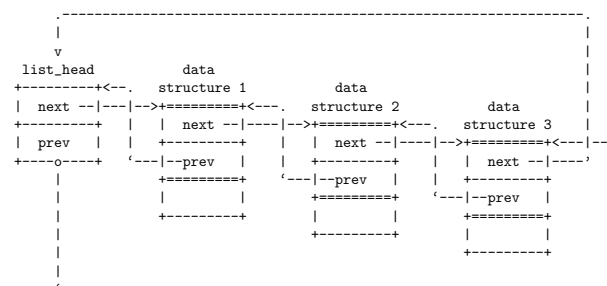
Linux, traditionally, has been somewhat different to the Unices, in that process #0 is the idle process, running `cpu_idle()`. It simply does nothing, in an infinite loop. It exists so that there's always a task ready to be scheduled.

Even this is an out-of-date description. The late 1980s and early 1990s was the advent of multi-threading operating systems, as a result of which process #0 became simply the system process. In the world of old single-threaded Unices, one could only get a separate flow of execution to do a continuous kernel task by `fork()`ing a process. So all of the kernel tasks (vmdaemon, pagedaemon, pagezero, bufdaemon, syncer, ktrace, and so forth on FreeBSD systems, for example) were low-numbered processes, `fork()`ed by process #0 after it `fork()`ed init. In multiple-threaded Unices, it makes no sense to create a whole new process context for something that runs entirely in kernel space, and doesn't need an address space, file descriptor table, and what not all to itself. So all of these tasks became (essentially) threads, sharing the address space of a single system process.

Along the way, several Unices finally lost the old swapper mechanism, that they were trying their utmost to avoid ever using anyway. OpenBSD's initialization code now simply drops into a `while(1) tsleep(...);` loop, for example, with no swapping mechanism anywhere.

So nowadays process #0 on a Unix is the system process, which effectively holds a number of kernel threads doing a number of things, ranging from page-out operations, through filesystem cache flushes and buffer zeroing, to idling when there's nothing else to run.

## Doubly Linked List



```

struct list_head {
    struct list_head *next, *prev;
};

struct task_struct {
    ...
    struct list_head tasks;
    ...
}

```



---

## List operations

**SET\_LINKS** insert into the list

**REMOVE\_LINKS** remove from the list

**for\_each\_process** scan the whole process list

```
#define for_each_process(p) \
    for (p = &init_task ; (p = next_task(p)) != &init_task ; )
```

**list\_for\_each** iterate over a list

```
#define list_for_each(pos, head) \
    for (pos = (head)->next; prefetch(pos->next), pos != (head); \
         pos = pos->next)
```

## Example: Iterate over a process' children

```
struct task_struct *task;
struct list_head *list;
list_for_each(list, &current->children) {
    task = list_entry(list, struct task_struct, sibling);
    /* task now points to one of current's children */
}
```



Expend the macro **list\_for\_each()**:

```
struct task_struct *task;
struct list_head *list;
for (list = (&current->children)->next;
     prefetch(list->next), list != (&current->children);
     list = list->next)
{
    task = list_entry(list, struct task_struct, sibling);
}
```

- (Linux Kernel Linked List Explained) we need to access *the item itself* not the variable "list" in the item! macro **list\_entry()** does just that.

**tmp = list\_entry(pos, struct kool\_list, list);**

Given

1. a pointer to **struct list\_head**,
2. type of data structure it is part of, and
3. it's name (**struct list\_head**'s name in the data structure)

it returns a pointer to the data structure in which the pointer is part of.

- Sec 3.2.6, *The Process Family Tree*, [Lov10].

---

## A task can be in multiple lists

```


struct task_struct {
    struct list_head run_list;
    struct list_head tasks;
    struct list_head ptrace_children;
    struct list_head ptrace_list;
    struct list_head children; /* list of my children */
    struct list_head sibling; /* linkage in my parent's children list */
}

```



## The List Of **TASK\_RUNNING** Processes

- Each CPU has its own runqueue
- Each runqueue has 140 lists
- One list per process priority
- Each list has zero to many tasks



```

struct task_struct {
    ...
    int prio, static_prio;
    struct list_head run_list;
    prio_array_t *array;
    ...
};

```



**prio** (Sec 3.2.2.1, *prio*, [RFS05]) the dynamic priority of a process is a value that depends on the processes' scheduling history and the specified *nice* value. It is updated at *sleep* time, which is when the process is not being executed and when timeslice is used up. This value, *prio*, is related to the value of the *static\_prio* field described next. The *prio* field holds  $\pm 5$  of the value of *static\_prio*, depending on the process' history; it will get a +5 bonus if it has slept a lot and a -5 handicap if it has been a processing hog and used up its timeslice.

**static\_prio** (Sec 3.2.2.2, *static\_prio*, [RFS05]) is equivalent to the *nice* value. The default value of *static\_prio* is  $\text{MAX\_PRIO} - 20$ . In our kernel, *MAX\_PRIO* defaults to 140.

**run\_list** (Sec 3.2.2.5, *The lists of TASK\_RUNNING processes*, [BC05]) The trick used to achieve the scheduler speedup consists of splitting the runqueue in many lists of runnable processes, one list per process priority. Each *task\_struct* descriptor includes a *run\_list* field of type *list\_head*. If the process priority is equal to *k* (a value ranging between 0 and 139), the *run\_list* field links the process descriptor into the list of runnable processes having priority *k*.

## Each Runqueue Has A `prio_array_t` Struct

```
typedef struct prio_array prio_array_t;

struct prio_array {
    unsigned int nr_active;
    unsigned long bitmap[BITMAP_SIZE];
    struct list_head queue[MAX_PRIO];
};
```

**nr\_active:** The number of process descriptors linked into the lists (the whole runqueue)

**bitmap:** A priority bitmap. Each flag is set if the priority list is not empty

**queue:** The 140 heads of the priority lists



- Each *runqueue* (NOT each priority) has a `prio_array_t`. (Sec 3.2.2.5, *The lists of TASK\_RUNNING processes*, [BC05]) As we'll see, the kernel must preserve a lot of data for every runqueue in the system; however, the main data structures of a runqueue are the lists of process descriptors belonging to the runqueue; all these lists are implemented by a single `prio_array_t` data structure.

---

## To Insert A Task Into A Runqueue List

```
static void enqueue_task(struct task_struct *p, prio_array_t *array)
{
    ...
    list_add_tail(&p->run_list, &array->queue[p->prio]);
    __set_bit(p->prio, array->bitmap);
    array->nr_active++;
    p->array = array;
}
```

**prio:** priority of this process

**array:** a pointer pointing to the `prio_array_t` of this runqueue

- To remove a process descriptor from a runqueue list, use `dequeue_task(p,array)` function.



## 2.3 Relationships Among Processes

---


### Relationships Among Processes

#### Family relationship

```
struct task_struct {
    ...
    struct list_head children; /* list of my children */
    struct list_head sibling; /* linkage in my parent's children list */
    ...
};
```

**children:** is the list head for the list of all child elements of the process

**sibling:** is used to link siblings with each other



---

## Other Relationships

A process can be:

- a leader of a process group or of a login session
- a leader of a thread group
- tracing the execution of other processes

```
struct task_struct {
    ...
    pid_t tgid;
    ...
    struct task_struct *group_leader; /* threadgroup leader */
    ...
    struct list_head ptrace_children;
    struct list_head ptrace_list;
    ...
};
```



---

## The Pid Hash Table And Chained Lists

**PID** ⇒ **process descriptor pointer?**

- Scanning the process list? — too slow
- Use hash tables

**Four hash tables have been introduced**

**Why 4?** For 4 types of PID

PID  
TGID  
PGID  
SID


} ⇒ task\_struct



Hash Table Type	Field Name	Description
PIDTYPE_PID	pid	PID of the process
PIDTYPE_TGID	tgid	PID of thread group leader process
PIDTYPE_PGID	pgrp	PID of the group leader process
PIDTYPE_SID	session	PID of the session leader process

## Collision

Multiple PIDs can be hashed into one table index



- [Chaining](#) is used to handle colliding PIDs
- No collision if the table is 32768 in size! But...



The PID is transformed into a table index using the `pid_hashfn` macro, which expands to:

```
#define pid_hashfn(x) hash_long((unsigned long) x, pidhash_shift)
```

The `pidhash_shift` variable stores the length in bits of a table index (11, in our example). The `hash_long()` function is used by many hash functions; on a 32-bit architecture it is essentially equivalent to:

```
unsigned long hash_long(unsigned long val, unsigned int bits)
{
    unsigned long hash = val * 0x9e370001UL;
    return hash >> (32 - bits);
}
```

Because in our example `pidhash_shift` is equal to 11, `pid_hashfn` yields values ranging between 0 and  $2^{11} - 1 = 2047$ .

## The pid data structure

```

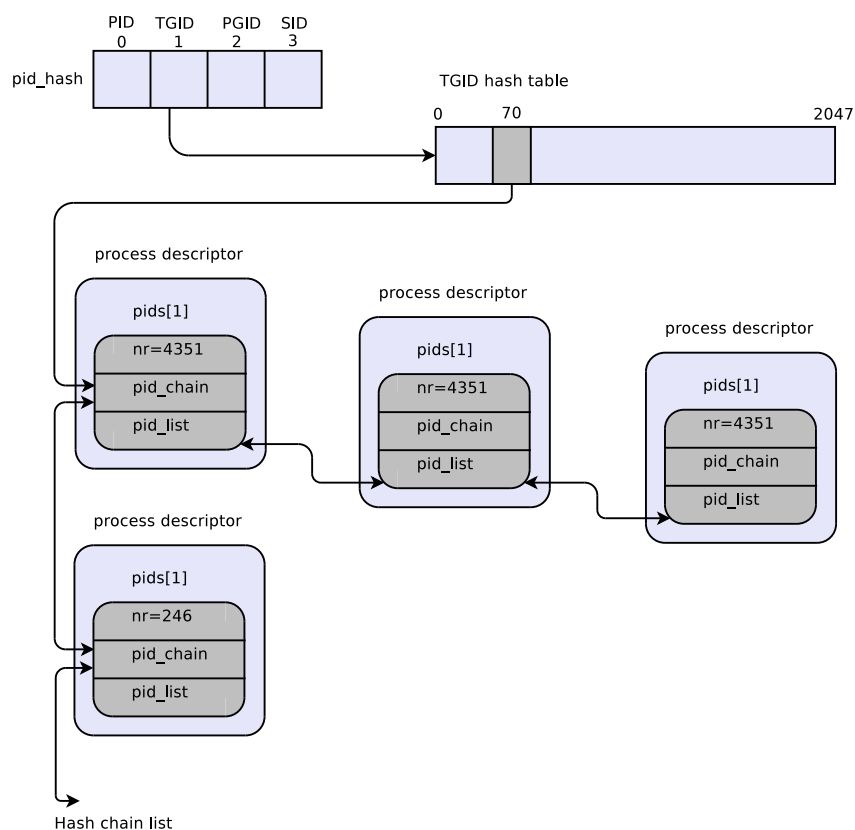
struct pid
{
    int nr;
    struct hlist_node pid_chain;
    struct list_head pid_list;
};

struct task_struct {
    ...
    struct pid pids[PIDTYPE_MAX];
    ...
}
```





## PID Hash Tables



## kernel/pid.c — Operations

- `do_each_task_pid(nr, type, task)`
- `while_each_task_pid(nr, type, task)`
- `find_task_by_pid_type(type, nr)`
- `find_task_by_pid(nr)`
- `attach_pid(task, type, nr)`
- `detach_pid(task, type)`
- `next_thread(task)`



## 2.4 How Processes Are Organized

## Wait Queues

- A wait queue represents a set of sleeping processes, which are woken up by the kernel when some condition becomes true.
- Wait queues are implemented as doubly linked lists whose elements include pointers to process descriptors.

Each wait queue is identified by a **\_\_wait\_queue\_head**

```
struct __wait_queue_head {
    spinlock_t lock;
    struct list_head task_list;
};
typedef struct __wait_queue_head wait_queue_head_t;
```

**lock:** avoid concurrent accesses.



---

Elements of a wait queue list are of type **wait\_queue\_t**:

```
struct __wait_queue {
    unsigned int flags;
    struct task_struct * task;
    wait_queue_func_t func;
    struct list_head task_list;
};
typedef struct __wait_queue wait_queue_t;
```

**task:** address of this sleeping process

**task\_list:** which wait queue are you in?

**flags:** 1 - exclusive; 0 - nonexclusive;

**func:** how it should be woken up?



## 2.5 Process Resource Limits

---

### Process Resource Limits

#### Limiting the resource use of a process

- The amount of system resources a process can use are stored in the **current->signal->rlim** field.
- **rlim** is an array of elements of type **struct rlimit**, one for each resource limit.

```
struct rlimit {
    unsigned long rlim_cur;
    unsigned long rlim_max;
};
```

**rlim\_cur:** the current resource limit for the resource

e.g. `current->signal->rlim[RLIMIT_CPU].rlim_cur` — the current limit on the CPU time of the running process.

**rlim\_max:** the maximum allowed value for the resource limit



---

## Resource Limits

RLIMIT_AS	The maximum size of process address space
RLIMIT_CORE	The maximum core dump file size
RLIMIT_CPU	The maximum CPU time for the process
RLIMIT_DATA	The maximum heap size
RLIMIT_FSIZE	The maximum file size allowed
RLIMIT_LOCKS	Maximum number of file locks
RLIMIT_MEMLOCK	The maximum size of nonswappable memory
RLIMIT_MSGQUEUE	Maximum number of bytes in POSIX message queues
RLIMIT_NOFILE	The maximum number of open file descriptors
RLIMIT_NPROC	The maximum number of processes of the user
RLIMIT_RSS	The maximum number of page frames owned by the process
RLIMIT_SIGPENDING	The maximum number of pending signals for the process
RLIMIT_STACK	The maximum stack size



## 3 Process Switch

---

### Process Switch

**Process execution context:** all information needed for the process execution

**Hardware context:** the set of registers used by a process

**Where is the hardware context stored?**

- partly in the process descriptor (PCB)
- partly in the Kernel Mode stack

### Process switch

- saving the hardware context of `prev`
- replacing it with the hardware context of `next`

Process switching occurs only in Kernel Mode.



---

### Task State Segment (TSS)

- For storing hardware contexts

- One TSS for each process (Intel's design)
- Hardware context switching
  - **far jmp** to the TSS of **next**

### Linux doesn't use hardware context switch

- One TSS for each CPU
  - The address of the kernel mode stack
  - I/O permission bitmap



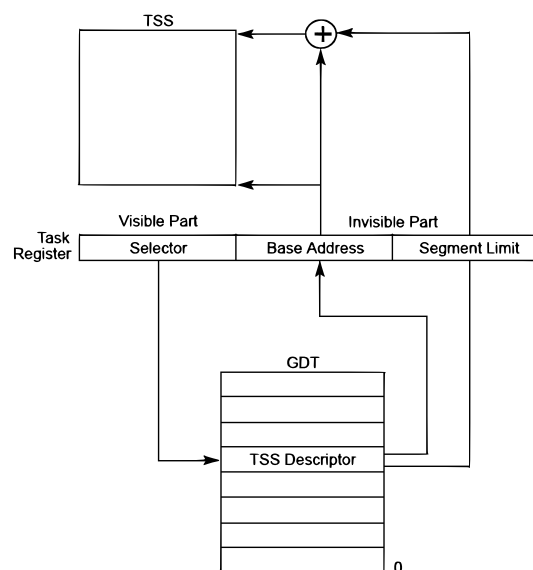
### Task State Segment Descriptor (TSSD)

31	24 23	20 19	16 15	8 7	0
7	BASE 31..24   1 1 0 0  LIMIT   1 DPL  0 1 0 1 1   BASE 23..16				4
3	SEGMENT BASE 15..0   SEGMENT LIMIT 15..0				0

- **S** bit set to 0;
- **Type** bits set to 9/11;
- **Busy** bit set to 1.



### The Task Register (**tr**)





---

## Where to save the hardware context?

```
struct task_struct{
    ...
    struct thread_struct thread;
    ...
}
```

- `thread_struct` includes fields for most of the CPU registers, except the general-purpose registers such as `eax`, `ebx`, etc., which are stored in the Kernel Mode stack.



---

## Performing The Process Switch

— `schedule()`

### Two steps:

1. Switching the Page Global Directory
2. Switching the Kernel Mode stack and the hardware context

### `switch_to(prev,next,last)`

- in any process switch three processes are involved, not just two
- 



Stackoverflow: **How does `schedule()` + `switch_to()` actually work?**

- When a process runs out of time-slice, the flag `TIF_NEED_RESCHED` is set by `scheduler_tick()`. ( Called from the timer interrupt handler, `set_tsk_need_resched()`)
- The kernel checks the flag, sees that it is set, and calls `schedule()` to switch to a new process. This flag is a message that schedule should be invoked as soon as possible because another process deserves to run. Upon returning to user-space or returning from an interrupt, the `TIF_NEED_RESCHED` flag is checked. If it is set, the kernel invokes the scheduler before continuing.

## 4 Creating Processes

REFs: [Mau08], [RFS05], [Lov10]

---

## Creating Processes

### The `clone()` system call

```
int clone(int (*fn)(void *), void *child_stack,  
          int flags, void *arg, ...  
          /* pid_t *ptid, struct user_desc *tls, pid_t *ctid */ );
```

### The traditional `fork()` system call

`clone(func, child_stack, SIGCHLD, NULL);`

- `child_stack`: parent stack pointer (copy-on-write)

### `vfork()`

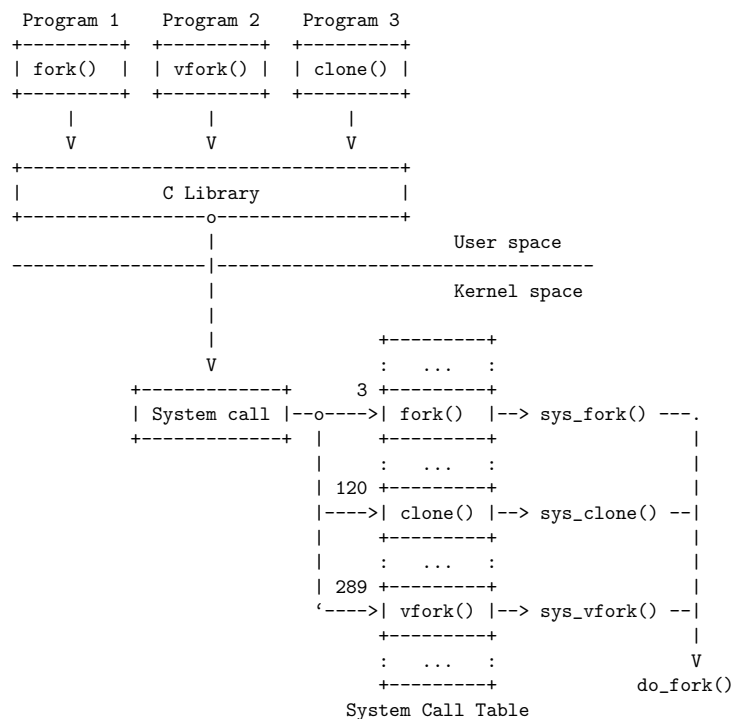
`clone(func, child_stack, CLONE_VM|CLONE_VFORK|SIGCHLD, NULL);`

- `child_stack`: parent stack pointer (copy-on-write)



---

### The `do_fork()` function does the real work



---

### `do_fork()` calls `copy_process()` to make a copy of process descriptor

```

long do_fork(unsigned long clone_flags,
             unsigned long stack_start,
             struct pt_regs *regs,
             unsigned long stack_size,
             int __user *parent_tidptr,
             int __user *child_tidptr)
{
    struct task_struct *p;
    ...
    long pid = alloc_pidmap();
    ...
    p = copy_process(clone_flags, stack_start, regs,
                    stack_size, parent_tidptr,
                    child_tidptr, pid);
    ...
    return pid;
}

```




---

## copy\_process()

1. dup\_task\_struct(): creates

- a new kernel mode stack
- thread\_info
- task\_struct

Values are identical to the parent

2. is current->signal->rlim[RLIMIT\_NPROC].rlim\_cur confirmed?
3. Update child's task\_struct
4. Set child's state to TASK\_UNINTERRUPTABLE
5. copy\_flags(): update flags in task\_struct
6. get\_pid() (check pidmap\_array bitmap)
7. Duplicate or share resources (opened files, FS info, signal, ...)
8. return p;



- Sec *Forking* in chap. 3 of [Lov10]
- 

## Creating A Kernel Thread

### kernel\_thread() is similar to clone()

```

int kernel_thread(int (*fn)(void *), void * arg, unsigned long flags)
{
    ...
    return do_fork(flags | CLONE_VM | CLONE_UNTRACED, 0, &regs, 0, NULL, NULL);
}

```



- Typically, a kernel thread continues executing its initial function forever (or at least until the system reboots, but with Linux you never know). The initial function usually implements a loop in which the kernel thread wakes up as needed, performs its duties, and then returns to sleep (Sec 3.4.2, *Kernel Threads*, [Lov10]).

---

## Process 0

**Process 0** is a kernel thread created from scratch during the initialization phase.

- Also called *idle process*, or *swapper process*
- Its data structures are *statically* allocated

### start\_kernel()

- Initializes all the data structures
- Enables interrupts
- Creates another kernel thread — *process 1*, the *init* process



---

## Call graph

```
start_kernel()
  '--> rest_init()
        |--> kernel_thread(init, NULL, CLONE_FS | CLONE_SIGHAND)
        '--> cpu_idle()
```

- After having created the *init* process, *process 0* executes the `cpu_idle()` function.
- Process 0 is selected by the scheduler only when there are no other processes in the `TASK_RUNNING` state.
- In multiprocessor systems there is a process 0 for each CPU.



---

## Process 1

- Created via `kernel_thread(init, NULL, CLONE_FS|CLONE_SIGHAND);`
- `PID` is 1
- shares all per-process kernel data structures with process 0
- starts executing the `init()` function
  - completes the initialization of the kernel
- `init()` invokes the `execve()` system call to load the executable program `init`



- As a result, the *init kernel thread* becomes a regular process having its own per-process kernel data structure
- The init process stays alive until the system is shut down



## 5 Destroying Processes

---

### Process Termination

- Usual way: call `exit()`
  - The C compiler places a call to `exit()` at the end of `main()`.
- Unusual way: Ctrl-C ...



---

### All process terminations are handled by `do_exit()`

- `tsk->flags |= PF_EXITING;` to indicate that the process is being eliminated
- `del_timer_sync(&tsk->real_timer);` to remove any kernel timers
- `exit_mm()`, `exit_sem()`, `__exit_files()`, `__exit_fs()`, `exit_namespace()`, `exit_thread()`: free pointers to the kernel data structures
- `tsk->exit_code = code;`
- `exit_notify()` to send signals to the task's parent
  - re-parents its children
  - sets the task's state to `TASK_ZOMBIE`
- `schedule()` to switch to a new process



---

### Process Removal

Cleaning up after a process and removing its process descriptor are separate.

#### Clean up

- done in `do_exit()`
- leaves a zombie
  - To provide information to its parent
  - The only memory it occupies is its kernel stack, the `thread_info` structure, and the `task_struct` structure.



---

## Removal

- `release_task()` is invoked by  
either `do_exit()` if the parent didn't wait  
or `wait4()/waitpid()`
- `free_uid()`
- `unhash_process`: to remove the process from the pidhash and from the task list
- `put_task_struct()`
  - free the pages containing the process's kernel stack and `thread_info` structure
  - de-allocate the slab cache containing the `task_struct`

