

# Precesses

Wang Xiaolin

June 13, 2013

# Processes

A process is

- ▶ an instance of a program in execution
- ▶ a dynamic entity (has lifetime)
- ▶ a collection of data structures describing the execution progress
- ▶ the unit of system resources allocation

The Linux kernel internally refers to processes as *tasks*.

# When A Process Is created

## The child

- ▶ is almost identical to the parent
  - ▶ has a logical copy of the parent's address space
  - ▶ executes the same code
- ▶ has its own data (stack and heap)

# Multithreaded Applications

## Threads

- ▶ are execution flows of a process
- ▶ share a large portion of the application data structures

## Lightweight processes (LWP) — Linux way of multithreaded applications

- ▶ each LWP is scheduled individually by the kernel
  - ▶ no nonblocking syscall is needed
- ▶ LWPs may share some resources, like the address space, the open files, and so on.

# Process Descriptor

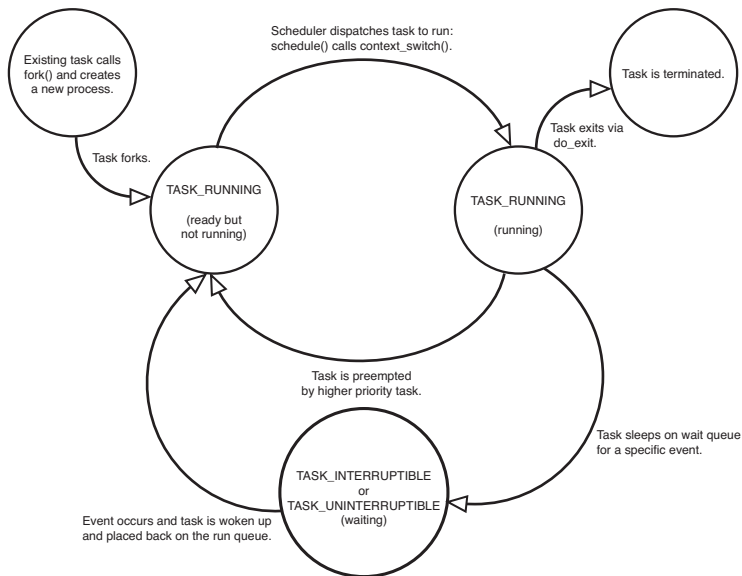
To manage processes, the kernel must have a clear picture of what each process is doing.

- ▶ the process's priority
- ▶ running or blocked
- ▶ its address space
- ▶ files it opened
- ▶ ...

**Process descriptor:** a `task_struct` type structure containing all the information related to a single process.

```
struct task_struct {  
    /* 160 lines of code in 2.6.11 */  
};
```

# Process State



# PID AND TGID

- ▶ kernel finds a process by its *process descriptor pointer* pointing to a `task_struct`
- ▶ users find a process by its `PID`
- ▶ all the threads of a multithreaded application share the same identifier  
    `tgid`: the PID of the thread group leader

```
struct task_struct {  
    ...  
    pid_t pid;  
    pid_t tgid;  
    ...  
};
```

```
~$ ps -eo pgid,ppid,pid,tgid,tid,nlwp,comm -sort pid
```

## How many PIDs can there be?

- ▶ `#define PID_MAX_DEFAULT 0x8000`
- ▶ Max PID number = `PID_MAX_DEFAULT - 1` = 32767
- ▶ `$ cat /proc/sys/kernel/pid_max`

## Which are the free PIDs?

```
static pidmap_t pidmap_array[PIDMAP_ENTRIES] =  
{  
    [ 0 ... PIDMAP_ENTRIES-1 ] =  
    { ATOMIC_INIT(BITS_PER_PAGE), NULL }  
};
```

`pidmap_array` consumes a single page.




# Process Descriptor Handling

**thread\_union**: 2 consecutive page frames (8K) containing

- ▶ a process kernel stack
- ▶ a **thread\_info** structure

```
union thread_union {  
    struct thread_info thread_info;  
    unsigned long stack[2048]; /* 1024 for 4KB stacks */  
};
```

# Kernel Stack



```

struct thread_info {
    struct task_struct    *task;           /* main task structure */
    struct exec_domain    *exec_domain;    /* execution domain */
    unsigned long         flags;           /* low level flags */
    unsigned long         status;          /* thread-synchronous flags */
    __u32                 cpu;            /* current CPU */
    __s32                  preempt_count;  /* 0 => preemptable, <0 => BUG */

    mm_segment_t          addr_limit;      /* thread address space:
                                           0-0xBFFFFFFF for user-thread
                                           0-0xFFFFFFFF for kernel-thread
                                           */
    struct restart_block   restart_block;

    unsigned long          previous_esp;    /* ESP of the previous stack in case
                                           of nested (IRQ) stacks
                                           */
    __u8                   supervisor_stack[0];
};

```

## Why both task\_struct and thread\_info?

- ▶ There wasn't a thread\_info in pre-2.6 kernel
- ▶ Size matters

## thread\_info and task\_struct are mutually linked

```
struct thread_info {  
    struct task_struct *task; /* main task structure */  
    ...  
};  
  
struct task_struct {  
    ...  
    struct thread_info *thread_info;  
    ...  
};
```

# Identifying The Current Process

Efficiency benefit from `thread_union`

- Easy get the base address of `thread_info` from `esp` register by masking out the 13 least significant bits of `esp`

## `current_thread_info()`

```
/* how to get the thread information struct from C */
static inline struct thread_info *current_thread_info(void)
{
    struct thread_info *ti;
    __asm__ ("andl %%esp, %%0;" : "=r" (ti) : "0" (~(THREAD_SIZE - 1)));
    return ti;
}
```

Can be seen as:

```
movl $0xffffe000, %ecx /* or 0xfffff000 for 4KB stacks */
andl %esp, %ecx
movl %ecx, p
```

## To get the process descriptor pointer

`current_thread_info()->task`

```
movl $0xffffe000,%ecx /* or 0xfffff000 for 4KB stacks */
andl %esp,%ecx
movl (%ecx),p
```

Because the `task` field is at offset 0 in `thread_info`, after executing these 3 instructions `p` contains the process descriptor pointer.

`current` — a marco pointing to the current running task

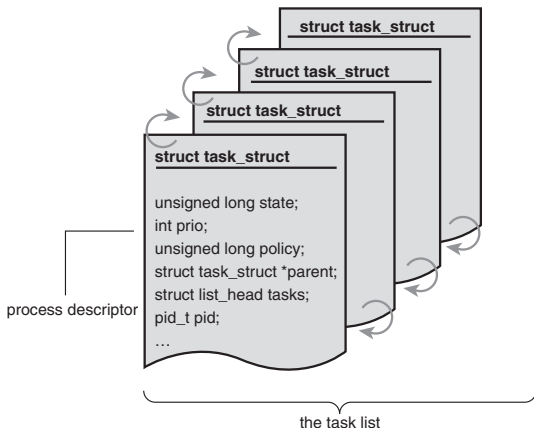
```
static inline struct task_struct * get_current(void)
{
    return current_thread_info()->task;
}

#define current get_current()
```


# Task List

The kernel stores the list of processes in a **circular doubly linked list** called the **task list**.

**Swapper** The head of this list, **init\_task**, process 0.



## Doubly Linked List



```
struct list_head {
    struct list_head *next, *prev;
};
```

```
struct task_struct {
    ...
    struct list_head tasks;
    ...
}
```



## List operations

**SET\_LINKS** insert into the list

**REMOVE\_LINKS** remove from the list

**for\_each\_process** scan the whole process list

```
#define for_each_process(p) \
    for (p = &init_task ; (p = next_task(p)) != &init_task ; )
```

**list\_for\_each** iterate over a list

```
#define list_for_each(pos, head) \
    for (pos = (head)->next; prefetch(pos->next), pos != (head); \
        pos = pos->next)
```

## Example: Iterate over a process' children


```
struct task_struct *task;
struct list_head *list;
list_for_each(list, &current->children) {
    task = list_entry(list, struct task_struct, sibling);
    /* task now points to one of current's children */
}
```

## A task can be in multiple lists

```
struct task_struct {  
    struct list_head run_list;  
    struct list_head tasks;  
    struct list_head ptrace_children;  
    struct list_head ptrace_list;  
    struct list_head children; /* list of my children */  
    struct list_head sibling; /* linkage in my parent's children list */  
}
```

# The List Of TASK\_RUNNING Processes

- ▶ Each CPU has its own runqueue
- ▶ Each runqueue has 140 lists
- ▶ One list per process priority
- ▶ Each list has zero to many tasks



```
struct task_struct {  
    ...  
    int prio, static_prio;  
    struct list_head run_list;  
    prio_array_t *array;  
    ...  
};
```

## Each Runqueue Has A prio\_array\_t Struct

```
typedef struct prio_array prio_array_t;  
  
struct prio_array {  
    unsigned int nr_active;  
    unsigned long bitmap[BITMAP_SIZE];  
    struct list_head queue[MAX_PRI0];  
};
```

**nr\_active:** The number of process descriptors linked into the lists (the whole runqueue)

**bitmap:** A priority bitmap. Each flag is set if the priority list is not empty

**queue:** The 140 heads of the priority lists

# To Insert A Task Into A Runqueue List

```
static void enqueue_task(struct task_struct *p, prio_array_t *array)
{
    ...
    list_add_tail(&p->run_list, &array->queue[p->prio]);
    __set_bit(p->prio, array->bitmap);
    array->nr_active++;
    p->array = array;
}
```

**prio:** priority of this process

**array:** a pointer pointing to the **prio\_array\_t** of this runqueue

- To remove a process descriptor from a runqueue list, use **dequeue\_task(p,array)** function.

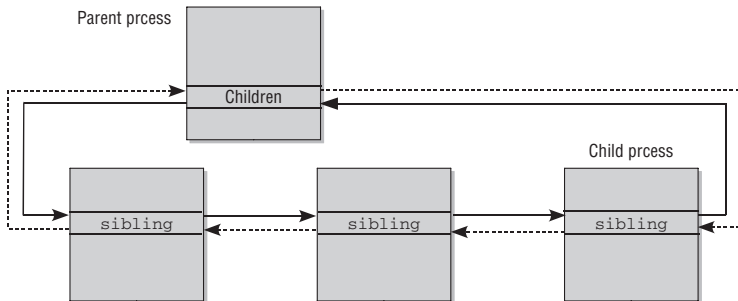
# Relationships Among Processes

## Family relationship

```
struct task_struct {  
    ...  
    struct list_head children; /* list of my children */  
    struct list_head sibling; /* linkage in my parent's children list */  
    ...  
};
```

**children:** is the list head for the list of all child elements of the process

**sibling:** is used to link siblings with each other



## Other Relationships

A process can be:

- ▶ a leader of a process group or of a login session
- ▶ a leader of a thread group
- ▶ tracing the execution of other processes

```
struct task_struct {  
    ...  
    pid_t tgid;  
    ...  
    struct task_struct *group_leader; /* threadgroup leader */  
    ...  
    struct list_head ptrace_children;  
    struct list_head ptrace_list;  
    ...  
};
```

# The Pid Hash Table And Chained Lists

PID  $\Rightarrow$  process descriptor pointer?

- ▶ Scanning the process list? — too slow
- ▶ Use hash tables

Four hash tables have been introduced

Why 4? For 4 types of PID

PID  
TGID  
PGID  
SID

}  $\Rightarrow$  task\_struct



## Collision

Multiple PIDs can be hashed into one table index

		Hash table		
+---+	PID	+-----+		
+---+		0		
		+-----+		
v				process descriptors
+-----+		+-----+		+---+ +---+
pid_hashfn()	-->m		<--- -->	885 <---> 169
+-----+		+-----+		+---+ +---+
		+-----+		+-----+
			<--- -->	29385
		+-----+		+-----+
	2047			
		+-----+		

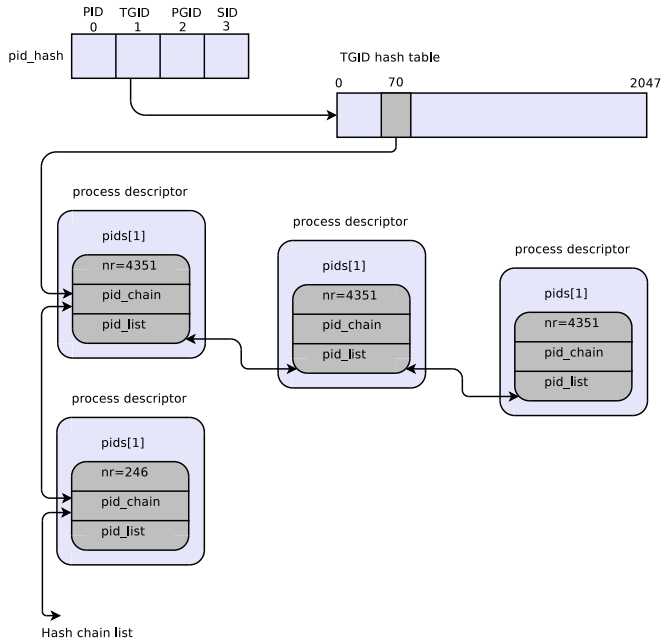
- **Chaining** is used to handle colliding PIDs
- No collision if the table is 32768 in size! But...

## The pid data structure

```
struct pid
{
    int nr;
    struct hlist_node pid_chain;
    struct list_head pid_list;
};
```

```
struct task_struct{
    ...
    struct pid pids[PIDTYPE_MAX];
    ...
}
```

# PID Hash Tables



## kernel/pid.c — Operations

- ▶ `do_each_task_pid(nr, type, task)`
- ▶ `while_each_task_pid(nr, type, task)`
- ▶ `find_task_by_pid_type(type, nr)`
- ▶ `find_task_by_pid(nr)`
- ▶ `attach_pid(task, type, nr)`
- ▶ `detach_pid(task, type)`
- ▶ `next_thread(task)`

# Wait Queues

- ▶ A wait queue represents a set of sleeping processes, which are woken up by the kernel when some condition becomes true.
- ▶ Wait queues are implemented as doubly linked lists whose elements include pointers to process descriptors.

Each wait queue is identified by a  
`__wait_queue_head`

```
struct __wait_queue_head {  
    spinlock_t lock;  
    struct list_head task_list;  
};  
typedef struct __wait_queue_head wait_queue_head_t;
```

`lock`: avoid concurrent accesses.

Elements of a wait queue list are of type  
`wait_queue_t`:

```
struct __wait_queue {  
    unsigned int flags;  
    struct task_struct * task;  
    wait_queue_func_t func;  
    struct list_head task_list;  
};  
typedef struct __wait_queue wait_queue_t;
```

`task`: address of this sleeping process

`task_list`: which wait queue are you in?

`flags`: 1 - exclusive; 0 - nonexclusive;

`func`: how it should be woken up?

# Process Resource Limits

## Limiting the resource use of a process

- ▶ The amount of system resources a process can use are stored in the `current->signal->rlim` field.
- ▶ `rlim` is an array of elements of type `struct rlimit`, one for each resource limit.

```
struct rlimit {  
    unsigned long    rlim_cur;  
    unsigned long    rlim_max;  
};
```

`rlim_cur`: the current resource limit for the resource  
e.g. `current->signal->rlim[RLIMIT_CPU].rlim_cur`  
— the current limit on the CPU time of the running process.

`rlim_max`: the maximum allowed value for the resource limit

# Resource Limits

RLIMIT_AS	The maximum size of process address space
RLIMIT_CORE	The maximum core dump file size
RLIMIT_CPU	The maximum CPU time for the process
RLIMIT_DATA	The maximum heap size
RLIMIT_FSIZE	The maximum file size allowed
RLIMIT_LOCKS	Maximum number of file locks
RLIMIT_MEMLOCK	The maximum size of nonswappable memory
RLIMIT_MSGQUEUE	Maximum number of bytes in POSIX message queues
RLIMIT_NOFILE	The maximum number of open file descriptors
RLIMIT_NPROC	The maximum number of processes of the user
RLIMIT_RSS	The maximum number of page frames owned by the process
RLIMIT_SIGPENDING	The maximum number of pending signals for the process
RLIMIT_STACK	The maximum stack size



# Process Switch

**Process execution context:** all information needed for the process execution

**Hardware context:** the set of registers used by a process

**Where is the hardware context stored?**

- ▶ partly in the process descriptor (PCB)
- ▶ partly in the Kernel Mode stack

**Process switch**

- ▶ saving the hardware context of **prev**
- ▶ replacing it with the hardware context of **next**

Process switching occurs only in Kernel Mode.


## Task State Segment (TSS)

- ▶ For storing hardware contexts
- ▶ One TSS for each process (Intel's design)
- ▶ Hardware context switching
  - ▶ `far jmp` to the TSS of `next`

## Linux doesn't use hardware context switch

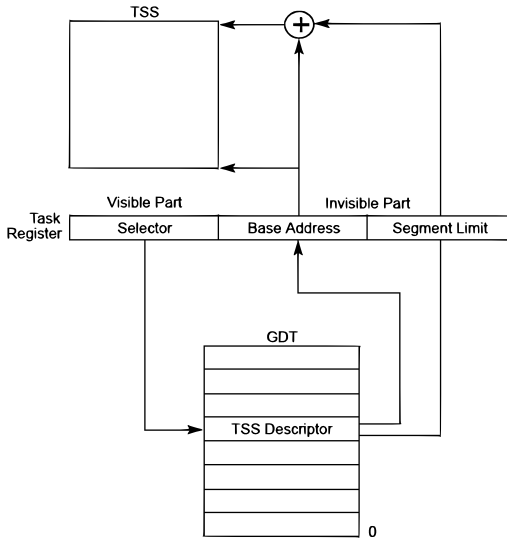
- ▶ One TSS for each CPU
  - ▶ The address of the kernel mode stack
  - ▶ I/O permission bitmap

## Task State Segment Descriptor (TSSD)



- ▶ S bit set to 0;
- ▶ Type bits set to 9/11;
- ▶ Busy bit set to 1.

# The Task Register (tr)



## Where to save the hardware context?

```
struct task_struct{  
    ...  
    struct thread_struct thread;  
    ...  
}
```

- ▶ `thread_struct` includes fields for most of the CPU registers, except the general-purpose registers such as `eax`, `ebx`, etc., which are stored in the Kernel Mode stack.

# Performing The Process Switch

— `schedule()`

## Two steps:

1. Switching the Page Global Directory
2. Switching the Kernel Mode stack and the hardware context

## `switch_to(prev,next,last)`

- ▶ in any process switch three processes are involved, not just two
- ▶

# Creating Processes

## The clone() system call

```
int clone(int (*fn)(void *), void *child_stack,  
          int flags, void *arg, ...  
          /* pid_t *ptid, struct user_desc *tls, pid_t *ctid */ );
```

## The traditional fork() system call

```
clone(func, child_stack, SIGCHLD, NULL);
```


- ▶ `child_stack`: parent stack pointer (copy-on-write)

## vfork()

```
clone(func, child_stack, CLONE_VM|CLONE_VFORK|SIGCHLD, NULL);
```

- ▶ `child_stack`: parent stack pointer (copy-on-write)

# The `do_fork()` function does the real work





**do\_fork()** calls **copy\_process()** to make a copy of process descriptor

```
long do_fork(unsigned long clone_flags,
             unsigned long stack_start,
             struct pt_regs *regs,
             unsigned long stack_size,
             int __user *parent_tidptr,
             int __user *child_tidptr)
{
    struct task_struct *p;
    ...
    long pid = alloc_pidmap();
    ...
    p = copy_process(clone_flags, stack_start, regs,
                    stack_size, parent_tidptr,
                    child_tidptr, pid);
    ...
    return pid;
}
```

## copy\_process()

1. `dup_task_struct()`: creates
  - ▶ a new kernel mode stack
  - ▶ `thread_info`
  - ▶ `task_struct`

Values are identical to the parent

2. is `current->signal->rlim[RLIMIT_NPROC].rlim_cur` confirmed?
3. Update child's `task_struct`
4. Set child's state to `TASK_UNINTERRUPTABLE`
5. `copy_flags()`: update flags in `task_struct`
6. `get_pid()` (check `pidmap_array` bitmap)
7. Duplicate or share resources (opened files, FS info, signal, ...)
8. `return p;`

# Creating A Kernel Thread

kernel\_thread() is similar to clone()

```
int kernel_thread(int (*fn)(void *), void * arg, unsigned long flags)
{
    ...
    return do_fork(flags | CLONE_VM | CLONE_UNTRACED, 0, &regs, 0, NULL, NULL);
}
```

# Process 0

Process 0 is a kernel thread created from scratch during the initialization phase.

- ▶ Also called *idle process*, or *swapper process*
- ▶ Its data structures are *statically* allocated

## start\_kernel()

- ▶ Initializes all the data structures
- ▶ Enables interrupts
- ▶ Creates another kernel thread — *process 1*, the *init process*

## Call graph

```
start_kernel()
  '--> rest_init()
        |--> kernel_thread(init, NULL, CLONE_FS | CLONE_SIGHAND)
        '--> cpu_idle()
```

- ▶ After having created the *init* process, *process 0* executes the `cpu_idle()` function.
- ▶ Process 0 is selected by the scheduler only when there are no other processes in the `TASK_RUNNING` state.
- ▶ In multiprocessor systems there is a process 0 for each CPU.

# Process 1

- ▶ Created via  
`kernel_thread(init, NULL, CLONE_FS|CLONE_SIGHAND);`
- ▶ PID is 1
- ▶ shares all per-process kernel data structures with process 0
- ▶ starts executing the `init()` function
  - ▶ completes the initialization of the kernel
- ▶ `init()` invokes the `execve()` system call to load the executable program `init`
  - ▶ As a result, the *init kernel thread* becomes a regular process having its own per-process kernel data structure
- ▶ The `init` process stays alive until the system is shut down

# Process Termination

- ▶ Usual way: call `exit()`
  - ▶ The C compiler places a call to `exit()` at the end of `main()`.
- ▶ Unusual way: `Ctrl-C ...`

## All process terminations are handled by `do_exit()`

- ▶ `tsk->flags |= PF_EXITING;` to indicate that the process is being eliminated
- ▶ `del_timer_sync(&tsk->real_timer);` to remove any kernel timers
- ▶ `exit_mm()`, `exit_sem()`, `__exit_files()`, `__exit_fs()`, `exit_namespace()`, `exit_thread()`: free pointers to the kernel data structures
- ▶ `tsk->exit_code = code;`
- ▶ `exit_notify()` to send signals to the task's parent
  - ▶ re-parents its children
  - ▶ sets the task's state to `TASK_ZOMBIE`
- ▶ `schedule()` to switch to a new process



# Process Removal

Cleaning up after a process and removing its process descriptor are separate.

## Clean up

- ▶ done in `do_exit()`
- ▶ leaves a zombie
  - ▶ To provide information to its parent
  - ▶ The only memory it occupies is its kernel stack, the `thread_info` structure, and the `task_struct` structure.

## Removal

- ▶ `release_task()` is invoked by  
either `do_exit()` if the parent didn't wait  
or `wait4()/waitpid()`
- ▶ `free_uid()`
- ▶ `unhash_process`: to remove the process from the  
pidhash and from the task list
- ▶ `put_task_struct()`
  - ▶ free the pages containing the process's kernel stack and  
`thread_info` structure
  - ▶ de-allocate the slab cache containing the `task_struct`