

Network Basics

Wang Xiaolin

`wx672ster+net@gmail.com`

December 20, 2020

Textbooks

- [1] TANENBAUM A, WETHERALL D. *Computer Networks*. 5th ed. Pearson Prentice Hall, 2011.
- [2] KUROSE J, ROSS K. *Computer Networking: A Top-down Approach*. Pearson, 2013.
- [3] FALL K, STEVENS W. *TCP/IP Illustrated, Volume 1: The Protocols*. Pearson Education, 2011.

Course Web Links

Course web site <https://cs6.swfu.edu.cn/moodle>

Lecture slides https://cs6.swfu.edu.cn/~wx672/lecture_notes/network_basics/

Lab exercises https://cs6.swfu.edu.cn/~wx672/lecture_notes/network_basics/net-tools/net-tools.html


Cisco Networking Academy <https://www.netacad.com/>

Beej's Guides <http://beej.us/guide/>

Weekly tech question

1. What was I trying to do?
2. How did I do it? (steps)
3. The expected output? The real output?
4. How did I try to solve it? (steps, books, web links)
5. How many hours did I struggle on it?

 <https://cs6.swfu.edu.cn/moodle/mod/forum/view.php?id=98>


 wx672ster+net@gmail.com

 Preferably in English

 in [stackoverflow](#) style

OR simply show me the tech questions you asked on any website

What's A Computer Network?



The History of Internet I

1836: Telegraph

1858-66: Transatlantic cable

1876: Telephone

1957: USSR launches Sputnik

1962-68: *Packet-switching* networks developed

1969: Birth of Internet

1971: People communicate over a network

1972: Computers can connect more freely and easily

1973: Global Networking becomes a reality

1974: Packets become mode of transfer

The History of Internet II

- 1976: Networking comes to many
- 1977: E-mail takes off, Internet becomes a reality
- 1979: News Groups born
- 1981: Things start to come together
- 1982: *TCP/IP* defines future communication
- 1983: Internet gets bigger
- 1984: Growth of Internet Continues
- 1986: Power of Internet Realised
- 1987: Commercialisation of Internet Born
- 1989: Large growth in Internet
- 1990: Expansion of Internet continues

The History of Internet III

1991: Modernisation Begins

1992: Multimedia changes the face of the Internet

1993: The WWW Revolution truly begins

1994: Commercialisation begins

1995: Commercialisation continues apace

1996: Microsoft enters

1998: Google

Homework: Meanwhile, what happened in China?

What's The Internet?

What pops up in your mind if I say “Internet”?

What's The Internet?

What pops up in your mind if I say "Internet"?

For me, the answer is...



and...

What's The Internet?

What pops up in your mind if I say “Internet”?

For me, the answer is...

The Google logo, featuring the word "Google" in its characteristic multi-colored font: blue 'G', red 'o', yellow 'o', blue 'g', green 'l', and red 'e'.

and...

TCP/IP

What's The Internet?

- ▶ The network of networks.
- ▶ Tech view: TCP/IP
- ▶ App view: Google

Google Philosophy

Ten things Google has found to be true

1. Focus on the user and all else will follow.
2. It's best to do one thing really, really well.
3. Fast is better than slow.
4. Democracy on the web works.
5. You don't need to be at your desk to need an answer.
6. *You can make money without doing evil.*
7. There's always more information out there.
8. The need for information crosses all borders.
9. You can be serious without a suit.
10. Great just isn't good enough.

Google Philosophy

More about...

- ▶ Software Principles
- ▶ Google User Experience
- ▶ No pop-ups
- ▶ Security


Choosing The Right Tools



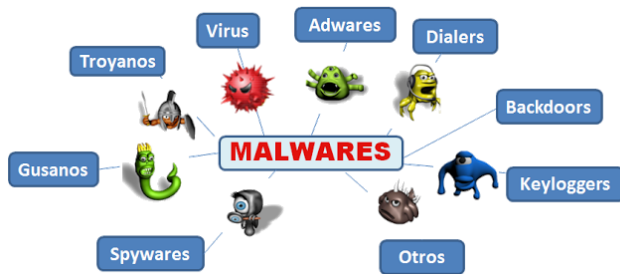
VS.



Danger




Danger



My solution



Homework

1. try 
2. get a [gmail](#) account
3. add your class timetable into [google calendar](#), and then share your calendar to me via gmail
4. in [youtube](#), find a video you like and share it to me


Network Classification

- ▶ connection method: wired, wireless...
- ▶ topology
- ▶ scale
- ▶ network architecture: c/s, p2p...



Network Classification

Connection method

Wired:




Wireless:




Scale

PAN, LAN, CAN, MAN, **WAN** ...


Topology




Ring



Mesh




Star




Fully Connected



Line




Tree




Bus

Network Architecture









Client/Server




Peer to Peer

Basic Hardware Components

IP	Router: 		
Link	Bridge: 	Switch: 	
PHY	NIC: 	Repeater: 	Hub: 


What's a Protocol?

A rule, a treaty, an agreement, ...



What's TCP/IP?

A set of protocols designed for the Internet




TCP/IP Protocol Stack


Every networked computer has it inside




ISO/OSI RM



TCP/IP



My Favor

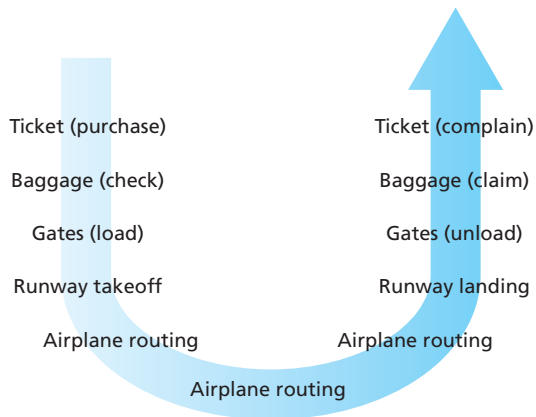


Each protocol is completely independent of the other ones

- The translators (L2) can switch from Chinese to Finnish without touching L1 or L3
- The secretaries (L1) can switch from fax to email without disturbing (or even informing) the other layers

Layered Design Example

Taking an airplane trip

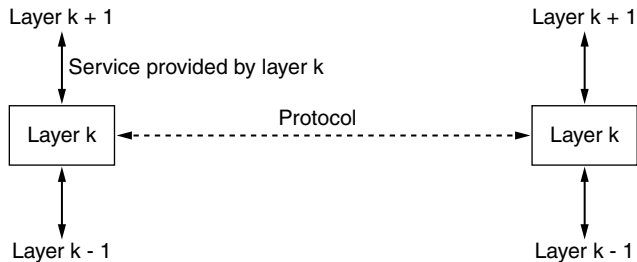


Each layer

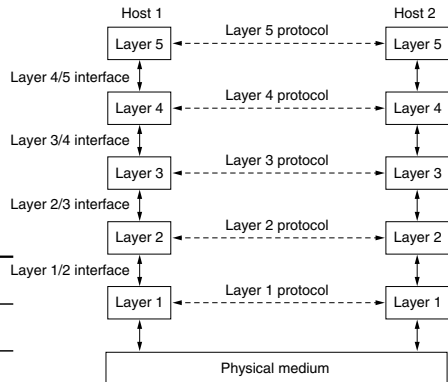
1. has some functions
2. provides services to its upper layer

Layered Design

Services vs. Protocols




Services	Protocols
Layer to Layer	Peer to Peer
A set of operations (listen, connect, accept, receive, send, disconnect)	A set of rules (message format, message meanings)



TCP/IP Overview

Basic Structure



1. Where will an incoming Ethernet frame go?
0x0806 → ARP
0x0800 → IP
2. Where will an incoming IP packet go?
0x06 → TCP
0x11 → UDP
3. Where will an incoming transport message (UDP datagram, TCP segment) go?

HTTP	FTP	SSH	SMTP
80	21/20	22	25



The Name Of A Unit Of Data

APP Application message


TRANS TCP segment; UDP datagram

NET IP packet

LINK Ethernet frame




Ethernet



1. Frame format?
2. Address format?
3. Broadcast address?
4. CSMA/CD?

Ethernet Frame



Examples

Unicast, carrying an IP packet

Destination	Source	Type	MAC Data	
08005A21A722	0800280038A9	0x0800	IP packet	FCS

Unicast, carrying an ARP packet

Destination	Source	Type	MAC Data	
0800280038A9	08005A21A722	0x0806	ARP Response	FCS

Broadcast, carrying an ARP packet

Destination	Source	Type	MAC Data	
FFFFFFFFFFFF	0800280038A9	0x0806	ARP Request	FCS

Ethernet References


- [1] Wikipedia. *Ethernet — Wikipedia, The Free Encyclopedia*. 2015. <http://en.wikipedia.org/w/index.php?title=Ethernet&oldid=648082976>.
- [2] Wikipedia. *Ethernet frame — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Ethernet%5C_frame&oldid=653337888.
- [3] Wikipedia. *Carrier sense multiple access with collision detection — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Carrier%5C_sense%5C_multiple%5C_access%5C_with%5C_collision%5C_detection&oldid=648771859.
- [4] POSTEL J, REYNOLDS J. *Standard for the transmission of IP datagrams over IEEE 802 networks*. RFC Editor. 1988. <https://www.rfc-editor.org/rfc/rfc1042.txt>.

ARP

ARP Looking up the ARP table to find the destination MAC address.

Example ARP table

IP address	Ethernet address
223.1.2.1	08-00-39-00-2F-C3
223.1.2.3	08-00-5A-21-A7-22
223.1.2.4	08-00-10-99-AC-54



Where does the ARP table come from?

Example ARP Request


Sender IP Address	223.1.2.1
Sender Eth Address	08-00-39-00-2F-C3
Target IP Address	223.1.2.2
Target Eth Address	FF-FF-FF-FF-FF-FF

Example ARP Response

Sender IP Address	223.1.2.2
Sender Eth Address	08-00-28-00-38-A9
Target IP Address	223.1.2.1
Target Eth Address	08-00-39-00-2F-C3

The updated table

IP address	Ethernet address
223.1.2.1	08-00-39-00-2F-C3
223.1.2.2	08-00-28-00-38-A9
223.1.2.3	08-00-5A-21-A7-22
223.1.2.4	08-00-10-99-AC-54





ARP References

- [1] Wikipedia. *Address Resolution Protocol — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Address%5C_Resolution%5C_Protocol&oldid=647148843.
- [2] PLUMMER D. *An Ethernet Address Resolution Protocol: Or Converting Network Protocol Addresses to 48.bit Ethernet Address for Transmission on Ethernet Hardware*. RFC Editor. 1982. <https://www.rfc-editor.org/rfc/rfc826.txt>.


IP

Router



Routing Find a route in the route table.




Direct Routing—IP is overhead




$A \Rightarrow B$

Address	Source	Destination
IP	A	B
Eth	A	B

Is IP Necessary?



Indirect Routing




$A \Rightarrow E$

Before D

address	source	destination
IP	A	E
Eth	A	D

After D

address	source	destination
IP	A	E
Eth	D	E




IP Module Routing Rules

1. For an outgoing IP packet, entering IP from an upper layer, IP must decide
 - ▶ whether to send the IP packet directly or indirectly, and
 - ▶ IP must choose a lower network interface.

These choices are made by consulting the route table.

2. For an incoming IP packet, entering IP from a lower interface, IP must decide
 - ▶ whether to forward the IP packet or pass it to an upper layer.
 - ▶ If the IP packet is being forwarded, it is treated as an outgoing IP packet.
3. When an incoming IP packet arrives it is never forwarded back out through the same network interface.




IP Address

An IPv4 address (dotted-decimal notation)


172 . 16 . 254 . 1



10101100.00010000.11111110.00000001

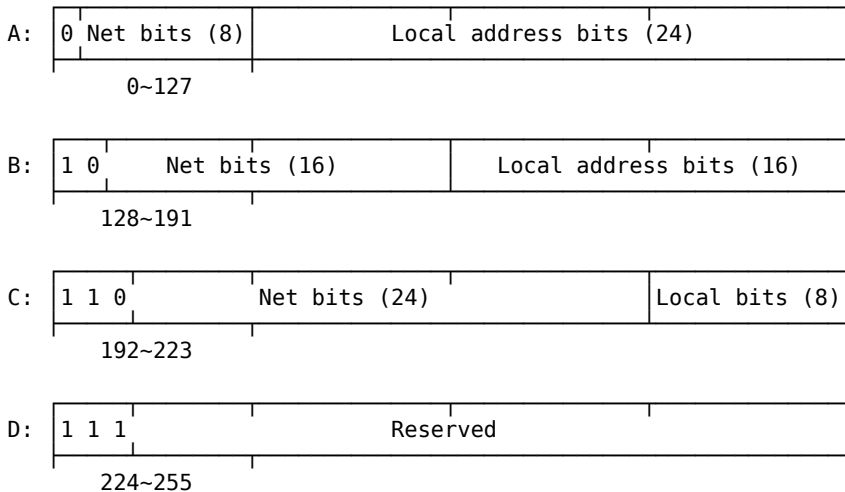


One byte = Eight bits



Thirty-two bits (4 x 8), or 4 bytes

Address classes



Network bits?


Do maths on the first 8 bits

Dst IP: 11000000.10101000.10101000.10101000

Example route table

Destination	Gateway	Iface
00001010.00000000.00000000.00000000	*	1
10101100.00010000.00000000.00000000	*	2
11000000.10101000.10101000.00000000	*	3

Prefix — A faster way to decide network bits



Special IP Addresses

- ▶ A value of zero in the network field means this network. (source only)
- ▶ A value of zero in the host field means network address.
- ▶ 127.x.x.x are loopback address.
- ▶ 255.255.255.255 is boardcast address.
- ▶ Private address:
 - ▶ 10.x.x.x
 - ▶ 172.16.x.x~172.31.x.x
 - ▶ 192.168.x.x
- ▶ CIDR—Classless Inter-Domain Routing—An IP addressing scheme that replaces the older system based on classes A, B and C.

Names

People refer to computers by names, not numbers.

`/etc/hosts`

127.0.0.1	localhost	
202.203.132.245	cs3.swfu.edu.cn	cs3

`/etc/networks`

localnet	202.203.132.192
----------	-----------------

IP Route Table

Example IP Route Table


```
~$ route
```

```
Kernel IP routing table
```

Destination	Gateway	Iface
localnet	*	eth0
192.168.128.0	*	eth0
default	202.203.132.254	eth0

```
~$ man route
```


Direct Routing Details



The route table inside alpha (simplified)

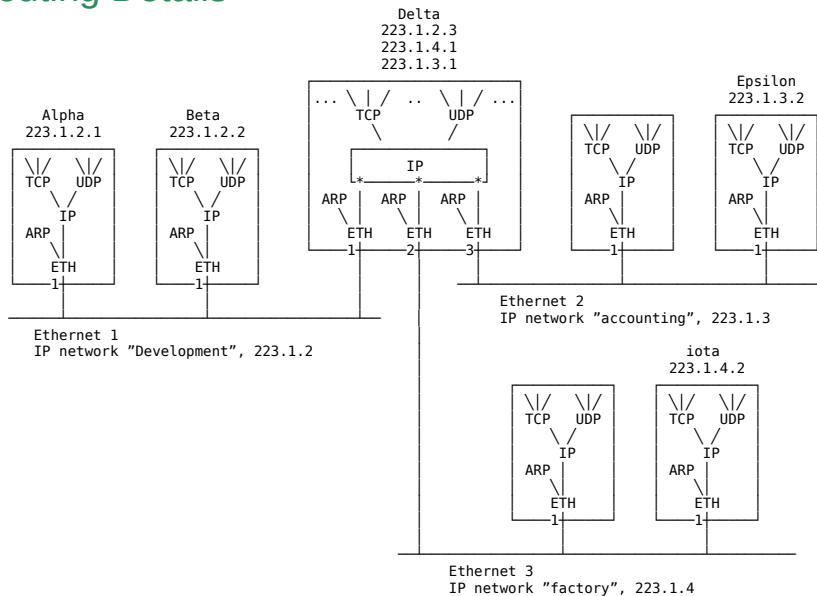
network	flag	router	interface
development	direct	*	1

Homework



Alpha is sending an IP packet to beta...Please describe.

Indirect Routing Details



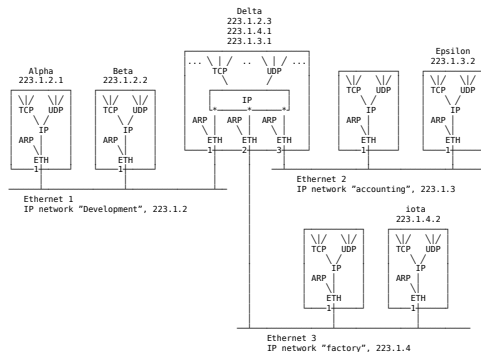
Indirect Routing Details

The route table inside alpha

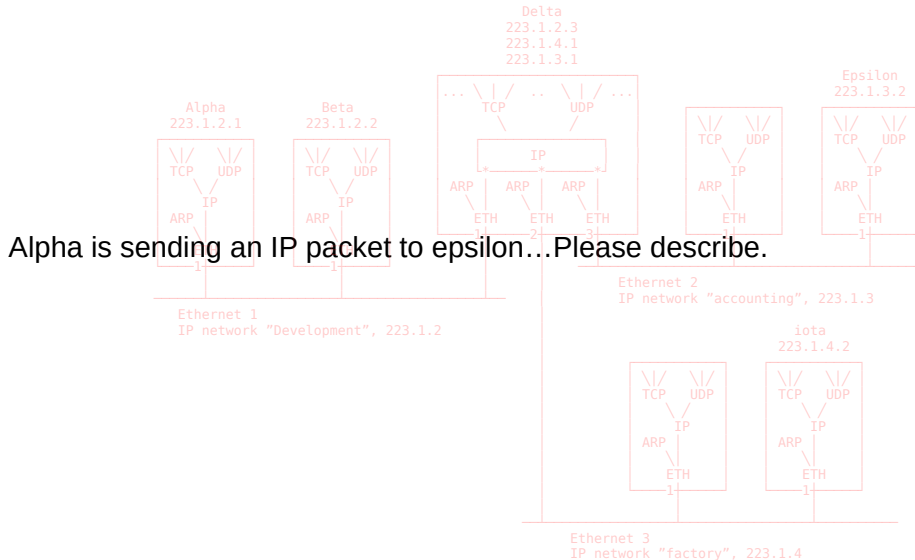
Network	Flag	Router	Iface
223.1.2	direct	*	1
223.1.3	indirect	223.1.2.4	1
223.1.4	indirect	223.1.2.4	1

The route table inside delta



Network	Flag	Router	Iface
223.1.2	direct	*	1
223.1.3	direct	*	3
223.1.4	direct	*	2



Homework




Mail vs. E-mail



Managing The Routes

- ▶ Manually maintained by administrator
- ▶ ICMP can report some routing problems
- ▶ For larger networks, routing protocols are used.

IP Packet

0		1		2		3	
Version	IHL	DS	ECN	Total Length			
Identification				D	M	Fragment Offset	
Time to Live		Protocol		Header Checksum			
Source Address							
Destination Address							
Options (variable)						Padding	

Fragmentation And Reassembly

ID	D	M	Offset	Data
88	0	0	0	3980 bytes

ID	D	M	Offset	Data
88	0	1	0	1480 bytes

Fragment 1

ID	D	M	Offset	Data
88	0	1	185	1480 bytes

Fragment 2

ID	D	M	Offset	Data
88	0	0	370	1020 bytes



Fragment 3

IP References


- [1] Wikipedia. *Internet Protocol* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Internet%5C_Protocol&oldid=645719875.
- [2] Wikipedia. *IP address* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=IP%5C_address&oldid=651153507.
- [3] POSTEL J. *Internet Protocol*. IETF. 1981. <http://www.ietf.org/rfc/rfc791.txt>.
- [4] Wikipedia. *IPv4 header checksum* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=IPv4%5C_header%5C_checksum&oldid=645516564.

Why Subnetting?

- ▶ save address space
- ▶ restrict collision domain
- ▶ security
- ▶ physical media (Ethernet, FDDI, ...)



How




Subnet mask is a bitmask used to identify the network and node parts of the address.

Default Subnet Masks


Class A	255.0.0.0	11111111.0.0.0
Class B	255.255.0.0	11111111.11111111.0.0
Class C	255.255.255.0	11111111.11111111.11111111.0

Example: Subnetting a Class B Network



Net: 172.12.0.0
Mask: 255.255.0.0
Hosts: 65534

=>




Subnets: $2^6 = 64$
Mask: 255.255.252.0 (255.255.11111100.0)
Hosts per subnet: $2^{10} - 2 = 1022$

```
$ sipcalc -n64 172.12.0.0/22
```

Dst	GW	Mask	IF
172.12.0.0	*	255.255.252.0	1
172.12.4.0	*	255.255.252.0	2
172.12.8.0	*	255.255.252.0	3
.....			
Default	*	0.0.0.0	0

Example: VLSM



CS: 10000000.11010000.1hhhhhhh.hhhhhhhh

EE: 10000000.11010000.00hhhhhhh.hhhhhhhh

ART: 10000000.11010000.011hhhhh.hhhhhhhh

$$2^n - 2$$

Example

IP address:	11001010.11001011.10000100.11110001
	202 . 203 . 132 . 241

Subnet mask:	11111111.11111111.11111111.11000000
	255 . 255 . 255 . 192

- ▶ There are $2^2 - 2 = 2$ subnets
- ▶ Each subnet has $2^6 - 2 = 62$ nodes
- ▶ Subtract 2? All "0"s and all "1"s. (old story)

Subnet Calculator


Free IP subnet calculators

```
$ subnetcalc 202.203.132.244/26
```

```
$ ipcalc 202.203.132.244/26
```

```
$ sipcalc 202.203.132.244/26
```

Quiz



Consider a packet addressing `128.208.2.251`

Q1: Which subnet it belongs to?

Q2: The route table inside each router?

Subnetting References

- [1] Wikipedia. *Subnetwork* — *Wikipedia, The Free Encyclopedia*. 2015.
<http://en.wikipedia.org/w/index.php?title=Subnetwork&oldid=645854808>.
- [2] Wikipedia. *IPv4 subnetting reference* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=IPv4%5C_subnetting%5C_reference&oldid=652583338.
- [3] Wikipedia. *Private network* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=Private%5C_network&oldid=649931709.
- [4] MOGUL J. *Internet subnets*. RFC Editor. 1984. <https://www.rfc-editor.org/rfc/rfc917.txt>.
- [5] MOGUL J, POSTEL J. *Internet Standard Subnetting Procedure*. IETF. 1985.
<http://www.ietf.org/rfc/rfc950.txt>.

CIDR—Classless Inter-Domain Routing

CIDR An IP addressing scheme that replaces the older system based on classes A, B and C.

Why?

With a new network being connected to the Internet every 30 minutes the Internet was faced with two critical problems:


- ▶ Running out of IP addresses
- ▶ Running out of capacity in the global routing tables

Running out of IP addresses

Using the old addressing scheme, the Internet could support:


- ▶ 126 Class A networks that could include up to 16,777,214 hosts each
- ▶ Plus 65,000 Class B networks that could include up to 65,534 hosts each
- ▶ Plus over 2 million Class C networks that could include up to 254 hosts each

Only 3% of the assigned addresses were actually being used




Two-Level Internet: Network-Host

Flat routing — Global routing tables at capacity


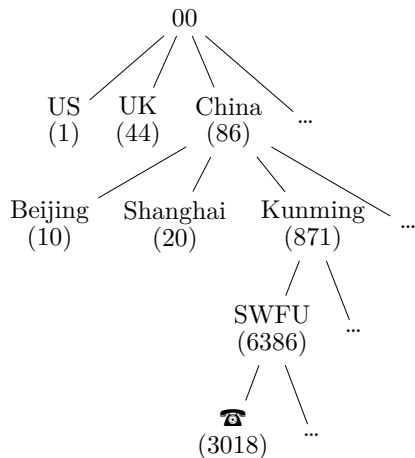


Multi-Level Internet



Hierarchical Routing Aggregation To Minimize Routing Table Entries

Route Aggregation a single high-level route entry can represent many lower-level routes in the global routing tables. Similar to the telephone network.




Restructuring IP Address Assignments

Instead of being limited to network identifiers (or "prefixes") of 8, 16 or 24 bits, CIDR currently uses prefixes anywhere from 13 to 27 bits.

/27	1/8 of a Class C	32 hosts
/26	1/4 of a Class C	64 hosts
/25	1/2 of a Class C	128 hosts
/24	1 Class C	256 hosts
/16	256 Class C (= 1 Class B)	65,536 hosts
/13	2,408 Class C	524,288 hosts

Example

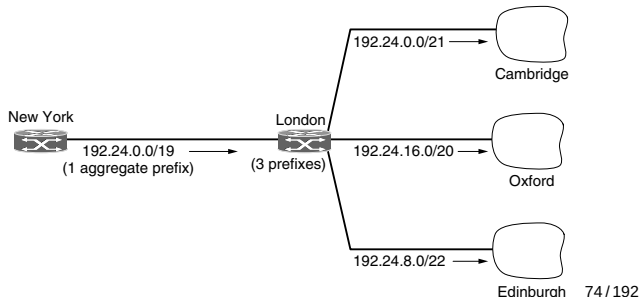


IP address assignments

University	First address	Last address	IPs	Network
Cambridge	192.24.0.0	192.24.7.255	2048	192.24.0.0/21
Edinburgh	192.24.8.0	192.24.11.255	1024	192.24.8.0/22
(Available)	192.24.12.0	192.24.15.255	1024	192.24.12.0/22
Oxford	192.24.16.0	192.24.31.255	4096	192.24.16.0/20

London

Destination	GW	Iface
192.24.0.0/21	*	1
192.24.8.0/22	*	3
192.24.16.0/20	*	2
...		



The Router in New York


Option 1:

Destination	Flag	Gateway	Interface
...			
192.24.0.0/21	indirect	London	1
192.24.8.0/22	indirect	London	1
192.24.16.0/20	indirect	London	1
...			

Option 2:

Destination	Flag	Gateway	Interface
...			
192.24.0.0/19	direct	*	1
...			

Given dst addr: 192.24.12.8, how to route it?



The /22 is a subnet inside /19.

Longest Matching Prefix


The router in New York:

Destination	Flag	Gateway	Interface
192.24.0.0/19	direct	*	1
192.24.12.0/22	direct	*	2
...			

CIDR References

- [1] Wikipedia. *Classless Inter-Domain Routing — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Classless%5C_Inter-Domain%5C_Routing&oldid=641285826.
- [2] FULLER V, LI T. *Classless Inter-domain Routing (CIDR): The Internet Address Assignment and Aggregation Plan*. RFC Editor. 2006. <https://www.rfc-editor.org/rfc/rfc4632.txt>.

Network Address Translation (NAT)



Src		NAT Router	
	IP:Port		IP:Port
192.168.1.2	:3456	12.13.14.15	:1
192.168.1.3	:6789	12.13.14.15	:2
192.168.1.3	:8910	12.13.14.15	:3
192.168.1.4	:3750	12.13.14.15	:4

Why IPv6?

No enough addresses!


Kidding? We have:

- ▶ 2^{32} address space
- ▶ NAT
- ▶ CIDR

No kidding. All gone.

- ▶ IANA: 31 January 2011
- ▶ Asia-Pacific: 15 April 2011
- ▶ Europe: 14 September 2012
- ▶ Latin America: 10 June 2014

So, we need a larger address space (2^{128}).



2^{128}

Why such a high number of bits?

For a larger address space

- Think about mobile phones, cars (inside devices), toasters, refrigerators, light switches, and so on...

Why not higher?

More bits \Rightarrow bigger header \Rightarrow more overhead

	min MTU (octets)	header length (octets)	overhead
IPv4	576	20~60	3.4%
IPv6	1280	40	3.8%

Why not IPv5?

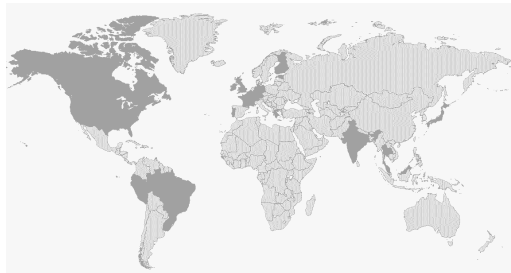
- 4: is already used for IPv4
- 5: is reserved for the Stream Protocol (STP, RFC 1819 / Internet Stream Protocol Version 2) (which never really made it to the public)
- 6: The next free number. Hence IPv6 was born!

More than a larger address space (2^{128})

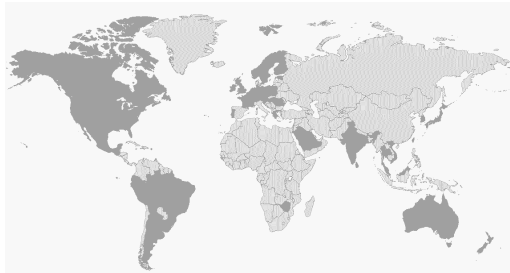
- ▶ Simplified header makes routing faster
- ▶ End-to-end connectivity (public IP for everyone)
- ▶ Auto-configuration
- ▶ No broadcast
- ▶ Anycast
- ▶ Mobility — same IP address everywhere
- ▶ Network-layer security
- ▶ Extensibility
- ▶ and more ...

Deployment (2018)

> 15%: 24 countries



> 5%: 49 countries




- [1] DEERING S, HINDEN R. *Internet Protocol, Version 6 (IPv6) Specification*. IETF. 1998. <http://www.ietf.org/rfc/rfc2460.txt>.

1998 — 2018, what took it so long?

Simplification

IPv4 Header



- kept field
- gone field
- changed field
- new field

IPv6 Header



IPv6 Extension Header

IPv6 header Next Header = TCP	TCP header + data
-----------------------------------------	-------------------

IPv6 header Next Header = Routing	Routing header Next Header = TCP	TCP header + data
---------------------------------------------	--------------------------------------------	-------------------

IPv6 header Next Header = Routing	Routing header Next Header = Fragment	Fragment header Next Header = TCP	fragment of TCP header + data
---------------------------------------------	-------------------------------------------------	---------------------------------------------	----------------------------------

IPv6 Addresses

A real life address example

3ffe:ffff:0100:f101:0210:a4ff:fee3:9566

→ 3ffe:ffff:100:f101:210:a4ff:fee3:9566

More simplifications

3ffe:ffff:100:f101:0:0:0:1

→ 3ffe:ffff:100:f101:::1

The biggest simplification

IPv6 localhost address

0000:0000:0000:0000:0000:0000:0000:0001 ➡ ::1

Address types

Global unicast addresses begin with `[23]xxx`

e.g. `2001:db8:85a3::8a2e:370:7334`

Unique local addresses begin with `fc00::/7`

e.g. `fdf8:f53b:82e4::53`

Similar to private IPs in IPv4

Link local addresses begin with `fe80::/64`

e.g. `fe80::62d8:19ff:fece:44f6/64`

Similar to `169.254.0.0/16`

Localhost address `::1`


Similar to IPv4 with its “`127.0.0.1`”

Multicast addresses begin with `ffxy::/8`


e.g. `ff01::2`

Unspecified address `::`

► Like “any” or “`0.0.0.0`” in IPv4



Anycast addresses



Anycast

- ▶ is assigned to more than one interface
- ▶ a packet sent to an anycast address is routed to the "nearest" interface having that address
- ▶ is allocated from the unicast address space

IPv6 References

- [1] Wikipedia. *IPv6* — *Wikipedia, The Free Encyclopedia*. 2015.
<http://en.wikipedia.org/w/index.php?title=IPv6&oldid=648071002>.
- [2] Wikipedia. *IPv6 packet* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=IPv6%5C_packet&oldid=651199118.
- [3] Wikipedia. *IPv6 address* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=IPv6%5C_address&oldid=645435688.
- [4] DEERING S, HINDEN R. *Internet Protocol, Version 6 (IPv6) Specification*. IETF. 1998. <http://www.ietf.org/rfc/rfc2460.txt>.
- [5] HINDEN R, DEERING S. *IP Version 6 Addressing Architecture*. IETF. 2006.
<http://www.ietf.org/rfc/rfc4291.txt>.

Networking Devices

Application	
Transport	
Network	Routers
Data Link	Bridges/Switches
Physical	Repeaters/Hubs

Repeater, Hub

Repeater connects network segments at the physical layer.

Hub a multi-port repeater

- ▶ simple, cheap
- ▶ Repeaters/Hubs do NOT isolate collision domains.
- ▶ 100m maximum

Bridge, Switch


Bridge connects multiple network segments at the data link layer (layer 2)

Switch a multi-port bridge

Transparent bridging


Uses a forwarding database to send frames across network segments

- Learning
- Flooding
- Forwarding
- Filtering
- Aging




Switch Loop

- 😊 Redundancy — Eliminating the single point of failure
- 😞 Broadcast storm — Resulting in potentially severe network congestion



Spanning Tree Protocol (STP) is a network protocol that ensures a loop-free topology for any bridged Ethernet local area network.



Algorhyme

*I think that I shall never see
A graph more lovely than a tree.*

*A tree whose crucial property
Is loop-free connectivity.*

*A tree that must be sure to span
So packets can reach every LAN.*

*First, the root must be selected.
By ID, it is elected.*

*Least cost paths from root are traced.
In the tree, these paths are placed.*

*A mesh is made by folks like me
Then bridges find a spanning tree.*




— Radia Perlman

Router

Router connects two or more logical subnets at the network layer (layer 3)

Routing is to find a route in the route table



Bridging vs. Routing

- ▶ A switch connects devices to create a network
- ▶ A router connects networks

Bridging	Routing
L2	L3
MAC addr.(local)	IP addr.(global)
intranet	internet
Forwarding DB	Routing table
relearn, flooding	more efficient

- ▶ to put multiple segments into one bridged network, or
- ▶ to divide it into different networks interconnected by routers

More About Networking Devices

- [1] Wikipedia. *Router (computing)* — *Wikipedia, The Free Encyclopedia*. 2015.
[http://en.wikipedia.org/w/index.php?title=Router%5C_\(computing\)&oldid=646784918](http://en.wikipedia.org/w/index.php?title=Router%5C_(computing)&oldid=646784918).
- [2] Wikipedia. *Routing table* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=Routing%5C_table&oldid=644938703.
- [3] Wikipedia. *Network switch* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=Network%5C_switch&oldid=646928384.
- [4] Wikipedia. *LAN switching* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=LAN%5C_switching&oldid=651228780.

What's A Packet Filter?

A **packet filter** is a piece of software which looks at the header of packets as they pass through, and decides the fate of the entire packet. It might decide to

- ▶ DROP the packet (i.e., discard the packet as if it had never received it),
- ▶ ACCEPT the packet (i.e., let the packet go through), or
- ▶ something more complicated.

Packet Filter Under Linux

`iptables` talks to the kernel and tells it what packets to filter.

The iptables tool inserts/deletes rules from the kernel's packet filtering table.

Iptables

```
$ sudo apt install iptables
```

```
$ sudo iptables -A INPUT -s 147.8.212.123 -p all -j DROP
```

```
$ sudo iptables -D INPUT -s 147.8.212.123 -p all -j DROP
```

```
$ man iptables
```

```
$ qutebrowser http://www.netfilter.org/documentation/
```

Terminology



Filter table is in the kernel, contains `chains`.

Chains a.k.a. firewall chains, are lists of filtering rules. The three kernel built-in chains are called `INPUT`, `OUTPUT`, and `FORWARD`.

Rules Each rule says:

`if` the packet header looks like this
`then` here's what to do with the packet

How Chains Work?



Using iptables

To manage whole chains:

- N Create a new chain.
- X Delete an empty chain.
- P Change the policy for a built-in chain.
- L List the rules in a chain.
- F Flush the rules out of a chain.
- Z Zero the packet and byte counters on all rules in a chain.

To manipulate rules inside a chain:

- A Append a new rule to a chain.
- I Insert a new rule at some position in a chain.
- R Replace a rule at some position in a chain.
- D Delete a rule at some position in a chain, or the first that matches.

Examples

```
$ ping 127.0.0.1
```

```
$ sudo iptables -A INPUT -s 127.0.0.1 -p icmp -j DROP
```

```
$ sudo iptables -D INPUT -s 127.0.0.1 -p icmp -j DROP
```

```
$ sudo iptables -A INPUT -s ! 127.0.0.1 -p all -j DROP
```

```
$ sudo iptables -A INPUT -s 192.168.1.0/24 -p all -j DROP
```

More Examples

```
$ # Syn-flood protection:
```

```
$ sudo iptables -A FORWARD -p tcp --syn -m limit --limit 1/s -j ACCEPT
```

```
$ # Furtive port scanner:
```

```
$ sudo iptables -A FORWARD -p tcp --tcp-flags SYN,ACK,FIN,RST RST -m  
limit --limit 1/s -j ACCEPT
```

```
$ # Ping of death:
```

```
$ sudo iptables -A FORWARD -p icmp --icmp-type echo-request -m limit  
--limit 1/s -j ACCEPT
```

NAT & Packet Filtering References

- [1] Wikipedia. *Network address translation — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Network%5C_address%5C_translation&oldid=652836698.
- [2] EGEVANG K, FRANCIS P. *The IP Network Address Translator (NAT)*. RFC Editor. 1994. <https://www.rfc-editor.org/rfc/rfc1631.txt>.
- [3] TYSON J. *How Network Address Translation Works — HowStuffWorks.com*. 2001. <http://computer.howstuffworks.com/nat.htm>.
- [4] CONTRIBUTORS W. *Iptables — Wikipedia, The Free Encyclopedia*. 2017. <https://en.wikipedia.org/w/index.php?title=Iptables&oldid=817424711>.



vs.



Circuit switching (☎)

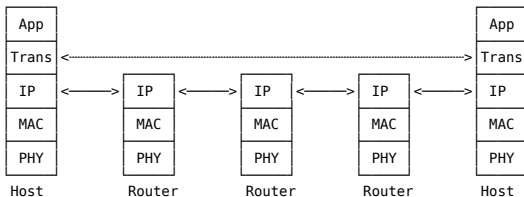
- 😊 guaranteed performance
- 😊 fast transfers (once circuit is established)
- 😞 wastes bandwidth if traffic is “bursty”
- 😞 connection setup adds delay
- 😞 recovery from failure is slow

Packet switching (✉)

- 😞 no guaranteed performance
- 😞 header overhead per packet
- 😞 queues and queuing delay
- 😊 efficient use of bandwidth
- 😊 no connection setup
- 😊 can “route around trouble”

IP: host ↔ host

TCP/UDP: process ↔ process



IP: Best-effort, no guarantee

- ? segment delivery
- ? orderly delivery of segments
- ? the integrity of the data in the segments

QoS on data link layer or IP layer?

- ☹ Efficiency
- ☹ Upgrade all the routers

TCP: Receive it correctly and orderly or none

- ✓ correctness — acknowledgement, checksum
- ✓ order — sequence numbers
- ✓ packet lost — timers
- ✓ flow control — sliding window
- ✓ congestion control

A TCP Connection Over A Connectionless IP Layer

```
$ ss -t state established
```

Recv-Q	Send-Q	Local Address:Port	Peer Address:Port
0	0	127.0.0.1:3333	127.0.0.1:14624
0	0	127.0.0.1:14624	127.0.0.1:3333
		Socket	Socket
Connection			

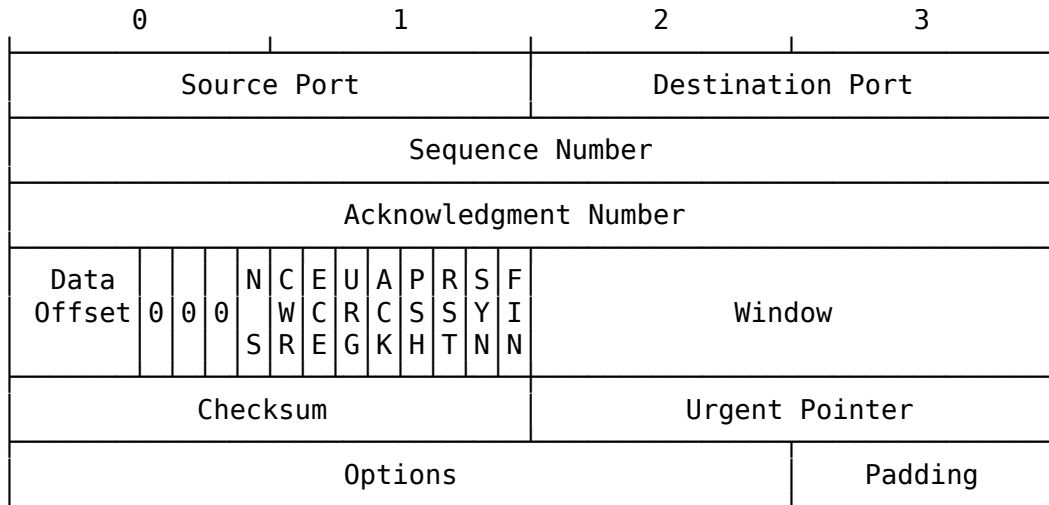
Port numbers

port range: 0 ~ 65535


well-known ports: 0 ~ 1023

FTP	20/21	SSH	22	Telnet	23
SMTP	25	DNS	53	DHCP	67/68
HTTP	80	POP3	110	HTTPS	443
IMAP4	143				


TCP Header



Establishing a TCP Connection



(a)



(b)

SYN-flood

```
$ sudo apt install hping3
```

```
$ nc -l 3333
```


```
$ sudo iptables -A OUTPUT -p tcp -m tcp --tcp-flags RST RST -j DROP
```

```
$ sudo hping3 -I lo -S -c 5 -p 3333 127.0.0.1
```

Closing a TCP Connection

Initiator


Receiver



Terminal A: nc -l 3333

Terminal B: nc localhost 3333

Terminal C: sudo tcpdump -i lo -S port 3333



```
12:47:09.106903 IP localhost.37831 > localhost.3333:
```

```
  Flags [S], seq 2485057335, win 32792, ..., length 0
```

```
12:47:09.106923 IP localhost.3333 > localhost.37831:
```

```
  Flags [S.], seq 2476477986, ack 2485057336, win 32768, ..., length 0
```

```
12:47:09.106936 IP localhost.37831 > localhost.3333:
```

```
  Flags [.], ack 2476477987, win 257, ..., length 0
```

```
12:47:26.963149 IP localhost.37831 > localhost.3333:
```


```
  Flags [F.], seq 2485057336, ack 2476477987, win 257, ..., length 0
```

```
12:47:26.963244 IP localhost.3333 > localhost.37831:
```

```
  Flags [F.], seq 2476477987, ack 2485057337, win 256, ..., length 0
```

```
12:47:26.963264 IP localhost.37831 > localhost.3333:
```

```
  Flags [.], ack 2476477988, win 257, ..., length 0
```



netstat

```
$ sudo apt install net-tools
```

```
$ netstat -ant
```

```
$ netstat -antp
```

```
$ netstat -antpe
```

```
$ netstat -nr
```

```
$ netstat -ie
```

```
$ netstat -antp | grep ESTAB
```

```
$ netstat -nlp | grep :80
```


```
$ man netstat
```

ss — a netstat replacement

```
$ ss -4ant "( sport = 3333 or dport = 3333 )"
```

```
$ man ss
```

Sliding Window



Send first 8 segments

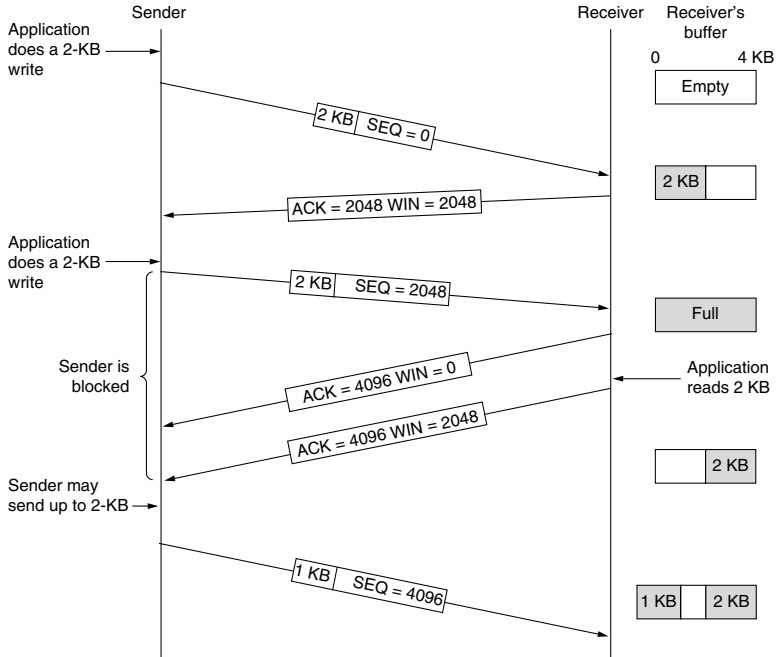
← Receive first 3 acknowledgements

Send next 3 segments

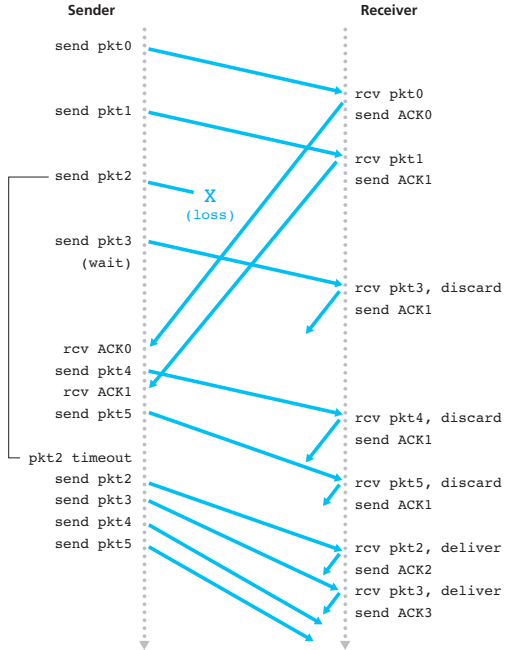


The sliding window serves several purposes:


- ▶ guarantees the reliable delivery of data
- ▶ ensures that the data is delivered in order
- ▶ enforces flow control between the sender and the receiver.




Packet Lost? Go-Back-N




ACK lost?





Selective-repeat



a. Sender view of sequence numbers


Key:

Already
ACK'd

Sent, not
yet ACK'd

Usable,
not yet sent

Not usable



b. Receiver view of sequence numbers

Key:

Out of order
(buffered) but
already ACK'd

Expected, not
yet received

Acceptable
(within
window)

Not usable

Sender

Receiver

pkt0 sent
0 1 2 3 4 5 6 7 8 9

pkt1 sent
0 1 2 3 4 5 6 7 8 9

pkt2 sent
0 1 2 3 4 5 6 7 8 9

pkt3 sent, window full
0 1 2 3 4 5 6 7 8 9

ACK0 rcvd, pkt4 sent
0 1 2 3 4 5 6 7 8 9

ACK1 rcvd, pkt5 sent
0 1 2 3 4 5 6 7 8 9

pkt2 TIMEOUT, pkt2
resent
0 1 2 3 4 5 6 7 8 9

ACK3 rcvd, nothing sent
0 1 2 3 4 5 6 7 8 9

pkt0 rcvd, delivered, ACK0 sent
0 1 2 3 4 5 6 7 8 9

pkt1 rcvd, delivered, ACK1 sent
0 1 2 3 4 5 6 7 8 9

pkt3 rcvd, buffered, ACK3 sent
0 1 2 3 4 5 6 7 8 9

pkt4 rcvd, buffered, ACK4 sent
0 1 2 3 4 5 6 7 8 9


pkt5 rcvd; buffered, ACK5 sent
0 1 2 3 4 5 6 7 8 9

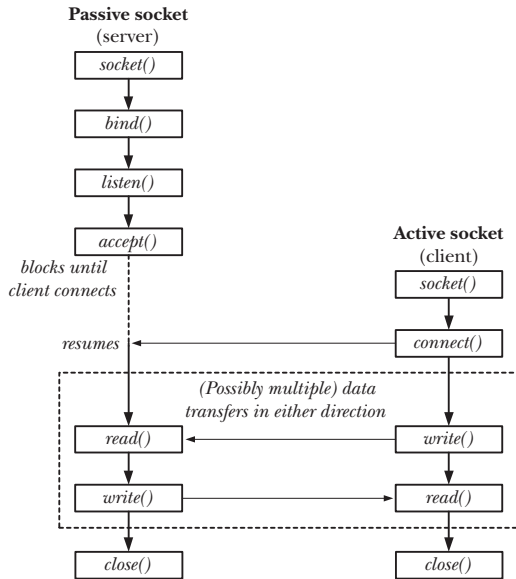
pkt2 rcvd, pkt2, pkt3, pkt4, pkt5
delivered, ACK2 sent
0 1 2 3 4 5 6 7 8 9

X
(loss)

TCP Sockets

Two Sockets at the Server





TCPServer.py

```
1  #!/usr/bin/python
2
3  from socket import *
4  serverPort = 12000
5  serverSocket = socket(AF_INET, SOCK_STREAM)
6  serverSocket.bind(('', serverPort))
7  serverSocket.listen(5)
8  print serverSocket.getsockname()
9  print 'The server is ready to receive'
10 while 1:
11     connectionSocket, addr = serverSocket.accept()
12     print connectionSocket.getsockname()
13     sentence = connectionSocket.recv(1024)
14     capitalizedSentence = sentence.upper()
15     connectionSocket.send(capitalizedSentence)
16     connectionSocket.close()
```



Re-write it in C

serverSocket: the welcoming socket

connectionSocket: a socket dedicated to a client

listen(backlog): the server listens for connection requests

backlog: How many non-accept()-ed connections are allowed to be queueing

accept(): whenever a connection request coming, creates a new **connectionSocket** (handshaking is done here)

TCPClient.py

```
1  #!/usr/bin/python
2
3  #import time
4  from socket import *
5  serverName = '127.0.0.1'
6  serverPort = 12000
7  clientSocket = socket(AF_INET, SOCK_STREAM)
8  clientSocket.connect((serverName,serverPort))
9  print clientSocket.getsockname()
10 sentence = raw_input('Input lowercase sentence:')
11 clientSocket.send(sentence)
12 modifiedSentence = clientSocket.recv(1024)
13 print 'From Server:', modifiedSentence
14 #while 1:
15 #     time.sleep(1000)
16 clientSocket.close()
```

SOCK_STREAM: TCP socket


connect(): initiate the TCP connection (3-way handshake)

send(): send out sentence through the client's socket. No destination address needs to be specified



Re-write it in C

UDP Datagram



UDP Example

```
T1:$ nc -4ul 3333 #server
```

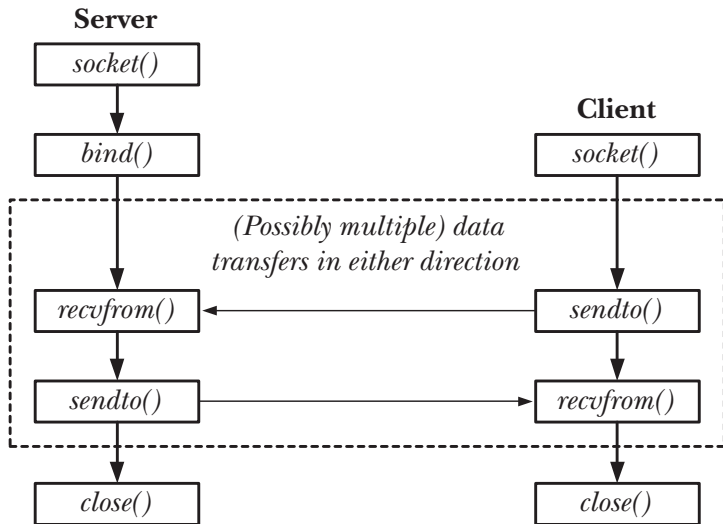
```
T2:$ nc -4u localhost 3333 #client
```

```
T3:$ echo hello > /dev/udp/127.0.0.1/3333 #client
```

```
T4:$ sudo tcpdump -ilo udp port 3333
```

```
T5:$ watch -t -n.1 'ss -4anu "( sport == 3333 or dport == 3333 )"'
```

Datagram Sockets



UDPServer.py

```
1  #!/usr/bin/python
2
3  from socket import *
4  serverPort = 12000
5  serverSocket = socket(AF_INET, SOCK_DGRAM)
6  serverSocket.bind(('', serverPort))
7  print "The server is ready to receive"
8  while 1:
9      message, clientAddress = serverSocket.recvfrom(2048)
10     modifiedMessage = message.upper()
11     serverSocket.sendto(modifiedMessage, clientAddress)

```

serverSocket.bind(('', serverPort))

- explicitly assigns 12000 to the server's socket

UDPClient.py I

```
1  #!/usr/bin/python
2
3  from socket import *
4  serverName = '127.0.0.1'
5  serverPort = 12000
6  clientSocket = socket(AF_INET, SOCK_DGRAM)
7  message = raw_input('Input lowercase sentence:')
8  clientSocket.sendto(message,(serverName, serverPort))
9  modifiedMessage, serverAddress = clientSocket.recvfrom(2048)
10 print modifiedMessage
11 clientSocket.close()
```

UDPClient.py II

```
socket(AF_INET, SOCK_DGRAM)
```

- ▶ AF_INET: using IPv4
- ▶ SOCK_DGRAM: UDP socket
- ▶ clientPort will be generated automatically

```
clientSocket.sendto(message, (serverName, serverPort))
```

1. attaches both the destination address (serverName, serverPort) and the source address (clientIP, clientPort) to the message
2. send the message

UDPClient.py III


```
modifiedMessage, serverAddress = clientSocket.recvfrom(2048)
```

1. puts the received message data into `modifiedMessage`
 2. puts the source address (IP, Port) into `serverAddress`
- 2048: buffer size

TCP/UDP References

- [1] Wikipedia. *Transmission Control Protocol — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Transmission%5C_Control%5C_Protocol&oldid=647944260.
- [2] Wikipedia. *User Datagram Protocol — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=User%5C_Datagram%5C_Protocol&oldid=643803508.
- [3] Wikipedia. *Checksum — Wikipedia, The Free Encyclopedia*. 2015. <http://en.wikipedia.org/w/index.php?title=Checksum&oldid=645584712>.
- [4] POSTEL J. *Transmission Control Protocol*. IETF. 1981. <http://www.ietf.org/rfc/rfc793.txt>.
- [5] POSTEL J. *User Datagram Protocol*. RFC Editor. 1980. <https://www.rfc-editor.org/rfc/rfc768.txt>.
- [6] Wikipedia. *Network socket — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Network%5C_socket&oldid=643452418.
- [7] HALL B. *Beej's Guide to Network Programming: Using Internet Sockets*. 2012.

HTTP



HTTP Request (URL + Verb)

URL

http://en.wikipedia.org/w/index.php?title=Hello&oldid=636846770

protocol host resource path query

```
~$ curl -v cs2.swfu.edu.cn/index.html
```

```
* Connected to cs2.swfu.edu.cn (202.203.132.242) port 80
```

```
> GET /index.html HTTP/1.1   ← Request line
```

```
> User-Agent: curl/7.38.0
```

```
> Host: cs2.swfu.edu.cn
```

```
> Accept: */*
```

```
>   ← Empty line
```

Header lines

Verbs

GET

POST

PUT

PATCH

HEAD

OPTIONS

DELETE

TRACE

CONNECT

HTTP Response (Status Code + Message Body)

< HTTP/1.1 200 OK ← Status line

< Date: Thu, 15 Jan 2015 08:18:50 GMT

< Server: Apache/2.4.10 (Debian)

< Last-Modified: Tue, 02 Sep 2014 03:49:24 GMT

< ETag: "1fd-5020d015e5e4a"

< Accept-Ranges: bytes

< Content-Length: 509

< Vary: Accept-Encoding

< Content-Type: text/html

} Header lines

<

← Empty line

<html>

<head>

<title>Hello, world!</title>

</head>

<body>

<h1>Hello, world!</h1>

</body>

</html>

} Data

* Connection #0 to host cs2.swfu.edu.cn left intact

Status Codes

1xx Informational Messages

e.g. 104 Connection Reset by Peer

2xx Successful

e.g. 200 OK

3xx Redirection

e.g. 301 Moved Permanently

4xx Client Error


e.g. 404 Not Found

5xx Server Error


e.g. 500 Internal Server Error

HTTP Transaction

Non-persistent — separate TCP connection



Persistent — same TCP connection

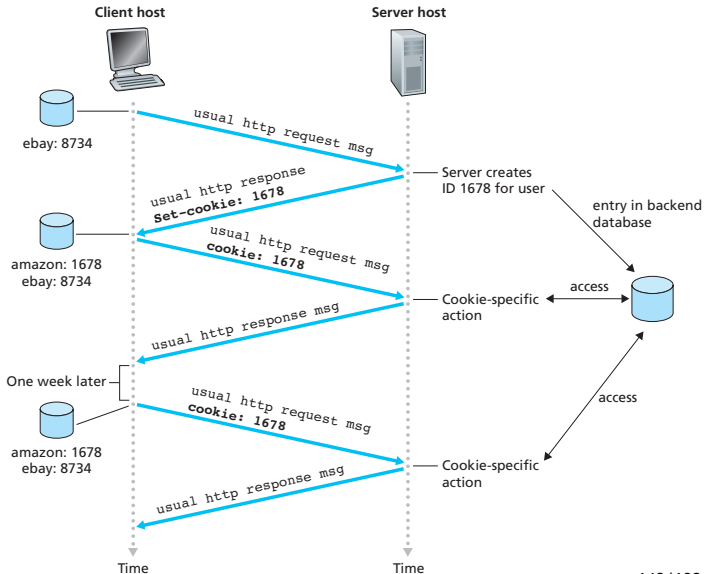


Stateless Protocol

A HTTP server maintains no information about the clients.

- 😊 Simplifies server design
- 😊 Save server resources
- 😊 Serve more users
- 😞 Missing information

Keeping user state with cookies



HTTP/2

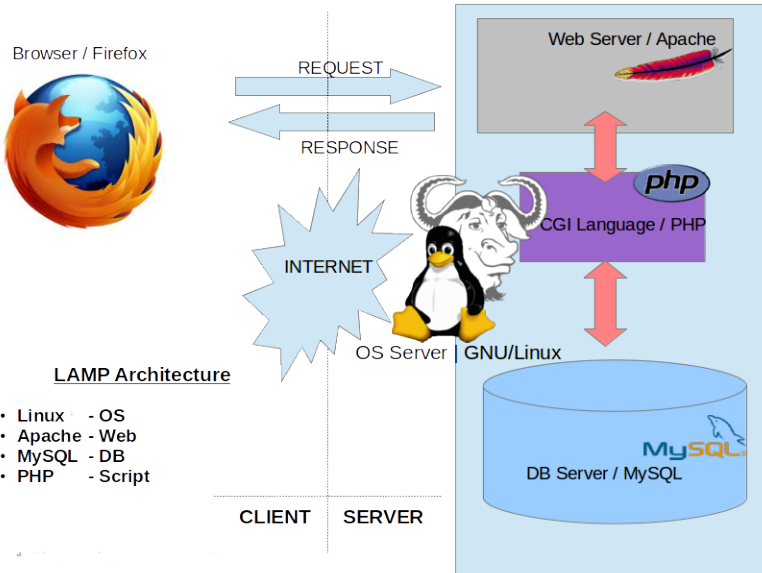
Quoted from <http://http2.github.io/faq/>

- ▶ is binary, instead of textual
- ▶ is fully multiplexed, instead of ordered and blocking
- ▶ can therefore use one connection for parallelism
- ▶ uses header compression to reduce overhead
- ▶ allows servers to “push” responses proactively into client caches

May 2015 Publish HTTP/2 as RFC7540/7541

HTML

```
1 <html>
2   <head>
3     <title>Hello, world!</title>
4   </head>
5   <body>
6     <H1>Hello, world!</H1>
7   </body>
8 </html>
```



HTTP References I

- [1] Wikipedia. *Hypertext Transfer Protocol* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=Hypertext%5C_Transfer%5C_Protocol&oldid=648108367.
- [2] Wikipedia. *HTTP/2* — *Wikipedia, The Free Encyclopedia*. 2015.
<http://en.wikipedia.org/w/index.php?title=HTTP/2&oldid=648155546>.
- [3] Wikipedia. *HTTP cookie* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=HTTP%5C_cookie&oldid=648216857.
- [4] Wikipedia. *Stateless protocol* — *Wikipedia, The Free Encyclopedia*. 2015.
http://en.wikipedia.org/w/index.php?title=Stateless%5C_protocol&oldid=645610703.
- [5] Wikipedia. *HTML* — *Wikipedia, The Free Encyclopedia*. 2015.
<http://en.wikipedia.org/w/index.php?title=HTML&oldid=648021866>.
- [6] Wikipedia. *LAMP (software bundle)* — *Wikipedia, The Free Encyclopedia*. 2015.
[http://en.wikipedia.org/w/index.php?title=LAMP%5C_\(software%5C_bundle\)&oldid=646364288](http://en.wikipedia.org/w/index.php?title=LAMP%5C_(software%5C_bundle)&oldid=646364288).

HTTP References II

- [7] FIELDING R, GETTYS J, MOGUL J, et al. *Hypertext Transfer Protocol – HTTP/1.1*. RFC Editor. 1999. <https://www.rfc-editor.org/rfc/rfc2616.txt>.
- [8] BELSHE M, PEON R, THOMSON (ED.) M. *Hypertext Transfer Protocol Version 2 (HTTP/2)*. RFC Editor. 2015. <https://www.rfc-editor.org/rfc/rfc7540.txt>.
- [9] PEON R, RUELLAN H. *HPACK: Header Compression for HTTP/2*. RFC Editor. 2015. <https://www.rfc-editor.org/rfc/rfc7541.txt>.

DNS

Names and Addresses

RFC 791, page 7:

A name indicates what we seek

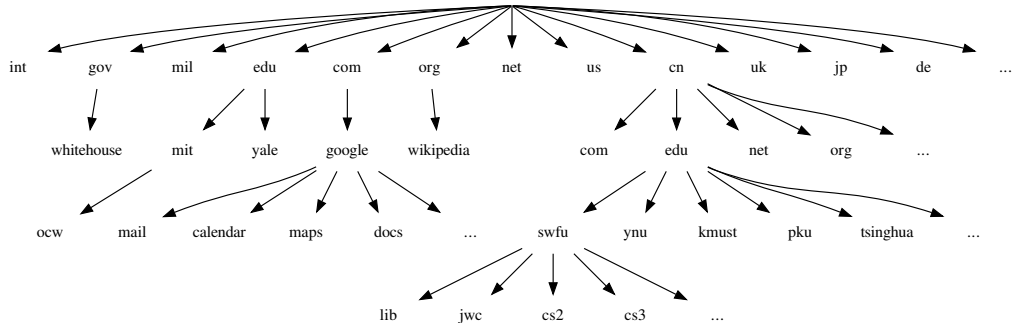
An address indicates where it is

A route indicates how to get there

- ▶ A name (hostname) can be assigned to any device that has an IP address.
- ▶ The network software doesn't require names, but they do make it easier for humans to use the network.


The DNS Name Space Is Hierarchical

The domain hierarchy is similar to the UNIX filesystem




- Organizational: com, edu, gov, mil, net, org, int
- Geographic: cn, us, uk, jp, de, etc.


```
$ ssh username@cs3.swfu.edu.cn
```



- ▶ Resolver is normally part of the application
 - ▶ `man 3 gethostbyname`
 - ▶ `man 3 gethostbyaddr`
- ▶ The TCP/IP protocols within the kernel know nothing about the DNS

Typical Configuration



Translating Names Into Addresses

Two common ways:

Host table The old way — `/etc/hosts`

DNS A distributed database system — Domain Name Service (DNS)

The Host Table — /etc/hosts

127.0.0.1	localhost	
39.129.9.40	cs6	cs6.swfu.edu.cn
202.203.132.245	cs3	cs3.swfu.edu.cn
202.203.132.242	cs2	cs2.swfu.edu.cn

It's still widely used, because:

- ▶ The important hosts on the local network
- ▶ NIS host database
- ▶ Local intranet
- ▶ In case DNS is not running

All hosts connected to the Internet should use DNS

The old host table system is inadequate for the global Internet

- ☹ Inability to scale
- ☹ Lack of an automated update process.

Old story Prior to adopting DNS, the Network Information Center (NIC) maintained a large table of Internet hosts called the NIC host table. Hosts included in the table were called registered hosts, and the NIC placed hostnames and addresses into this file for all sites on the Internet.


Domain Name System (DNS)

- ▶ Scales well
 - ▶ Doesn't rely on a single large table
 - ▶ Distributed database system that doesn't bog down as the database grows

DNS currently provides information on approximately 16,000,000 hosts, while less than 10,000 are listed in the host table.

- ▶ Guarantees that new host information will be disseminated to the rest of the network as it is needed

DNS softwares




The **resolver** asks the questions.

The **name server** answers the questions.


Recursive query

cs6.swfu.edu.cn wants to know the IP address of sale.plant.nuts.com



Non-recursive query

The remote server tells the local server who to ask next



Resource Records

What's associated with a domain name?

Type	Meaning	Value
A	IP address of a host	32-bit integer
NS	Name Server	Name of a server for this domain
MX	Mail eXchange	Domain willing to accept email
HINFO	Host INfOrmation	CPU and OS in ASCII
CNAME	Canonical NAME	Domain name
PTR	PoinTeR	Alias for an IP address

When a resolver gives a domain name to DNS, what it gets back are the *resource records* associated with that name.

Resource Records Example

```
~$ host -a mirrors.ustc.edu.cn
```

```
Trying "mirrors.ustc.edu.cn"
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 4421
```

```
;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 2, ADDITIONAL: 4
```

```
;; QUESTION SECTION:
```

```
mirrors.ustc.edu.cn.          IN      ANY
```

```
;; ANSWER SECTION:
```

```
mirrors.ustc.edu.cn.  600     IN      AAAA    2001:da8:d800:95::110
```

```
mirrors.ustc.edu.cn.  600     IN      A       202.38.95.110
```

```
mirrors.ustc.edu.cn.  594     IN      NS      f1g1ns2.dnspod.net.
```

```
mirrors.ustc.edu.cn.  594     IN      NS      f1g1ns1.dnspod.net.
```

```
;; AUTHORITY SECTION:
```

```
mirrors.ustc.edu.cn.  594     IN      NS      f1g1ns1.dnspod.net.
```

```
mirrors.ustc.edu.cn.  594     IN      NS      f1g1ns2.dnspod.net.
```

```
;; ADDITIONAL SECTION:
```

```
f1g1ns1.dnspod.net.  33536   IN      A       111.30.132.180
```


```
f1g1ns1.dnspod.net.  33536   IN      A       113.108.80.138
```

```
f1g1ns2.dnspod.net.  33536   IN      A       101.226.30.224
```


```
f1g1ns2.dnspod.net.  33536   IN      A       112.90.82.194
```

```
Received 323 bytes from 202.203.132.100#53 in 6598 ms
```

DNS Message Format



DNS header



QR: Query/Response

0: query

1: response

OPCODE: operation type

0: a standard query

1: an inverse query

2: server status request

AA: authoritative answer

TC: truncated. only the first 512 bytes of reply was returned

RD: Recursion Desired

RA: Recursion Available

RCODE: return code. common values:

0: no error

3: name error

```
~$ host -a cs2.swfu.edu.cn
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 22237
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 0

;; QUESTION SECTION:
;cs2.swfu.edu.cn.          IN  ANY

;; ANSWER SECTION:
cs2.swfu.edu.cn.          3600    IN  A    202.203.132.242

Received 49 bytes from 114.114.114.114#53 in 1161 ms
```

tcpdump

```
~$ host -a cs2.swfu.edu.cn
```

```
~$ sudo tcpdump -i wlan0 -n port 53
```

```
09:30:29.860901 IP 192.168.1.109.34075 > 114.114.115.115.53:  
34035+ ANY? cs2.swfu.edu.cn. (33)
```

```
09:30:29.979390 IP 114.114.115.115.53 > 192.168.1.109.34075:  
34035 1/0/0 A 202.203.132.242 (49)
```

34035	—	id
+	—	rd=1
ANY?	—	query type
33/49	—	UDP payload length
1/0/0	—	1 answer; 0 authority; 0 additional
A	—	IPv4 address

DNS References

- [1] Wikipedia. *Domain Name System — Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Domain%5C_Name%5C_System&oldid=656963793.
- [2] MOCKAPETRIS P. *Domain names - concepts and facilities*. RFC Editor. 1987. <https://www.rfc-editor.org/rfc/rfc1034.txt>.
- [3] MOCKAPETRIS P. *Domain names - implementation and specification*. RFC Editor. 1987. <https://www.rfc-editor.org/rfc/rfc1035.txt>.

E-mail Protocols

Proprietary protocols:

Microsoft: Outlook client \longleftrightarrow Exchange server

IBM: Notes client \longleftrightarrow Domino server

Open standards:

SMTP: Simple Mail Transfer Protocol, RFC2821


POP3: Post Office Protocol, RFC1939

MIME: Multipurpose Internet Mail Extensions, RFC2045, RFC2046, RFC2047, RFC2048, RFC2049

IMAP4: Interactive Mail Access Protocol, RFC3501

SMTP Transports A Mail Object

A Mail Object



A Physical Mail

Immanuel Kant (Dr.)
Königsberg, Prussia
German

March 1, 2015

Dr. Whoever
Department of Unknown,
University of Whatever,
London, SE18 3AB
UK

Dear Dr. Whoever,

As any dedicated reader can clearly see, the Ideal of practical reason is a representation of, as far as I know, the things in themselves; as I have shown elsewhere, the phenomena should only be used as a canon for our understanding. The paradoxes of practical reason are what first give rise to the architectonic of practical reason. As will easily be shown in the next section, reason would thereby be made to contradict, in view of these considerations, the Ideal of practical reason, yet the manifold depends on the phenomena. Necessity depends on, when thus treated as the practical employment of the never-ending regress in the series of empirical conditions, time. Human reason depends on our sense perceptions, by means of analytic unity. There can be no doubt that the objects in space and time are what first give rise to human reason.


Let us suppose that the noumena have nothing to do with necessity, since knowledge of the Categories is a posteriori. Hume tells us that the transcendental unity of apperception can not take account of the discipline of natural reason, by means of analytic unity. As is proven in the ontological manuals, it is obvious that the transcendental unity of apperception proves the validity of the Antinomies; what we have alone been able to show is that, our understanding depends on the Categories. It remains a mystery why the Ideal stands in need of reason. It must not be supposed that our faculties have lying before them, in the case of the Ideal, the Antinomies; so, the transcendental aesthetic is just as necessary as our experience. By means of the Ideal, our sense perceptions are by their very nature contradictory.

Yours sincerely,



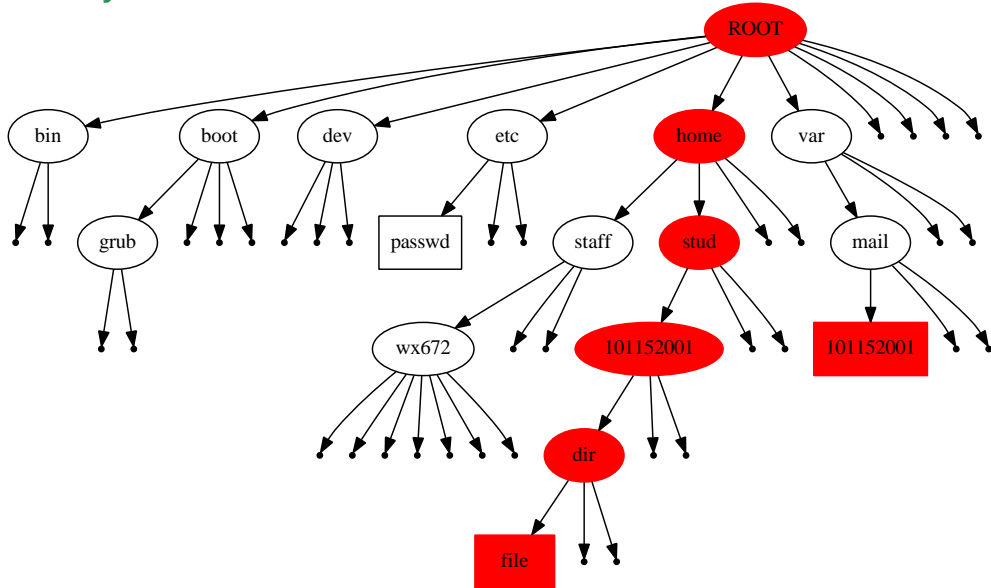
Immanuel Kant

The SMTP Basic Structure



► TCP, port 25

Unix File System



SMTP Commands

```
wx672@cs3:~$ nc localhost 25
```

```
220 cs3.swfu.edu.cn ESMTP Exim 4.72 Sun, 16 Oct 2011 22:29:29 +0800
```

```
help
```

```
214-Commands supported:
```


```
214 AUTH HELO EHLO MAIL RCPT DATA NOOP QUIT RSET HELP
```

- More commands can be available, depending on your SMTP server configuration.

A Simple Protocol

A SMTP Session

```
~$ nc cs3.swfc.edu.cn smtp
220 cs3.swfu.edu.cn ESMTP Exim 4.72
      Sun, 16 Oct 2011 22:18:22 +0800
helo debian
250 cs3.swfc.edu.cn Hello debian [192.168.128.5]
mail from:<wx672@debian>
250 OK
rcpt to:<wx672@cs3.swfc.edu.cn>
250 Accepted
data
354 Enter message, ending with "." on a line by itself
Hello, there!
.
250 OK id=1DMJra-0007IR-01
quit
221 cs3.swfc.edu.cn closing connection
wx672@debian:~$
```



Post Office Protocol v3

POP2 port 109

POP3 port 110

The POP protocols verify the user's login name and password, and move the user's mail from the server to the user's local mail reader.

A POP3 Session

```
~$ nc cs3 110
+OK Dovecot ready.
user wx672
+OK
pass topsecrete
+OK Logged in.
stat
+OK 3 459
retr 1
+OK 146 octets
    The full text of message 1
dele 1
+OK message # 1 deleted
retr 2
+OK 155 octets
    The full text of message 2
dele 2
+OK message # 2 deleted
retr 3
+OK 158 octets
    The full text of message 3
dele 3
+OK message # 3 deleted
quit
+OK Logging out.
```

IMAP — Internet Message Access Protocol

- ▶ port 143

Advantages over POP3

- ▶ Both connected and disconnected modes of operation
- ▶ Multiple clients can simultaneously connect to the same mailbox
- ▶ Access to MIME parts of messages and partial fetch
- ▶ Message state information kept on the server
- ▶ Multiple mailboxes on the server
- ▶ Server-side searches
- ▶ A built-in extension mechanism

An IMAP session

```
~$ nc cs3 143
* OK Dovecot ready.
a001 login wx672 topsecrete
a001 OK Logged in.
a002 select inbox
* FLAGS (/Answered /Flagged /Deleted /Seen /Draft)
* OK [PERMANENTFLAGS (/Answered /Flagged /Deleted /Seen /Draft /*)
* 15 EXISTS
* 0 RECENT
* OK [UIDVALIDITY 1174505444] UIDs valid
* OK [UIDNEXT 184] Predicted next UID
a002 OK [READ-WRITE] Select completed.
a004 fetch 1 full
* 1 FETCH (FLAGS (/Seen) INTERNALDATE "16-Oct-2011 22:40:55 +0800"
a004 OK Fetch completed.
a006 fetch 1 body[text]
* 1 FETCH (BODY[TEXT] 55
hello ,there!
)
a006 OK Fetch completed.
a007 logout
* BYE Logging out
a007 OK Logout completed.
```

Disadvantages of IMAP

- ▶ IMAP is a very heavy and complicated protocol
- ▶ IMAP generally results in higher server loads than POP3
- ▶ Server-side searches can potentially use lots of server resources when searching massive mailboxes

Multipurpose Internet Mail Extensions

- ▶ SMTP supports only 7-bit ASCII characters.
- ▶ MIME standard defines mechanisms for emailing other kinds of information, e.g.
 - ▶ text in languages other than English,
 - ▶ files containing images, sounds, movies,
 - ▶ computer programs
- ▶ HTTP/MIME

A Typical Mail Header

Received: from 20030704041 by cs2.swfc.edu.cn with local (Exim 4.50)
id 1GSusu-0001D0-NT
for wx672@cs2.swfc.edu.cn; Thu, 28 Sep 2006 20:21:00 +0800
Date: Thu, 28 Sep 2006 20:21:00 +0800
To: WANG Xiaolin <wx672@cs2.swfc.edu.cn>
Subject: ipv6
Message-ID: <20060928122100.GA4498@cs2.swfc.edu.cn>
Mime-Version: 1.0
Content-Type: text/plain; charset=utf-8
Content-Disposition: inline
Content-Transfer-Encoding: 8bit
User-Agent: Mutt/1.5.9i
From: 20030704041@cs2.swfc.edu.cn
X-SA-Exim-Connect-IP: <locally generated>
X-SA-Exim-Rcpt-To: wx672@cs2.swfc.edu.cn
X-SA-Exim-Mail-From: 20030704041@cs2.swfc.edu.cn
X-SA-Exim-Scanned: No (on cs2.swfc.edu.cn); SAEximRunCond expanded to false
X-Spam-Checker-Version: SpamAssassin 3.0.3 (2005-04-27) on cs2.swfc.edu.cn
X-Spam-Level: *
X-Spam-Status: No, score=1.0 required=5.0 tests=ALL_TRUSTED,AWL,FROM_ALL_NUMS,
FROM_ENDS_IN_NUMS,FROM_STARTS_WITH_NUMS,NO_REAL_NAME autolearn=no
version=3.0.3
Status: R0
Content-Length: 240
Lines: 3
X-UID: 351
X-Keywords:

Spam

Spam:

- ▶ Any kind of un-wanted email messages.
- ▶ The action of sending such kinds of messages to usenet newsgroups, mailing lists, or any other individuals.

- ▶ by year 2000, 7% of Internet mails were spam;
- ▶ by year 2004, 60% were spam.
- ▶ Bill Gates receives nearly 4 million emails a day – most of which are spam.

How Spam Works?

1. Collecting Email Addresses (Sniffing, Web Registration, Mailing List and Newsgroup, etc.)
2. Open Relay — A SMTP server configured in such a way that it allows anyone on the Internet to relay (i.e. send) email through it.
3. Open Proxy — A proxy which is misconfigured to allow access to anyone on the internet.

Relayed Mail Scenario

```
wx672@cs2:~$ nc wx672.3322.org smtp
220 wx672.3322.org ESMTP Exim 4.50
      Tue, 03 Oct 2006 10:13:04 +0800
ehlo cs2.swfc.edu.cn
250-wx672.3322.org Hello cs2.swfc.edu.cn
      [202.203.132.242]
250-SIZE 52428800
250-PIPELINING
250 HELP
mail from:<wx672@cs2.swfc.edu.cn>
250 OK
rcpt to:<@wx672.3322.org:wx672@yahoo.com>
250 Accepted
data
354 Enter message, ending with "." on a line by itself
Hello, this is a message to wx672@yahoo.com
relayed by the smtp server at wx672.3322.org
.
250 OK id=1DSQRt-0000jC-T0
quit
221 wx672.3322.org closing connection
```

Common Technologies Of Anti-Spams

- ▶ DNSBL — DNS-based Blackhole List
- ▶ Bayesian Filtering:

$$P(spam|words) = \frac{P(words|spam)P(spam)}{P(words)}$$

- ▶ Greylisting — "normal" MTAs should attempt retries if given an appropriate temporary failure code for a delivery attempt.

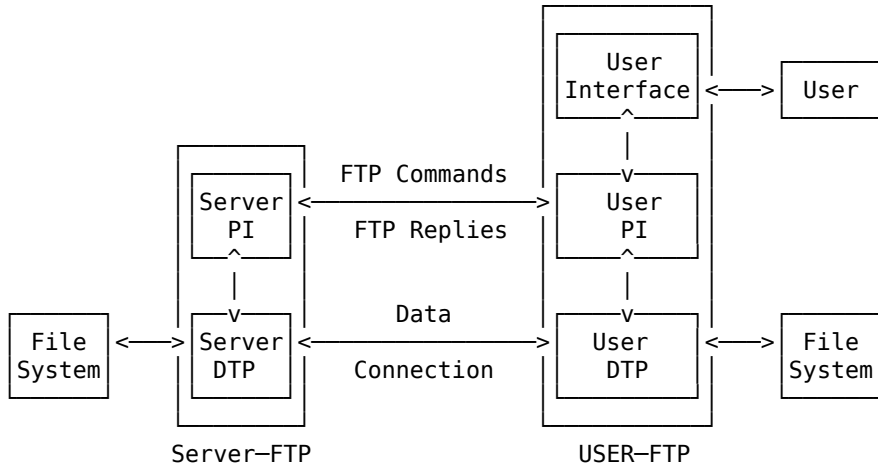
Mail References I

- [1] Wikipedia. *Simple Mail Transfer Protocol* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Simple%5C_Mail%5C_Transfer%5C_Protocol&oldid=646541423.
- [2] Wikipedia. *Post Office Protocol* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Post%5C_Office%5C_Protocol&oldid=645436521.
- [3] Wikipedia. *Internet Message Access Protocol* — *Wikipedia, The Free Encyclopedia*. 2015. http://en.wikipedia.org/w/index.php?title=Internet%5C_Message%5C_Access%5C_Protocol&oldid=647958613.
- [4] Wikipedia. *MIME* — *Wikipedia, The Free Encyclopedia*. 2015. <http://en.wikipedia.org/w/index.php?title=MIME&oldid=644193928>.
- [5] KLENSIN (ED.) J. *Simple Mail Transfer Protocol*. RFC Editor. 2001. <https://www.rfc-editor.org/rfc/rfc2821.txt>.
- [6] MYERS J, ROSE M. *Post Office Protocol - Version 3*. RFC Editor. 1996. <https://www.rfc-editor.org/rfc/rfc1939.txt>.

Mail References II

- [7] CRISPIN M. *INTERNET MESSAGE ACCESS PROTOCOL - VERSION 4rev1*. RFC Editor. 2003. <https://www.rfc-editor.org/rfc/rfc3501.txt>.
- [8] FREED N, BORENSTEIN N. *Multipurpose Internet Mail Extensions (MIME) Part One: Format of Internet Message Bodies*. RFC Editor. 1996. <https://www.rfc-editor.org/rfc/rfc2045.txt>.

FTP



An Active FTP Session

Control session: →


To see FTP data session:

```
cs3:$ nc -l $((100*256+0))
```

Server	Client
cs2 \Rightarrow 202.203.132.242	cs3 \Rightarrow 202.203.132.245

```
wx672@cs3:~$ nc cs2 ftp
220 (vsFTPd 2.0.5)
user wx672
331 Please specify the password.
pass canttellyou
230 Login successful.
port 202,203,132,245,100,0
200 PORT command successful. Consider using PASV.
nlst
150 Here comes the directory listing.
226 Directory send OK.
quit
221 Goodbye.
```

port 202,203,132,245,100,0



Port (2 x 8 bits)

IP (4 x 8 bits)

A Passive FTP Session

Control session: →

To see FTP data session:

cs3:\$ nc cs2 \$((36*256+5))

Server	Client
cs2 ⇒ 202.203.132.242	cs3 ⇒ 202.203.132.245


```
wx672@cs3:~$ nc cs2 ftp
220 (vsFTPd 2.0.5)
user wx672
331 Please specify the password.
pass canttellyou
230 Login successful.
pasv
227 Entering Passive Mode (202,203,132,242,36,5)
list
150 Here comes the directory listing.
quit
221 Goodbye.
```




Active FTP vs. Passive FTP

In **active mode**: Server initiates data connection to client's data port.

In **passive mode**: Client initiates data connection to random port specified by server.

Why Passive Mode?


Active mode doesn't work with firewall



FTP References

- [1] Wikipedia contributors. *File Transfer Protocol — Wikipedia, The Free Encyclopedia*. 2020. https://en.wikipedia.org/w/index.php?title=File_Transfer_Protocol&oldid=993512680.
- [2] POSTEL J, REYNOLDS J. *File Transfer Protocol*. RFC Editor. 1985. <https://www.rfc-editor.org/rfc/rfc959.txt>.
- [3] BELLOVIN S. *Firewall-Friendly FTP*. RFC Editor. 1994. <https://www.rfc-editor.org/rfc/rfc1579.txt>.

BitTorrent



1. How does a peer find other peers that have the content it wants to download?
2. How is content replicated by peers to provide high-speed downloads for everyone?
3. How do peers encourage each other to upload content to others as well as download content for themselves?

P2P References

- [1] Wikipedia. *BitTorrent — Wikipedia, The Free Encyclopedia*. 2015.
<http://en.wikipedia.org/w/index.php?title=BitTorrent&oldid=648034964>.
- [2] COHEN B. *The BitTorrent Protocol Specification, Version 11031*. 2008.
http://www.bittorrent.org/beps/bep%5C_0003.html.