

# Operating Systems

Lecture Handouts

Wang Xiaolin

wx672ster+os@gmail.com

May 18, 2019

## Contents

<b>I Introduction</b>	<b>4</b>
1 What's an Operating System	4
2 OS Services	7
3 Hardware	7
4 Bootstrapping	9
5 Interrupt	10
6 System Calls	11
<b>II Process And Thread</b>	<b>14</b>
7 Processes	15
7.1 What's a Process . . . . .	15
7.2 PCB . . . . .	15
7.3 Process Creation . . . . .	16
7.4 Process State . . . . .	17
7.5 CPU Switch From Process To Process . . . . .	17
8 Threads	18
8.1 Processes vs. Threads . . . . .	18
8.2 Why Thread? . . . . .	18
8.3 Thread Characteristics . . . . .	19
8.4 POSIX Threads . . . . .	20
8.5 User-Level Threads vs. Kernel-level Threads . . . . .	22
8.6 Linux Threads . . . . .	23
9 Process Synchronization	26
9.1 IPC . . . . .	26
9.2 Shared Memory . . . . .	26
9.3 Race Condition and Mutual Exclusion . . . . .	27
9.4 Semaphores . . . . .	31
9.5 Monitors . . . . .	36
9.6 Message Passing . . . . .	37
9.7 Classical IPC Problems . . . . .	38

<b>10 CPU Scheduling</b>	<b>42</b>
10.1 Process Scheduling Queues . . . . .	42
10.2 Scheduling . . . . .	43
10.3 Process Classification . . . . .	44
10.4 Process Behavior . . . . .	44
10.5 Process Schedulers . . . . .	44
10.6 Scheduling In Batch Systems . . . . .	45
10.7 Scheduling In Interactive Systems . . . . .	45
10.8 Thread Scheduling . . . . .	46
10.9 Linux Scheduling . . . . .	46
<b>11 Deadlock</b>	<b>48</b>
11.1 Resources . . . . .	48
11.2 Introduction to Deadlocks . . . . .	49
11.3 Deadlock Modeling . . . . .	50
11.4 Deadlock Detection and Recovery . . . . .	51
11.5 Deadlock Avoidance . . . . .	53
11.6 Deadlock Prevention . . . . .	54
11.7 The Ostrich Algorithm . . . . .	55
<b>III Memory Management</b>	<b>55</b>
<b>12 Background</b>	<b>55</b>
<b>13 Contiguous Memory Allocation</b>	<b>61</b>
<b>14 Virtual Memory</b>	<b>63</b>
14.1 Paging . . . . .	63
14.2 Demand Paging . . . . .	68
14.3 Copy-on-Write . . . . .	69
14.4 Memory mapped files . . . . .	69
14.5 Page Replacement Algorithms . . . . .	69
14.6 Allocation of Frames . . . . .	72
14.7 Thrashing And Working Set Model . . . . .	72
14.8 Other Issues . . . . .	74
14.9 Segmentation . . . . .	76
<b>IV File Systems</b>	<b>82</b>
<b>15 File System Structure</b>	<b>82</b>
<b>16 Files</b>	<b>83</b>
<b>17 Directories</b>	<b>86</b>
<b>18 File System Implementation</b>	<b>88</b>
18.1 Basic Structures . . . . .	88
18.2 Implementing Files . . . . .	88
18.3 Implementing Directories . . . . .	96
18.4 Shared Files . . . . .	98
18.5 Disk Space Management . . . . .	100
<b>19 Ext2 File System</b>	<b>101</b>
19.1 Ext2 File System Layout . . . . .	101
19.2 Ext2 Block groups . . . . .	102
19.3 Ext2 Inode . . . . .	103
19.4 Ext2 Superblock . . . . .	104
19.5 Ext2 Directory . . . . .	104
<b>20 Virtural File Systems</b>	<b>105</b>

<b>V Input/Output</b>	<b>111</b>
<b>21 Principles of I/O Hardware</b>	<b>111</b>
21.1 Programmed I/O . . . . .	115
21.2 Interrupt-Driven I/O . . . . .	116
21.3 Direct Memory Access (DMA) . . . . .	116
<b>22 I/O Software Layers</b>	<b>118</b>
22.1 Interrupt Handlers . . . . .	118
22.2 Device Drivers . . . . .	118
22.3 Device-Independent I/O Software . . . . .	119
22.4 User-Space I/O Software . . . . .	120
<b>23 Disks</b>	<b>120</b>
23.1 Disk Scheduling . . . . .	120
23.2 RAID Structure . . . . .	123

## References

- [1] BRYANT R E, O'HALLARON D R. Computer Systems: A Programmer's Perspective. 2nd ed. USA: Addison-Wesley, 2010.
- [2] TANENBAUM A S. Modern Operating Systems. 3rd ed. Prentice Hall Press, 2007.
- [3] BOVET D, CESATI M. Understanding The Linux Kernel. 3rd ed. O'Reilly, 2005.
- [4] Wikipedia. Process (computing) — Wikipedia, The Free Encyclopedia. 2014. [http://en.wikipedia.org/w/index.php?title=Process\\_\(computing\)&oldid=639847817](http://en.wikipedia.org/w/index.php?title=Process_(computing)&oldid=639847817).
- [5] DOWNEY A B. The Little Book of Semaphores. greenteapress.com, 2008.
- [6] Silberschatz, Galvin, Gagne. Operating System Concepts Essentials. 1st ed. John Wiley & Sons, 2011.
- [7] LOVE R. Linux Kernel Development. Addison-Wesley, 2010.
- [8] Wikipedia. Compiler — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Compiler&oldid=661266598>.
- [9] Wikipedia. Assembly language — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Assembly\\_language&oldid=661185928](http://en.wikipedia.org/w/index.php?title=Assembly_language&oldid=661185928).
- [10] Wikipedia. Linker (computing) — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Linker\\_\(computing\)&oldid=652892136](http://en.wikipedia.org/w/index.php?title=Linker_(computing)&oldid=652892136).
- [11] Wikipedia. Loader (computing) — Wikipedia, The Free Encyclopedia. 2012. [http://en.wikipedia.org/w/index.php?title=Loader\\_\(computing\)&oldid=520743198](http://en.wikipedia.org/w/index.php?title=Loader_(computing)&oldid=520743198).
- [12] Wikipedia. Dynamic linker — Wikipedia, The Free Encyclopedia. 2012. [http://en.wikipedia.org/w/index.php?title=Dynamic\\_linker&oldid=517400345](http://en.wikipedia.org/w/index.php?title=Dynamic_linker&oldid=517400345).
- [13] LEVINE J. Linkers and Loaders. Morgan Kaufmann, 2000.
- [14] Wikipedia. Page replacement algorithm — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Page\\_replacement\\_algorithm&oldid=656936753](http://en.wikipedia.org/w/index.php?title=Page_replacement_algorithm&oldid=656936753).
- [15] Intel. INTEL 80386 Programmer's Reference Manual. <http://pdos.csail.mit.edu/6.828/2004/readings/i386/toc.htm>. 1986.
- [16] Wikipedia. Executable and Linkable Format — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Executable\\_and\\_Linkable\\_Format&oldid=659380509](http://en.wikipedia.org/w/index.php?title=Executable_and_Linkable_Format&oldid=659380509).
- [17] Wikipedia. Open (system call) — Wikipedia, The Free Encyclopedia. 2014. [http://en.wikipedia.org/w/index.php?title=Open\\_\(system\\_call\)&oldid=611838618](http://en.wikipedia.org/w/index.php?title=Open_(system_call)&oldid=611838618).
- [18] Wikipedia. File descriptor — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=File\\_descriptor&oldid=661810398](http://en.wikipedia.org/w/index.php?title=File_descriptor&oldid=661810398).
- [19] STALLINGS W. Operating Systems: Internals and Design Principles. 7th ed. Prentice Hall, 2011.
- [20] Wikipedia. File Allocation Table — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=File\\_Allocation\\_Table&oldid=661104239](http://en.wikipedia.org/w/index.php?title=File_Allocation_Table&oldid=661104239).
- [21] Wikipedia. Inode — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Inode&oldid=647736522>.

- [22] BACH M. The design of the UNIX operating system. Prentice-Hall, 1986.
- [23] THOMPSON K. Unix Implementation. Bell System Technical Journal, 1978, 57: 1931-1946.
- [24] Wikipedia. Directed acyclic graph — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Directed\\_acyclic\\_graph&oldid=654052259](http://en.wikipedia.org/w/index.php?title=Directed_acyclic_graph&oldid=654052259).
- [25] POIRIER D. The Second Extended File System Internal Layout. Web, 2011.
- [26] RUSLING D A. The Linux Kernel. Linux Documentation Project, 1999.
- [27] CARD R, TS'O T, TWEEDIE S. Design and Implementation of the Second Extended Filesystem. Dutch International Symposium on Linux, 1996.
- [28] MORGAN D. Analyzing a filesystem. [http://homepage.smc.edu/morgan\\_david/cs40/analyze-ext2.htm](http://homepage.smc.edu/morgan_david/cs40/analyze-ext2.htm). 2012.
- [29] MAUERER W. Professional Linux Kernel Architecture. John Wiley & Sons, 2008.

#### **Course Web Links**

**Course web site:** <https://cs6.swfu.edu.cn/moodle>

**Lecture slides:** [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/os/slides/](https://cs2.swfu.edu.cn/~wx672/lecture_notes/os/slides/)

**Source code:** [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/os/src/](https://cs2.swfu.edu.cn/~wx672/lecture_notes/os/src/)

**Lab instructions** [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/os/lab.html](https://cs2.swfu.edu.cn/~wx672/lecture_notes/os/lab.html)

**Sample lab report:** [https://cs2.swfu.edu.cn/~wx672/lecture\\_notes/os/sample-report/](https://cs2.swfu.edu.cn/~wx672/lecture_notes/os/sample-report/)

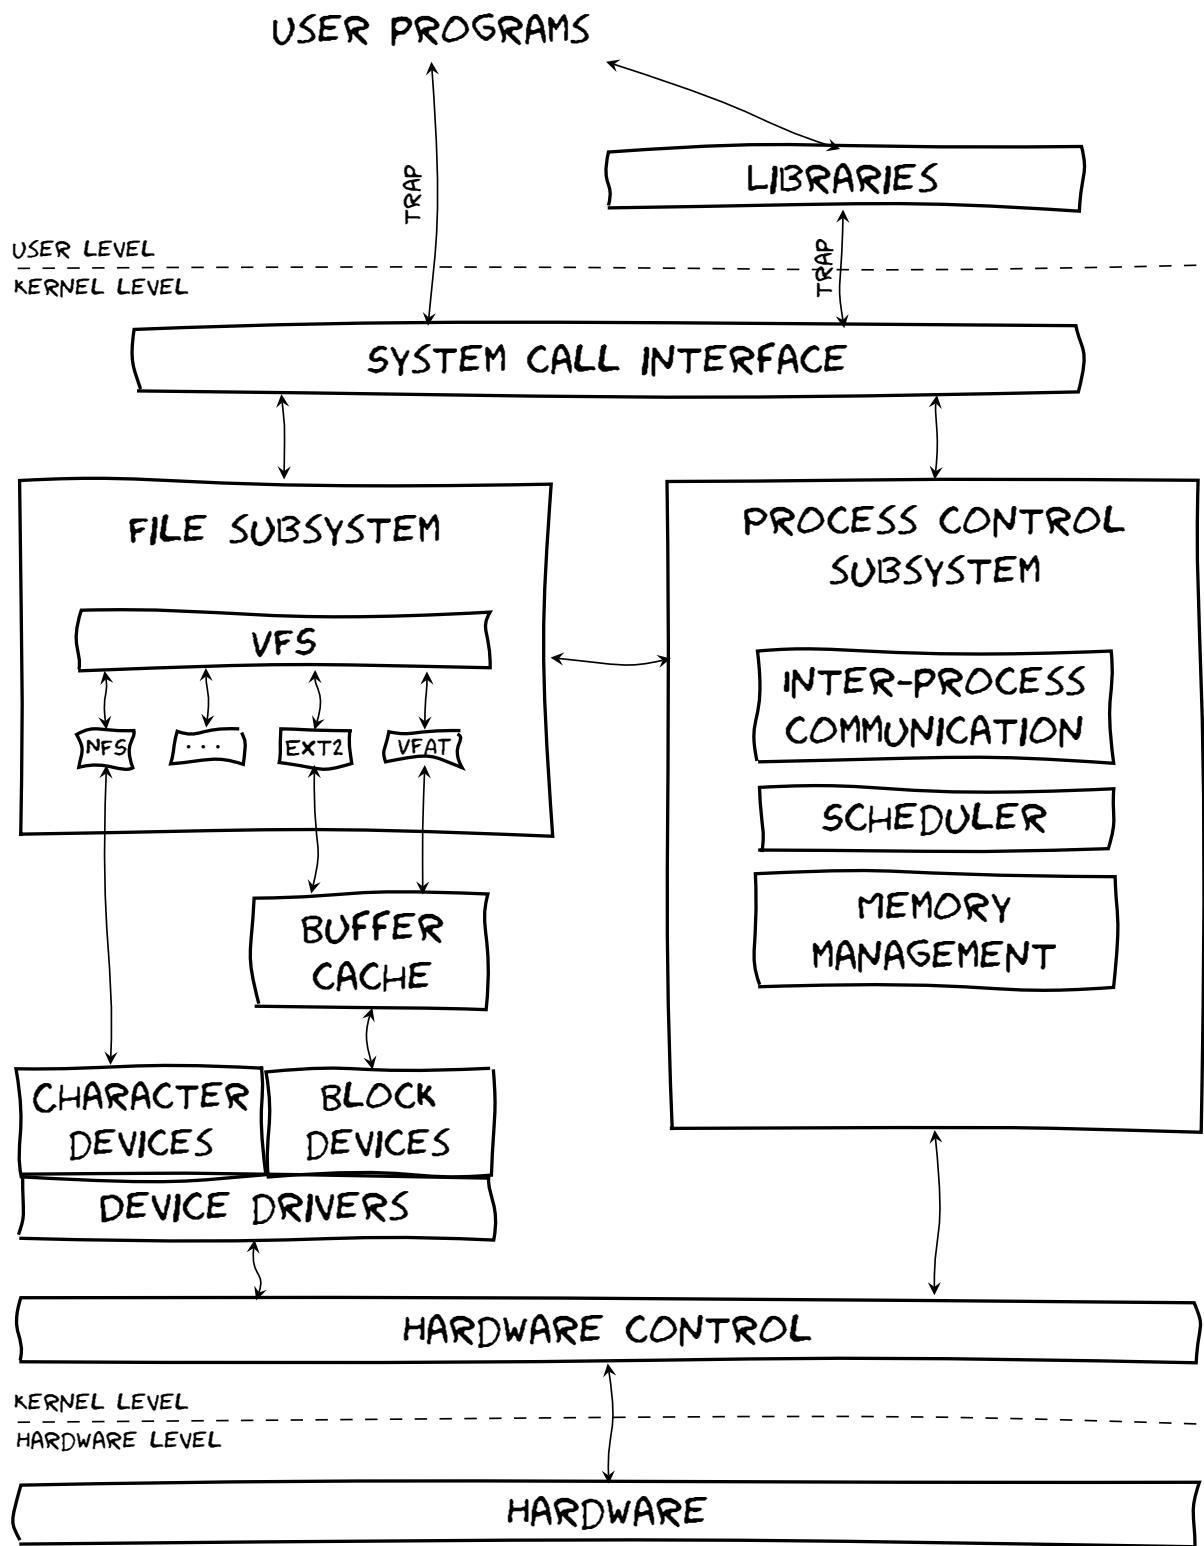
## **Part I**

# **Introduction**

## **1 What's an Operating System**

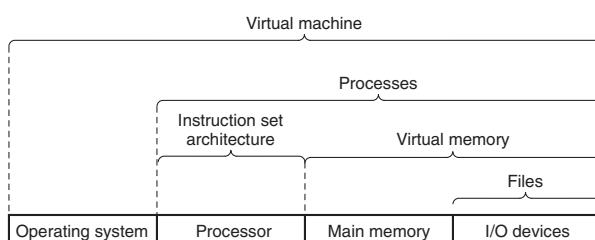
### **What's an Operating System?**

- “*Everything a vendor ships when you order an operating system*”
- It’s the program that runs all the time
- It’s a *resource manager*
  - Each program gets time with the resource
  - Each program gets space on the resource
- It’s a *control program*
  - Controls execution of programs to prevent errors and improper use of the computer
- No universally accepted definition



### Abstractions

To hide the complexity of the actual implementations



See also: [Computer Systems: A Programmer's Perspective, Sec. 1.9.2, The Importance of Abstractions in Computer Systems].

## System Goals

### Convenient vs. Efficient

- Convenient for the user — for PCs
- Efficient — for mainframes, multiusers
- UNIX
  - Started with keyboard + printer, none paid to convenience
  - Now, still concentrating on efficiency, with GUI support

## History of Operating Systems

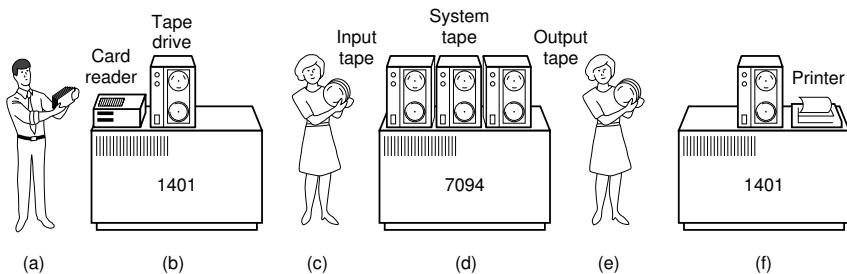


Fig. 1-2. An early batch system. (a) Programmers bring cards to 1401. (b) 1401 reads batch of jobs onto tape. (c) Operator carries input tape to 7094. (d) 7094 does computing. (e) Operator carries output tape to 1401. (f) 1401 prints output.

### 1945 - 1955 First generation

- vacuum tubes, plug boards

### 1955 - 1965 Second generation

- transistors, batch systems

### 1965 - 1980 Third generation

- ICs and multiprogramming

### 1980 - present Fourth generation

- personal computers

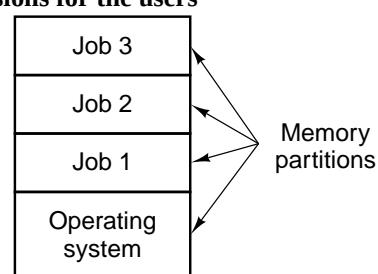
## Multi-programming is the first instance where the OS must make decisions for the users

**Job scheduling** — decides which job should be loaded into the memory.

**Memory management** — because several programs in memory at the same time

**CPU scheduling** — choose one job among all the jobs are ready to run

**Process management** — make sure processes don't offend each other



## The Operating System Zoo

- Mainframe OS
- Server OS
- Multiprocessor OS
- Personal computer OS
- Real-time OS
- Embedded OS
- Smart card OS

## 2 OS Services

### OS Services

*Like a government*

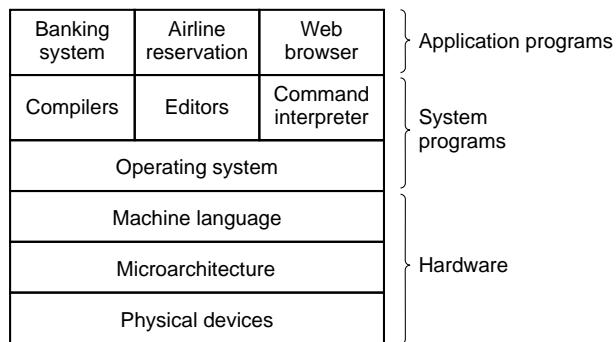
#### Helping the users:

- User interface
- Program execution
- I/O operation
- File system manipulation
- Communication
- Error detection

#### Keeping the system efficient:

- Resource allocation
- Accounting
- Protection and security

## A Computer System



## 3 Hardware

### CPU Working Cycle



1. Fetch the first instruction from memory
2. Decode it to determine its type and operands
3. execute it

### Special CPU Registers

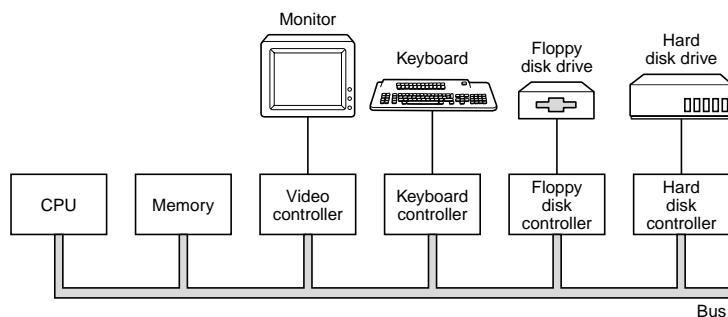
**Program counter (PC):** keeps the memory address of the next instruction to be fetched

**Stack pointer (SP):** → the top of the current stack in memory

**Program status (PS):** holds

- condition code bits
- processor state

### System Bus



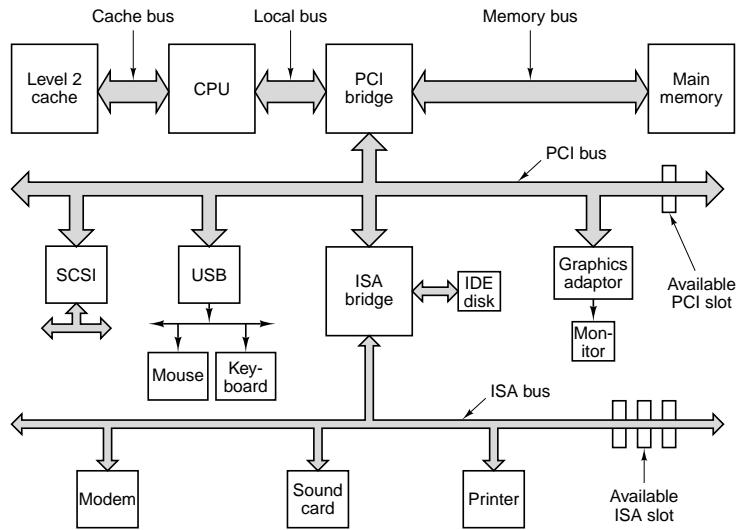
**Address Bus:** specifies the memory locations (addresses) for the data transfers

**Data Bus:** holds the data transferred. Bidirectional

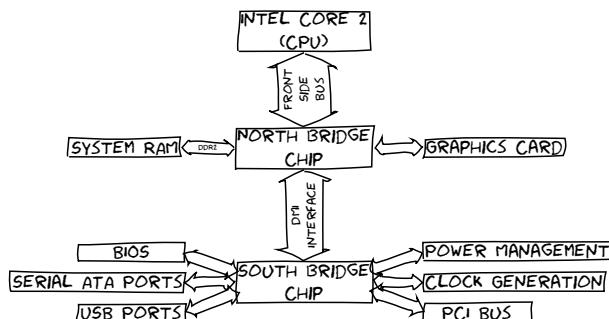
**Control Bus:** contains various lines used to route timing and control signals throughout the system

## Controllers and Peripherals

- Peripherals are real devices controlled by controller chips
- Controllers are processors like the CPU itself, have control registers
- Device driver writes to the registers, thus control it
- Controllers are connected to the CPU and to each other by a variety of buses



## Motherboard Chipsets



See also: *Motherboard Chipsets And The Memory Map*<sup>1</sup>.

- The CPU doesn't know what it's connected to
  - CPU test bench? network router? toaster? brain implant?
- The CPU talks to the outside world through its pins
  - some pins to transmit the physical memory address
  - other pins to transmit the values
- The CPU's gateway to the world is the *front-side bus*

## Intel Core 2 QX6600

- 33 pins to transmit the physical memory address
  - so there are  $2^{33}$  choices of memory locations
- 64 pins to send or receive data
  - so data path is 64-bit wide, or 8-byte chunks

This allows the CPU to physically address 64GB of memory ( $2^{33} \times 8B$ )

<sup>1</sup><http://duartes.org/gustavo/blog/post/motherboard-chipsets-memory-map>

See also: *Datasheet for Intel Core 2 Quad-Core Q6000 Sequence*<sup>2</sup>.

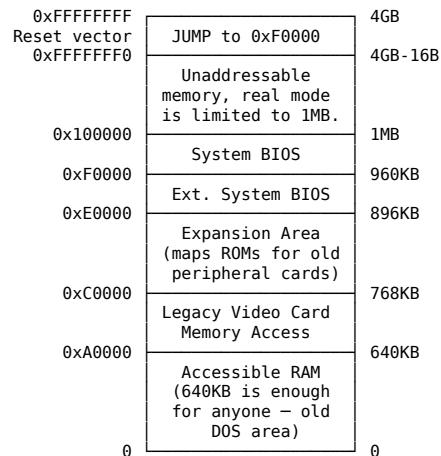
### Some physical memory addresses are mapped away!

- only the addresses, not the spaces
- Memory holes
  - $640KB \sim 1MB$
  - `/proc/iomem`

### Memory-mapped I/O

- BIOS ROM
- video cards
- PCI cards
- ...

This is why 32-bit OSes have problems using 4 gigs of RAM.



### the northbridge

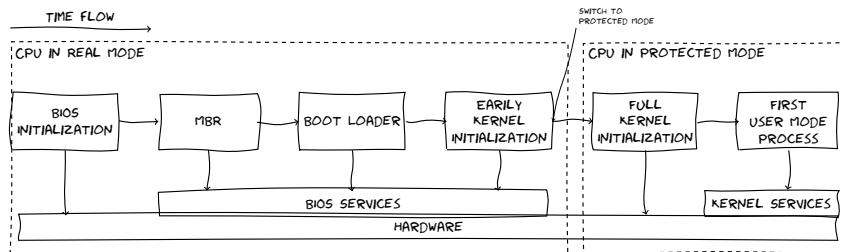
1. receives a physical memory request
2. decides where to route it
  - to RAM? to video card? to ...?
  - decision made via the *memory address map*
- When is the memory address map built? `setup()`.

## 4 Bootstrapping

### Bootstrapping

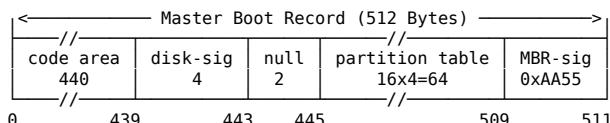
#### Can you pull yourself up by your own bootstraps?

A computer cannot run without first loading software but must be running before any software can be loaded.



### Intel x86 Bootstrapping

1. BIOS (0xfffffff0)
  - ➡ POST ➡ HW init ➡ Find a boot device (FD,CD,HD...) ➡ Copy sector zero (MBR) to RAM (0x00007c00)
2. MBR – the first 512 bytes, contains
  - Small code (< 446 Bytes), e.g. GRUB stage 1, for loading GRUB stage 2
  - the primary partition table ( $16 \times 4 = 64$  Bytes)
  - its job is to load the second-stage boot loader.
3. GRUB stage 2 — load the OS kernel into RAM
4. ⚙ startup
5. init — the first user-space program



\$ sudo dd -n512 /dev/sda

<sup>2</sup><http://download.intel.com/design/processor/datasheets/31559205.pdf>

## 5 Interrupt

Why Interrupt?

While a process is reading a disk file, can we do...

```
1 | while(!done_reading_a_file())
2 | {
3 |     let_CPU_wait();
4 |     // or...
5 |     lend_CPU_to_others();
6 | }
7 | operate_on_the_file();
```

Modern OS are Interrupt Driven

HW INT by sending a signal to CPU

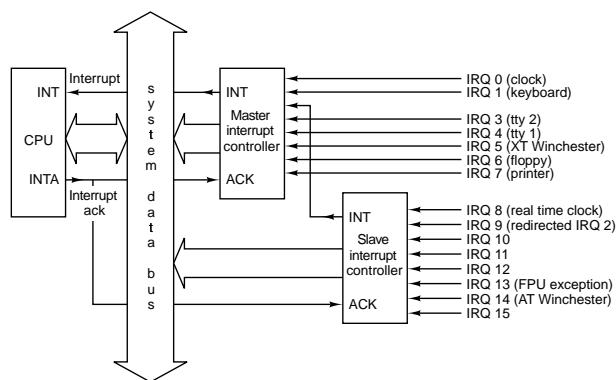
SW INT by executing a *system call*

Trap (exception) is a software-generated INT caused by an error or by a specific request from an user program

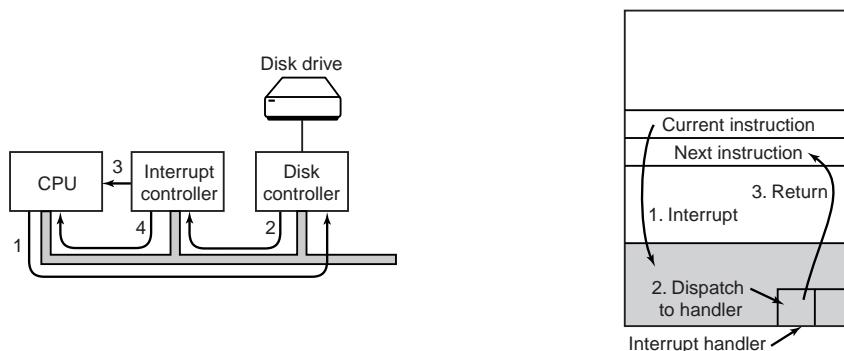
Interrupt vector is an array of pointers → the memory addresses of *interrupt handlers*. This array is indexed by a unique device number

```
$ less /proc/devices
$ less /proc/interrupts
```

Programmable Interrupt Controllers

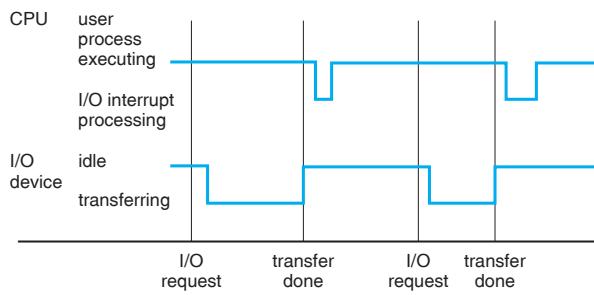


Interrupt Processing



Detailed explanation: in [Modern Operating Systems, Sec. 1.3.5, I/O Devices].

Interrupt Timeline

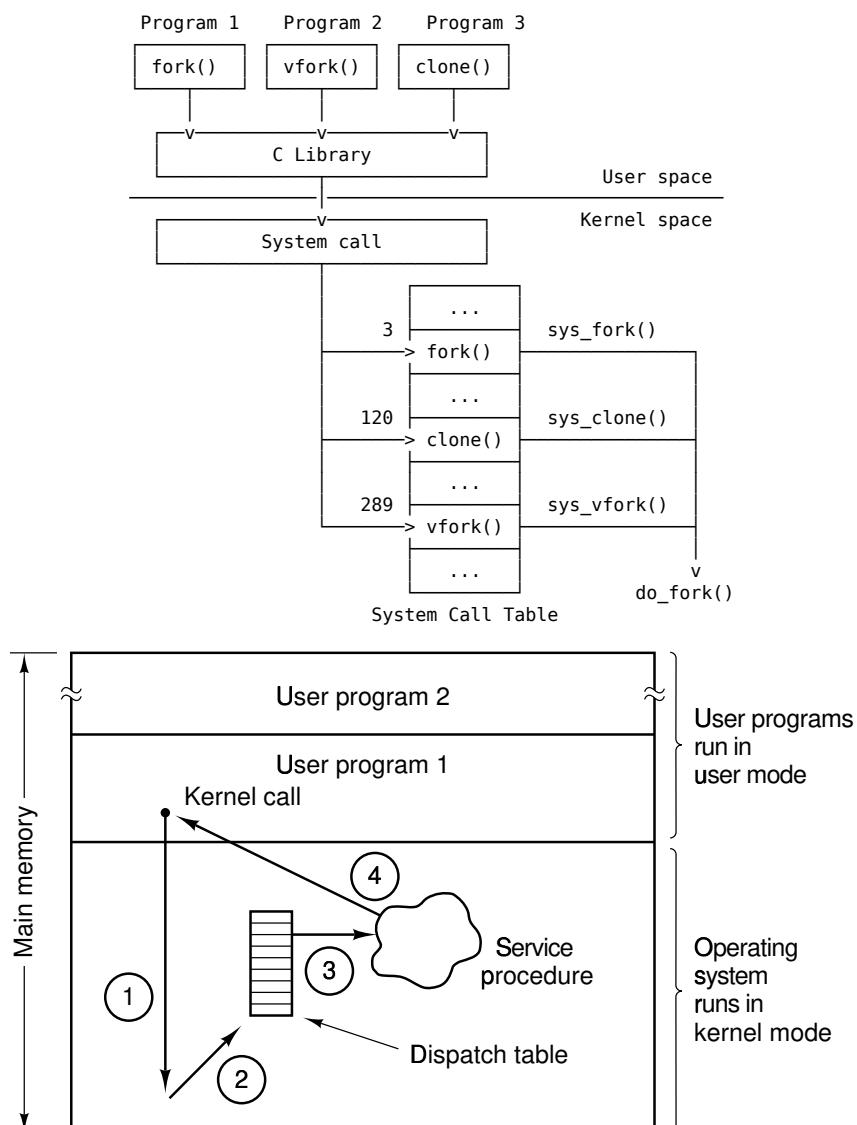


## 6 System Calls

### System Calls

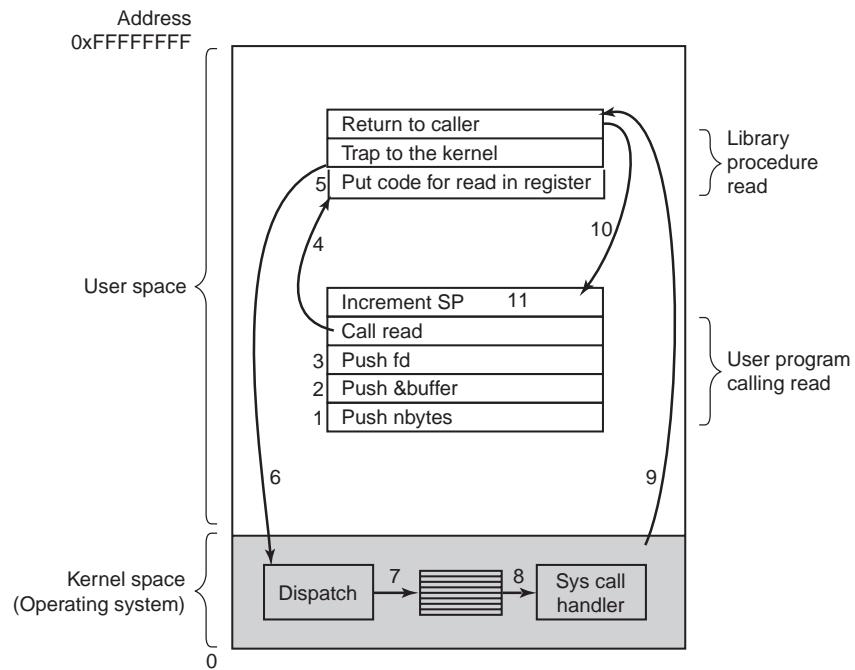
#### A System Call

- is how a program requests a service from an OS kernel
- provides the interface between a process and the OS



**Figure 1-16.** How a system call can be made: (1) User program traps to the kernel. (2) Operating system determines service number required. (3) Operating system calls service procedure. (4) Control is returned to user program.

The 11 steps in making the system call `read(fd, buffer, nbytes)`



### Example

*Linux INT80h*

**Interrupt Vector Table:** The very first 1KiB of x86 memory.

- 256 entries × 4B = 1KiB
- Each entry is a complete memory address (segment:offset)
- It's populated by Linux and BIOS
- Slot 80h: address of the kernel services dispatcher (☞ sys-call table)

### Example

Msg: db "Hello, world"

MsgLen: equ \$-Msg

```

mov eax,4      ; Specify sys_write syscall
mov ebx,1      ; Specify File Descriptor 1 (STDOUT)
mov ecx,Msg    ; Pass offset of the message
mov edx,MsgLen ; Pass the length of the message
int 80H        ; Make syscall to output the text to STDOUT

```

Process management	
Call	Description
pid = fork()	Create a child process identical to the parent
pid = waitpid(pid, &statloc, options)	Wait for a child to terminate
s = execve(name, argv, environp)	Replace a process' core image
exit(status)	Terminate process execution and return status

File management	
Call	Description
fd = open(file, how, ...)	Open a file for reading, writing or both
s = close(fd)	Close an open file
n = read(fd, buffer, nbytes)	Read data from a file into a buffer
n = write(fd, buffer, nbytes)	Write data from a buffer into a file
position = lseek(fd, offset, whence)	Move the file pointer
s = stat(name, &buf)	Get a file's status information

Directory and file system management	
Call	Description
s = mkdir(name, mode)	Create a new directory
s = rmdir(name)	Remove an empty directory
s = link(name1, name2)	Create a new entry, name2, pointing to name1
s = unlink(name)	Remove a directory entry
s = mount(special, name, flag)	Mount a file system
s = umount(special)	Unmount a file system

Miscellaneous	
Call	Description
s = chdir(dirname)	Change the working directory
s = chmod(name, mode)	Change a file's protection bits
s = kill(pid, signal)	Send a signal to a process
seconds = time(&seconds)	Get the elapsed time since Jan. 1, 1970

Fig. 1-18. Some of the major POSIX system calls. The return code *s* is  $-1$  if an error has occurred. The return codes are as follows: *pid* is a process id, *fd* is a file descriptor, *n* is a byte count, *position* is an offset within the file, and *seconds* is the elapsed time. The parameters are explained in the text.

## System Call Examples

`fork()`

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main ()
5 {
6     printf ("Hello World!\n");
7     fork();
8     printf ("Goodbye Cruel World!\n");
9     return 0;
10 }
```

\$ man 2 fork

`exec()`

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main()
5 {
6     printf("Hello World!\n");
7     if( fork() != 0 )
8         printf("I am the parent process.\n");
9     else {
10        printf("A child is listing the directory contents... \n");
11        execl("/bin/ls", "ls", "-al", NULL);
12    }
13    return 0;
14 }
```

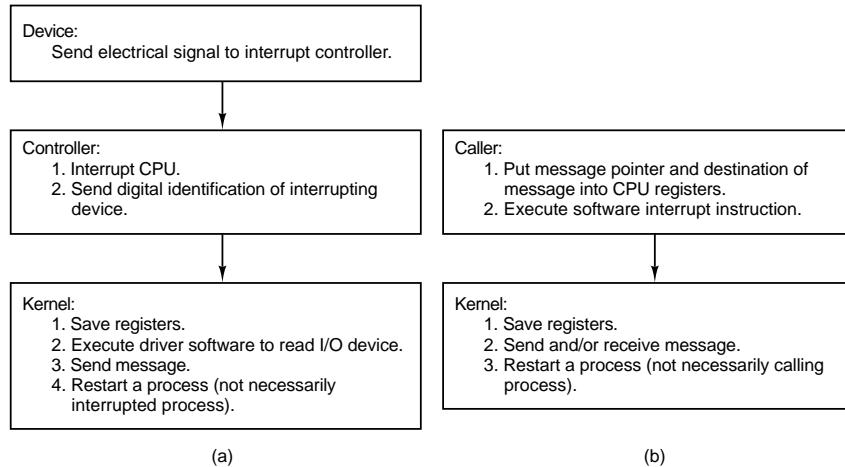
\$ man 3 exec

Quoted from stackoverflow: What is the difference between the functions of the *exec* family of system calls:

There is no *exec* system call — this is usually used to refer to all the *execXX* calls as a group. They all do essentially the same thing: loading a new program into the current process, and provide it with arguments and environment variables. The differences are in how the program is found, how the arguments are specified, and where the environment comes from.

- The calls with *v* in the name take an array parameter to specify the *argv[]* array (*vector*) of the new program.
- The calls with *l* in the name take the arguments of the new program as a variable-length argument *list* to the function itself.
- The calls with *e* in the name take an extra argument to provide the *environment* of the new program; otherwise, the program inherits the current process's environment.
- The calls with *p* in the name search the *PATH* environment variable to find the program if it doesn't have a directory in it (i.e. it doesn't contain a / character). Otherwise, the program name is always treated as a path to the executable.

## Hardware INT vs. Software INT



## References

- [1] Wikipedia. Interrupt — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Interrupt&oldid=646521061>.
- [2] Wikipedia. System call — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=System\\_call&oldid=647910319](http://en.wikipedia.org/w/index.php?title=System_call&oldid=647910319).

# Part II

# Process And Thread

## 7 Processes

### 7.1 What's a Process

#### Process

A **process** is an instance of a program in execution

#### Processes are like human beings:

- » they are generated
- » they have a life
- » they optionally generate one or more child processes, and
- » eventually they die

A small difference:

- sex is not really common among processes
- each process has just one parent

The term "process" is often used with several different meanings. In this book, we stick to the usual OS textbook definition: a process is an instance of a program in execution. You might think of it as *the collection of data structures that fully describes how far the execution of the program has progressed.* [Understanding The Linux Kernel, Sec. 3.1, Processes, Lightweight Processes, and Threads]

Processes are like human beings: they are generated, they have a more or less significant life, they optionally generate one or more child processes, and eventually they die. A small difference is that sex is not really common among processes each process has just one parent.

From the kernel's point of view, the purpose of a process is to act as an entity to which system resources (CPU time, memory, etc.) are allocated.

In general, a computer system process consists of (or is said to 'own') the following resources [Process (computing) — Wikipedia, The Free Encyclopedia]:

- An image of the executable machine code associated with a program.
- Memory (typically some region of virtual memory); which includes the executable code, process-specific data (input and output), a call stack (to keep track of active subroutines and/or other events), and a heap to hold intermediate computation data generated during run time.
- Operating system descriptors of resources that are allocated to the process, such as file descriptors (Unix terminology) or handles (Windows), and data sources and sinks.
- Security attributes, such as the process owner and the process' set of permissions (allowable operations).
- Processor state (context), such as the content of registers, physical memory addressing, etc. The state is typically stored in computer registers when the process is executing, and in memory otherwise.

The operating system holds most of this information about active processes in data structures called process control blocks.

Any subset of resource, but typically at least the processor state, may be associated with each of the process' threads in operating systems that support threads or 'daughter' processes.

The operating system keeps its processes separated and allocates the resources they need, so that they are less likely to interfere with each other and cause system failures (e.g., deadlock or thrashing). The operating system may also provide mechanisms for inter-process communication to enable processes to interact in safe and predictable ways.

### 7.2 PCB

#### Process Control Block (PCB)

##### Implementation

A process is *the collection of data structures* that fully describes how far the execution of the program has progressed.

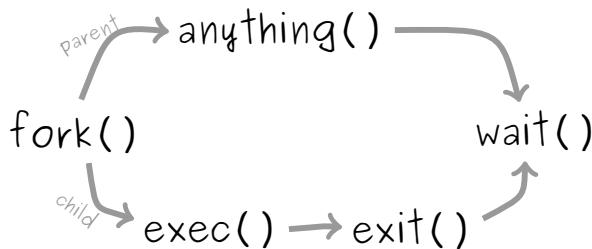
- Each process is represented by a *PCB*
- `task_struct` in

process state
PID
program counter
registers
memory limits
list of open files
...

To manage processes, the kernel must have a clear picture of what each process is doing. It must know, for instance, the process's priority, whether it is running on a CPU or blocked on an event, what address space has been assigned to it, which files it is allowed to address, and so on. This is the role of the process descriptor a `task_struct` type structure whose fields contain all the information related to a single process. As the repository of so much information, the process descriptor is rather complex. In addition to a large number of fields containing process attributes, the process descriptor contains several pointers to other data structures that, in turn, contain pointers to other structures[*Understanding The Linux Kernel*, Sec. 3.2, *Process Descriptor*].

## 7.3 Process Creation

### Process Creation



- When a process is created, it is almost identical to its parent
  - It receives a (logical) copy of the parent's address space, and
  - executes the same code as the parent
- The parent and child have separate copies of the data (stack and heap)

When a process is created, it is almost identical to its parent. It receives a (logical) copy of the parent's address space and executes the same code as the parent, beginning at the next instruction following the process creation system call. Although the parent and child may share the pages containing the program code (text), they have separate copies of the data (stack and heap), so that changes by the child to a memory location are invisible to the parent (and vice versa)[*Understanding The Linux Kernel*, Sec. 3.1, *Processes, Lightweight Processes, and Threads*].

While earlier Unix kernels employed this simple model, modern Unix systems do not. They support *multi-threaded applications* user programs having many relatively independent execution flows sharing a large portion of the application data structures. In such systems, a process is composed of several *user threads* (or simply *threads*), each of which represents an execution flow of the process. Nowadays, most multi-threaded applications are written using standard sets of library functions called *pthread (POSIX thread) libraries*.

Traditional Unix systems treat all processes in the same way: resources owned by the parent process are duplicated in the child process. This approach makes process creation very slow and inefficient, because it requires copying the entire address space of the parent process. The child process rarely needs to read or modify all the resources inherited from the parent; in many cases, it issues an immediate `execve()` and wipes out the address space that was so carefully copied[*Understanding The Linux Kernel*, Sec. 3.4, *Creating Processes*].

Modern Unix kernels solve this problem by introducing three different mechanisms:

- Copy On Write
- Lightweight processes
- The `vfork()` system call

### Forking in C

```

1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main ()
5 {
6     printf("Hello World!\n");
7     fork();
8     printf("Goodbye Cruel World!\n");
9     return 0;
10 }
  
```

\$ man fork

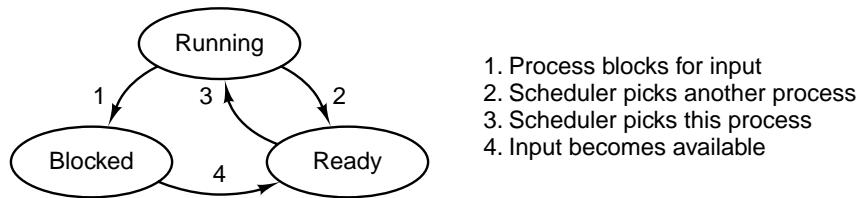
```
exec()
```

```
1 int main()
2 {
3     pid_t pid;
4     /* fork another process */
5     pid = fork();
6     if (pid < 0) { /* error occurred */
7         fprintf(stderr, "Fork Failed");
8         exit(-1);
9     }
10    else if (pid == 0) { /* child process */
11        execlp("/bin/ls", "ls", NULL);
12    }
13    else { /* parent process */
14        /* wait for the child to complete */
15        wait(NULL);
16        printf ("Child Complete");
17        exit(0);
18    }
19    return 0;
20 }
```

```
$ man 3 exec
```

## 7.4 Process State

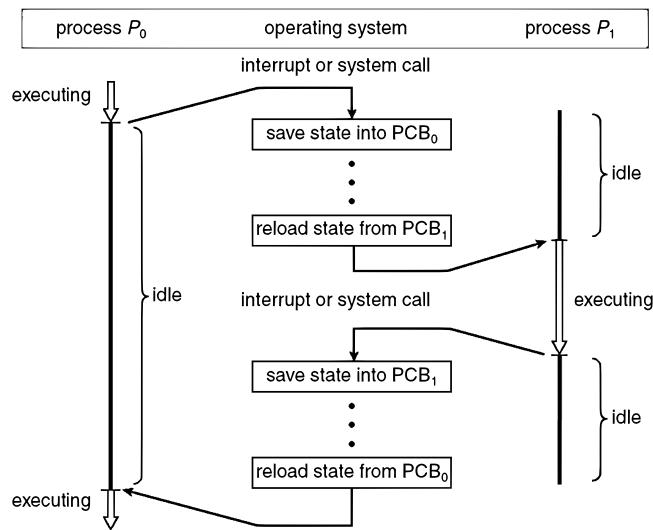
### Process State Transition



See also: [Understanding The Linux Kernel, Sec. 3.2.1, Process State].

## 7.5 CPU Switch From Process To Process

### CPU Switch From Process To Process



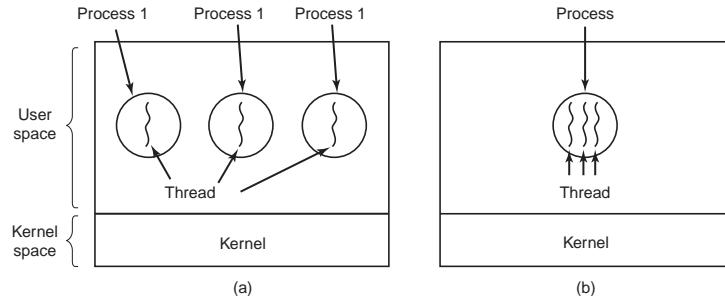
See also: [Understanding The Linux Kernel, Sec. 3.3, Process Switch].

## 8 Threads

### 8.1 Processes vs. Threads

#### Process vs. Thread

a single-threaded process = resource + execution  
a multi-threaded process = resource + executions



**A process** = a unit of resource ownership, used to group resources together;

**A thread** = a unit of scheduling, scheduled for execution on the CPU.

#### Process vs. Thread

**multiple threads running in one process:**  
share an address space and other resources

**multiple processes running in one computer:**  
share physical memory, disk, printers ...

#### No protection between threads

**impossible** — because process is the minimum unit of resource management

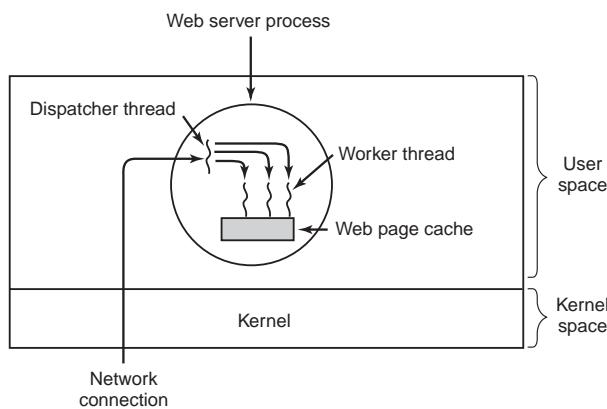
**unnecessary** — a process is owned by a single user

#### Threads

code, data, open files, signals...		
thread ID	thread ID	thread ID
program counter	program counter	program counter
register set	register set	register set
stack	stack	stack

### 8.2 Why Thread?

#### A Multi-threaded Web Server



```

while (TRUE) {
    get_next_request(&buf);
    handoff_work(&buf);
}

```

(a)

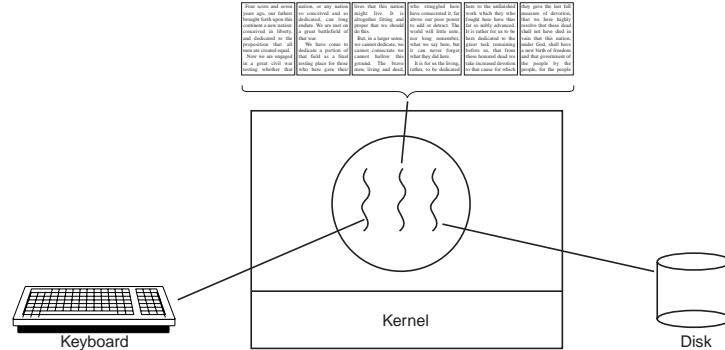
```

while (TRUE) {
    wait_for_work(&buf)
    look_for_page_in_cache(&buf, &page);
    if (page_not_in_cache(&page))
        read_page_from_disk(&buf, &page);
    return_page(&page);
}

```

(b)

### A Word Processor With 3 Threads



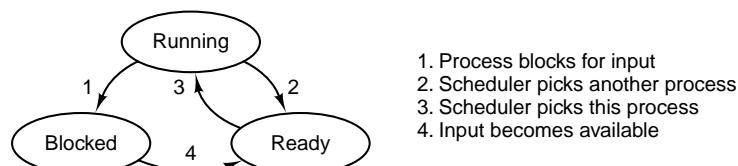
### Why Having a Kind of Process Within a Process?

- Responsiveness
  - Good for interactive applications.
  - A process with multiple threads makes a great server (e.g. a web server):
    - Have one server process, many "worker" threads – if one thread blocks (e.g. on a read), others can still continue executing
- Economy – Threads are cheap!
  - Cheap to create – only need a stack and storage for registers
  - Use very little resources – don't need new address space, global data, program code, or OS resources
    - switches are fast – only have to save/restore PC, SP, and registers
- Resource sharing – Threads can pass data via shared memory; no need for IPC
- Can take advantage of multiprocessors

## 8.3 Thread Characteristics

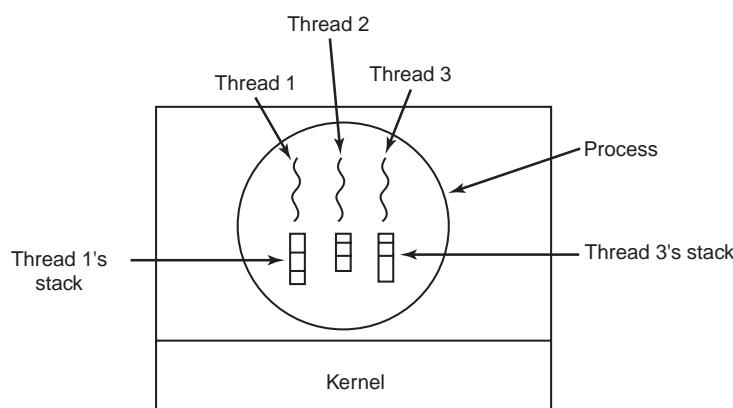
### Thread States Transition

*Same as process states transition*

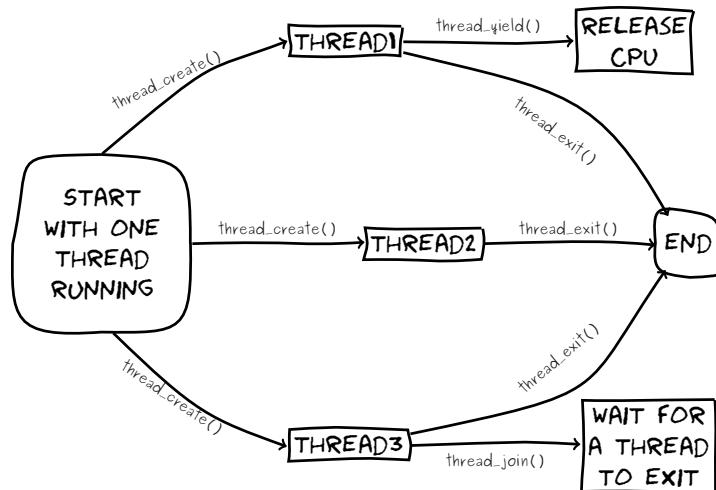


### Each Thread Has Its Own Stack

- A typical stack stores local data and call information for (usually nested) procedure calls.
- Each thread generally has a different execution history.



## Thread Operations



## 8.4 POSIX Threads

### POSIX Threads

**IEEE 1003.1c** The standard for writing portable threaded programs. The threads package it defines is called *Pthreads*, including over 60 function calls, supported by most UNIX systems.

#### Some of the Pthreads function calls

Thread call	Description
<code>pthread_create</code>	Create a new thread
<code>pthread_exit</code>	Terminate the calling thread
<code>pthread_join</code>	Wait for a specific thread to exit
<code>pthread_yield</code>	Release the CPU to let another thread run
<code>pthread_attr_init</code>	Create and initialize a thread's attribute structure
<code>pthread_attr_destroy</code>	Remove a thread's attribute structure

### Pthreads

*Example 1*

```

1 #include <pthread.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4 #include <stdio.h>
5
6 void *thread_function(void *arg){
7     int i;
8     for( i=0; i<20; i++ ){
9         printf("Thread says hi!\n");
10        sleep(1);
11    }
12    return NULL;
13 }
14
15 int main(void){
16     pthread_t mythread;
17     if(pthread_create(&mythread, NULL, thread_function, NULL)){
18         printf("error creating thread.");
19         abort();
20     }
21
22     if(pthread_join(mythread, NULL)){
23         printf("error joining thread.");
24         abort();
25     }
26     exit(0);
27 }
```

See also:

- IBM Developworks: POSIX threads explained<sup>3</sup>.
- stackoverflow.com: What is the difference between `exit()` and `abort()`?<sup>4</sup>.

## Pthreads

`pthread_t` defined in `pthread.h`, is often called a "thread id" (`tid`);  
`pthread_create()` returns zero on success and a non-zero value on failure;  
`pthread_join()` returns zero on success and a non-zero value on failure;

### How to use pthread?

```

• #include<pthread.h>
$ gcc thread1.c -o thread1 -pthread
$ ./thread1
```

## Pthreads

### Example 2

---

<sup>3</sup><http://www.ibm.com/developerworks/linux/library/l-posix1/index.html>

<sup>4</sup><http://stackoverflow.com/questions/397075/what-is-the-difference-between-exit-and-abort>

```

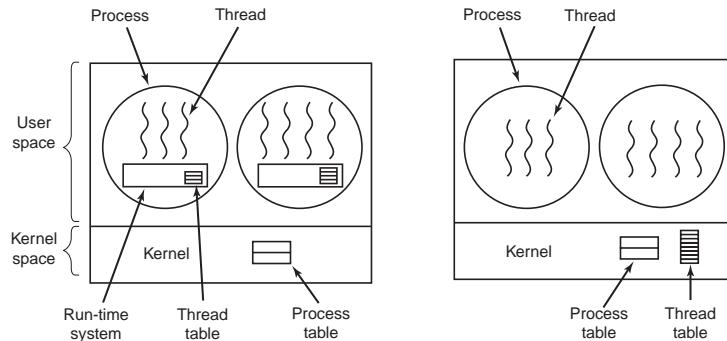
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 #define NUMBER_OF_THREADS 10
6
7 void *print_hello_world(void *tid)
8 {
9     /* prints the thread's identifier, then exits.*/
10    printf ("Thread %d: Hello World!\n", tid);
11    pthread_exit(NULL);
12 }
13
14 int main(int argc, char *argv[])
15 {
16     pthread_t threads[NUMBER_OF_THREADS];
17     int status, i;
18     for (i=0; i<NUMBER_OF_THREADS; i++)
19     {
20         printf ("Main: creating thread %d\n",i);
21         status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);
22
23         if(status != 0){
24             printf ("Oops. pthread_create returned error code %d\n",status);
25             exit(-1);
26         }
27     }
28     exit(NULL);
29 }
```

## Pthreads

With or without `pthread_join()`? Check it by yourself.

## 8.5 User-Level Threads vs. Kernel-level Threads

### User-Level Threads vs. Kernel-Level Threads



### User-Level Threads

**User-level threads** provide a library of functions to allow user processes to create and manage their own threads.

- ⌚ No need to modify the OS
- ⌚ Simple representation
  - each thread is represented simply by a PC, regs, stack, and a small TCB, all stored in the user process' address space
- ⌚ Simple Management
  - creating a new thread, switching between threads, and synchronization between threads can all be done without intervention of the kernel
- ⌚ Fast

- thread switching is not much more expensive than a procedure call
- ⌚ Flexible
  - CPU scheduling (among threads) can be customized to suit the needs of the algorithm – each process can use a different thread scheduling algorithm

### User-Level Threads

- ⌚ Lack of coordination between threads and OS kernel
    - Process as a whole gets one time slice
    - Same time slice, whether process has 1 thread or 1000 threads
    - Also – up to each thread to relinquish control to other threads in that process
  - ⌚ Requires non-blocking system calls (i.e. a multithreaded kernel)
    - Otherwise, entire process will block in the kernel, even if there are runnable threads left in the process
    - part of motivation for user-level threads was not to have to modify the OS
  - ⌚ If one thread causes a page fault(interrupt!), the entire process blocks
- See also: *More about blocking and non-blocking calls*<sup>5</sup>.

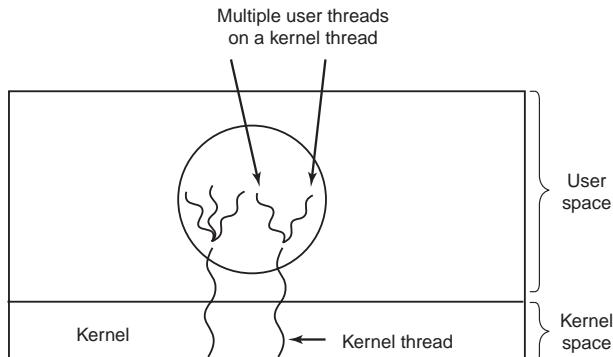
### Kernel-Level Threads

**Kernel-level threads** kernel provides system calls to create and manage threads

- ⌚ Kernel has full knowledge of all threads
  - Scheduler may choose to give a process with 10 threads more time than process with only 1 thread
- ⌚ Good for applications that frequently block (e.g. server processes with frequent interprocess communication)
- ⌚ Slow – thread operations are 100s of times slower than for user-level threads
- ⌚ Significant overhead and increased kernel complexity – kernel must manage and schedule threads as well as processes
  - Requires a full thread control block (TCB) for each thread

### Hybrid Implementations

*Combine the advantages of two*



### Programming Complications

- `fork()`: shall the child has the threads that its parent has?
- What happens if one thread closes a file while another is still reading from it?
- What happens if several threads notice that there is too little memory?

*And sometimes, threads fix the symptom, but not the problem.*

## 8.6 Linux Threads

### Linux Threads

#### To the Linux kernel, there is no concept of a thread

- Linux implements all threads as standard processes
- To Linux, a thread is merely a process that shares certain resources with other processes
- Some OS (MS Windows, Sun Solaris) have cheap threads and expensive processes.
- Linux processes are already quite lightweight

<sup>5</sup><http://www.daniweb.com/software-development/computer-science/threads/384575/synchronous-vs-asynchronous-blocking-vs-non-blocking>

On a 75MHz Pentium      thread:  $1.7\mu s$   
                                 fork:  $1.8\mu s$

[*Understanding The Linux Kernel*, Sec. 3.1, *Processes, Lightweight Processes, and Threads*] Older versions of the Linux kernel offered no support for multithreaded applications. From the kernel point of view, a multithreaded application was just a normal process. The multiple execution flows of a multithreaded application were created, handled, and scheduled entirely in User Mode, usually by means of a POSIX-compliant *pthread* library.

However, such an implementation of multithreaded applications is not very satisfactory. For instance, suppose a chess program uses two threads: one of them controls the graphical chessboard, waiting for the moves of the human player and showing the moves of the computer, while the other thread ponders the next move of the game. While the first thread waits for the human move, the second thread should run continuously, thus exploiting the thinking time of the human player. However, if the chess program is just a single process, the first thread cannot simply issue a blocking system call waiting for a user action; otherwise, the second thread is blocked as well. Instead, the first thread must employ sophisticated nonblocking techniques to ensure that the process remains runnable.

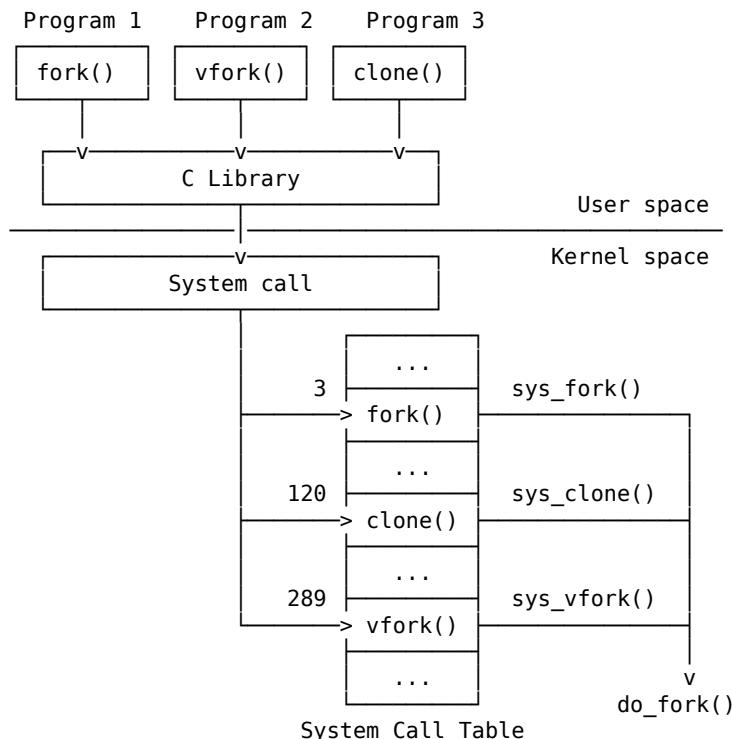
Linux uses *lightweight processes* to offer better support for multithreaded applications. Basically, two lightweight processes may share some resources, like the address space, the open files, and so on. Whenever one of them modifies a shared resource, the other immediately sees the change. Of course, the two processes must synchronize themselves when accessing the shared resource.

A straightforward way to implement multithreaded applications is to associate a light-weight process with each thread. In this way, the threads can access the same set of application data structures by simply sharing the same memory address space, the same set of open files, and so on; at the same time, each thread can be scheduled independently by the kernel so that one may sleep while another remains runnable. Examples of POSIX-compliant *pthread* libraries that use Linux's lightweight processes are *LinuxThreads*, *Native POSIX Thread Library (NPTL)*, and IBM's *Next Generation Posix Threading Package (NGPT)*.

POSIX-compliant multithreaded applications are best handled by kernels that support "thread groups". In Linux a *thread group* is basically a set of lightweight processes that implement a multithreaded application and act as a whole with regards to some system calls such as `getpid()`, `kill()`, and `_exit()`.

## Linux Threads

`clone()` creates a separate process that shares the address space of the calling process. The cloned task behaves *much like a separate thread*.



`clone()`

```

1 #include <sched.h>
2 int clone(int (*fn) (void *), void *child_stack,
3           int flags, void *arg, ...);
  
```

- arg 1** the function to be executed, i.e. `fn(arg)`, which returns an `int`;
  - arg 2** a pointer to a (usually malloced) memory space to be used as the stack for the new thread;
  - arg 3** a set of flags used to indicate how much the calling process is to be shared. In fact,  
`clone(0) == fork()`
  - arg 4** the arguments passed to the function.
- It returns the PID of the child process or -1 on failure.
- ```
$ man clone
```

## The `clone()` System Call

Some flags:

| flag                       | Shared                |
|----------------------------|-----------------------|
| <code>CLONE_FS</code>      | File-system info      |
| <code>CLONE_VM</code>      | Same memory space     |
| <code>CLONE_SIGHAND</code> | Signal handlers       |
| <code>CLONE_FILES</code>   | The set of open files |

### In practice, one should try to avoid calling `clone()` directly

Instead, use a threading library (such as pthreads) which use `clone()` when starting a thread (such as during a call to `pthread_create()`)

## `clone()` Example

```

1 #include <unistd.h>
2 #include <sched.h>
3 #include <sys/types.h>
4 #include <stdlib.h>
5 #include <string.h>
6 #include <stdio.h>
7 #include <fcntl.h>
8
9 int variable;
10
11 int do_something()
12 {
13     variable = 42;
14     _exit(0);
15 }
16
17 int main(void)
18 {
19     void *child_stack;
20     variable = 9;
21
22     child_stack = (void *) malloc(16384);
23     printf("The variable was %d\n", variable);
24
25     clone(do_something, child_stack,
26           CLONE_FS | CLONE_VM | CLONE_FILES, NULL);
27     sleep(1);
28
29     printf("The variable is now %d\n", variable);
30     return 0;
31 }
```

## Stack Grows Downwards

```
child_stack = (void**)malloc(8192) + 8192/sizeof(*child_stack);
```

## References

- [1] Wikipedia. Process (computing) — Wikipedia, The Free Encyclopedia. 2014. [http://en.wikipedia.org/w/index.php?title=Process\\_\(computing\)&oldid=639847817](http://en.wikipedia.org/w/index.php?title=Process_(computing)&oldid=639847817).
- [2] Wikipedia. Thread (computing) — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Thread\\_\(computing\)&oldid=648172980](http://en.wikipedia.org/w/index.php?title=Thread_(computing)&oldid=648172980).
- [3] Wikipedia. Process control block — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Process\\_control\\_block&oldid=646933587](http://en.wikipedia.org/w/index.php?title=Process_control_block&oldid=646933587).

## 9 Process Synchronization

### 9.1 IPC

#### Interprocess Communication

##### Example:

```
$ ps | head -2 | tail -1 | cut -f2 -d' '
```

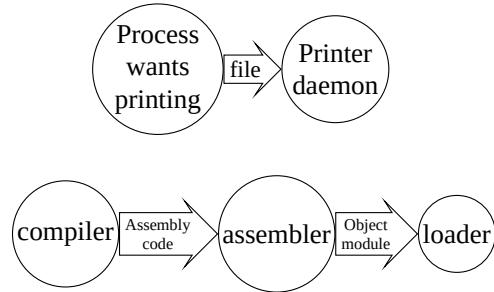
##### IPC issues:

1. How one process can pass information to another
2. Be sure processes do not get into each other's way
  - e.g. in an airline reservation system, two processes compete for the last seat
3. Proper sequencing when dependencies are present
  - e.g. if A produces data and B prints them, B has to wait until A has produced some data

##### Two models of IPC:

- Shared memory
- Message passing (e.g. sockets)

#### Producer-Consumer Problem



### 9.2 Shared Memory

#### Process Synchronization

##### Producer-Consumer Problem

- Consumers don't try to remove objects from Buffer when it is empty.
- Producers don't try to add objects to the Buffer when it is full.

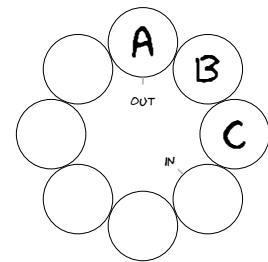
```
1 while(TRUE){           1 while(TRUE){  
2     while(FULL);      2     while(EMPTY);  
3     item = produceItem(); 3     item = removeItem();  
4     insertItem(item);   4     consumeItem(item);  
5 }                     5 }
```

How to define *full/empty*?

## Producer-Consumer Problem

— Bounded-Buffer Problem (Circular Array)

**Front(out):** the first full position  
**Rear(in):** the next free position



Full or empty when “ $front == rear$ ”?

## Producer-Consumer Problem

**Common solution:**

**Full:** when “ $(in + 1) \% BUFFER\_SIZE == out$ ”

Actually, this is “full – 1”

**Empty:** when “ $in == out$ ”

Can only use “ $BUFFER\_SIZE - 1$ ” elements

**Shared data:**

```
1 #define BUFFER_SIZE 6
2 typedef struct {
3     ...
4 } item;
5 item buffer[BUFFER_SIZE];
6 int in = 0; //the next free position
7 int out = 0; //the first full position
```

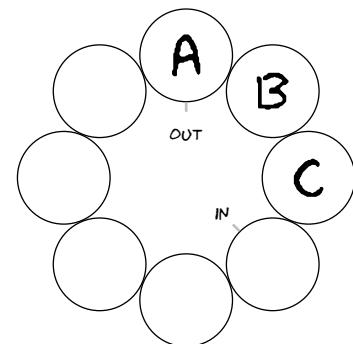
## Bounded-Buffer Problem

**Producer:**

```
1 while (true) {
2     /* do nothing -- no free buffers */
3     while (((in + 1) % BUFFER_SIZE) == out);
4
5     buffer[in] = item;
6     in = (in + 1) % BUFFER_SIZE;
7 }
```

**Consumer:**

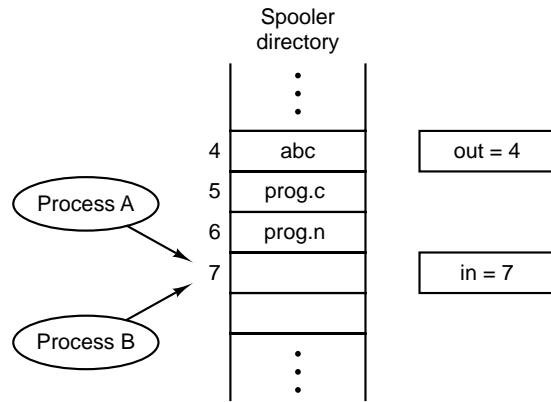
```
1 while (true) {
2     while (in == out); // do nothing
3     // remove an item from the buffer
4     item = buffer[out];
5     out = (out + 1) % BUFFER_SIZE;
6     return item;
7 }
```



## 9.3 Race Condition and Mutual Exclusion

**Race Conditions**

Now, let's have two producers



## Race Conditions

*Two producers*

```

1 #define BUFFER_SIZE 100
2 typedef struct {
3     ...
4 } item;
5 item buffer[BUFFER_SIZE]
6 int in = 0;
7 int out = 0;
```

**Process A and B do the same thing:**

```

1 while (true) {
2     while (((in + 1) % BUFFER_SIZE) == out);
3     buffer[in] = item;
4     in = (in + 1) % BUFFER_SIZE;
5 }
```

## Race Conditions

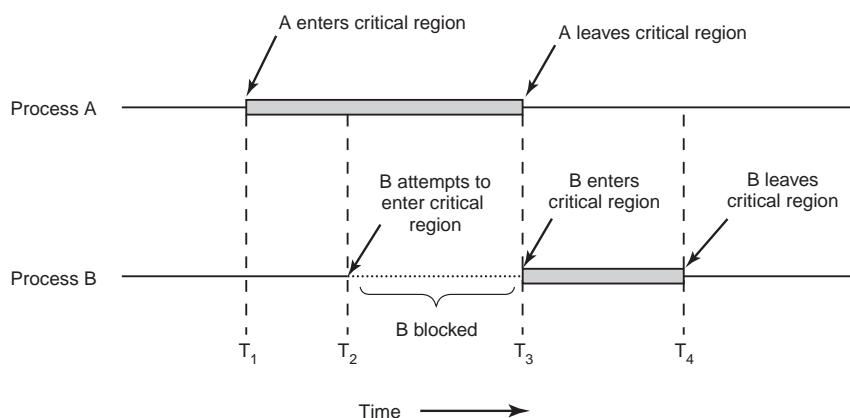
**Problem:** Process B started using one of the *shared variables* before Process A was finished with it.

**Solution:** *Mutual exclusion*. If one process is using a shared variable or file, the other processes will be excluded from doing the same thing.

## Critical Regions

*Mutual Exclusion*

**Critical Region:** is a piece of code accessing a common resource.



## Critical Region

A solution to the critical region problem must satisfy three conditions:

**Mutual Exclusion:** No two process may be simultaneously inside their critical regions.

**Progress:** No process running outside its critical region may block other processes.

**Bounded Waiting:** No process should have to wait forever to enter its critical region.

## Mutual Exclusion With Busy Waiting

*Disabling Interrupts*

```
1 {  
2 ...  
3 disableINT();      Problems:  
4 critical_code();    • It's not wise to give user process the power of turning off INTs.  
5 enableINT();        – Suppose one did it, and never turned them on again  
6 ...  
7 }
```

Disabling INTs is often a useful technique within the kernel itself but is not a general mutual exclusion mechanism for user processes.

## Mutual Exclusion With Busy Waiting

*Lock Variables*

```
1 int lock=0; //shared variable  
2 {  
3 ...  
4 while(lock); //busy waiting   Problem:  
5 lock=1;  
6 critical_code();  
7 lock=0;  
8 ...  
9 }
```

- What if an interrupt occurs right at line 5?
- Checking the lock again while backing from an interrupt?

## Mutual Exclusion With Busy Waiting

*Strict Alternation*

```
1 while(TRUE){  
2   while(turn != 0);  
3   critical_region();  
4   turn = 1;  
5   noncritical_region();  
6 }           1 while(TRUE){  
2   while(turn != 1);  
3   critical_region();  
4   turn = 0;  
5   noncritical_region();  
6 }
```

### Problem: violates condition-2

- One process can be blocked by another not in its critical region.
- Requires the two processes strictly alternate in entering their critical region.

## Mutual Exclusion With Busy Waiting

*Peterson's Solution*

|                                                                        |                                                                                                                                                         |
|------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>P0</b>                                                              | <b>P1</b>                                                                                                                                               |
| <pre> 1 int interest[0] = 0; 2 int interest[1] = 0; 3 int turn; </pre> | <pre> 1 interest[1] = 1; 2 turn = 0; 3 while(interest[0] == 1 &amp;&amp; turn == 0); 4 critical_section(); 5 interest[0] = 0; 6 interest[1] = 0; </pre> |

## References

- [1] Wikipedia. Peterson's algorithm — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Peterson%27s\\_algorithm&oldid=646078826](http://en.wikipedia.org/w/index.php?title=Peterson%27s_algorithm&oldid=646078826).

### Mutual Exclusion With Busy Waiting

*Hardware Solution: The TSL Instruction*

#### Lock the memory bus

|                                                                                   |                                                                                                                                                           |
|-----------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre> enter_region: TSL REGISTER,LOCK CMP REGISTER,#0 JNE enter_region RET </pre> | copy lock to register and set lock to 1<br>  was lock zero?<br>  if it was non zero, lock was set, so loop<br>  return to caller; critical region entered |
| <pre> leave_region: MOVE LOCK,#0 RET </pre>                                       | store a 0 in lock<br>  return to caller                                                                                                                   |

See also: [Modern Operating Systems, Sec. 2.3.3, Mutual Exclusion With Busy Waiting, p. 124].

### Mutual Exclusion Without Busy Waiting

*Sleep & Wakeup*

```

1 #define N 100 /* number of slots in the buffer */
2 int count = 0; /* number of items in the buffer */

1 void producer(){
2     int item;
3     while(TRUE){
4         item = produce_item();
5         if(count == N)
6             sleep();
7         insert_item(item);
8         count++;
9         if(count == 1)
10            wakeup(consumer);
11    }
12 }

1 void consumer(){
2     int item;
3     while(TRUE){
4         if(count == 0)
5             sleep();
6         item = rm_item();
7         count--;
8         if(count == N - 1)
9             wakeup(producer);
10        consume_item(item);
11    }
12 }

```

### Producer-Consumer Problem

*Race Condition*

#### Problem

1. Consumer is going to sleep upon seeing an empty buffer, but INT occurs;
2. Producer inserts an item, increasing count to 1, then call `wakeup(consumer)`;
3. But the consumer is not asleep, though count was 0. So the `wakeup()` signal is lost;
4. Consumer is back from INT remembering count is 0, and goes to sleep;
5. Producer sooner or later will fill up the buffer and also goes to sleep;
6. Both will sleep forever, and waiting to be waken up by the other process. Deadlock!

## Producer-Consumer Problem

Race Condition

### Solution: Add a *wakeup waiting bit*

1. The bit is set, when a wakeup is sent to an awaken process;
  2. Later, when the process wants to sleep, it checks the bit first. Turns it off if it's set, and stays awake.
- What if many processes try going to sleep?

## 9.4 Semaphores

### What is a Semaphore?

- A locking mechanism
- An integer or ADT

| Atomic Operations |             |
|-------------------|-------------|
| P()               | V()         |
| Wait()            | Signal()    |
| Down()            | Up()        |
| Decrement()       | Increment() |
| ...               | ...         |

```
1 down(S){           1 up(S){  
2   while(S<=0);    2   S++;  
3   S--;             3 }  
4 }
```

### More meaningful names:

- increment\_and\_wake\_a\_waiting\_process\_if\_any()
- decrement\_and\_block\_if\_the\_result\_is\_negative()

### Semaphore

#### How to ensure atomic?

1. For single CPU, implement up() and down() as system calls, with the OS disabling all interrupts while accessing the semaphore;
2. For multiple CPUs, to make sure only one CPU at a time examines the semaphore, a lock variable should be used with the TSL instructions.

### Semaphore is a Special Integer

#### A semaphore is like an integer, with three differences:

1. You can initialize its value to any integer, but after that the only operations you are allowed to perform are *increment* ( $S++$ ) and *decrement* ( $S--$ ).
2. When a thread decrements the semaphore, if the result is negative ( $S \leq 0$ ), the thread blocks itself and cannot continue until another thread increments the semaphore.
3. When a thread increments the semaphore, if there are other threads waiting, one of the waiting threads gets unblocked.

### Why Semaphores?

We don't need semaphores to solve synchronization problems, but there are some advantages to using them:

- Semaphores impose deliberate constraints that help programmers avoid errors.
- Solutions using semaphores are often clean and organized, making it easy to demonstrate their correctness.
- Semaphores can be implemented efficiently on many systems, so solutions that use semaphores are portable and usually efficient.

## The Simplest Use of Semaphore

*Signaling*

- One thread sends a signal to another thread to indicate that something has happened
- it solves the serialization problem

Signaling makes it possible to guarantee that a section of code in one thread will run before a section of code in another thread

```
1 statement a1      1 sem.wait()  
2 sem.signal()     2 statement b1
```

What's the initial value of sem?

## Semaphore

*Rendezvous Puzzle*

|                                                |                                                |
|------------------------------------------------|------------------------------------------------|
| <pre> 1   statement a1 2   statement a2 </pre> | <pre> 1   statement b1 2   statement b2 </pre> |
|------------------------------------------------|------------------------------------------------|

Q: How to guarantee that

- 1. a1 happens before b2, and
- 2. b1 happens before a2

$a1 \rightarrow b2; b1 \rightarrow a2$

Hint: Use two semaphores initialized to 0.

|             | <b>Thread A:</b>                                                 | <b>Thread B:</b>                                                 |
|-------------|------------------------------------------------------------------|------------------------------------------------------------------|
| Solution 1: | <pre> statement a1 sem1.wait() sem2.signal() statement a2 </pre> | <pre> statement b1 sem1.signal() sem2.wait() statement b2 </pre> |
| Solution 2: | <pre> statement a1 sem2.signal() sem1.wait() statement a2 </pre> | <pre> statement b1 sem1.signal() sem2.wait() statement b2 </pre> |
| Solution 3: | <pre> statement a1 sem2.wait() sem1.signal() statement a2 </pre> | <pre> statement b1 sem1.wait() sem2.signal() statement b2 </pre> |

*Solution 3 has deadlock!*

## Mutex

- A second common use for semaphores is to enforce mutual exclusion
- It guarantees that only one thread accesses the shared variable at a time
- A mutex is like a token that passes from one thread to another, allowing one thread at a time to proceed

Q: Add semaphores to the following example to enforce mutual exclusion to the shared variable *i*.

**Thread A:** i++

**Thread B:** i++

**Why?** Because *i++* is not atomic.

### i++ is not atomic in assembly language

```

1 LOAD    [i], r0      ;load the value of 'i' into
2                      ;a register from memory
3 ADD     r0, 1        ;increment the value
4                      ;in the register
5 STORE   r0, [i]      ;write the updated
6                      ;value back to memory

```

Interrupts might occur in between. So, *i++* needs to be protected with a mutex.

## Mutex Solution

### Create a semaphore named mutex that is initialized to 1

- 1: a thread may proceed and access the shared variable
- 0: it has to wait for another thread to release the mutex

|                                                              |                                                              |
|--------------------------------------------------------------|--------------------------------------------------------------|
| <pre> 1   mutex.wait() 2       i++ 3   mutex.signal() </pre> | <pre> 1   mutex.wait() 2       i++ 3   mutex.signal() </pre> |
|--------------------------------------------------------------|--------------------------------------------------------------|

## Multiplex — Without Busy Waiting

```

1| typedef struct{
2|     int space;           //number of free resources
3|     struct process *P; //a list of queueing producers
4|     struct process *C; //a list of queueing consumers
5| } semaphore;
6| semaphore S;
7| S.space = 5;

1| void down(S){
2|     S.space--;
3|     if(S.space == 4){
4|         rmFromQueue(S.C);
5|         wakeup(S.C);
6|     }
7|     if(S.space < 0){
8|         addToQueue(S.P);
9|         sleep();
10}
11}

1| void up(S){
2|     S.space++;
3|     if(S.space > 5){
4|         addToQueue(S.C);
5|         sleep();
6|     }
7|     if(S.space >= 0){
8|         rmFromQueue(S.P);
9|         wakeup(S.P);
10}
11}

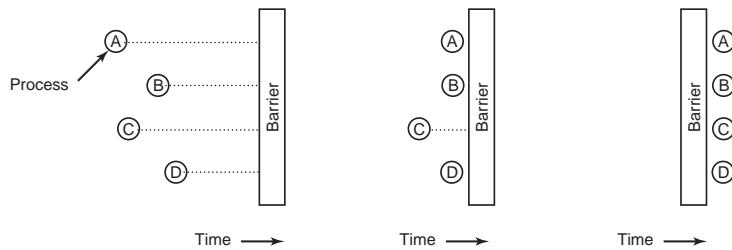
```

if S.space < 0,  
S.space == Number of queuing producers  
if S.space > 5,  
S.space == Number of queuing consumers + 5

**The work flow** There are several processes running simultaneously. They all need to access some common resources.

1. Assuming S.space == 3 in the beginning
2. Process P1 comes and take one resource away. S.space == 2 now.
3. Process P2 comes and take the 2nd resource away. S.space == 1 now.
4. Process P3 comes and take the last resource away. S.space == 0 now.
5. Process P4 comes and sees nothing left. It has to sleep. S.space == -1 now.
6. Process P5 comes and sees nothing left. It has to sleep. S.space == -2 now.
7. At this moment, there are 2 processes (P4 and P5) sleeping. In another word, they are queuing for resources.
8. Now, P1 finishes using the resource, and released it. After it does a S.space++, it finds out that S.space <= 0. So it wakes up a Process (say P4) in the queue.
9. P4 wakes up, and back to execute the instruction right after sleep().
10. P4 (or P2|P3) finishes using the resource, and releases it. After it does a S.space++, it finds out that S.space <= 0. So it wakes up P5 in the queue.
11. the queue is empty now.

## Barrier



1. Processes approaching a barrier
2. All processes but one blocked at the barrier
3. When the last process arrives at the barrier, all of them are let through

### Synchronization requirement:

```

specific_task()
critical_point()

```

No thread executes `critical_point()` until after all threads have executed `specific_task()`.

## Barrier Solution

```
1 n = the number of threads
2 count = 0
3 mutex = Semaphore(1)
4 barrier = Semaphore(0)
```

count: keeps track of how many threads have arrived

mutex: provides exclusive access to count

barrier: is locked ( $\leq 0$ ) until all threads arrive

When `barrier.value < 0`,

```
barrier.value == Number of queueing processes
```

```
1 #define CHAIRS 5
2 semaphore customers = 0;
3 semaphore bber = ?;
4 semaphore mutex = 1;
5 int waiting = 0;
6
7 void barber(void)
8 {
9     while (TRUE) {
10         wait(&customers);
11         cutHair();
12     }
13 }
```

```
1 void customer(void)
2 {
3     if (waiting == CHAIRS)
4         goHome();
5     else {
6         wait(&mutex);
7         waiting++;
8         signal(&mutex);
9         signal(&customers);
10        wait(&bber);
11        getHairCut();
12        signal(&bber);
13        wait(&mutex);
14        waiting--;
15        signal(&mutex);
16    }
17 }
```

*Only one thread can pass the barrier!*

## Barrier Solution

```
1 specific_task();
2
3 mutex.wait();
4     count++;
5 mutex.signal();
6
7 if (count == n)
8     barrier.signal();
9
10 barrier.wait();
11 barrier.signal();
12
13 critical_point();
```

```
1 specific_task();
2
3 mutex.wait();
4     count++;
5
6 if (count == n)
7     barrier.signal();
8
9 barrier.wait();
10 barrier.signal();
11 mutex.signal();
12
13 critical_point();
```

✿ Blocking on a semaphore while holding a mutex!✿

```
barrier.wait();
barrier.signal();
```

## Turnstile

This pattern, a `wait` and a `signal` in rapid succession, occurs often enough that it has a name called a *turnstile*, because

- it allows one thread to pass at a time, and
- it can be locked to bar all threads

## Semaphores

*Producer-Consumer Problem*

### Whenever an event occurs

- a producer thread creates an *event object* and adds it to the event buffer. Concurrently,
- consumer threads take events out of the buffer and process them. In this case, the consumers are called “*event handlers*”.

```

1| event = waitForEvent()      1| event = buffer.get()
2| buffer.add(event)          2| event.process()

```

Q: Add synchronization statements to the producer and consumer code to enforce the synchronization constraints

1. Mutual exclusion
2. Serialization

See also: [The Little Book of Semaphores, Sec. 4.1, Producer-Consumer Problem].

## Semaphores — Producer-Consumer Problem

*Solution*

**Initialization:**

```

mutex = Semaphore(1)
items = Semaphore(0)

```

- mutex provides exclusive access to the buffer
- items:
  - + number of items in the buffer
  - number of consumer threads in queue

## Semaphores — Producer-Consumer Problem

*Solution*

```

1| event = waitForEvent()      1| items.wait()
2| mutex.wait()                2| mutex.wait()
3|     buffer.add(event)       3|     event = buffer.get()
4|     items.signal()          4|     mutex.signal()
5| mutex.signal()              5| event.process()

```

or,

```

1| event = waitForEvent()      1| mutex.wait()
2| mutex.wait()                2|     items.wait()
3|     buffer.add(event)       3|     event = buffer.get()
4|     mutex.signal()          4|     mutex.signal()
5| items.signal()              5| event.process()

```

⌘ any time you wait for a semaphore while holding a mutex!

```

items.signal()                      items.wait()
{
  items++;
  if(items == 0)
    wakeup(consumer);
}
{                                     {
  items--;
  if(items < 0)
    sleep();
}

```

## Producer-Consumer Problem With Bounded-Buffer

**Given:**

```

semaphore items = 0;
semaphore spaces = BUFFER_SIZE;

```

**Can we?**

```

if (items >= BUFFER_SIZE)
  producer.block();

```

**if:** the buffer is full

**then:** the producer blocks until a consumer removes an item

*No!* We can't check the current value of a semaphore, because

*! the only operations are wait and signal.*

*? But...*

**Why can't check the current value of a semaphore?** We DO have seen:

```
void S.down(){  
    S.value--;  
    if(S.value < 0){  
        addToQueue(S.L);  
        sleep();  
    }  
}  
}  
void S.up(){  
    S.value++;  
    if(S.value <= 0){  
        rmFromQueue(S.L);  
        wakeup(S.L);  
    }  
}
```

Notice that the checking is within down() and up(), and is not available to user process to use it directly.

## Producer-Consumer Problem With Bounded-Buffer

```
1 | semaphore items = 0;  
2 | semaphore spaces = BUFFER_SIZE;  
  
1 | void producer() {  
2 |     while (true) {  
3 |         item = produceItem();  
4 |         down(spaces);  
5 |         putIntoBuffer(item);  
6 |         up(items);  
7 |     }  
8 | }  
1 | void consumer() {  
2 |     while (true) {  
3 |         down(items);  
4 |         item = rmFromBuffer();  
5 |         up(spaces);  
6 |         consumeItem(item);  
7 |     }  
8 | }
```

works fine when there is only one producer and one consumer, because putIntoBuffer() is not atomic.

**putIntoBuffer() could contain two actions:**

1. determining the next available slot
2. writing into it

**Race condition:**

1. Two producers decrement spaces
2. One of the producers determines the next empty slot in the buffer
3. Second producer determines the next empty slot and gets the same result as the first producer
4. Both producers write into the same slot

putIntoBuffer() needs to be protected with a mutex.

## With a mutex

```
1 | semaphore mutex = 1; //controls access to c.r.  
2 | semaphore items = 0;  
3 | semaphore spaces = BUFFER_SIZE;  
  
1 | void producer() {  
2 |     while (true) {  
3 |         item = produceItem();  
4 |         down(spaces);  
5 |         down(mutex);  
6 |         putIntoBuffer(item);  
7 |         up(mutex);  
8 |         up(items);  
9 |     }  
10 | }  
1 | void consumer() {  
2 |     while (true) {  
3 |         down(items);  
4 |         down(mutex);  
5 |         item = rmFromBuffer();  
6 |         up(mutex);  
7 |         up(spaces);  
8 |         consumeItem(item);  
9 |     }  
10 | }
```

## 9.5 Monitors

### Monitors

**Monitor** a high-level synchronization object for achieving mutual exclusion.

- It's a language concept, and C does not have it.
- Only one process can be active in a monitor at any instant.
- It is up to the compiler to implement mutual exclusion on monitor entries.
  - The programmer just needs to know that by turning all the critical regions into monitor procedures, no two processes will ever execute their critical regions at the same time.

```

1 monitor example
2   integer i;
3   condition c;
4
5   procedure producer();
6   ...
7   end;
8
9   procedure consumer();
10  ...
11  end;
12 end monitor;
```

## Monitor

*The producer-consumer problem*

```

1 monitor ProducerConsumer
2   condition full, empty;
3   integer count;
4
5   procedure insert(item: integer);
6   begin
7     if count = N then wait(full);
8     insert_item(item);
9     count := count + 1;
10    if count = 1 then signal(empty)
11  end;
12
13  function remove: integer;
14  begin
15    if count = 0 then wait(empty);
16    remove = remove_item;
17    count := count - 1;
18    if count = N - 1 then signal(full)
19  end;
20  count := 0;
21 end monitor;
```

```

1 procedure producer;
2 begin
3   while true do
4     begin
5       item = produce_item;
6       ProducerConsumer.insert(item)
7     end;
8   end;
9
10 procedure consumer;
11 begin
12   while true do
13     begin
14       item = ProducerConsumer.remove;
15       consume_item(item)
16     end;
17   end;
```

## 9.6 Message Passing

### Message Passing

- Semaphores are too low level
- Monitors are not usable except in a few programming languages
- Neither monitor nor semaphore is suitable for distributed systems
- No conflicts, easier to implement

Message passing uses two primitives, `send` and `receive` system calls:

- `send(destination, &message);`
- `receive(source, &message);`

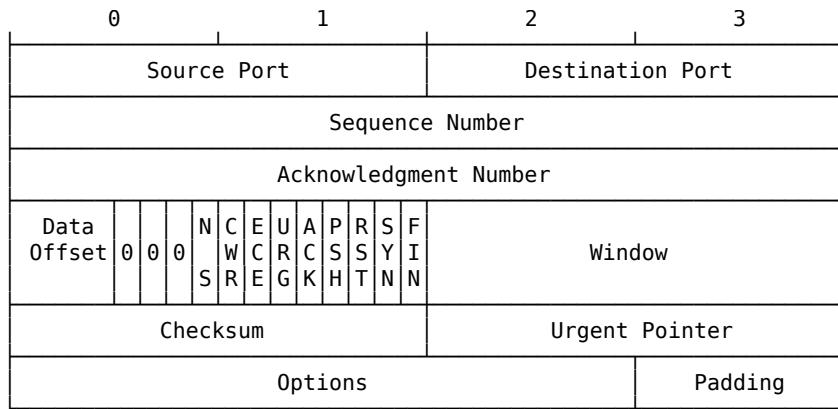
### Message Passing

#### Design issues

- Message can be lost by network; — ACK
- What if the ACK is lost? — SEQ
- What if two processes have the same name? — socket
- Am I talking with the right guy? Or maybe a MIM? — authentication
- What if the sender and the receiver on the same machine? — Copying messages is always slower than doing a semaphore operation or entering a monitor.

### Message Passing

*TCP Header Format*



## Message Passing

*The producer-consumer problem*

```

1 #define N 100 /* number of slots in the buffer */
2 void producer(void)
3 {
4     int item;
5     message m;           /* message buffer */
6     while (TRUE) {
7         item = produce_item(); /* generate something to put in buffer */
8         receive(consumer, &m); /* wait for an empty to arrive */
9         build_message(&m, item); /* construct a message to send */
10        send(consumer, &m);   /* send item to consumer */
11    }
12 }
13
14 void consumer(void)
15 {
16     int item, i;
17     message m;
18     for (i=0; i<N; i++) send(producer, &m); /* send N empties */
19     while (TRUE) {
20         receive(producer, &m); /* get message containing item */
21         item = extract_item(&m); /* extract item from message */
22         send(producer, &m);   /* send back empty reply */
23         consume_item(item); /* do something with the item */
24     }
25 }
```

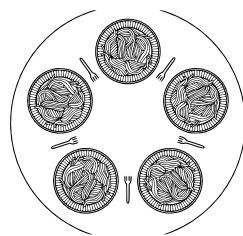
## 9.7 Classical IPC Problems

### 9.7.1 The Dining Philosophers Problem

**The Dining Philosophers Problem**

```

1 while True:
2     think()
3     get_forks()
4     eat()
5     put_forks()
```



How to implement `get_forks()` and `put_forks()` to ensure

1. No deadlock
2. No starvation
3. Allow more than one philosopher to eat at the same time

**The Dining Philosophers Problem**

*Deadlock*

```

#define N 5          /* number of philosophers */

void philosopher(int i)    /* i: philosopher number, from 0 to 4 */
{
    while (TRUE) {
        think();      /* philosopher is thinking */
        take_fork(i); /* take left fork */
        take_fork((i+1) % N); /* take right fork; % is modulo operator */
        eat();         /* yum-yum, spaghetti */
        put_fork(i);   /* put left fork back on the table */
        put_fork((i+1) % N); /* put right fork back on the table */
    }
}

```

- Put down the left fork and wait for a while if the right one is not available? Similar to CSMA/CD — Starvation

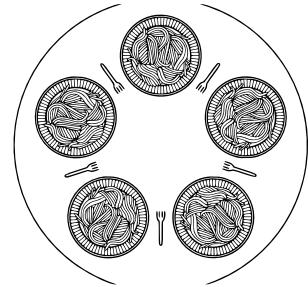
### The Dining Philosophers Problem

With One Mutex

```

1 #define N 5
2 semaphore mutex=1;
3
4 void philosopher(int i)
5 {
6     while (TRUE) {
7         think();
8         wait(&mutex);
9         take_fork(i);
10        take_fork((i+1) % N);
11        eat();
12        put_fork(i);
13        put_fork((i+1) % N);
14        signal(&mutex);
15    }
16 }

```



- Only one philosopher can eat at a time.
- How about 2 mutexes? 5 mutexes?

### The Dining Philosophers Problem

AST Solution (Part 1)

A philosopher may only move into eating state if neither neighbor is eating

```

1 #define N 5          /* number of philosophers */
2 #define LEFT (i+N-1)%N /* number of i's left neighbor */
3 #define RIGHT (i+1)%N /* number of i's right neighbor */
4 #define THINKING 0    /* philosopher is thinking */
5 #define HUNGRY 1       /* philosopher is trying to get forks */
6 #define EATING 2       /* philosopher is eating */
7 typedef int semaphore;
8 int state[N];           /* state of everyone */
9 semaphore mutex = 1;    /* for critical regions */
10 semaphore s[N];        /* one semaphore per philosopher */
11
12 void philosopher(int i) /* i: philosopher number, from 0 to N-1 */
13 {
14     while (TRUE) {
15         think();
16         take_forks(i); /* acquire two forks or block */
17         eat();
18         put_forks(i);  /* put both forks back on table */
19     }
20 }

```

## The Dining Philosophers Problem

AST Solution (Part 2)

```
1 void take_forks(int i)           /* i: philosopher number, from 0 to N-1 */
2 {
3     down(&mutex);
4     state[i] = HUNGRY;
5     test(i);
6     up(&mutex);
7     down(&s[i]);
8 }
9 void put_forks(i)               /* i: philosopher number, from 0 to N-1 */
10 {
11     down(&mutex);
12     state[i] = THINKING;
13     test(LEFT);
14     test(RIGHT);
15     up(&mutex);
16 }
17 void test(i)                  /* i: philosopher number, from 0 to N-1 */
18 {
19     if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
20         state[i] = EATING;
21         up(&s[i]);
22     }
23 }
```

- ✿ Starvation can happen!

### Step by step

1. If 5 philosophers `take_forks(i)` at the same time, only one can get `mutex`.
2. The one who gets `mutex` sets his `state` to `HUNGRY`. And then,
3. `test(i);` try to get 2 forks.
  - If his `LEFT` and `RIGHT` are not `EATING`, success to get 2 forks.
    - i. sets his `state` to `EATING`
    - ii. `up(&s[i]);` The initial value of `s(i)` is 0.
  - Now, his `LEFT` and `RIGHT` will fail to get 2 forks, even if they could grab `mutex`.
  - If either `LEFT` or `RIGHT` are `EATING`, fail to get 2 forks.
4. `release mutex`
5. `down(&s[i]);`
  - (a) block if forks are not acquired
  - (b) `eat()` if 2 forks are acquired
6. After `eat()`ing, the philosopher doing `put_forks(i)` has to get `mutex` first.
  - because `state[i]` can be changed by more than one philosopher.
7. After getting `mutex`, set his `state` to `THINKING`
8. `test(LEFT);` see if `LEFT` can now eat?
  - If `LEFT` is `HUNGRY`, and `LEFT`'s `LEFT` is not `EATING`, and `LEFT`'s `RIGHT` (me) is not `EATING`
    - i. set `LEFT`'s `state` to `EATING`
    - ii. `up(&s[LEFT]);`
  - If `LEFT` is not `HUNGRY`, or `LEFT`'s `LEFT` is `EATING`, or `LEFT`'s `RIGHT` (me) is `EATING`, `LEFT` fails to get 2 forks.
9. `test(RIGHT);` see if `RIGHT` can now eat?
10. `release mutex`

## The Dining Philosophers Problem

More Solutions

- If there is at least one leftie and at least one rightie, then deadlock is not possible
- Wikipedia: Dining philosophers problem

### 9.7.2 The Readers-Writers Problem

#### The Readers-Writers Problem

**Constraint:** no process may access the shared data for reading or writing while another process is writing to it.

```
1 void reader(void)
2 {
3     semaphore mutex = 1;
4     semaphore noOther = 1;
5     int readers = 0;
6
7     void writer(void)
8     {
9         while (TRUE) {
10             wait(&noOther);
11             writing();
12             signal(&noOther);
13         }
14     }
15
16     while (TRUE) {
17         wait(&mutex);
18         readers++;
19         if (readers == 1)
20             wait(&noOther);
21         signal(&mutex);
22         reading();
23         wait(&mutex);
24         readers--;
25         if (readers == 0)
26             signal(&noOther);
27         signal(&mutex);
28         anything();
29     }
30 }
```

**Starvation** The writer could be blocked forever if there are always someone reading.

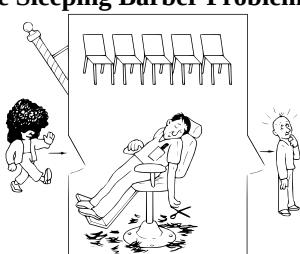
### The Readers-Writers Problem

*No starvation*

```
1 void reader(void)
2 {
3     semaphore mutex = 1;
4     semaphore noOther = 1;
5     semaphore turnstile = 1;
6     int readers = 0;
7
8     void writer(void)
9     {
10         while (TRUE) {
11             turnstile.wait();
12             turnstile.signal();
13
14             wait(&mutex);
15             readers++;
16             if (readers == 1)
17                 wait(&noOther);
18             signal(&mutex);
19             reading();
20             wait(&mutex);
21             readers--;
22             if (readers == 0)
23                 signal(&noOther);
24             signal(&mutex);
25             anything();
26         }
27     }
28 }
```

### 9.7.3 The Sleeping Barber Problem

#### The Sleeping Barber Problem



#### Where's the problem?

- the barber saw an empty room right before a customer arrives the waiting room;
- Several customer could race for a single chair;

## Solution

```
1 #define CHAIRS 5
2 semaphore customers = 0; // any customers or not?
3 semaphore bber = 0;      // barber is busy
4 semaphore mutex = 1;
5 int waiting = 0;         // queueing customers

1 void barber(void)           1 void customer(void)
2 {                           2 {
3     while (TRUE) {          3     if(waiting == CHAIRS)
4         wait(&customers);   4         goHome();
5         wait(&mutex);       5     else {
6             waiting--;        6         wait(&mutex);
7             signal(&mutex);    7         waiting++;
8             cutHair();         8         signal(&mutex);
9             signal(&bber);    9         signal(&customers);
10 }                         10        wait(&bber);
11 }                         11        getHairCut();
12 }                         12 }
13 }
```

## Solution2

```
1 #define CHAIRS 5
2 semaphore customers = 0;
3 semaphore bber = ???;
4 semaphore mutex = 1;
5 int waiting = 0;
6
7 void customer(void)
8 {
9     wait(&mutex);
10    if (waiting == CHAIRS){
11        signal(&mutex);
12        goHome();
13    } else {
14        waiting++;
15        signal(&customers);
16        signal(&mutex);
17        wait(&bber);
18        getHairCut();
19        wait(&mutex);
20        waiting--;
21        signal(&mutex);
22        signal(&bber);
23    }
24 }
```

## References

- [1] Wikipedia. Inter-process communication — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Inter-process\\_communication&oldid=645037874](http://en.wikipedia.org/w/index.php?title=Inter-process_communication&oldid=645037874).
- [2] Wikipedia. Semaphore (programming) — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?titleSemaphore\\_\(programming\)&oldid=647556304](http://en.wikipedia.org/w/index.php?titleSemaphore_(programming)&oldid=647556304).

# 10 CPU Scheduling

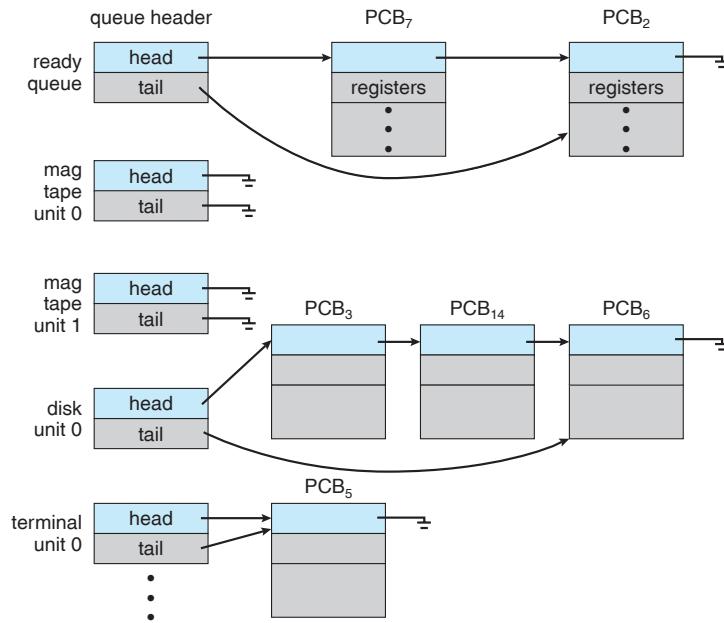
## 10.1 Process Scheduling Queues

### Scheduling Queues

**Job queue** consists all the processes in the system

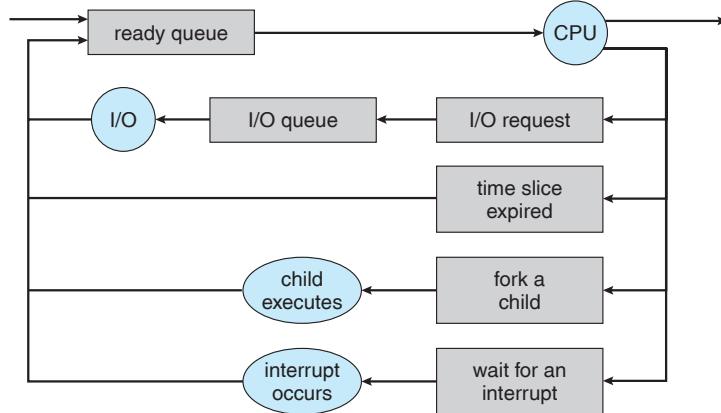
**Ready queue** A linked list consists processes in the main memory ready for execute

**Device queue** Each device has its own device queue



- The *tail* pointer — When adding a new process to the queue, don't have to find the tail by traversing the list

### Queueing Diagram



## 10.2 Scheduling

### Scheduling

- Scheduler uses *scheduling algorithm* to choose a process from the ready queue
- Scheduling doesn't matter much on simple PCs, because
  1. Most of the time there is only one active process
  2. The CPU is too fast to be a scarce resource any more
- Scheduler has to make efficient use of the CPU because process switching is expensive
  1. User mode → kernel mode
  2. Save process state, registers, memory map...
  3. Selecting a new process to run by running the scheduling algorithm
  4. Load the memory map of the new process
  5. The process switch usually invalidates the entire memory cache

### Scheduling Algorithm Goals

#### All systems

**Fairness** giving each process a fair share of the CPU

**Policy enforcement** seeing that stated policy is carried out

**Balance** keeping all parts of the system busy

### Batch systems

**Throughput** maximize jobs per hour

**Turnaround time** minimize time between submission and termination

**CPU utilization** keep the CPU busy all the time

#### Interactive systems

**Response time** respond to requests quickly

**Proportionality** meet users' expectations

#### Real-time systems

**Meeting deadlines** avoid losing data

**Predictability** avoid quality degradation in multimedia systems

See also: [Modern Operating Systems, Sec. 2.4.1.5, *Scheduling Algorithm Goals*, p. 150].

## 10.3 Process Classification

### Process Classification

#### Traditionally

CPU-bound processes vs. I/O-bound processes

#### Alternatively

**Interactive processes** responsiveness

- command shells, editors, graphical apps

**Batch processes** no user interaction, run in background, often penalized by the scheduler

- programming language compilers, database search engines, scientific computations

**Real-time processes** video and sound apps, robot controllers, programs that collect data from physical sensors

- should never be blocked by lower-priority processes
- should have a short guaranteed response time with a minimum variance

The two classifications we just offered are somewhat independent. For instance, a batch process can be either I/O-bound (e.g., a database server) or CPU-bound (e.g., an image-rendering program).

While real-time programs are explicitly recognized as such by the scheduling algorithm in Linux, there is no easy way to distinguish between interactive and batch programs. The Linux 2.6 scheduler implements a sophisticated heuristic algorithm based on the past behavior of the processes to decide whether a given process should be considered as interactive or batch. Of course, the scheduler tends to favor interactive processes over batch ones.

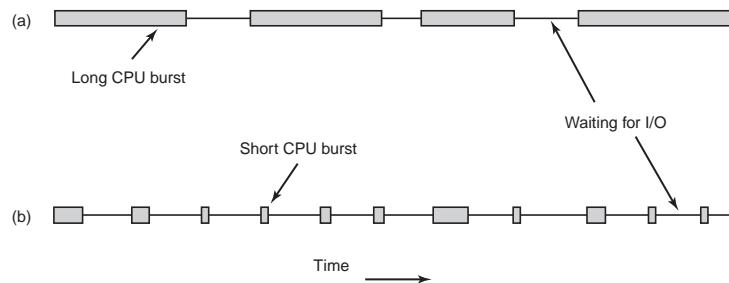
## 10.4 Process Behavior

### Process Behavior

#### CPU-bound vs. I/O-bound

Types of CPU bursts:

- long bursts – CPU bound (i.e. batch work)
- short bursts – I/O bound (i.e. emacs)



As CPUs get faster, processes tend to get more I/O-bound.

## 10.5 Process Schedulers

### Schedulers

**Long-term scheduler** (or job scheduler) - selects which processes should be brought into the ready queue.

**Short-term scheduler** (or CPU scheduler) - selects which process should be executed next and allocates CPU.

**Medium-term scheduler** swapping.

- LTS is responsible for a good *process mix* of I/O-bound and CPU-bound process leading to best performance.
- Time-sharing systems, e.g. UNIX, often have no long-term scheduler.

## Nonpreemptive vs. preemptive

A **nonpreemptive scheduling algorithm** lets a process run as long as it wants until it blocks (I/O or waiting for another process) or until it voluntarily releases the CPU.

A **preemptive scheduling algorithm** will forcibly suspend a process after it runs for sometime. — clock interruptable

## 10.6 Scheduling In Batch Systems

### Scheduling In Batch Systems

#### *First-Come First-Served*

- nonpreemptive
- simple
- also has a disadvantage

What if a CPU-bound process (e.g. runs 1s at a time) followed by many I/O-bound processes (e.g. 1000 disk reads to complete)?

\* In this case, a preemptive scheduling is preferred.

### Scheduling In Batch Systems

#### *Shortest Job First*



### Average turnaround time

$$(a) (8 + 12 + 16 + 20) \div 4 = 14$$

$$(b) (4 + 8 + 12 + 20) \div 4 = 11$$

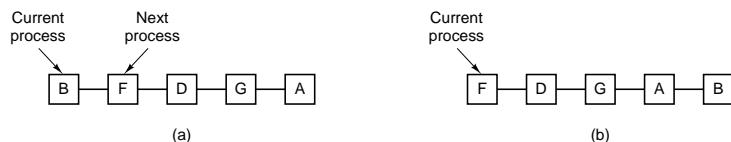
How to know the length of the next CPU burst?

- For long-term (job) scheduling, user provides
- For short-term scheduling, no way

## 10.7 Scheduling In Interactive Systems

### Scheduling In Interactive Systems

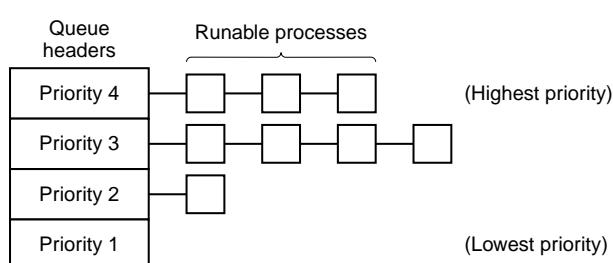
#### *Round-Robin Scheduling*



- Simple, and most widely used;
- Each process is assigned a time interval, called its *quantum*;
- How long should the quantum be?
  - too short — too many process switches, lower CPU efficiency;
  - too long — poor response to short interactive requests;
  - usually around 20 ~ 50ms.

### Scheduling In Interactive Systems

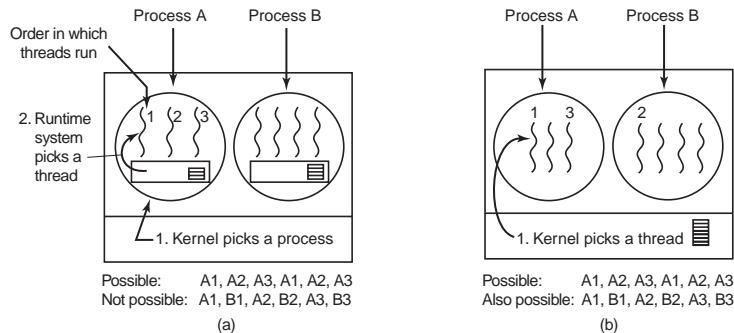
#### *Priority Scheduling*



- SJF is a priority scheduling;
  - *Starvation* — low priority processes may never execute;
    - *Aging* — as time progresses increase the priority of the process;
- ```
$ man nice
```

## 10.8 Thread Scheduling

### Thread Scheduling

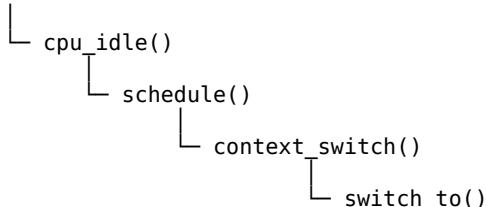


- With kernel-level threads, sometimes a full context switch is required
- Each process can have its own application-specific thread scheduler, which usually works better than kernel can

## 10.9 Linux Scheduling

- [Operating System Concepts Essentials, Sec. 5.6.3, Example: Linux Scheduling].
- [Operating System Concepts Essentials, Sec. 15.5, Scheduling].
- [Understanding The Linux Kernel, Chap. 7, Process Scheduling].
- [Linux Kernel Development, Chap. 4, Process Scheduling].

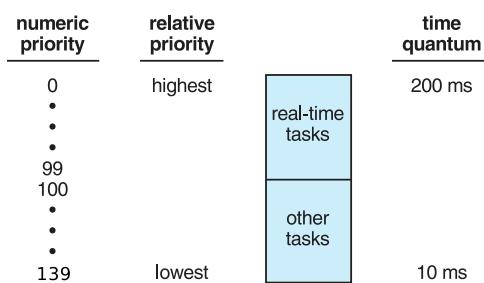
### Call graph:



### Process Scheduling In Linux

A preemptive, priority-based algorithm with two separate priority ranges:

1. *real-time* range (0 ~ 99), for tasks where absolute priorities are more important than fairness
2. *nice value* range (100 ~ 139), for fair preemptive scheduling among multiple processes



### In Linux, Process Priority is Dynamic

#### The scheduler keeps track of what processes are doing and adjusts their priorities periodically

- Processes that have been denied the use of a CPU for a long time interval are boosted by dynamically increasing their priority (usually I/O-bound)
- Processes running for a long time are penalized by decreasing their priority (usually CPU-bound)
- Priority adjustments are performed only on user tasks, not on real-time tasks

## Tasks are determined to be I/O-bound or CPU-bound based on an interactivity heuristic

A task's interactiveness metric is calculated based on how much time the task executes compared to how much time it sleeps

## Problems With The Pre-2.6 Scheduler

- ⌚ an algorithm with  $O(n)$  complexity
  - a single runqueue for all processors
    - ⌚ good for load balancing
    - ⌚ bad for CPU caches, when a task is rescheduled from one CPU to another
  - ⌚ a single runqueue lock — only one CPU working at a time

The scheduling algorithm used in earlier versions of Linux was quite simple and straightforward: at every process switch the kernel scanned the list of runnable processes, computed their priorities, and selected the "best" process to run. The main drawback of that algorithm is that the time spent in choosing the best process depends on the number of runnable processes; therefore, the algorithm is too costly, that is, it spends too much time in high-end systems running thousands of processes [*Understanding The Linux Kernel*, Sec. 7.2, *The Scheduling Algorithm*].

## Scheduling In Linux 2.6 Kernel

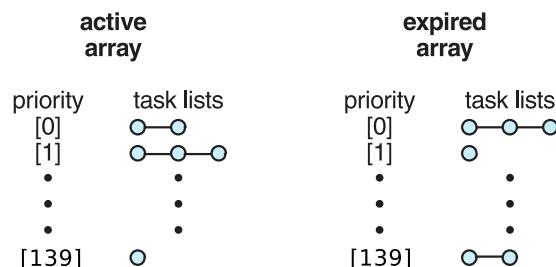
- ⌚  $O(1)$  — Time for finding a task to execute depends not on *the number of active tasks* but instead on *the number of priorities*
- ⌚ Each CPU has its own *runqueue*, and schedules itself independently; better cache efficiency
  - The job of the scheduler is simple — Choose the task on the highest priority list to execute

### How to know there are processes waiting in a priority list?

A priority bitmap (5 32-bit words for 140 priorities) is used to define when tasks are on a given priority list.

- `find-first-bit-set` instruction is used to find the highest priority bit.

### Each runqueue has two priority arrays



### 10.9.1 Completely Fair Scheduling

#### Completely Fair Scheduling (CFS)

##### Linux's Process Scheduler

**up to 2.4:** simple, scaled poorly

- $O(n)$
- non-preemptive in kernel mode
- single run queue (cache? SMP?)

**from 2.5 on:**  $O(1)$  scheduler

- ⌚ 140 priority lists — scaled well
- ⌚ one run queue per CPU — true SMP support
- ⌚ preemptive
- ⌚ ideal for large server workloads
- ⌚ showed latency on desktop systems

**from 2.6.23 on:** Completely Fair Scheduler (CFS)

- ⌚ improved interactive performance

## Completely Fair Scheduler (CFS)

For a perfect (unreal) multitasking CPU

- $n$  runnable processes can run at the same time
- each process should receive  $\frac{1}{n}$  of CPU power

For a real world CPU

- can run only a single task at once — unfair
  - ◎ while one task is running
  - ◎ the others have to wait
- $p \rightarrow \text{wait\_runtime}$  is the amount of time the task should now run on the CPU for it becomes completely fair and balanced.
  - ◎ on ideal CPU, the  $p \rightarrow \text{wait\_runtime}$  value would always be zero
- CFS always tries to run the task with the largest  $p \rightarrow \text{wait\_runtime}$  value

See also: *Discussing the Completely Fair Scheduler*<sup>6</sup>.

## CFS

In practice it works like this:

- While a task is using the CPU, its  $\text{wait\_runtime}$  decreases

```
wait_runtime = wait_runtime - time_running
```

if: its  $\text{wait\_runtime} \neq \text{MAX}_{\text{wait\_runtime}}$  (among all processes)

then: it gets preempted

- Newly woken tasks ( $\text{wait\_runtime} = 0$ ) are put into the tree more and more to the right
- slowly but surely giving a chance for every task to become the “leftmost task” and thus get on the CPU within a deterministic amount of time

## References

- [1] Wikipedia. Scheduling (computing) — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Scheduling\\_\(computing\)&oldid=647892123](http://en.wikipedia.org/w/index.php?title=Scheduling_(computing)&oldid=647892123).

# 11 Deadlock

## 11.1 Resources

### A Major Class of Deadlocks Involve Resources

#### Processes need access to resources in reasonable order

Suppose...

- a process holds resource A and requests resource B. At same time,
- another process holds B and requests A

Both are blocked and remain so

#### Examples of computer resources

- printers
- memory space
- data (e.g. a locked record in a DB)
- semaphores

#### Resources

---

<sup>6</sup><http://kerneltrap.org/node/8208>

```

typedef int semaphore;
    semaphore resource_1;
    semaphore resource_2;

    void process_A(void) {
        down(&resource_1);
        down(&resource_2);
        use_both_resources( );
        up(&resource_2);
        up(&resource_1);
    }

    void process_B(void) {
        down(&resource_1);
        down(&resource_2);
        use_both_resources( );
        up(&resource_2);
        up(&resource_1);
    }

```

(a)

```

semaphore resource_1;
semaphore resource_2;

void process_A(void) {
    down(&resource_1);
    down(&resource_2);
    use_both_resources( );
    up(&resource_2);
    up(&resource_1);
}

void process_B(void) {
    down(&resource_2);
    down(&resource_1);
    use_both_resources( );
    up(&resource_1);
    up(&resource_2);
}

```

(b)

## Resources

### Deadlocks occur when ...

processes are granted exclusive access to *resources*  
e.g. devices, data records, files, ...

**Preemptable resources** can be taken away from a process with no ill effects

e.g. memory

**Nonpreemptable resources** will cause the process to fail if taken away

e.g. CD recorder

In general, deadlocks involve nonpreemptable resources.

## Resources

### Sequence of events required to use a resource

request	use	release
open()		close()
allocate()		free()

### What if request is denied?

Requesting process

- may be blocked
- may fail with error code

## 11.2 Introduction to Deadlocks

### The Best Illustration of a Deadlock

#### A law passed by the Kansas legislature early in the 20th century

“When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.”

### Introduction to Deadlocks

#### Deadlock

A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.

- Usually the event is release of a currently held resource

- None of the processes can ...
  - run
  - release resources
  - be awakened

#### Four Conditions For Deadlocks

**Mutual exclusion condition** each resource can only be assigned to one process or is available

**Hold and wait condition** process holding resources can request additional

**No preemption condition** previously granted resources cannot forcibly taken away

**Circular wait condition**

- must be a circular chain of 2 or more processes
- each is waiting for resource held by next member of the chain

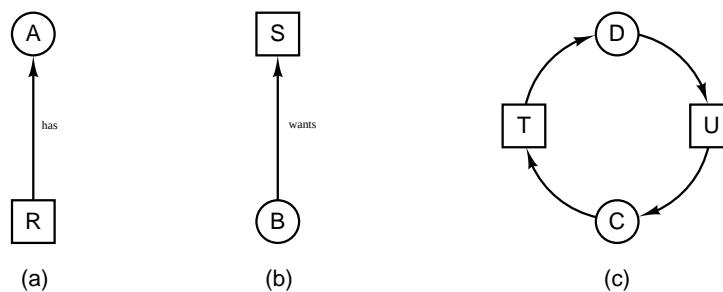
#### Four Conditions For Deadlocks

**Unlocking a deadlock is to answer 4 questions:**

1. Can a resource be assigned to more than one process at once?
2. Can a process hold a resource and ask for another?
3. can resources be preempted?
4. Can circular waits exits?

### 11.3 Deadlock Modeling

#### Resource-Allocation Graph



#### Strategies for dealing with Deadlocks

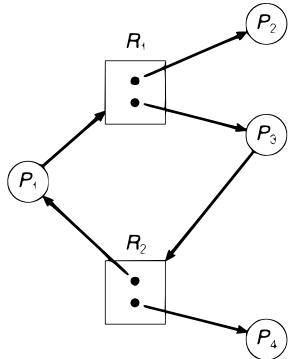
1. detection and recovery
2. dynamic avoidance — careful resource allocation
3. prevention — negating one of the four necessary conditions
4. just ignore the problem altogether

#### Resource-Allocation Graph



- The right graph has deadlock

### Resource-Allocation Graph



#### Basic facts:

- No cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock

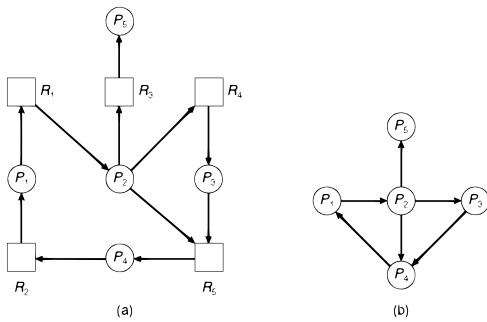
## 11.4 Deadlock Detection and Recovery

### Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

### Deadlock Detection

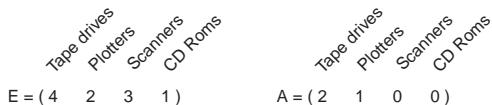
*Single Instance of Each Resource Type — Wait-for Graph*



- $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ ;
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.

### Deadlock Detection

*Several Instances of a Resource Type*



Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

Row n:

C: current allocation to process n

R: current requirement of process n

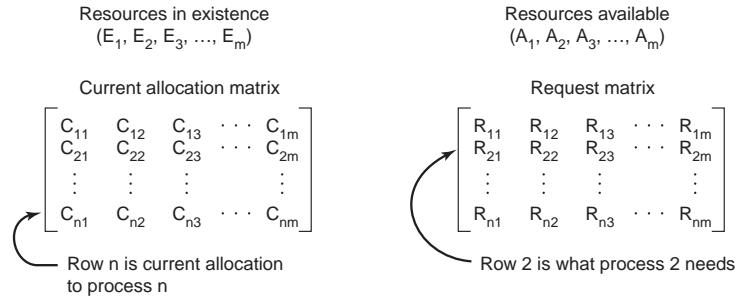
Column m:

C: current allocation of resource class m

R: current requirement of resource class m

## Deadlock Detection

*Several Instances of a Resource Type*



$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

e.g.

$$(C_{13} + C_{23} + \dots + C_{n3}) + A_3 = E_3$$

**n:** number of processes;

**m:** number of resource classes;

**E:** a vector of existing resources

- $E = [E_1, E_2, \dots, E_m]$
- $E_i = 2$  means system has 2 resources of class  $i$ , ( $1 \leq i \leq m$ );

**A:** a vector of available resources;

- $A = [A_1, A_2, \dots, A_m]$
- $A_i = 2$  means system has 2 resources of class  $i$  left unassigned;

$C_{ij}$ : is the number of instances of resource  $j$  that process  $i$  holds;

e.g.  $C_{31} = 2$  means  $P_3$  has 2 resources of class 1;

$R_{ij}$ : is the number of instances of resource  $j$  that process  $i$  wants;

e.g.  $R_{43} = 2$  means  $P_4$  wants 2 resources of class 3;

### Maths recall: vectors comparison

For two vectors,  $X$  and  $Y$

$$X \leq Y \quad \text{iff} \quad X_i \leq Y_i \quad \text{for } 0 \leq i \leq m$$

e.g.

$$\begin{bmatrix} 1 & 2 & 3 & 4 \end{bmatrix} \leq \begin{bmatrix} 2 & 3 & 4 & 4 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 3 & 4 \end{bmatrix} \not\leq \begin{bmatrix} 2 & 3 & 2 & 4 \end{bmatrix}$$

## Deadlock Detection

*Several Instances of a Resource Type*

$E = (4 \quad 2 \quad 3 \quad 1)$ Tape drives Plotters Scanners CD Roms	$A = (2 \quad 1 \quad 0 \quad 0)$ Tape drives Plotters Scanners CD Roms
---	---

Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

$A$

$$(2 \ 1 \ 0 \ 0) \geq R_3, (2 \ 1 \ 0 \ 0)$$

$$(2 \ 2 \ 2 \ 0) \geq R_2, (1 \ 0 \ 1 \ 0)$$

$$(4 \ 2 \ 2 \ 1) \geq R_1, (2 \ 0 \ 0 \ 1)$$

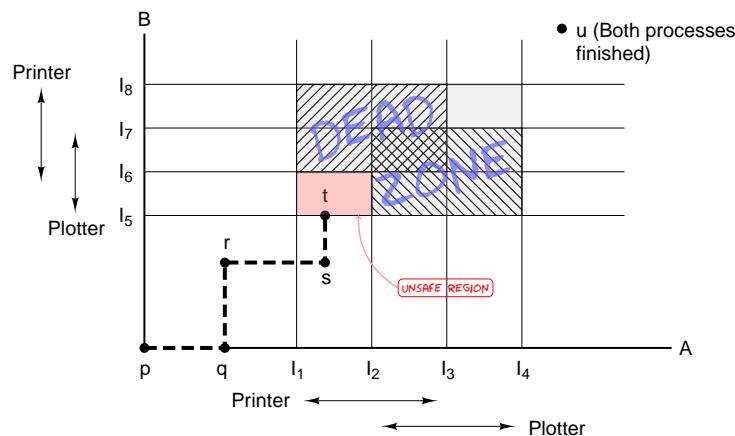
## Recovery From Deadlock

- Recovery through preemption
  - take a resource from some other process
  - depends on nature of the resource
- Recovery through rollback
  - checkpoint a process periodically
  - use this saved state
  - restart the process if it is found deadlocked
- Recovery through killing processes

## 11.5 Deadlock Avoidance

### Deadlock Avoidance

*Resource Trajectories*



- B is requesting a resource at point t. The system must decide whether to grant it or not.
- Deadlock is unavoidable if you get into *unsafe region*.

### Deadlock Avoidance

*Safe and Unsafe States*

Assuming  $E = 10$

#### Unsafe

Has Max			Has Max			Has Max			Has Max		
A	3	9	A	4	9	A	4	9	A	4	9
B	2	4	B	2	4	B	4	4	B	—	—
C	2	7	C	2	7	C	2	7	C	2	7
Free: 3			Free: 2			Free: 0			Free: 4		
(a)	(b)	(c)	(d)								

#### Safe

Has Max			Has Max			Has Max			Has Max		
A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	—	B	0	—
C	2	7	C	2	7	C	2	7	C	7	7
Free: 3			Free: 1			Free: 5			Free: 0		
(a)	(b)	(c)	(d)	(e)							

- Given totally 10 resources, for process A, B, C:

- A has 3, and need 6 more
- B has 2, and need 2 more
- C has 2, and need 5 more
- 3 left available

- allocate 1 to A,
  - A has 4, and need 5 more
  - B unchange
  - C unchange
  - 2 left available
- ...

## Deadlock Avoidance

*The Banker's Algorithm for a Single Resource*

The banker's algorithm considers each request as it occurs, and sees if granting it leads to a safe state.

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10  
(a)

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2  
(b)

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1  
(c)

*unsafe ≠ deadlock*

## Deadlock Avoidance

*The Banker's Algorithm for Multiple Resources*

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

$E = (6342)$   
 $P = (5322)$   
 $A = (1020)$

$$\begin{aligned} D \rightarrow A \rightarrow B, C, E \\ D \rightarrow E \rightarrow A \rightarrow B, C \end{aligned}$$

## Deadlock Avoidance

*Mission Impossible*

### In practice

- processes rarely know in advance their max future resource needs;
- the number of processes is not fixed;
- the number of available resources is not fixed.

**Conclusion:** Deadlock avoidance is essentially a mission impossible.

## 11.6 Deadlock Prevention

### Deadlock Prevention

*Break The Four Conditions*

#### Attacking the Mutual Exclusion Condition

- For example, using a *printing daemon* to avoid exclusive access to a printer.
- Not always possible
  - Not required for sharable resources;
  - must hold for nonsharable resources.
- The best we can do is to avoid mutual exclusion as much as possible
  - Avoid assigning a resource when not really necessary
  - Try to make sure as few processes as possible may actually claim the resource

See also: [Modern Operating Systems, Sec. 6.6.1, Attacking the Mutual Exclusion Condition, p. 452] for the *printer daemon* example.

#### Attacking the Hold and Wait Condition

Must guarantee that whenever a process requests a resource, it does not hold any other resources.

**Try:** the processes must request all their resources before starting execution  
if everything is available

```

then can run
  if one or more resources are busy
    then nothing will be allocated (just wait)

```

**Problem:**

- many processes don't know what they will need before running
- Low resource utilization; starvation possible

**Attacking the No Preemption Condition**

```

if a process that is holding some resources requests another resource that cannot be immediately allocated to it
then 1. All resources currently being held are released
      2. Preempted resources are added to the list of resources for which the process is waiting
      3. Process will be restarted only when it can regain its old resources, as well as the new ones that it is
          requesting

```

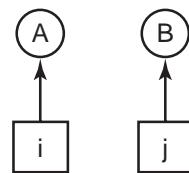
Low resource utilization; starvation possible

**Attacking Circular Wait Condition**

Impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD Rom drive

(a)



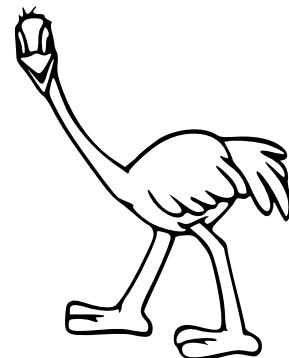
(b)

It's hard to find an ordering that satisfies everyone.

## 11.7 The Ostrich Algorithm

**The Ostrich Algorithm**

- Pretend there is no problem
- Reasonable if
  - deadlocks occur very rarely
  - cost of prevention is high
- UNIX and Windows takes this approach
- It is a trade off between
  - convenience
  - correctness



## References

- [1] Wikipedia. Deadlock — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Deadlock&oldid=645602687>.

## Part III

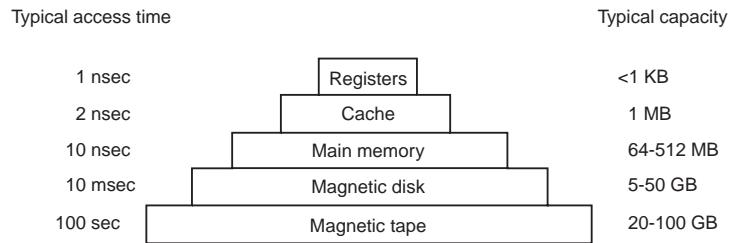
# Memory Management

## 12 Background

**Memory Management**

**In a perfect world** Memory is *large, fast, non-volatile*

**In real world ...**



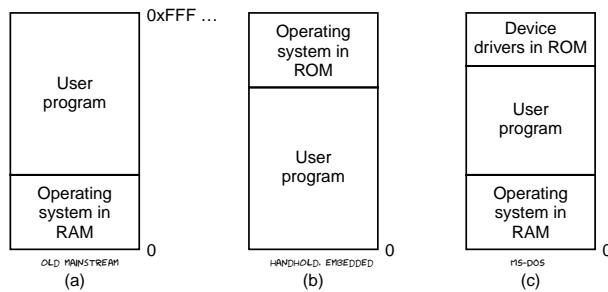
Memory manager handles the memory hierarchy.

## Basic Memory Management

### Real Mode

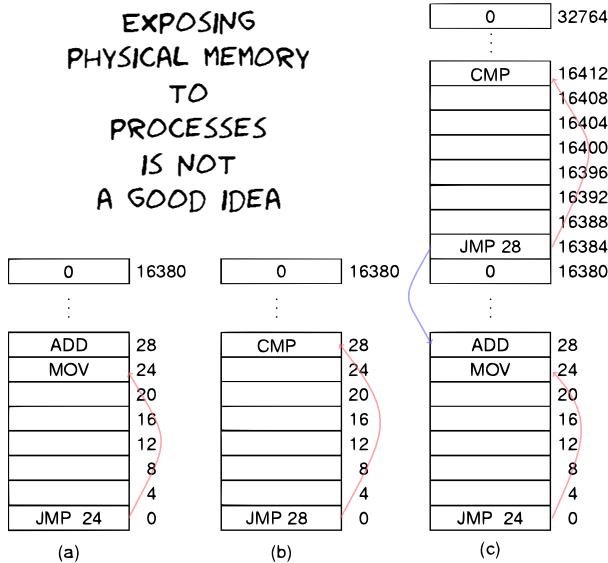
In the old days ...

- Every program simply saw the physical memory
- mono-programming without swapping or paging



## Basic Memory Management

### Relocation Problem



- (a) only one program in memory
- (b) only another program in memory
- (c) both in memory

## Memory Protection

### Protected mode

We need

- Protect the OS from access by user programs
- Protect user programs from one another

**Protected mode** is an operational mode of x86-compatible CPU.

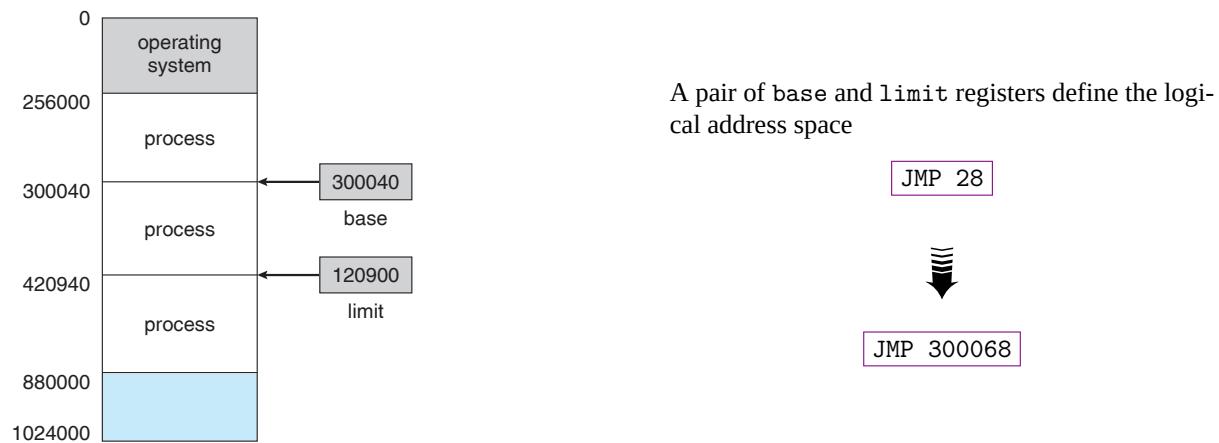
- The purpose is to protect everyone else (including the OS) from your program.

## Memory Protection

### Logical Address Space

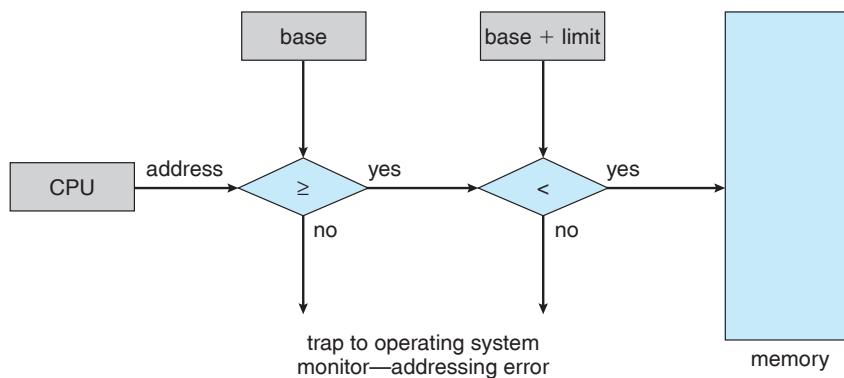
**Base register** holds the smallest legal physical memory address

**Limit register** contains the size of the range

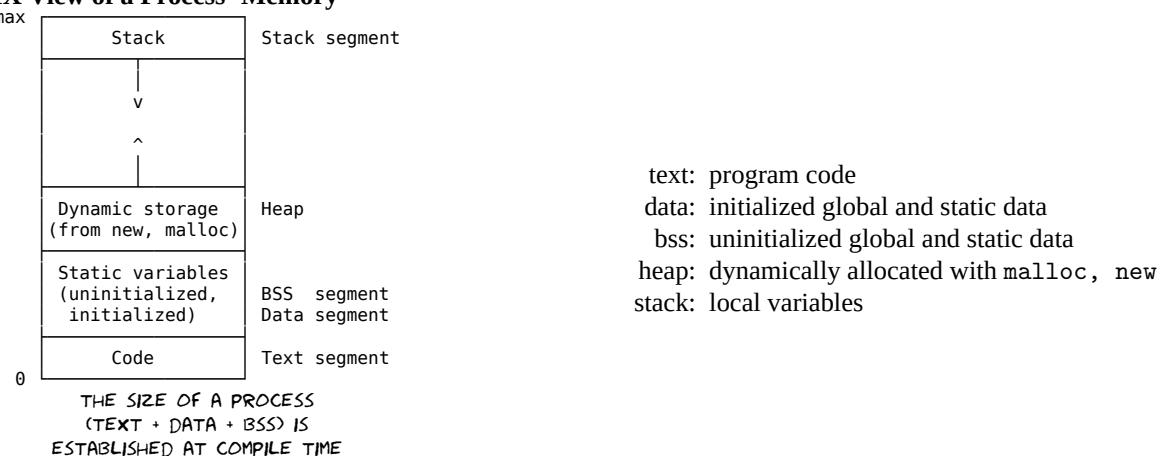


## Memory Protection

### Base and limit registers



## UNIX View of a Process' Memory



## Stack vs. Heap

Stack	Heap
compile-time allocation	run-time allocation
auto clean-up	you clean-up
inflexible	flexible
smaller	bigger
quicker	slower

## How large is the ...

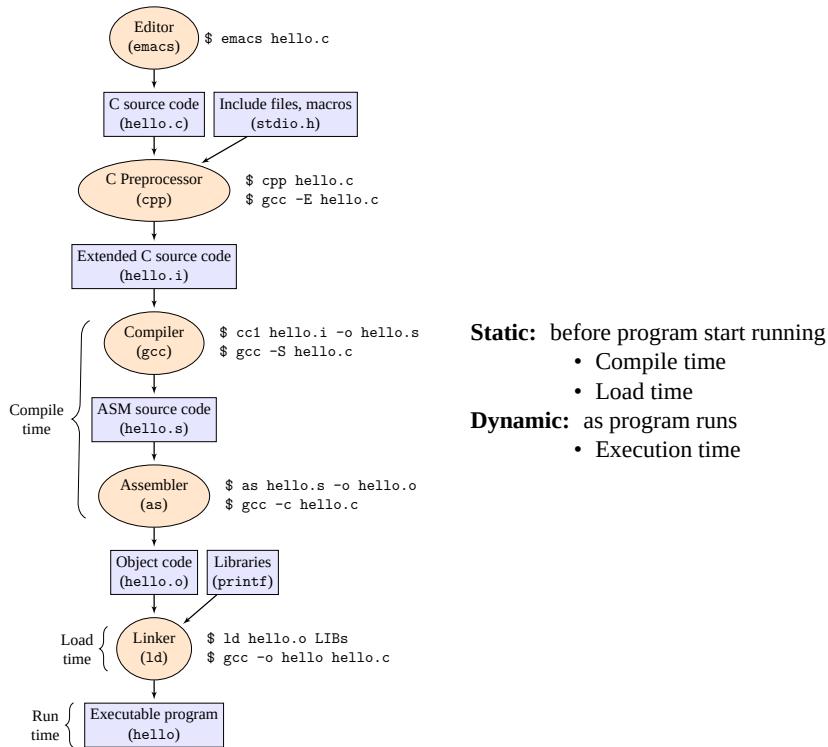
**stack:** ulimit -s

**heap:** could be as large as your virtual memory

**text|data|bss:** size a.out

## Multi-step Processing of a User Program

When is space allocated?



**Static:** before program start running

- Compile time
- Load time

**Dynamic:** as program runs

- Execution time

**Compiler** The name "compiler" is primarily used for programs that translate source code from a high-level programming language to a lower level language (e.g., assembly language or machine code)[*Compiler — Wikipedia, The Free Encyclopedia*].

**Assembler** An assembler creates object code by translating assembly instruction mnemonics into opcodes, and by resolving symbolic names for memory locations and other entities[*Assembly language — Wikipedia, The Free Encyclopedia*].

**Linker** Computer programs typically comprise several parts or modules; all these parts/modules need not be contained within a single object file, and in such case refer to each other by means of symbols[*Linker (computing) — Wikipedia, The Free Encyclopedia*].

When a program comprises multiple object files, the linker combines these files into a unified executable program, resolving the symbols as it goes along.

Linkers can take objects from a collection called a library. Some linkers do not include the whole library in the output; they only include its symbols that are referenced from other object files or libraries. Libraries exist for diverse purposes, and one or more system libraries are usually linked in by default.

The linker also takes care of arranging the objects in a program's address space. This may involve relocating code that assumes a specific base address to another base. Since a compiler seldom knows where an object will reside, it often assumes a fixed base location (for example, zero).

**Loader** An assembler creates object code by translating assembly instruction mnemonics into opcodes, and by resolving symbolic names for memory locations and other entities. ... Loading a program involves reading the contents of executable file, the file containing the program text, into memory, and then carrying out other required preparatory tasks to prepare the executable for running. Once loading is complete, the operating system starts the program by passing control to the loaded program code[*Loader (computing) — Wikipedia, The Free Encyclopedia*]

**Dynamic linker** A dynamic linker is the part of an operating system (OS) that loads (copies from persistent storage to RAM) and links (fills jump tables and relocates pointers) the shared libraries needed by an executable at run time, that is, when it is executed. The specific operating system and executable format determine how the dynamic linker functions and how it is implemented. Linking is often referred to as a process that is performed at compile time of the executable while a dynamic linker is in actuality a special loader that loads external shared libraries into a running process and then binds those shared libraries dynamically to the

running process. The specifics of how a dynamic linker functions is operating-system dependent [*Dynamic linker — Wikipedia, The Free Encyclopedia*]

Linkers and Loaders allow programs to be built from modules rather than as one big monolith.

See also:

- [Computer Systems: A Programmer's Perspective, Chap. 7, Linking].
- COMPILER, ASSEMBLER, LINKER AND LOADER: A BRIEF STORY<sup>7</sup>.
- Linkers and Loaders<sup>8</sup>.
- [Linkers and Loaders, Links and loaders].
- Linux Journal: Linkers and Loaders<sup>9</sup>. Discussing how compilers, links and loaders work and the benefits of shared libraries.

## Address Binding

*Who assigns memory to segments?*

### Static-binding: before a program starts running

**Compile time:** Compiler and assembler generate an object file for each source file

**Load time:**

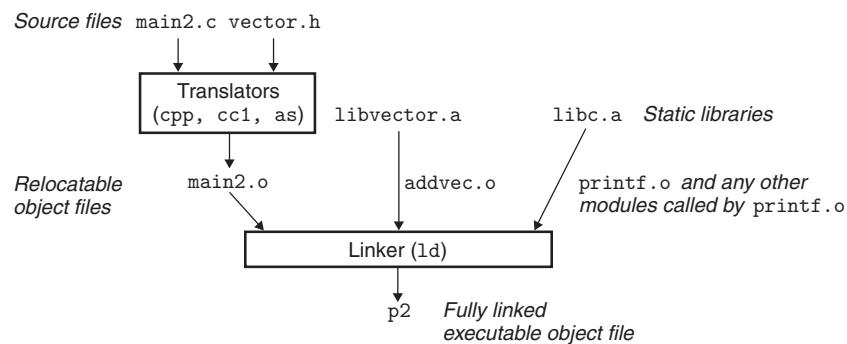
- Linker combines all the object files into a single executable object file
- Loader (part of OS) loads an executable object file into memory at location(s) determined by the OS
  - invoked via the execve system call

### Dynamic-binding: as program runs

- Execution time:
  - uses new and malloc to dynamically allocate memory
  - gets space on stack during function calls
- Address binding has nothing to do with physical memory (RAM). It determines the addresses of objects in the address space (virtual memory) of a process.

## Static loading

- The entire program and all data of a process must be in physical memory for the process to execute
- The size of a process is thus limited to the size of physical memory



## Dynamic Linking

A *dynamic linker* is actually a special loader that loads external shared libraries into a running process

- Small piece of code, stub, used to locate the appropriate memory-resident library routine
- Only one copy in memory
- Don't have to re-link after a library update

```

1 if ( stub_is_executed ){
2     if ( !routine_in_memory )
3         load_routine_into_memory();
4     stub_replaces_itself_with_routine();
5     execute_routine();
6 }

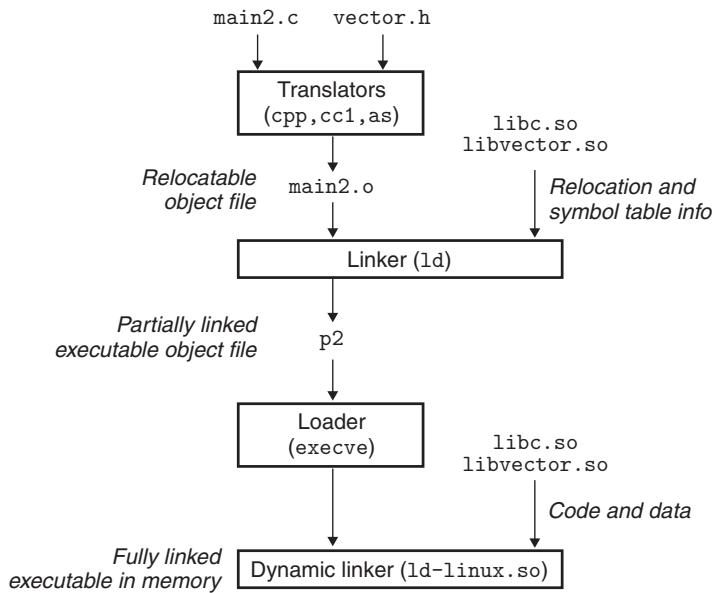
```

<sup>7</sup><http://www.tenouk.com/ModuleW.html>

<sup>8</sup><http://www.iecc.com/linker/>

<sup>9</sup><http://www.linuxjournal.com/article/6463>

## Dynamic Linking



**More about dynamic linking** Many operating system environments allow dynamic linking, that is the postponing of the resolving of some undefined symbols until a program is run. That means that the executable code still contains undefined symbols, plus a list of objects or libraries that will provide definitions for these. Loading the program will load these objects/libraries as well, and perform a final linking. Dynamic linking needs no linker[*Linker (computing)* — Wikipedia, The Free Encyclopedia].

This approach offers two advantages:

- Often-used libraries (for example the standard system libraries) need to be stored in only one location, not duplicated in every single binary.
- If an error in a library function is corrected by replacing the library, all programs using it dynamically will benefit from the correction after restarting them. Programs that included this function by static linking would have to be re-linked first.

There are also disadvantages:

- Known on the Windows platform as "DLL Hell", an incompatible updated DLL will break executables that depended on the behavior of the previous DLL.
- A program, together with the libraries it uses, might be certified (e.g. as to correctness, documentation requirements, or performance) as a package, but not if components can be replaced. (This also argues against automatic OS updates in critical systems; in both cases, the OS and libraries form part of a qualified environment.)

## Logical vs. Physical Address Space

- Mapping logical address space to physical address space is central to MM

**Logical address** generated by the CPU; also referred to as *virtual address*

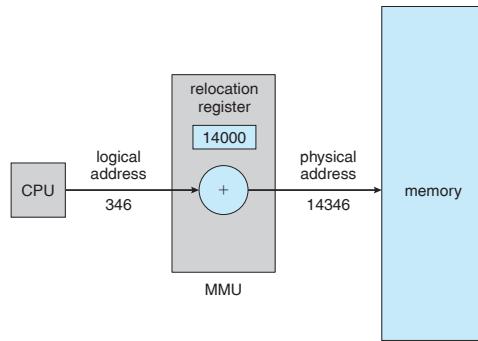
**Physical address** address seen by the memory unit

- In compile-time and load-time address binding schemes, LAS and PAS are identical in size
- In execution-time address binding scheme, they are differ.

## Logical vs. Physical Address Space

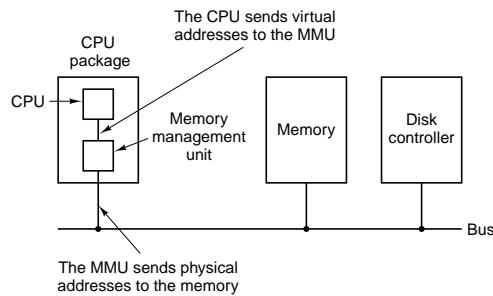
The user program

- deals with logical addresses
- never sees the real physical addresses

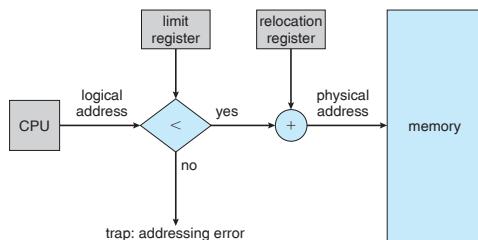


## MMU

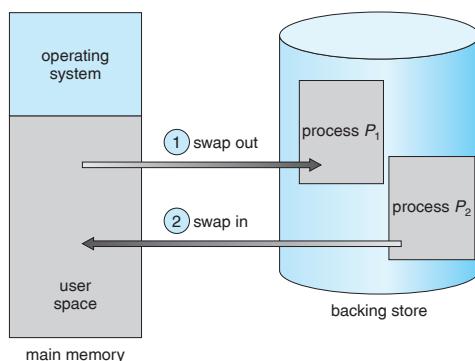
*Memory Management Unit*



## Memory Protection



## Swapping



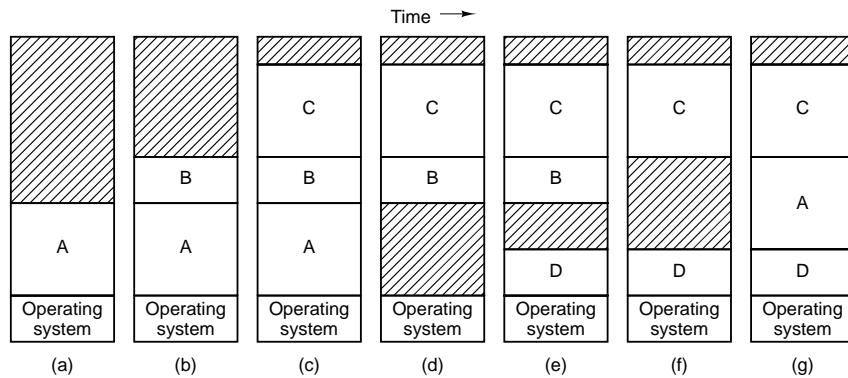
### Major part of swap time is transfer time

Total transfer time is directly proportional to the amount of memory swapped

## 13 Contiguous Memory Allocation

### Contiguous Memory Allocation

*Multiple-partition allocation*

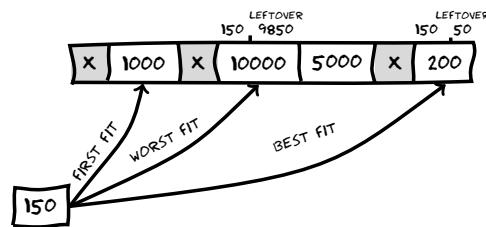


Operating system maintains information about:

- a allocated partitions
- b free partitions (hole)

### Dynamic Storage-Allocation Problem

*First Fit, Best Fit, Worst Fit*



**First-fit:** The first hole that is big enough

**Best-fit:** The smallest hole that is big enough

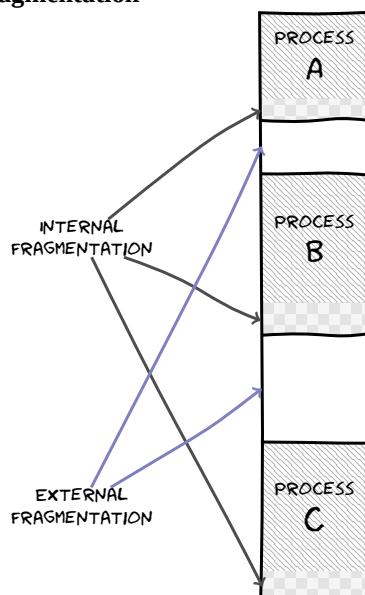
- Must search entire list, unless ordered by size
- Produces the smallest leftover hole

**Worst-fit:** The largest hole

- Must also search entire list
- Produces the largest leftover hole

- First-fit and best-fit better than worst-fit in terms of speed and storage utilization
- First-fit is generally faster

### Fragmentation



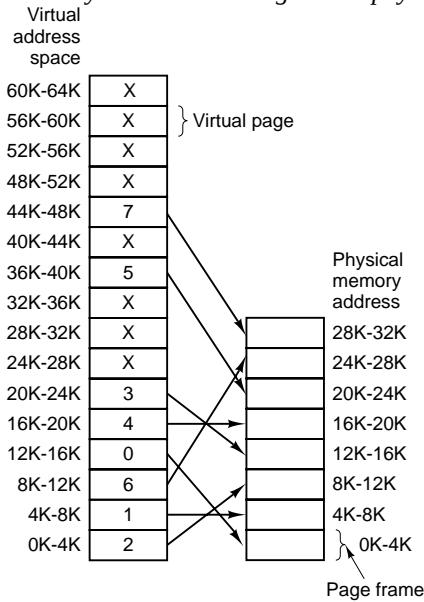
#### Reduce external fragmentation by

- *Compaction* is possible only if relocation is dynamic, and is done at execution time
- *Noncontiguous memory allocation*
  - Paging
  - Segmentation

# 14 Virtual Memory

## Virtual Memory

*Logical memory can be much larger than physical memory*



## Address translation

$$\begin{array}{ccc}
 \text{virtual} & \xrightarrow{\text{page table}} & \text{physical} \\
 \text{address} & & \text{address} \\
 \\ 
 \text{Page 0} & \xrightarrow{\text{map to}} & \text{Frame 2} \\
 \\ 
 0_{\text{virtual}} & \xrightarrow{\text{map to}} & 8192_{\text{physical}} \\
 \\ 
 20500_{\text{vir}} & \xrightarrow{\text{map to}} & 12308_{\text{phy}} \\
 (20k + 20)_{\text{vir}} & & (12k + 20)_{\text{phy}}
 \end{array}$$

## 14.1 Paging

### Paging

*Address Translation Scheme*

**Address generated by CPU is divided into:**

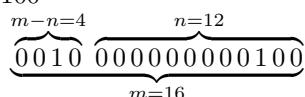
**Page number(p):** an index into a *page table*

**Page offset(d):** to be copied into memory

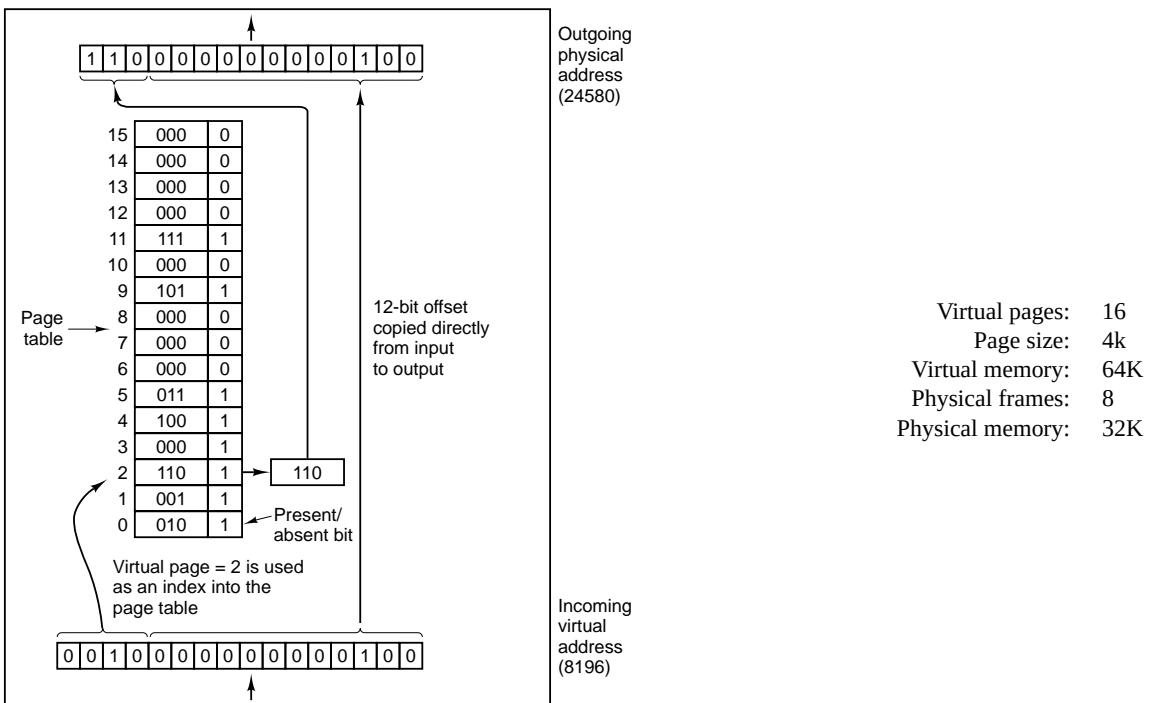
Given *logical address space* ( $2^m$ ) and *page size* ( $2^n$ ),

$$\text{number of pages} = \frac{2^m}{2^n} = 2^{m-n}$$

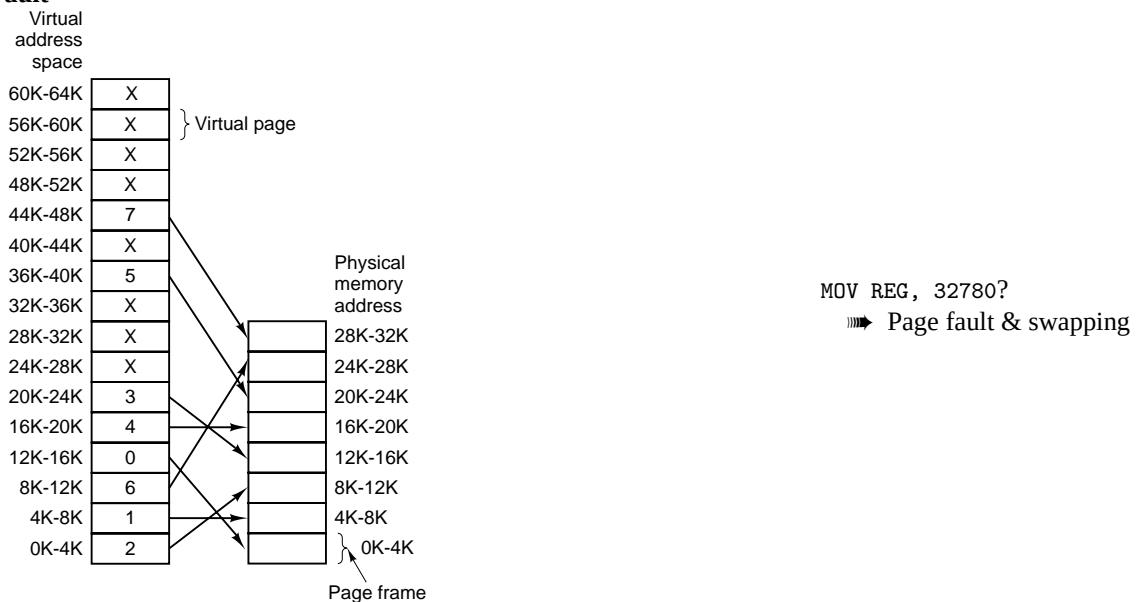
**Example: addressing to 0010000000000100**



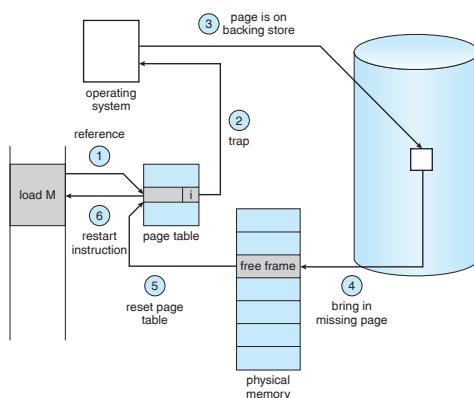
$$\text{page number} = 0010 = 2, \quad \text{page offset} = 000000000100$$



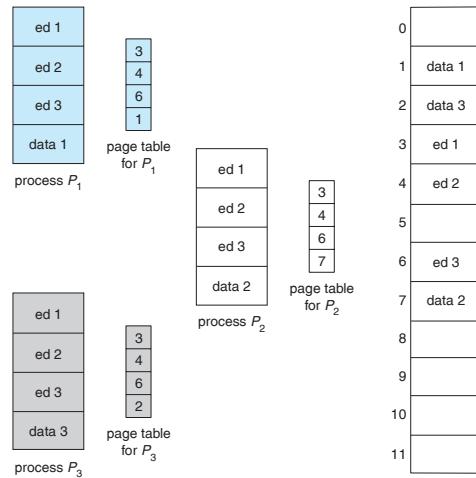
## Page Fault



## Page Fault Handling



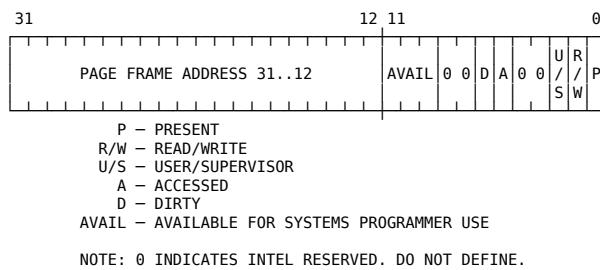
## Shared Pages



## Page Table Entry

### Intel i386 Page Table Entry

- Commonly 4 bytes (32 bits) long
- Page size is usually 4k ( $2^{12}$  bytes). OS dependent  
    \$ getconf PAGESIZE
- Could have  $2^{32-12} = 2^{20} = 1M$  pages  
    Could address  $1M \times 4KB = 4GB$  memory



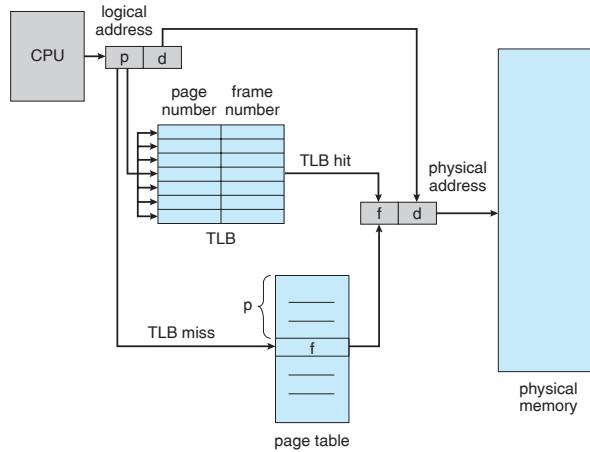
## Page Table

- Page table is kept in main memory
- Usually one page table for each process
- *Page-table base register (PTBR)*: A pointer to the page table is stored in PCB
- *Page-table length register (PRLR)*: indicates size of the page table
- Slow
  - Requires two memory accesses. One for the page table and one for the data/instruction.
- TLB

## Translation Lookaside Buffer (TLB)

### Fact: 80-20 rule

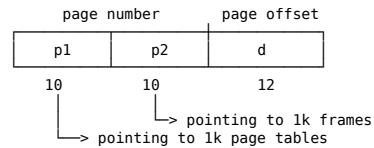
- Only a small fraction of the PTEs are heavily read; the rest are barely used at all



### Multilevel Page Tables

- a 1M-entry page table eats 4M memory
- while 100 processes running, 400M memory is gone for page tables
- avoid keeping all the page tables in memory all the time

### A two-level scheme



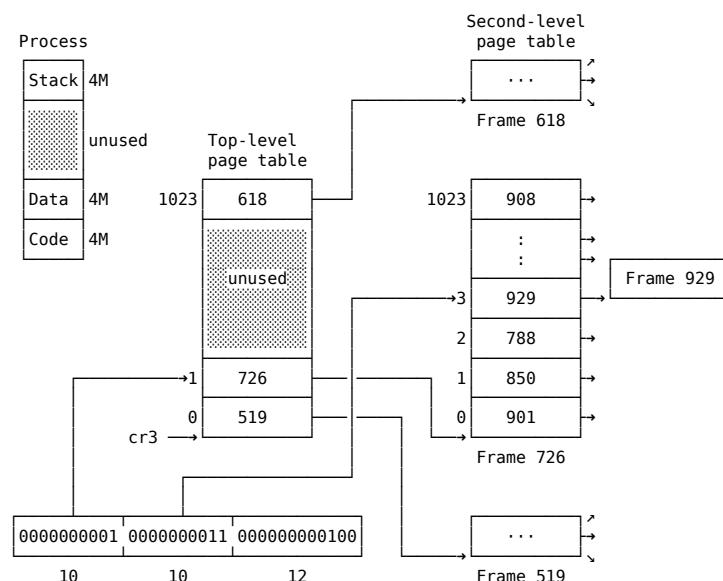
**p1:** is an index into the outer page table

**p2:** is the displacement within the page of the outer page table

- Split one huge page table into 1k small page tables
  - i.e. the huge page table has 1k entries.
  - Each entry keeps a page frame number of a small page table.
- Each small page table has 1k entries
  - Each entry keeps a page frame number of a physical frame.

### Two-Level Page Tables

*Example*



### Problem With 64-bit Systems

*Given:*

- virtual address space = 64 bits

- page size =  $4 KB = 2^{12} B$

? How much space would a simple single-level page table take?

if Each page table entry takes 4 Bytes

then The whole page table ( $2^{64-12}$  entries) will take

$$2^{64-12} \times 4 B = 2^{54} B = 16 PB \quad (\text{peta} \Rightarrow \text{tera} \Rightarrow \text{giga})!$$

And this is for ONE process!

**Multi-level?**

if 10 bits for each level

then  $\frac{64-12}{10} = 5$  levels are required

5 memory access for each address translation!

## Inverted Page Tables

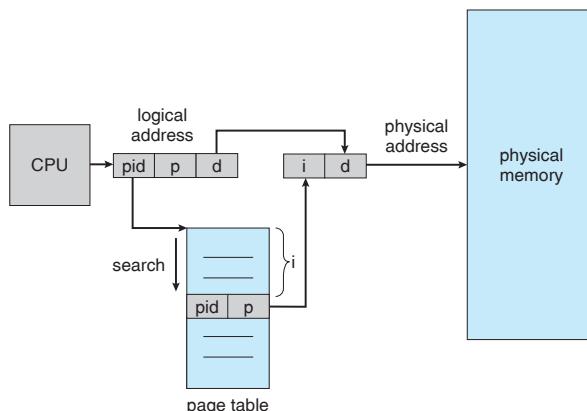
*Index with frame number*

### Inverted Page Table:

- One entry for each *physical* frame
  - The physical frame number is the table index
- A single global page table for all processes
  - The table is shared — PID is required
- Physical pages are now mapped to virtual — each entry contains a virtual page number instead of a physical one
- Information bits, e.g. protection bit, are as usual

### Find index according to entry contents

$(pid, p) \Rightarrow i$



Std. PTE (32-bit sys.):

page	frame address	info
20	12	

indexed by page number

if  $2^{20}$  entries, 4 B each  
then  $SIZE_{page\ table} = 2^{20} \times 4 = 4 MB$   
(for each process)

Inverted PTE (64-bit sys.):

pid	virtual page number	info
16	52	12

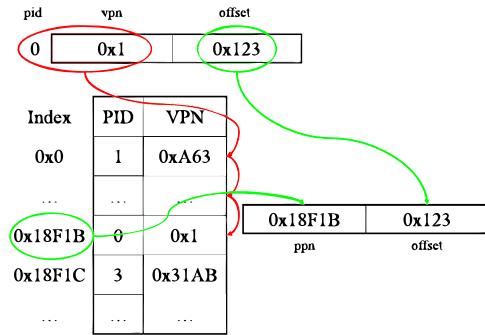
indexed by frame number

if assuming

- 16 bits for PID
- 52 bits for virtual page number
- 12 bits of information

then each entry takes  $16 + 52 + 12 = 80 \text{ bits} = 10 \text{ bytes}$   
if physical mem =  $1G (2^{30} B)$ , and page size =  $4K (2^{12} B)$ ,  
we'll have  $2^{30-12} = 2^{18} \text{ pages}$   
then  $SIZE_{page\ table} = 2^{18} \times 10 B = 2.5 MB$   
(for all processes)

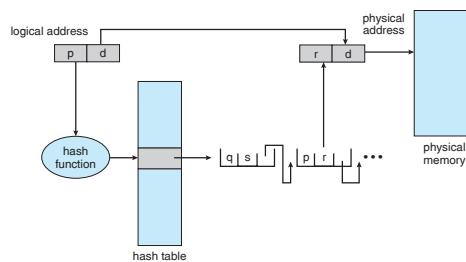
**Inefficient:** Require searching the entire table



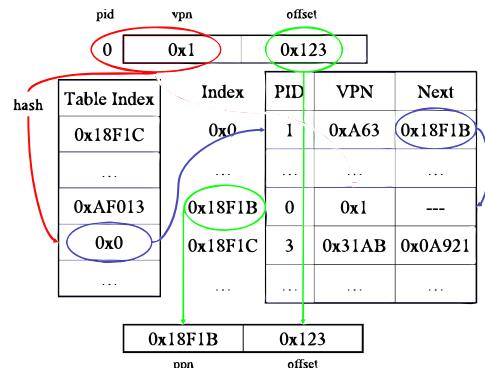
### Hashed Inverted Page Tables

A **hash anchor table** — an extra level before the actual page table

- maps process IDs virtual page numbers  $\Rightarrow$  page table entries
- Since collisions may occur, the page table must do chaining



### Hashed Inverted Page Table



## 14.2 Demand Paging

### Demand Paging

With demand paging, the size of the LAS is no longer constrained by physical memory

- Bring a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
- Lazy swapper* - never swaps a page into memory unless page will be needed
  - Swapper* deals with entire processes
  - Pager (Lazy swapper)* deals with pages

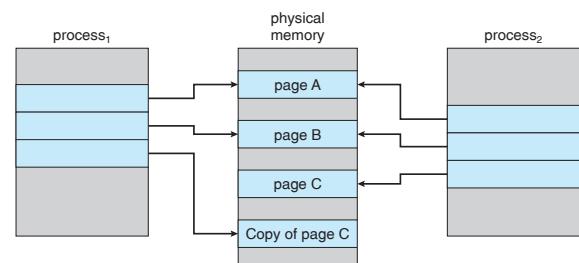
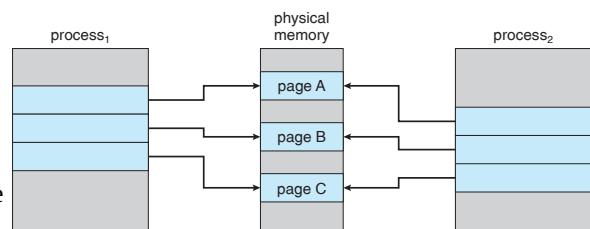
**More about demand paging** In the purest form of paging, processes are started up with none of their pages in memory. As soon as the CPU tries to fetch the first instruction, it gets a page fault, causing the operating system to bring in the page containing the first instruction. Other page faults for global variables and the stack usually follow quickly. After a while, the process has most of the pages it needs and settles down to run with relatively few page faults. This strategy is called *demand paging* because pages are loaded only on demand, not in advance [Modern Operating Systems, Sec. 3.4.8, P. 207].

### 14.3 Copy-on-Write

#### Copy-on-Write

More efficient process creation

- Parent and child processes initially share the same pages in memory
- Only the modified page is copied upon modification occurs
- Free pages are allocated from a *pool* of zeroed-out pages

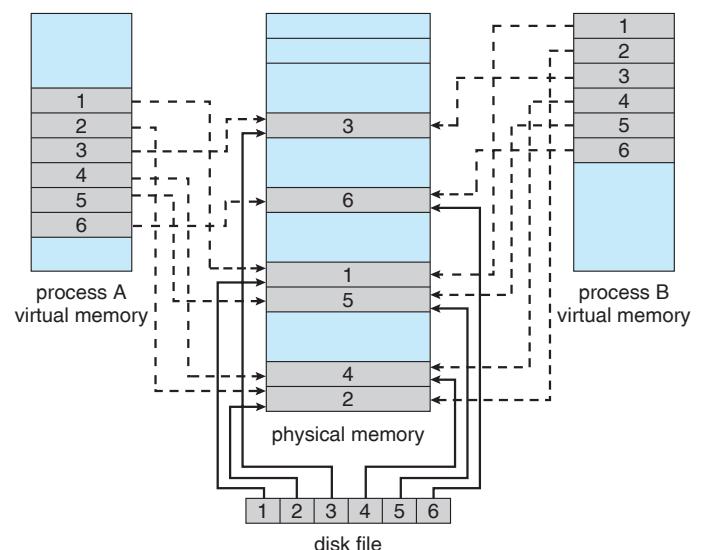


### 14.4 Memory mapped files

#### Memory Mapped Files

Mapping a file (disk block) to one or more memory pages

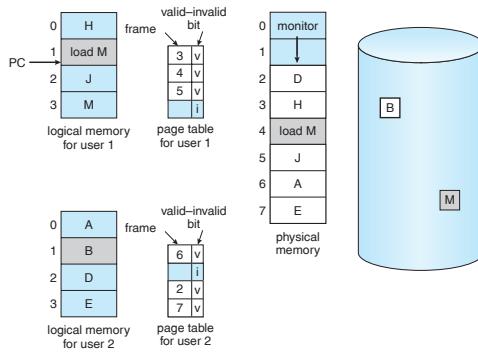
- Improved I/O performance — much faster than *read()* and *write()* system calls
- Lazy loading (*demand paging*) — only a small portion of file is loaded initially
- A mapped file can be shared, like *shared library*



### 14.5 Page Replacement Algorithms

#### Need For Page Replacement

**Page replacement:** find some page in memory, but not really in use, swap it out



Linux calls it the *Page Frame Reclaiming Algorithm*<sup>10</sup>, it's basically LRU with a bias towards non-dirty pages.  
See also:

- [Page replacement algorithm — Wikipedia, The Free Encyclopedia, Wikipedia: Page replacement algorithm].
- [Understanding The Linux Kernel, Chap. 17, Page Frame Reclaiming].
- PageReplacementDesign<sup>11</sup>.

### Performance Concern

Because disk I/O is so expensive, we must solve two major problems to implement demand paging.

**Frame-allocation algorithm** If we have multiple processes in memory, we must decide how many frames to allocate to each process.

**Page-replacement algorithm** When page replacement is required, we must select the frames that are to be replaced.

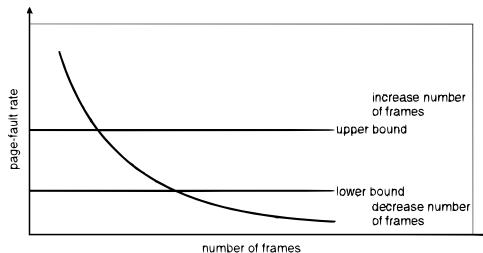
### Performance

We want an algorithm resulting in lowest page-fault rate

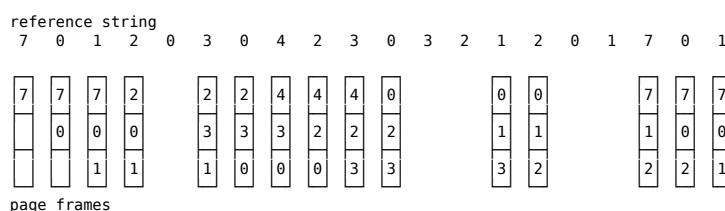
- Is the victim page modified?
- Pick a random page to swap out?
- Pick a page from the faulting process' own pages? Or from others?

### Page-Fault Frequency Scheme

Establish "acceptable" page-fault rate



### FIFO Page Replacement Algorithm



- Maintain a linked list (FIFO queue) of all pages
  - in order they came into memory
- Page at beginning of list replaced
- Disadvantage
  - The oldest page may be often used
  - Belady's anomaly

<sup>10</sup><http://stackoverflow.com/questions/5889825/page-replacement-algorithm>

<sup>11</sup><http://linux-mm.org/PageReplacementDesign>

## FIFO Page Replacement Algorithm

### *Belady's Anomaly*

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

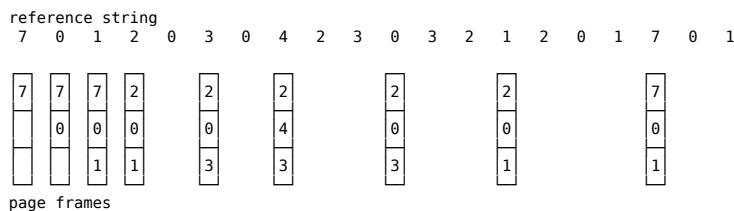
1	4	5	3																
2		1	4	4															
3		2	2	5															

- 4 frames

1	4	2																	
2	2	3																	
3	5	4																	
4	1	5																	

- Belady's Anomaly: *more frames  $\Rightarrow$  more page faults*

## Optimal Page Replacement Algorithm (OPT)



- Replace page needed at the farthest point in future
  - Optimal but not feasible
- Estimate by ...
  - logging page use on previous runs of process
  - although this is impractical, similar to SJF CPU-scheduling, it can be used for comparison studies

## Least Recently Used (LRU) Algorithm

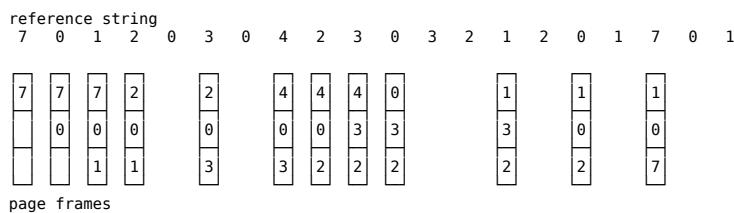
FIFO uses the time when a page was brought into memory

OPT uses the time when a page is to be used

LRU uses the recent past as an approximation of the near future

Replace the page that *has not been used* for the longest period of time

**Assume recently-used-pages will be used again soon**



## LRU Implementations

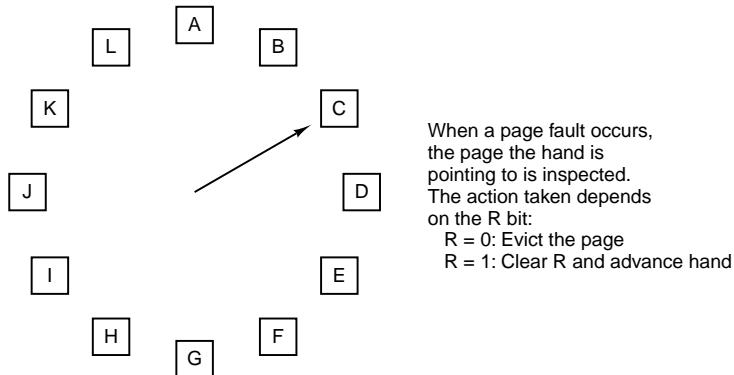
**Counters:** record the time of the last reference to each page

- choose page with lowest value counter
- Keep counter in each page table entry
- counter overflow — periodically zero the counter
- require a *search* of the page table to find the LRU page
- update time-of-use field in the page table *every memory reference!*

**Stack:** keep a linked list (stack) of pages

- most recently used at top, least (LRU) at bottom
  - no search for replacement
- whenever a page is referenced, it's removed from the stack and put on the top
  - update this list *every memory reference!*

## Second Chance Page Replacement Algorithm



## 14.6 Allocation of Frames

### Allocation of Frames

- Each process needs minimum number of pages
- Fixed Allocation
  - Equal allocation — e.g., 100 frames and 5 processes, give each process 20 frames.
  - Proportional allocation — Allocate according to the size of process

$$a_i = \frac{s_i}{\sum s_i} \times m$$

$s_i$ : size of process  $p_i$

$m$ : total number of frames

$a_i$ : frames allocated to  $p_i$

- Priority Allocation — Use a proportional allocation scheme using priorities rather than size

$$\frac{\text{priority}_i}{\sum \text{priority}_i} \quad \text{or} \quad \left( \frac{s_i}{\sum s_i}, \frac{\text{priority}_i}{\sum \text{priority}_i} \right)$$

### Global vs. Local Allocation

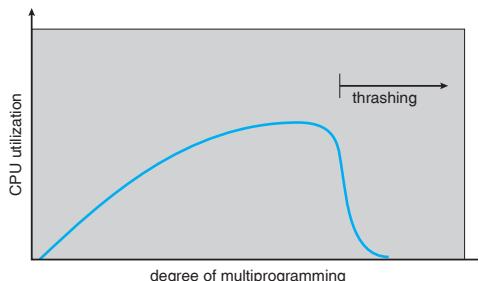
If process  $P_i$  generates a page fault, it can select a replacement frame

- from its own frames — *Local replacement*
- from the set of all frames; one process can take a frame from another — *Global replacement*
  - from a process with lower priority number

*Global replacement generally results in greater system throughput.*

## 14.7 Thrashing And Working Set Model

### Thrashing



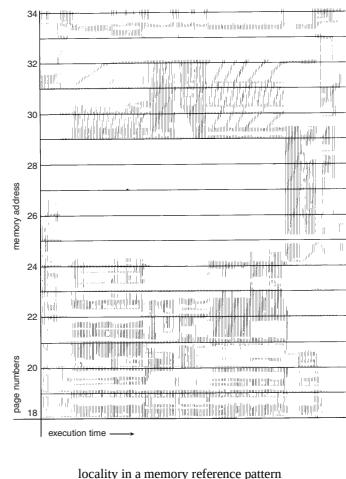
### Thrashing

1. CPU not busy  $\Rightarrow$  add more processes
2. a process needs more frames  $\Rightarrow$  faulting, and taking frames away from others
3. these processes also need these pages  $\Rightarrow$  also faulting, and taking frames away from others  $\Rightarrow$  chain reaction
4. more and more processes queueing for the paging device  $\Rightarrow$  ready queue is empty  $\Rightarrow$  CPU has nothing to do  $\Rightarrow$  add more processes  $\Rightarrow$  more page faults
5. MMU is busy, but no work is getting done, because processes are busy paging — *thrashing*

## Demand Paging and Thrashing

### Locality Model

- A *locality* is a set of pages that are actively used together
- Process migrates from one locality to another

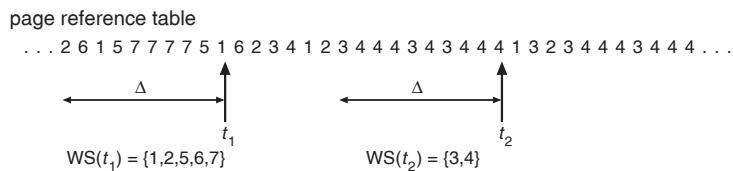


### Why does thrashing occur?

$$\sum_{i=(0,n)} Locality_i > total\ memory\ size$$

### Working-Set Model

**Working Set (WS)** The set of pages that a process is currently( $\Delta$ ) using. ( $\approx$  locality)



$\Delta$ : Working-set window. In this example,

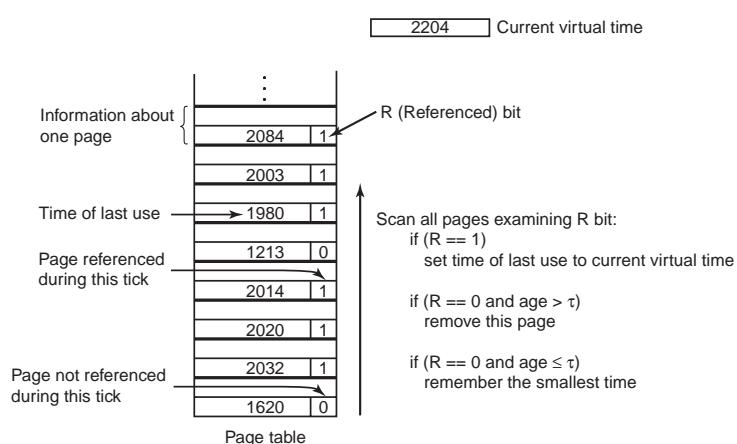
$$\Delta = 10 \text{ memory access}$$

**WSS:** Working-set size.  $WS(t_1) = \{\overbrace{1, 2, 5, 6, 7}\}$

- The accuracy of the working set depends on the selection of  $\Delta$
- Thrashing, if  $\sum WSS_i > SIZE_{total\ memory}$

### The Working-Set Page Replacement Algorithm

To evict a page that is not in the working set



$$age = Current\ virtual\ time - Time\ of\ last\ use$$

**Current virtual time** The amount of CPU time a process has actually used since it started is often called its *current virtual time*[Modern Operating Systems, sec 3.4.8, p 209].

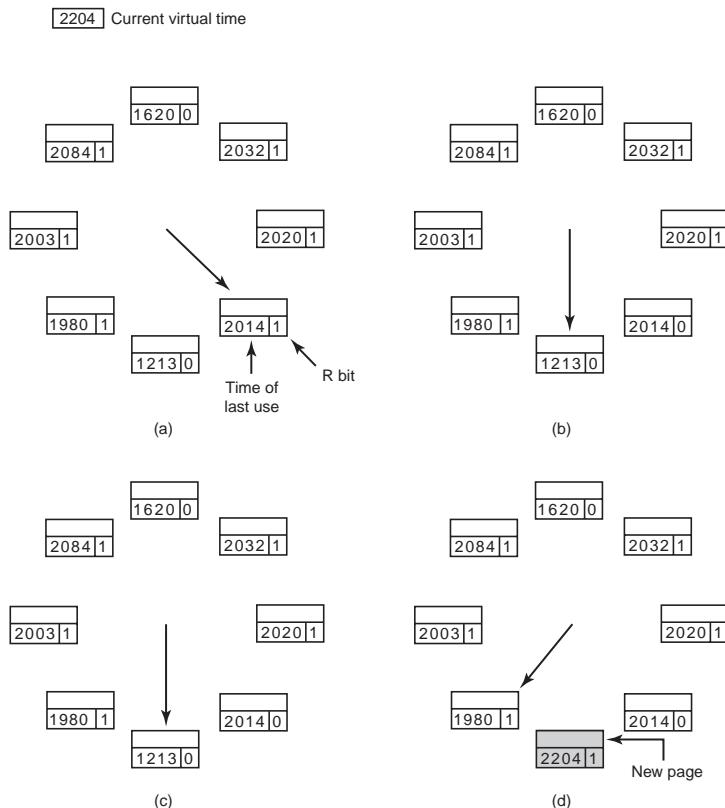
The algorithm works as follows. The hardware is assumed to set the R and M bits, as discussed earlier. Similarly, a periodic clock interrupt is assumed to cause software to run that clears the Referenced bit on every clock tick. On every page fault, the page table is scanned to look for a suitable page to evict[*Modern Operating Systems*, sec 3.4.8, p 210].

As each entry is processed, the R bit is examined. If it is 1, the current virtual time is written into the *Time of last use* field in the page table, indicating that the page was in use at the time the fault occurred. Since the page has been referenced during the current clock tick, it is clearly in the working set and is not a candidate for removal ( $\tau$  is assumed to span multiple clock ticks).

If R is 0, the page has not been referenced during the current clock tick and may be a candidate for removal. To see whether or not it should be removed, its age (the current virtual time minus its *Time of last use*) is computed and compared to  $\tau$ . If the age is greater than  $\tau$ , the page is no longer in the working set and the new page replaces it. The scan continues updating the remaining entries.

However, if R is 0 but the age is less than or equal to  $\tau$ , the page is still in the working set. The page is temporarily spared, but the page with the greatest age (smallest value of *Time of last use*) is noted. If the entire table is scanned without finding a candidate to evict, that means that all pages are in the working set. In that case, if one or more pages with R = 0 were found, the one with the greatest age is evicted. In the worst case, all pages have been referenced during the current clock tick (and thus all have R = 1), so one is chosen at random for removal, preferably a clean page, if one exists.

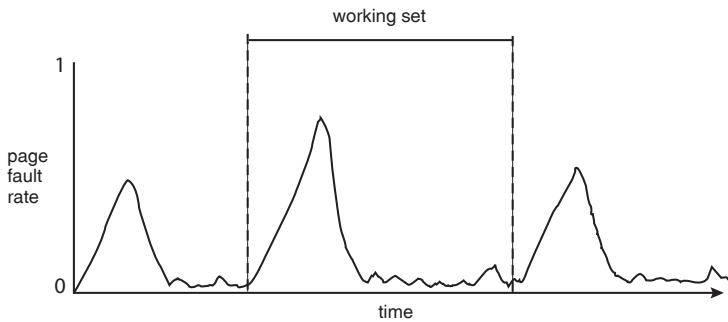
### The WSClock Page Replacement Algorithm Combine Working Set Algorithm With Clock Algorithm



The basic working set algorithm is cumbersome, since the entire page table has to be scanned at each page fault until a suitable candidate is located. An improved algorithm, that is based on the clock algorithm but also uses the working set information, is called *WSClock* (Carr and Hennessey, 1981). Due to its simplicity of implementation and good performance, it is widely used in practice[*Modern Operating Systems*, Sec 3.4.9, P. 211].

## 14.8 Other Issues

### Other Issues — Prepaging



- reduce faulting rate at (re)startup
  - remember working-set in PCB
- Not always work
  - if prepaged pages are unused, I/O and memory was wasted

### Other Issues — Page Size

Larger page size

- ⊗ Bigger internal fragmentation
- ⊗ longer I/O time

Smaller page size

- ⊗ Larger page table
- ⊗ more page faults
  - one page fault for each byte, if page size = 1 Byte
  - for a  $200K$  process, with page size =  $200K$ , only one page fault

No best answer

`$ getconf PAGESIZE`

### Other Issues — TLB Reach

- Ideally, the working set of each process is stored in the TLB
  - Otherwise there is a high degree of page faults
- *TLB Reach* — The amount of memory accessible from the TLB

$$TLB\ Reach = (TLB\ Size) \times (Page\ Size)$$

- Increase the page size
  - Internal fragmentation may be increased
- Provide multiple page sizes
  - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation
    - \* UltraSPARC supports page sizes of  $8KB$ ,  $64KB$ ,  $512KB$ , and  $4MB$
    - \* Pentium supports page sizes of  $4KB$  and  $4MB$

### Other Issues — Program Structure

Careful selection of data structures and programming structures can increase locality, i.e. lower the page-fault rate and the number of pages in the working set.

#### Example

- A stack has a good locality, since access is always made to the top
- A hash table has a bad locality, since it's designed to scatter references
- Programming language
  - Pointers tend to randomize access to memory
  - OO programs tend to have a poor locality

### Other Issues — Program Structure

#### Example

- `int[i][j] = int[128][128]`
- Assuming page size is 128 words, then
- Each row (128 words) takes one page

If the process has fewer than 128 frames...

#### Program 1:

```
for(j=0;j<128;j++)
    for(i=0;i<128;i++)
        data[i][j] = 0;
```

Worst case:

$$128 \times 128 = 16,384 \text{ page faults}$$

See also: [Operating System Concepts Essentials, Sec. 8.9.5, Program Structure].

#### Program 2:

```
for(i=0;i<128;i++)
    for(j=0;j<128;j++)
        data[i][j] = 0;
```

Worst case:

$$128 \text{ page faults}$$

### Other Issues — I/O interlock

Sometimes it is necessary to lock pages in memory so that they are not paged out.

#### Example

- The OS
- I/O operation — the frame into which the I/O device was scheduled to write should not be replaced.
- New page that was just brought in — looks like the best candidate to be replaced because it was not accessed yet, nor was it modified.

### Other Issues — I/O interlock

#### Case 1

##### Be sure the following sequence of events does not occur:

1. A process issues an I/O request, and then queueing for that I/O device
2. The CPU is given to other processes
3. These processes cause page faults
4. The waiting process' page is unluckily replaced
5. When its I/O request is served, the specific frame is now being used by another process

### Other Issues — I/O interlock

#### Case 2

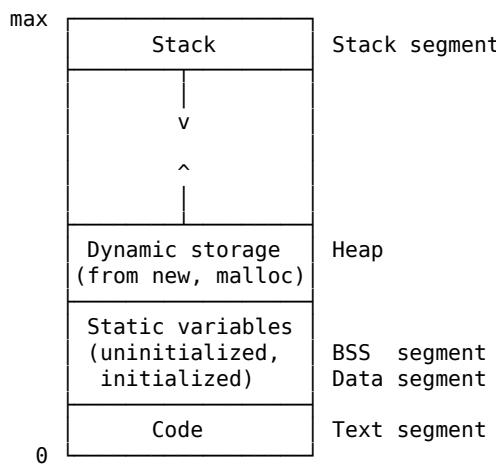
##### Another bad sequence of events:

1. A low-priority process faults
2. The paging system selects a replacement frame. Then, the necessary page is loaded into memory
3. The low-priority process is now ready to continue, and waiting in the ready queue
4. A high-priority process faults
5. The paging system looks for a replacement
  - (a) It sees a page that is in memory but not been referenced nor modified: perfect!
  - (b) It doesn't know the page is just brought in for the low-priority process

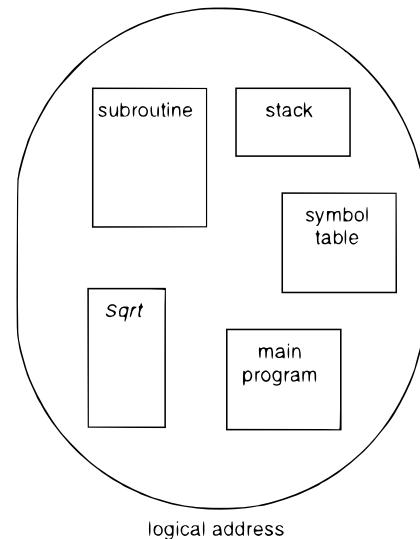
## 14.9 Segmentation

### Two Views of A Virtual Address Space

**One-dimensional**  
a linear array of bytes



**Two-dimensional**  
a collection of variable-sized segments

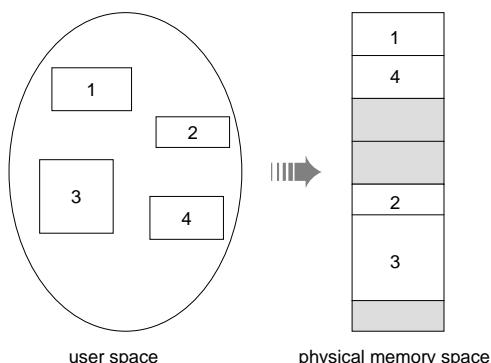


### User's View

- A program is a collection of segments
- A segment is a logical unit such as:

main program	procedure	function
method	object	local variables
global variables	common block	stack
symbol table	arrays	

### Logical And Physical View of Segmentation



### Segmentation Architecture

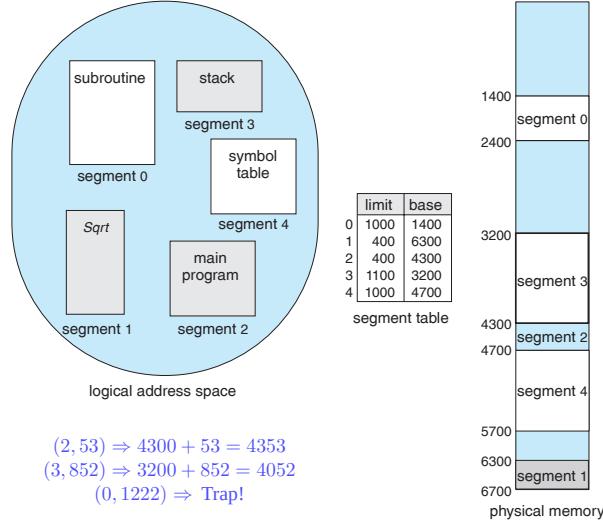
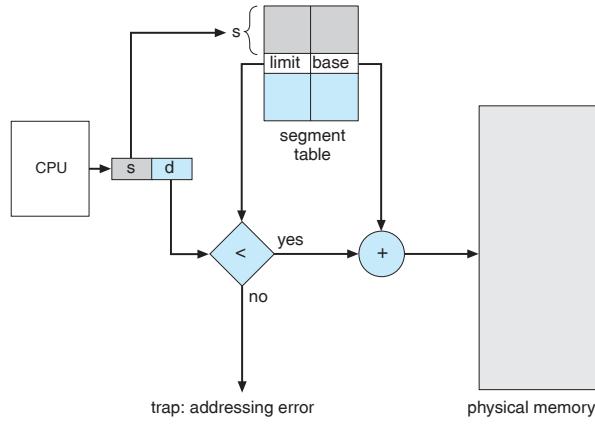
- Logical address consists of a *two tuple*:

<segment-number, offset>

- *Segment table* maps 2D virtual addresses into 1D physical addresses; each table entry has:
  - *base* contains the starting physical address where the segments reside in memory
  - *limit* specifies the length of the segment
- *Segment-table base register (STBR)* → the segment table's location in memory
- *Segment-table length register (STLR)* indicates number of segments used by a program;

segment number s is legal if s < STLR

Segmentation hardware



## Advantages of Segmentation

- Each segment can be
  - located independently
  - separately protected
  - grow independently
- Segments can be shared between processes

## Problems with Segmentation

- Variable allocation
- Difficult to find holes in physical memory
- Must use one of non-trivial placement algorithm
  - first fit, best fit, worst fit
- External fragmentation

See also: <http://cseweb.ucsd.edu/classes/fa03/cse120/Lec08.pdf>.

## Linux prefers paging to segmentation

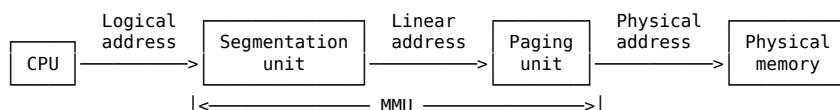
### Because

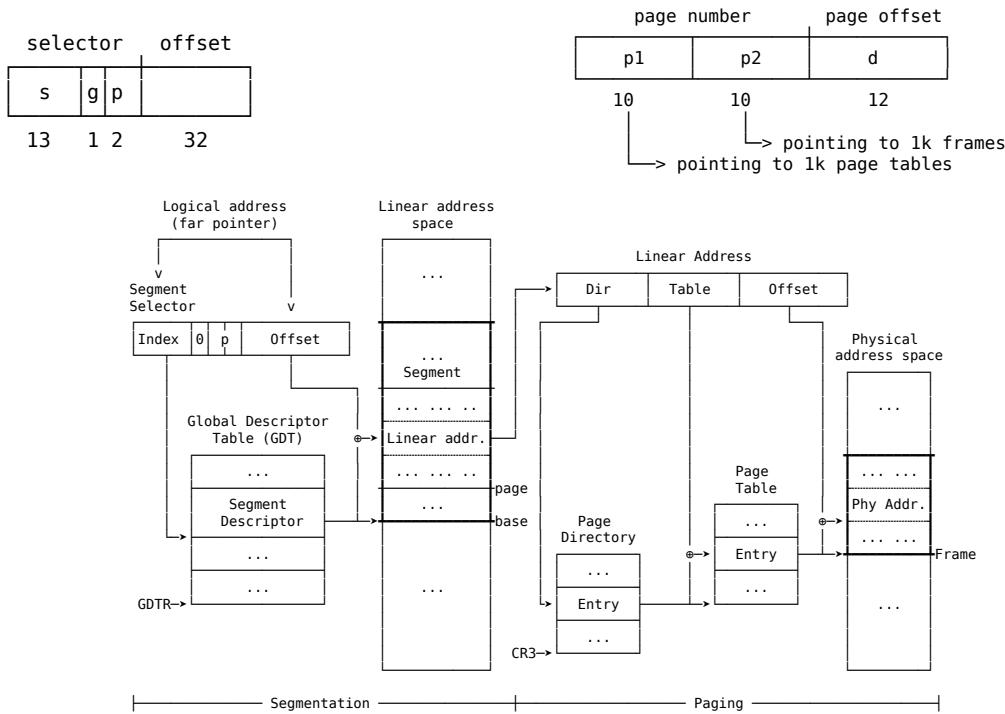
- Segmentation and paging are somewhat redundant
- Memory management is simpler when all processes share the same set of linear addresses
- Maximum portability. RISC architectures in particular have limited support for segmentation

The Linux 2.6 uses segmentation only when required by the 80x86 architecture.

## Case Study: The Intel Pentium

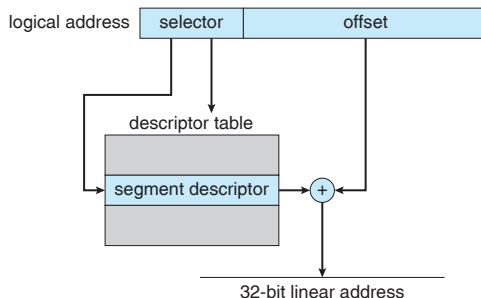
### Segmentation With Paging





## Segmentation

**Logical Address  $\Rightarrow$  Linear Address**

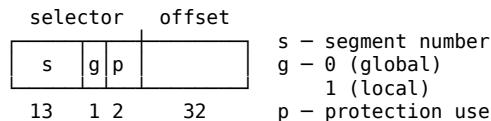


## Segment Selectors

A logical address consists of two parts:

segment selector	:	offset
16 bits		32 bits

**Segment selector** is an index into GDT/LDT



## Segment Descriptor Tables

All the segments are organized in 2 tables:

**GDT Global Descriptor Table**

- shared by all processes
- GDTR stores address and size of the GDT

**LDT Local Descriptor Table**

- one process each
- LDTR stores address and size of the LDT

**Segment descriptors** are entries in either GDT or LDT, 8-byte long

## Analogy

Process	$\iff$	Process Descriptor(PCB)
File	$\iff$	Inode
Segment	$\iff$	Segment Descriptor

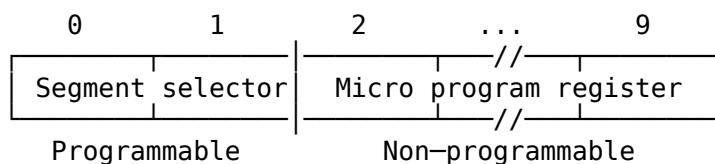
See also:

- Memory Translation And Segmentation<sup>12</sup>.
- The GDT<sup>13</sup>.
- The GDT and IDT<sup>14</sup>.

## Segment Registers

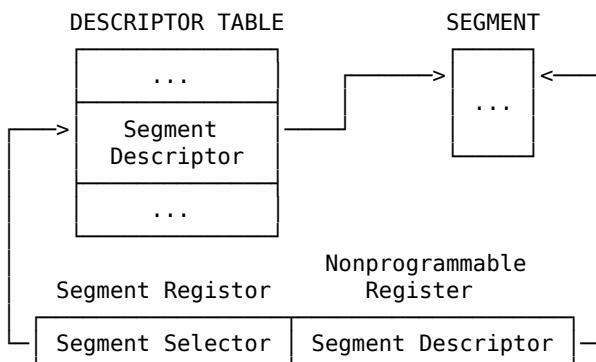
### The Intel Pentium has

- 6 *segment registers*, allowing 6 segments to be addressed at any one time by a process
  - Each segment register  $\lhd$  an entry in LDT/GDT
- 6 8-byte *micro program registers* to hold descriptors from either LDT or GDT
  - avoid having to read the descriptor from memory for every memory reference



### Fast access to segment descriptors

#### An additional nonprogrammable register for each segment register



### Segment registers hold segment selectors

**cs** code segment register

CPL 2-bit, specifies the Current Privilege Level of the CPU

00 - Kernel mode

11 - User mode

**ss** stack segment register

**ds** data segment register

**es/fs/gs** general purpose registers, may refer to arbitrary data segments

See also:

- [INTEL 80386 Programmer's Reference Manual, Sec. 6.3.2, Restricting Access to Data].
- CPU Rings, Privilege, and Protection<sup>15</sup>.

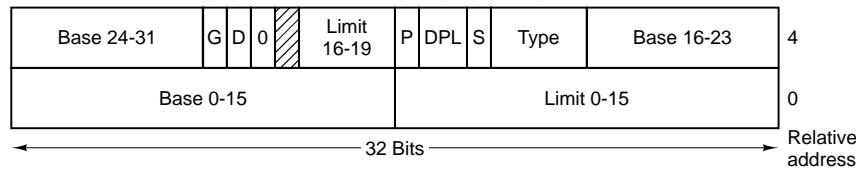
### Example: A LDT entry for code segment

<sup>12</sup><http://duartes.org/gustavo/blog/post/memory-translation-and-segmentation>

<sup>13</sup><http://www.osdever.net/bkerndev/Docs/gdt.htm>

<sup>14</sup>[http://www.jamesmolloy.co.uk/tutorial\\_html/4.-The%20GDT%20and%20IDT.html](http://www.jamesmolloy.co.uk/tutorial_html/4.-The%20GDT%20and%20IDT.html)

<sup>15</sup><http://duartes.org/gustavo/blog/post/cpu-rings-privilege-and-protection>



<b>Base:</b>	Where the seg starts	<b>Limit:</b>	$20 \text{ bit, } \Rightarrow 2^{20} \text{ in size}$
<b>G:</b>	Granularity flag 0 - seg size in bytes 1 - in 4096 bytes	<b>S:</b>	System flag 0 - system seg, e.g. LDT 1 - normal code/data seg
<b>D/B:</b>	0 - 16-bit seg 1 - 32-bit seg	<b>Type:</b>	seg type (cs/ds/tss)
<b>DPL:</b>	Descriptor Privilege Level. 0 or 3	<b>P:</b>	Seg-Present flag 0 - not in memory 1 - in memory
<b>AVL:</b>	ignored by Linux		

## The Four Main Linux Segments

Every process in Linux has these 4 segments

Segment	Base	G	Limit	S	Type	DPL	D/B	P
user code	0x00000000	1	0xfffffff	1	10	3	1	1
user data	0x00000000	1	0xfffffff	1	2	3	1	1
kernel code	0x00000000	1	0xfffffff	1	10	0	1	1
kernel data	0x00000000	1	0xfffffff	1	2	0	1	1

All linear addresses start at 0, end at 4G-1

- All processes share the same set of linear addresses
- Logical addresses coincide with linear addresses

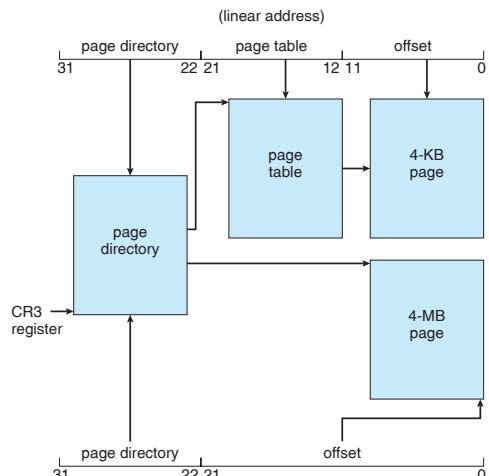
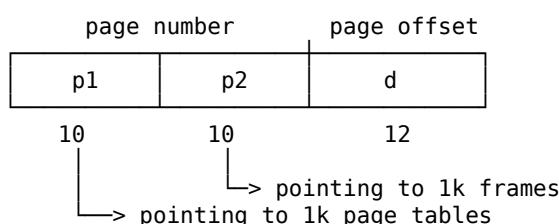
## Pentium Paging

Linear Address  $\Rightarrow$  Physical Address

Two page size in Pentium:

4K: 2-level paging

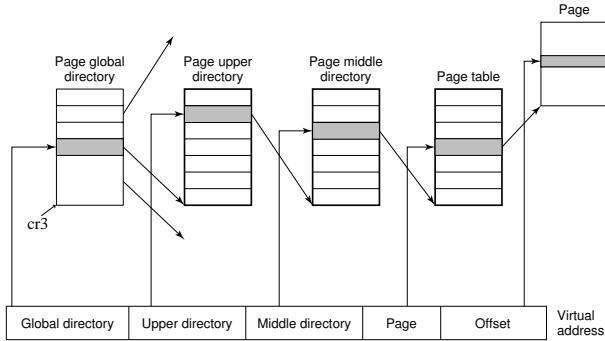
4M: 1-level paging



- The CR3 register  $\rightarrow$  the top level page table for the current process.

## Paging In Linux

4-level paging for both 32-bit and 64-bit



#### 4-level paging for both 32-bit and 64-bit

- 64-bit: *four-level paging*
  1. Page Global Directory
  2. Page Upper Directory
  3. Page Middle Directory
  4. Page Table
- 32-bit: *two-level paging*
  1. Page Global Directory
  2. Page Upper Directory — 0 bits; 1 entry
  3. Page Middle Directory — 0 bits; 1 entry
  4. Page Table

The same code can work on 32-bit and 64-bit architectures

Arch	Page size	Address bits	Paging levels	Address splitting
x86	4KB(12bits)	32	2	10 + 0 + 0 + 10 + 12
x86-PAE	4KB(12bits)	32	3	2 + 0 + 9 + 9 + 12
x86-64	4KB(12bits)	48	4	9 + 9 + 9 + 9 + 12

## References

- [1] Wikipedia. Memory management — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Memory\\_management&oldid=647917932](http://en.wikipedia.org/w/index.php?title=Memory_management&oldid=647917932).
- [2] Wikipedia. Virtual memory — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Virtual\\_memory&oldid=644430956](http://en.wikipedia.org/w/index.php?title=Virtual_memory&oldid=644430956).

## Part IV

# File Systems

## 15 File System Structure

### Long-term Information Storage Requirements

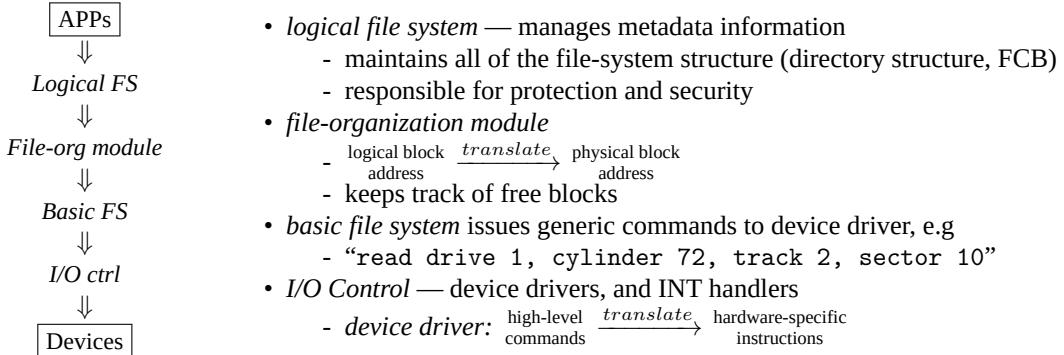
- Must store large amounts of data
- Information stored must survive the termination of the process using it
- Multiple processes must be able to access the information concurrently

### File-System Structure

#### File-system design addressing two problems:

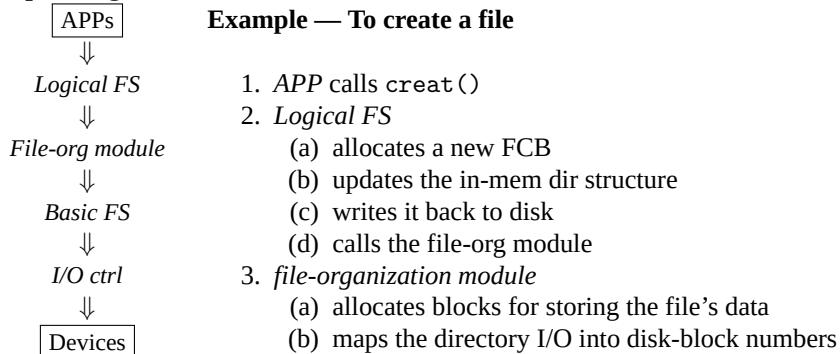
1. defining how the FS should look to the user
  - defining a file and its attributes
  - the operations allowed on a file
  - directory structure
2. creating algorithms and data structures to map the logical FS onto the physical disk

## File-System — A Layered Design



See also [Modern Operating Systems, Sec. 1.3.5, I/O Devices].

## The Operating Structure



### Benefit of layered design

The I/O control and sometimes the basic file system code can be used by multiple file systems.

## 16 Files

### File

*A Logical View Of Information Storage*

#### User’s view

- A file is the smallest storage unit on disk.
  - Data cannot be written to disk unless they are within a file

#### UNIX view

- Each file is a sequence of 8-bit bytes
  - It’s up to the application program to interpret this byte stream.

### File

*What Is Stored In A File?*

Source code, object files, executable files, shell scripts, PostScript...

#### Different type of files have different structure

- UNIX looks at contents to determine type
- **Shell scripts** start with “#!”
- **PDF** start with “%PDF...”
- **Executables** start with *magic number*
- Windows uses file naming conventions
  - **executables** end with “.exe” and “.com”
  - **MS-Word** end with “.doc”
  - **MS-Excel** end with “.xls”

## File Naming

### Vary from system to system

- Name length?
- Characters? Digits? Special characters?
- Extension?
- Case sensitive?

## File Types

**Regular files:** ASCII, binary

**Directories:** Maintaining the structure of the FS

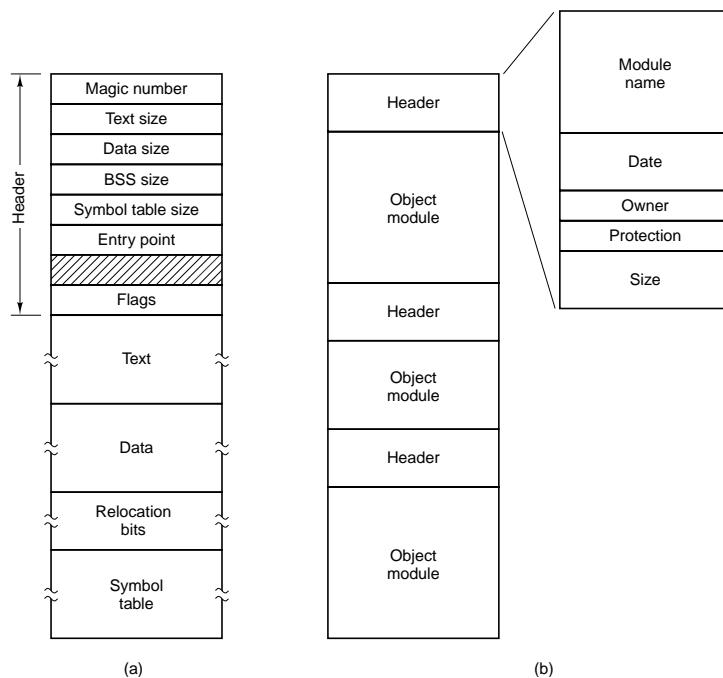
**In UNIX, everything is a file**

**Character special files:** I/O related, such as terminals, printers ...

**Block special files:** Devices that can contain file systems, i.e. disks

disks — logically, linear collections of blocks; disk driver translates them into physical block addresses

## Binary files



(a) An UNIX executable file

(b) An UNIX archive file

See also:

- [Executable and Linkable Format — Wikipedia, The Free Encyclopedia, Wikipedia:ELF].
- OSDev: ELF<sup>16</sup>.

## File Attributes — Metadata

- *Name* - only information kept in human-readable form
- *Identifier* - unique tag (number) identifies file within file system
- *Type* - needed for systems that support different types
- *Location* - pointer to file location on device
- *Size* - current file size
- *Protection* - controls who can do reading, writing, executing
- *Time, date, and user identification* - data for protection, security, and usage monitoring

<sup>16</sup><http://wiki.osdev.org/ELF>

## File Operations

### POSIX file system calls

creat(name, mode)	read(fd, buffer, byte_count)
open(name, flags)	write(fd, buffer, byte_count)
close(fd)	lseek(fd, offset, whence)
link(oldname, newname)	chown(name, owner, group)
unlink(name)	fchown(fd, owner, group)
truncate(name, size)	chmod(name, mode)
ftruncate(fd, size)	fchmod(fd, mode)
stat(name, buffer)	utimes(name, times)
fstat(fd, buffer)	

### An Example Program Using File System Calls

```
/* File copy program. Error checking and reporting is minimal. */

#include <sys/types.h>                      /* include necessary header files */
#include <fcntl.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]);           /* ANSI prototype */

#define BUF_SIZE 4096                         /* use a buffer size of 4096 bytes */
#define OUTPUT_MODE 0700                        /* protection bits for output file */

int main(int argc, char *argv[])
{
    int in_fd, out_fd, rd_count, wt_count;
    char buffer[BUF_SIZE];

    if (argc != 3) exit(1);                   /* syntax error if argc is not 3 */

    /* Open the input file and create the output file */
    in_fd = open(argv[1], O_RDONLY);          /* open the source file */
    if (in_fd < 0) exit(2);                  /* if it cannot be opened, exit */
    out_fd = creat(argv[2], OUTPUT_MODE);     /* create the destination file */
    if (out_fd < 0) exit(3);                 /* if it cannot be created, exit */

    /* Copy loop */
    while (TRUE) {
        rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
        if (rd_count <= 0) break;                /* if end of file or error, exit loop */
        wt_count = write(out_fd, buffer, rd_count); /* write data */
        if (wt_count <= 0) exit(4);              /* wt_count <= 0 is an error */
    }

    /* Close the files */
    close(in_fd);
    close(out_fd);
    if (rd_count == 0)                       /* no error on last read */
        exit(0);
    else
        exit(5);                           /* error on last read */
}
```

`open()`

`fd open(pathname, flags)`

A per-process *open-file table* is kept in the OS

- upon a successful `open()` syscall, a new entry is added into this table
- indexed by *file descriptor (fd)*

To see files opened by a process, e.g. `init`

`$ lsof -p 1`

### Why `open()` is needed?

To avoid constant searching

- Without `open()`, every file operation involves searching the directory for the file.

The purpose of the `open()` call is to allow the system to fetch the attributes and list of disk addresses into main memory for rapid access on later calls [*Modern Operating Systems*, Sec. 4.1.6, *File Operations*].

See also:

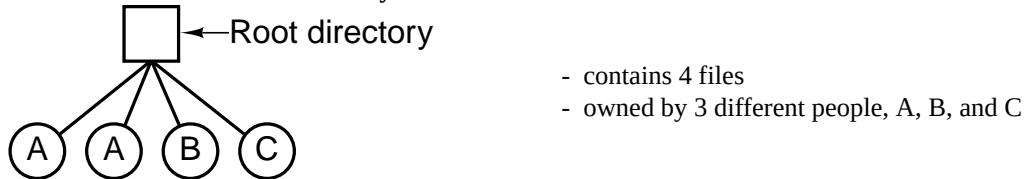
- [Open (system call) — Wikipedia, The Free Encyclopedia, Wikipedia:open() syscall].
- [File descriptor — Wikipedia, The Free Encyclopedia, Wikipedia:File descriptor].

## 17 Directories

### Directories

#### *Single-Level Directory Systems*

All files are contained in the same directory



### Limitations

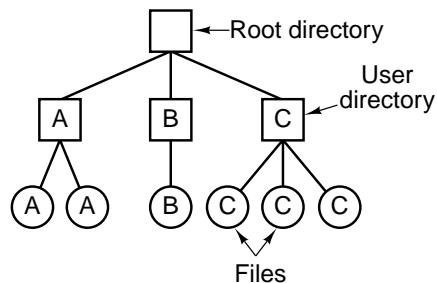
- name collision
- file searching

Often used on simple embedded devices, such as telephone, digital cameras...

### Directories

#### *Two-level Directory Systems*

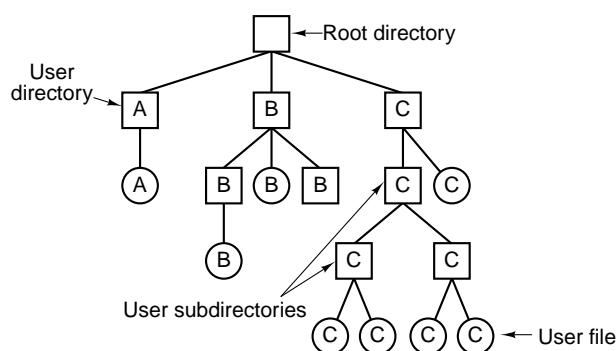
A separate directory for each user



**Limitation:** hard to access others files

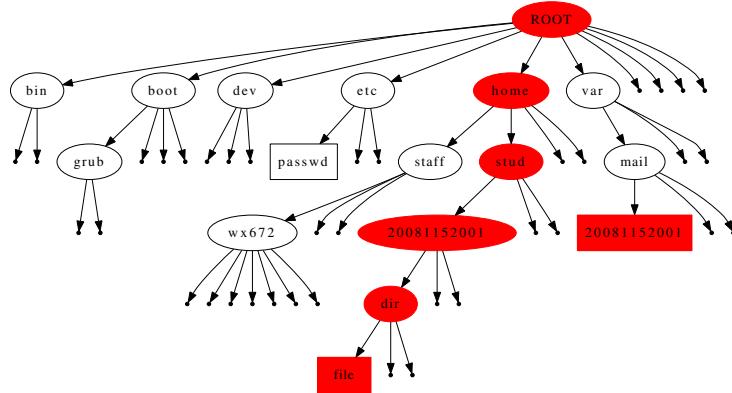
### Directories

#### *Hierarchical Directory Systems*



### Directories

#### *Path Names*



## Directories

### Directory Operations

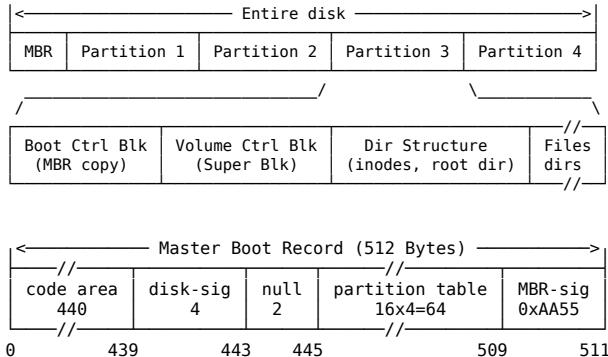
Create	Delete	Rename	Link
Opendir	Closedir	Readdir	Unlink

## 18 File System Implementation

### 18.1 Basic Structures

#### File System Implementation

##### A typical file system layout



**MBR, partition table, and booting** File systems are stored on disks. Most disks can be divided up into one or more partitions, with independent file systems on each partition. Sector 0 of the disk is called the MBR (Master Boot Record ) and is used to boot the computer. The end of the MBR contains the partition table. This table gives the starting and ending addresses of each partition. One of the partitions in the table is marked as active. When the computer is booted, the BIOS reads in and executes the MBR. The first thing the MBR program does is locate the active partition, read in its first block, called the boot block , and execute it. The program in the boot block loads the operating system contained in that partition. For uniformity, every partition starts with a boot block, even if it does not contain a bootable operating system. Besides, it might contain one in the future, so reserving a boot block is a good idea anyway[*Modern Operating Systems*, Sec 4.3.1, *File System Layout*].

The *superblock* is read into memory when the computer is booted or the file system is first touched.

#### On-Disk Information Structure

##### Boot control block a MBR copy

UFS: Boot block

NTFS: Partition boot sector

##### Volume control block Contains volume details

number of blocks	size of blocks
free-block count	free-block pointers
free FCB count	free FCB pointers

UFS: Superblock

NTFS: Master File Table

**Directory structure** Organizes the files *FCB*, *File control block*, contains file details (metadata).

UFS: I-node

NTFS: Stored in MFT using a relational database structure, with one row per file

### Each File-System Has a Superblock

#### Superblock

Keeps information about the file system

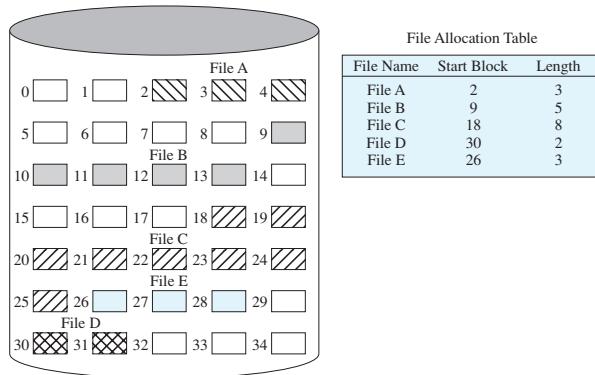
- Type — ext2, ext3, ext4...
- Size
- Status — how it's mounted, free blocks, free inodes, ...
- Information about other metadata structures

```
$ sudo dumpe2fs /dev/sda1 | less
```

## 18.2 Implementing Files

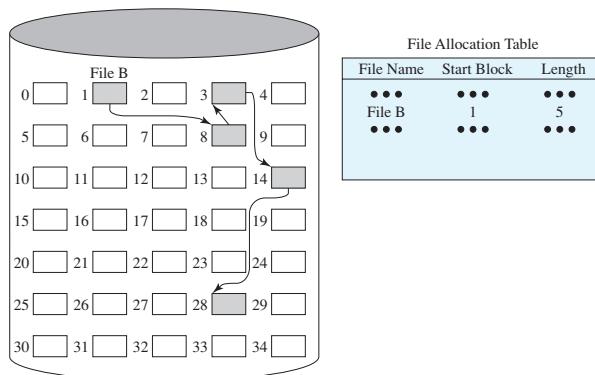
### Implementing Files

#### Contiguous Allocation



- simple;
- good for read only;
- fragmentation

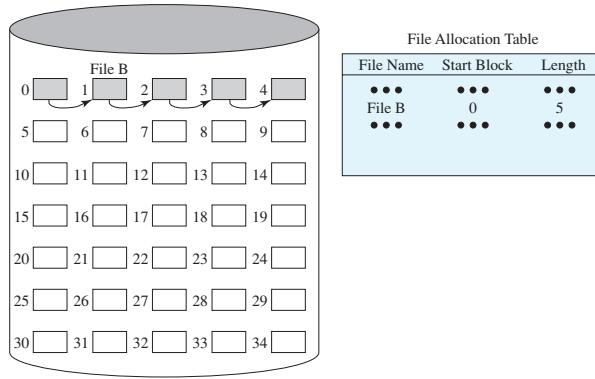
**Linked List (Chained) Allocation** A pointer in each disk block



- no waste block;
- slow random access;
- not  $2^n$

**Consolidation** One consequence of chaining, as described so far, is that there is *no accommodation of the principle of locality*. Thus, if it is necessary to bring in several blocks of a file at a time, as in sequential processing, then a series of accesses to different parts of the disk are required. This is perhaps a more significant effect on a single-user system but may also be of concern on a shared system. To overcome this problem, some systems periodically consolidate files (fig. 298)[*Operating Systems: Internals and Design Principles*, Sec. 12.7, Secondary Storage Management, P. 547]

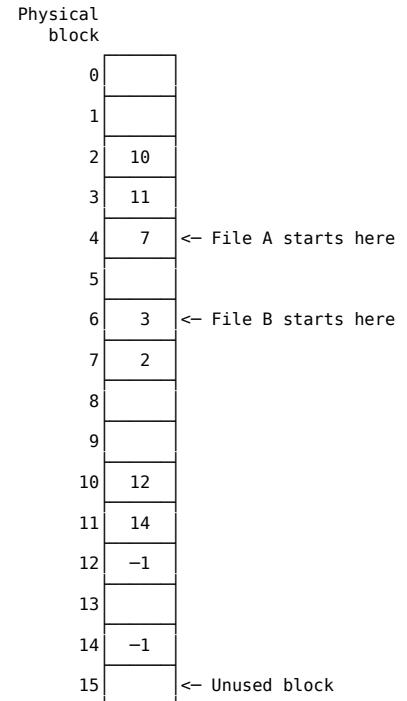
**Linked List (Chained) Allocation** Though there is no external fragmentation, consolidation is still preferred.



FAT: Linked list allocation with a table in RAM

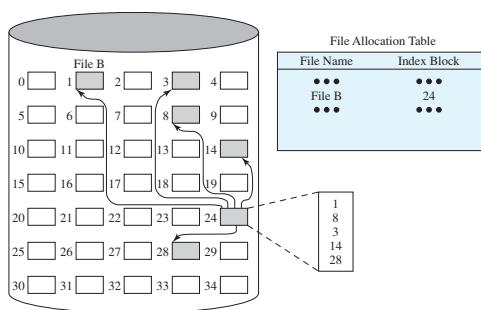
- Taking the pointer out of each disk block, and putting it into a table in memory
- fast random access (chain is in RAM)
- is  $2^n$
- the entire table must be in RAM

$$disk \nearrow \Rightarrow FAT \nearrow \Rightarrow RAM_{used} \nearrow$$



See also: [File Allocation Table — Wikipedia, The Free Encyclopedia, Wikipedia:FAT].

### Indexed Allocation

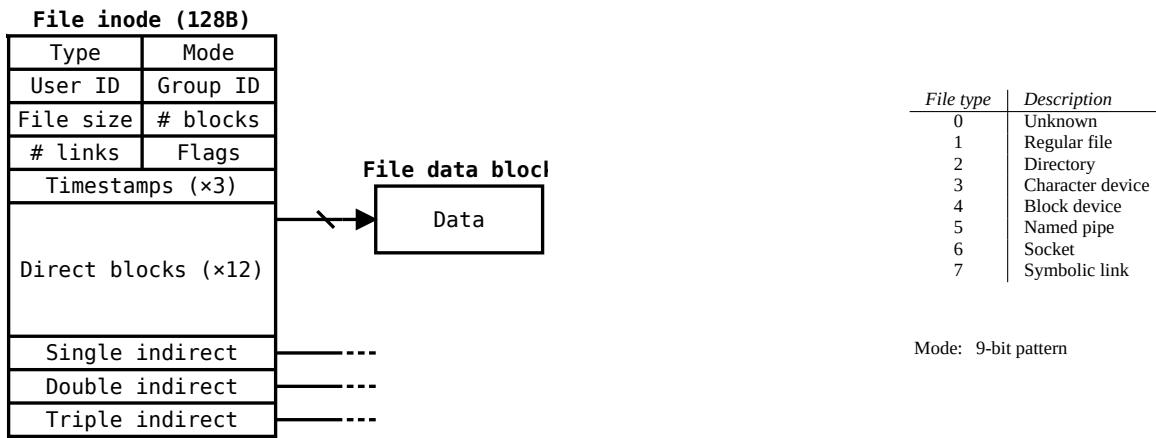


**I-node** A data structure for each file. An i-node is in memory *only if* the file is open

$$files_{opened} \nearrow \Rightarrow RAM_{used} \nearrow$$

See also: [Inode — Wikipedia, The Free Encyclopedia, Wikipedia:inode].

### I-node — FCB in UNIX



### A little quiz

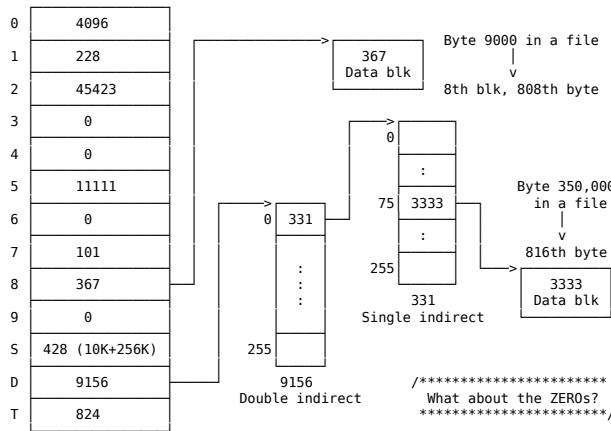
- in one terminal, to create a file “a”, do:  
\$ echo hello > /tmp/a
- to track its contents and keep it open, do:  
\$ tail -f /tmp/a
- in another terminal, delete this file “a”, do:  
\$ rm -f /tmp/a
- make sure it’s gone, do:  
\$ ls -l /tmp/a  
\$ ls -li /proc/`pidof tail`/fd  
\$ lsof -p `pidof tail` | grep deleted  
as you can see, /tmp/a is marked as “deleted”. Now, do:  
\$ echo "another a" >> /tmp/a  
\$ ls -li /tmp/a

Is this /tmp/a same as the deleted one? (check the inodes)

### Inode Quiz

Given: block size is 1KB  
pointer size is 4B

Addressing: byte offset 9000  
byte offset 350,000



Several block entries in the inode are 0, meaning that the logical block entries contain no data. This happens if no process ever wrote data into the file at any byte offsets corresponding to those blocks and hence the block numbers remain at their initial value, 0. No disk space is wasted for such blocks. Process can cause such a block layout in a file by using the `lseek()` and `write()` system calls [*The design of the UNIX operating system*, Sec. 4.2, *Structure of a Regular File*].

### UNIX In-Core Data Structure

**mount table** — Info about each mounted FS

**directory-structure cache** — Dir-info of recently accessed dirs

**inode table** — An in-core version of the on-disk inode table

**file table**

- global
- keeps inode of each open file
- keeps track of
  - how many processes are associated with each open file
  - where the next read and write will start
  - access rights

### **user file descriptor table**

- per process
- identifies all open files for a process

Find them in the kernel source

**user file descriptor table** — struct `fdtable` in `include/linux/fdtable.h`

**open file table** — struct `files_struct` in `include/linux/fdtable.h`

**inode** — struct `inode` in `include/linux/fs.h`

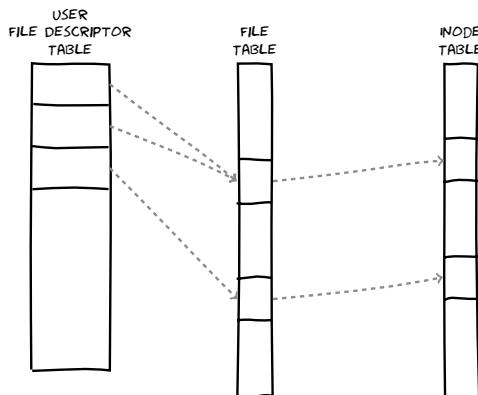
**inode table** — ?

**superblock** — struct `super_block` in `include/linux/fs.h`

**dentry** — struct `dentry` in `include/linux/dcache.h`

**file** — struct `file` in `include/linux/fs.h`

## **UNIX In-Core Data Structure**



**open() / creat()**

1. add entry in each table
2. returns a *file descriptor* — an index into the user file descriptor table

**Open file descriptor table** a global table, whose address is contained in the `files` field of the process descriptor, specifies which files are currently opened by the process. It is a `files_struct` structure whose fields are illustrated in Table 1 [*Understanding The Linux Kernel*, Sec. 12.2.6, *Files Associated with a Process*].

**Tbl. 1:** The fields of the `files_struct` structure

Type	Field	Description
<code>atomic_t</code>	<code>count</code>	Number of processes sharing this table
<code>rwlock_t</code>	<code>file_lock</code>	Read/write spin lock for the table fields
<code>int</code>	<code>max_fds</code>	Current maximum number of file objects
<code>int</code>	<code>max_fdset</code>	Current maximum number of file descriptors
<code>int</code>	<code>next_fd</code>	Maximum file descriptors ever allocated plus 1
<code>struct file **</code>	<code>fd</code>	Pointer to array of file object pointers
<code>fd_set *</code>	<code>close_on_exec</code>	Pointer to file descriptors to be closed on <code>exec()</code>
<code>fd_set *</code>	<code>open_fds</code>	Pointer to open file descriptors
<code>fd_set</code>	<code>close_on_exec_init</code>	Initial set of file descriptors to be closed on <code>exec()</code>
<code>fd_set</code>	<code>open_fds_init</code>	Initial set of file descriptors
<code>struct file *[]</code>	<code>fd_array</code>	Initial array of file object pointers

The `fd` field points to an array of pointers to file objects. The size of the array is stored in the `max_fds` field. Usually, `fd` points to the `fd_array` field of the `files_struct` structure, which includes 32 file object pointers. If the process opens more than 32 files, the kernel allocates a new, larger array of file pointers and stores its address in the `fd` fields; it also updates the `max_fds` field.

For every file with an entry in the `fd` array, the array index is the file descriptor. Usually, the first element (index 0) of the array is associated with the standard input of the process, the second with the standard output, and the third with the standard error (See fig. 12-3<sup>17</sup>). Unix processes use the file descriptor as the main file identifier. Notice that, thanks to the `dup()`, `dup2()`, and `fcntl(2) F_SETFL` system calls, two file descriptors may refer to the same opened file, that is, two elements of the array could point to the same file object. Users see this all the time when they use shell constructs such as `2>&1` to redirect the standard error to the standard output.

**open()** A call to `open()` creates a new *open file description*, an entry in the system-wide table of open files. This entry records the file offset and the file status flags (modifiable via the `fcntl(2) F_SETFL` operation). A *file descriptor* is a reference to one of these entries; this reference is unaffected if *pathname* is subsequently removed or modified to refer to a different file. The new open file description is initially not shared with any other process, but sharing may arise via `fork(2)` [*man 2 open*].

The internal representation of a file is given by an *inode*, which contains a description of the disk layout of the file data and other information such as the file owner, access permissions, and access times. The term *inode* is a contraction of the term *index node* and is commonly used in literature on the UNIX system. Every file has one *inode*, but it may have several names, all of which map into the *inode*. Each name is called a *link*. When a process refers to a file by name, the kernel parses the file name one component at a time, checks that the process has permission to search the directories in the path, and eventually retrieves the *inode* for the file. For example, if a process calls

```
open("/fs2/mjb/rje/sourcefile", 1);
```

the kernel retrieves the *inode* for “/fs2/mjb/rje/sourcefile”. When a process creates a new file, the kernel assigns it an unused *inode*. *Inodes* are stored in the file system, as will be seen shortly, but the kernel reads them into an in-core *inode* table when manipulating files.

The kernel contains two other data structures, the *file table* and the *user file descriptor table*. The *file table* is a global kernel structure, but the *user file descriptor table* is allocated per process. When a process opens or creates a file, the kernel allocates an entry from each table, corresponding to the file’s *inode*. *Entries in the three structures — user file descriptor table, file table, and inode table — maintain the state of the file and the user’s access to it.* The *file table* keeps track of the byte offset in the file where the user’s next `read` or `write` will start, and the access rights allowed to the opening process. The *user file descriptor table* identifies all open files for a process. Fig. 304 shows the tables and their relationship to each other. The kernel returns a *file descriptor* for the `open` and `creat` system calls, which is an index into the *user file descriptor table*. When executing `read` and `write` system calls, the kernel uses the *file descriptor* to access the *user file descriptor table*, follows pointers to the *file table* and *inode table* entries, and, from the *inode*, finds the data in the file. Chapters 4 and 5 describe these data structures in great detail. For now, suffice it to say that use of three tables allows various degrees of sharing access to a file [*The design of the UNIX operating system*, Sec. 2.2.1, *An Overview of the File Subsystem*].

The `open` system call is the first step a process must take to access the data in a file. The syntax for the `open` system call is

```
fd = open(pathname, flags, modes);
```

where *pathname* is a file name, *flags* indicate the type of open (such as for reading or writing), and *modes* give the file permissions if the file is being created. The `open` system call returns an integer called the *user file descriptor*. Other file operations, such as reading, writing, seeking, duplicating the file descriptor, setting file I/O parameters, determining file status, and closing the file, use the *file descriptor* that the `open` system call returns [*The design of the UNIX operating system*, Sec. 5.1, *Open*].

The kernel searches the file system for the file name parameter using algorithm *namei* (see fig. 1). It checks permissions for opening the file after it finds the in-core *inode* and allocates an entry in the *file table* for the open file. The *file table* entry contains a pointer to the *inode* of the open file and a field that indicates the byte offset in the file where the kernel expects the next `read` or `write` to begin. The kernel initializes the offset to 0 during the `open` call, meaning that the initial `read` or `write` starts at the beginning of a file by default. Alternatively, a process can open a file in *write-append* mode, in which case the kernel initializes the offset to the size of the file. The kernel allocates an entry in a private table in the process *u area*, called the *user file descriptor table*, and notes the index of this entry. The index is the *file descriptor* that is returned to the user. The entry in the *user file table* points to the entry in the global *file table*.

**Q1:** Can `open()` return an *inode number*?

**Q2:** Can we keep the I/O pointers in the *inode*? Possibly adding a few pointers in the *inode* data structure in a similar fashion of those block pointers (direct/single-indirect/double-indirect/triple-indirect). Each pointer pointing to an I/O pointer record.

## The Tables

---

<sup>17</sup>In book Understanding The Linux Kernel

```

algorithm open
inputs: file name
        type of open
        file permissions (for creation type of open)
output: file descriptor
{
    convert file name to inode (algorithm namei);
    if (file does not exist or not permitted access)
        return (error);
    allocate file table entry for inode, initialize count, offset;
    allocate user file descriptor entry, set pointer to file table entry;
    if (type of open specifies truncate file)
        free all file blocks (algorithm free);
    unlock (inode); /* locked above in namei */
    return (user file descriptor);
}

```

**Fig. 1:** Algorithm for opening a file

### Two levels of internal tables in the OS

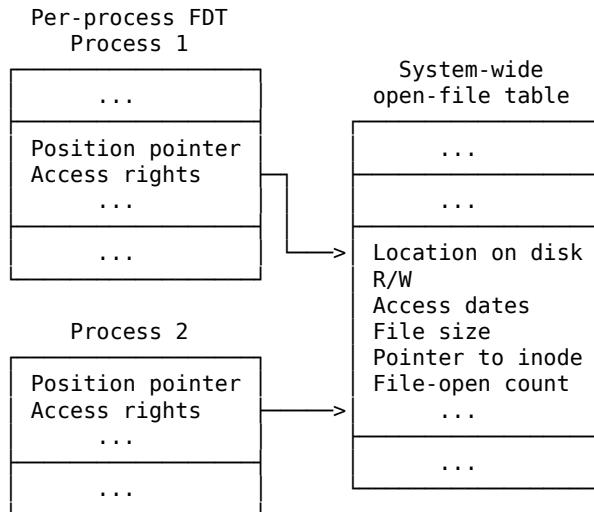
**A per-process table** tracks all files that a process has open. Stores

- the current-file-position pointer (not really)
- access rights
- more...

a.k.a file descriptor table

**A system-wide table** keeps process-independent information, such as

- the location of the file on disk
- access dates
- file size
- file open count — the number of processes opening this file

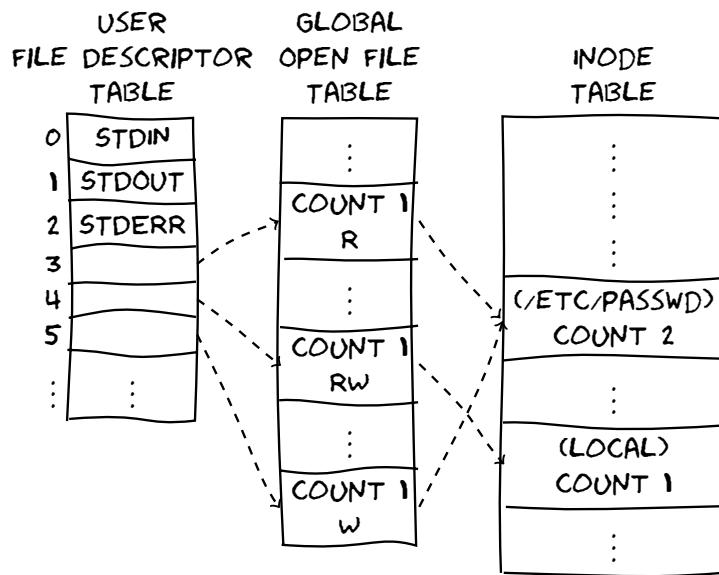


**A process executes the following code:**

```

fd1 = open("/etc/passwd", O_RDONLY);
fd2 = open("local", O_RDWR);
fd3 = open("/etc/passwd", O_WRONLY);

```

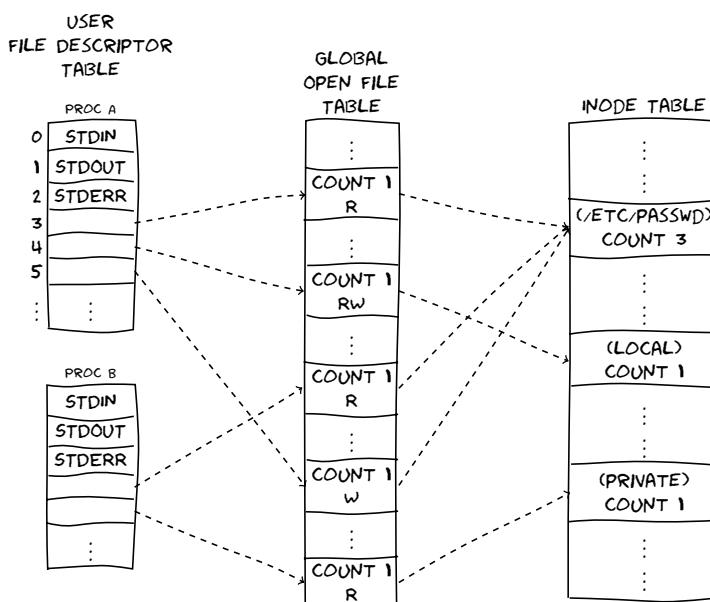


See also:

- [The design of the UNIX operating system, Sec. 5.1, Open].
- File descriptor manipulation<sup>18</sup>.

#### One more process B:

```
fd1 = open("/etc/passwd", O_RDONLY);
fd2 = open("private", O_RDONLY);
```

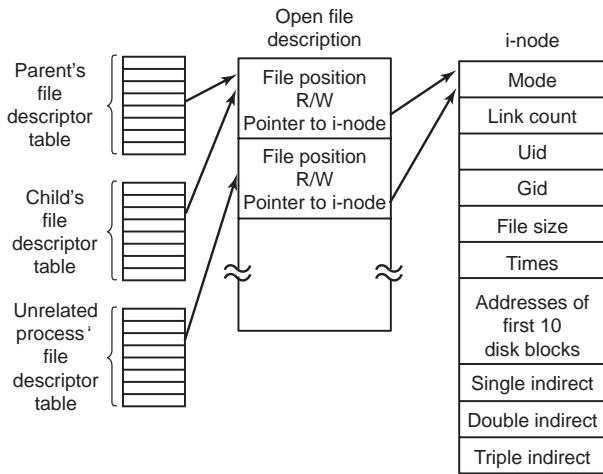


#### Why File Table?

To allow a parent and child to share a file position, but to provide unrelated processes with their own values.

---

<sup>18</sup><http://www.cim.mcgill.ca/~franco/OpSys-304-427/lecture-notes/node27.html#SECTION00063000000000000000>



## Why File Table?

### Where To Put File Position Info?

**Inode table?** No. Multiple processes can open the same file. Each one has its own file position.

**User file descriptor table?** No. Trouble in file sharing.

### Example

```

1 #!/bin/bash
2 echo hello > A
3 echo world
? Where should the "world" be?

```

Why file table?

**File system implementation** With file sharing, it is necessary to allow related processes to share a common I/O pointer and yet have separate I/O pointers for independent processes that access the same file. With these two conditions, the I/O pointer cannot reside in the i-node table nor can it reside in the list of open files for the process. A new table (the open file table) was invented for the sole purpose of holding the I/O pointer. Processes that share the same open file (the result of `forks`) share a common open file table entry. A separate open of the same file will only share the i-node table entry, but will have distinct open file table entries[“Unix Implementation”, Sec. 4.1].

**Open** The user file descriptor table entry could conceivably contain the file offset for the position of the next I/O operation and point directly to the in-core inode entry for the file, eliminating the need for a separate kernel file table. The examples above show a one-to-one relationship between user file descriptor entries and kernel file table entries. Thompson notes, however, that he implemented the file table as a separate structure to allow sharing of the offset pointer between several user file descriptors[“Unix Implementation”, *Thompson 78, p. 1943*]. The `dup` and `fork` system calls, explained in [*The design of the UNIX operating system*, Sec. 5.13] and [*The design of the UNIX operating system*, Sec. 7.1], manipulate the data structures to allow such sharing [*The design of the UNIX operating system*, Sec. 5.1].

**The Linux File System** ... The idea is to start with this file descriptor and end up with the corresponding i-node. Let us consider one possible design: just put a pointer to the i-node in the file descriptor table. Although simple, unfortunately this method does not work. The problem is as follows. Associated with every file descriptor is a file position that tells at which byte the next read (or write) will start. Where should it go? One possibility is to put it in the i-node table. However, this approach fails if two or more unrelated processes happen to open the same file at the same time because each one has its own file position[*Modern Operating Systems*, Sec. 10.6].

A second possibility is to put the file position in the file descriptor table. In that way, every process that opens a file gets its own private file position. Unfortunately this scheme fails too, but the reasoning is more subtle and has to do with the nature of file sharing in Linux. Consider a shell script, `s`, consisting of two commands, `p1` and `p2`, to be run in order. If the shell script is called by the command line

`S >x`

it is expected that `p1` will write its output to `x`, and then `p2` will write its output to `x` also, starting at the place where `p1` stopped.

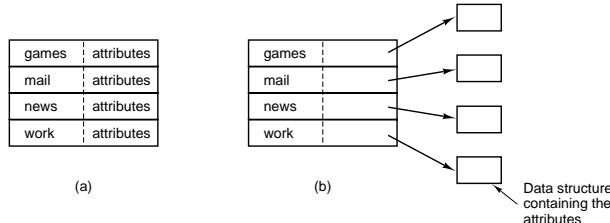
When the shell forks off `p1`, `x` is initially empty, so `p1` just starts writing at file position 0. However, when `p1` finishes, some mechanism is needed to make sure that the initial file position that `p2` sees is not 0 (which it would be if the file position were kept in the file descriptor table), but the value `p1` ended with.

The way this is achieved is shown in Fig 309. The trick is to introduce a new table, the *open file description table*, between the file descriptor table and the i-node table, and put the file position (and read/write bit) there.

In this figure, the parent is the shell and the child is first p1 and later p2. When the shell forks off p1, its user structure (including the file descriptor table) is an exact copy of the shell's, so both of them point to the same open file description table entry. When p1 finishes, the shell's file descriptor is still pointing to the open file description containing p1's file position. When the shell now forks off p2, the new child automatically inherits the file position, without either it or the shell even having to know what that position is.

## 18.3 Implementing Directories

### Implementing Directories



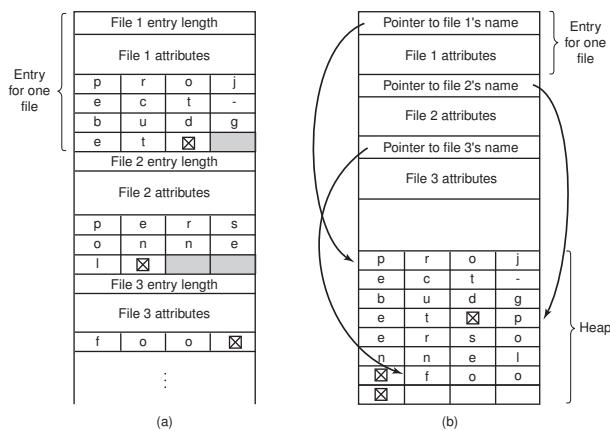
(a) A simple directory (Windows)

- fixed size entries
- disk addresses and attributes in directory entry

(b) Directory in which each entry just refers to an i-node (UNIX)

The maximum possible size for a file on a FAT32 volume is 4 GiB minus 1 byte or 4,294,967,295 ( $2^{32}-1$ ) bytes. This limit is a consequence of the file length entry in the directory table and would also affect huge FAT16 partitions with a sufficient sector size[*File Allocation Table — Wikipedia, The Free Encyclopedia*].

### How Long A File Name Can Be?



How long file name is implemented?

1. The simplest approach is to set a limit on file name length, typically 255 characters, and then use one of the designs of fig. 311 with 255 characters reserved for each file name. This approach is simple, but wastes a great deal of directory space, since few files have such long names. For efficiency reasons, a different structure is desirable.
2. One alternative is to give up the idea that all directory entries are the same size. With this method, each directory entry contains a fixed portion, typically starting with the length of the entry, and then followed by data with a fixed format, usually including the owner, creation time, protection information, and other attributes. This fixed-length header is followed by the actual file name, however long it may be, as shown in fig. 312(a) in big-endian format (e.g., SPARC). In this example we have three files, project-budget, personnel, and foo. Each file name is terminated by a special character (usually 0), which is represented in the figure by a box with a cross in it. To allow each directory entry to begin on a word boundary, each file name is filled out to an integral number of words, shown by shaded boxes in the figure.

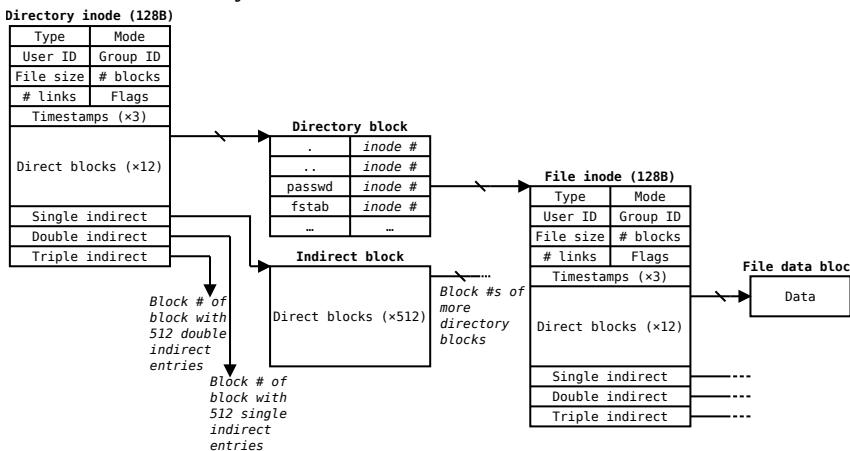
A disadvantage of this method is that when a file is removed, a variable-sized gap is introduced into the directory into which the next file to be entered may not fit. This problem is the same one we saw with contiguous disk files, only now compacting the directory is feasible because it is entirely in memory. Another problem is that a single directory entry may span multiple pages, so a page fault may occur while reading a file name.

3. Another way to handle variable-length names is to make the directory entries themselves all fixed length and keep the file names together in a heap at the end of the directory, as shown in fig. 312(b). This method has the advantage that when an entry is removed, the next file entered will always fit there. Of course, the heap must be managed and page faults can still occur while processing file names. One minor win here is that there is no longer any real need for file names to begin at word boundaries, so no filler characters are needed after file names in fig. 312(b) as they are in fig. 312(a).

4. Ext2's approach is a bit different. See Sec. 19.5.

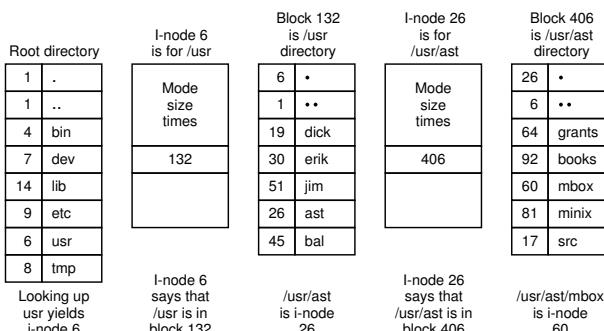
See also: [Modern Operating Systems, Sec. 4.3.3, Implementing Directories].

### UNIX Treats a Directory as a File



- A directory is a file whose data is a sequence of entries, each consisting of an inode number and the name of a file contained in the directory.
- Each (disk) block in a directory file consists of a linked list of entries; each entry contains the length of the entry, the name of a file, and the inode number of the inode to which that entry refers [Operating System Concepts Essentials, Sec. 15.7.2, The Linux ext2fs File System].

### The steps in looking up /usr/ast/mbox



- First the file system locates the root directory. In UNIX its i-node is located at a fixed place on the disk. From this i-node, it locates the root directory, which can be anywhere on the disk, but say block 1 [Modern Operating Systems, Sec. 4.5, Example File Systems].

Then it reads the root directory and looks up the first component of the path, `usr`, in the root directory to find the i-node number of the file `/usr`. Locating an i-node from its number is straightforward, since each one has a fixed location on the disk. From this i-node, the system locates the directory for `/usr` and looks up the next component, `ast`, in it. When it has found the entry for `ast`, it has the i-node for the directory `/usr/ast`. From this i-node it can find the directory itself and look up `mbox`. The i-node for this file is then read into memory and kept there until the file is closed. The lookup process is illustrated in fig. 314.

Relative path names are looked up the same way as absolute ones, only starting from the working directory instead of starting from the root directory. Every directory has entries for `.` and `..` which are put there when the directory is created. The entry `.` has the i-node number for the current directory, and the entry for `..` has the i-node number for the parent directory. Thus, a procedure looking up `../dick/prog.c` simply looks up `..` in the working directory, finds the i-node number for the parent directory, and searches that directory for `dick`. No special mechanism is needed to handle these names. As far as the directory system is concerned, they are just ordinary ASCII strings, just the same as any other names. The only bit of trickery here is that `..` in the root directory points to itself.

## 18.4 Shared Files

### File Sharing

#### Multiple Users

User IDs identify users, allowing permissions and protections to be per-user

Group IDs allow users to be in groups, permitting group access rights

#### Example: 9-bit pattern

rwxr-x--- means:

user	group	other
rwx	r-x	---
111	1-1	000
7	5	0

### File Sharing

#### Remote File Systems

Networking — allows file system access between systems

- Manually via programs like FTP
- Automatically, seamlessly using distributed file systems
- Semi automatically, via the world wide web

C/S model — allows clients to *mount* remote file systems from servers

- NFS — standard UNIX client-server file sharing protocol
- CIFS — standard Windows protocol
- Standard system calls are translated into remote calls

Distributed Information Systems (distributed naming services)

- such as LDAP, DNS, NIS, Active Directory implement unified access to information needed for remote computing

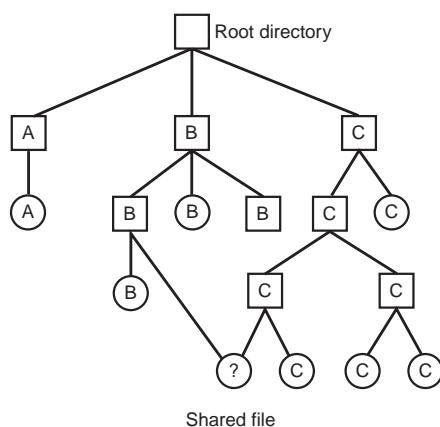
### File Sharing

#### Protection

- File owner/creator should be able to control:
  - what can be done
  - by whom
- Types of access
  - Read
  - Write
  - Execute
  - Append
  - Delete
  - List

### Shared Files

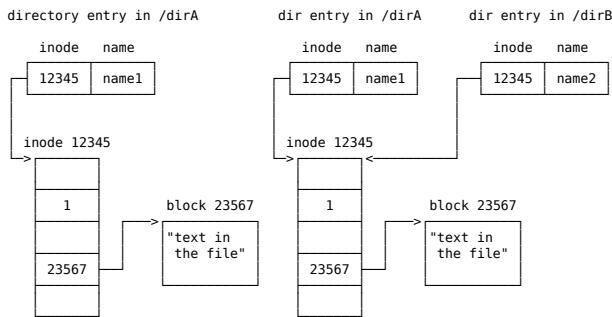
#### Hard Links vs. Soft Links



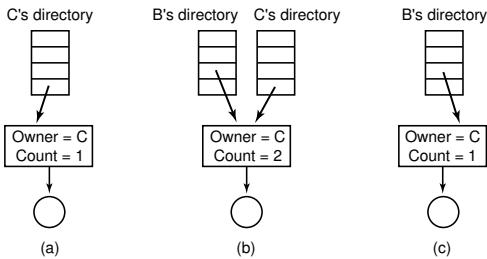
See also: [Directed acyclic graph — Wikipedia, The Free Encyclopedia, Wikipedia:dag].

## Hard Links

### Hard links → the same inode



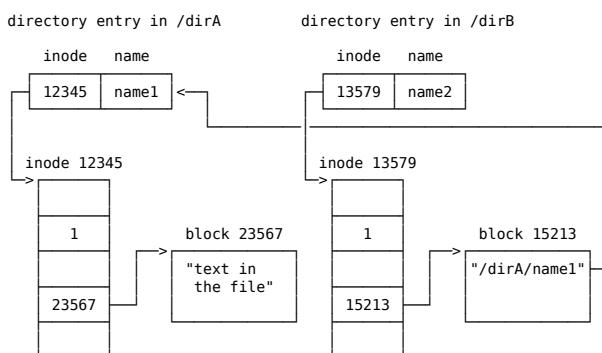
### Drawback



- Both hard and soft links have drawbacks[*Modern Operating Systems*, Sec. 4.3.4, Shared Files].
- # echo 0 > /proc/sys/fs/protected\_hardlinks
- Why hard links not allowed to directories in UNIX/Linux<sup>19</sup>?
  - ... if you were allowed to do this for directories, two different directories in different points in the filesystem could point to the same thing. In fact, a subdir could point back to its grandparent, creating a loop.
  - Why is this loop a concern? Because when you are traversing, there is no way to detect you are looping (without keeping track of inode numbers as you traverse). Imagine you are writing the du command, which needs to recurse through subdirs to find out about disk usage. How would du know when it hit a loop? It is error prone and a lot of bookkeeping that du would have to do, just to pull off this simple task.
- The Ultimate Linux Soft and Hard Link Guide (10 Ln Command Examples)<sup>20</sup>

## Symbolic Links

### A symbolic link has its own inode → a directory entry



## 18.5 Disk Space Management

### Disk Space Management

#### Statistics

<sup>19</sup><http://unix.stackexchange.com/questions/22394/why-hard-links-not-allowed-to-directories-in-unix-linux>

<sup>20</sup><http://www.thegeekstuff.com/2010/10/linux-ln-command-examples/>

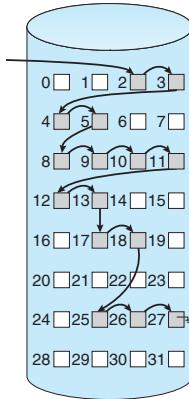
Length	VU 1984	VU 2005	Web
1	1.79	1.38	6.67
2	1.88	1.53	7.67
4	2.01	1.65	8.33
8	2.31	1.80	11.30
16	3.32	2.15	11.46
32	5.13	3.15	12.33
64	8.71	4.98	26.10
128	14.73	8.03	28.49
256	23.09	13.29	32.10
512	34.44	20.62	39.94
1 KB	48.05	30.91	47.82
2 KB	60.87	46.09	59.44
4 KB	75.31	59.13	70.64
8 KB	84.97	69.96	79.69
Length	VU 1984	VU 2005	Web
16 KB	92.53	78.92	86.79
32 KB	97.21	85.87	91.65
64 KB	99.18	90.84	94.80
128 KB	99.84	93.73	96.93
256 KB	99.96	96.12	98.48
512 KB	100.00	97.73	98.99
1 MB	100.00	98.87	99.62
2 MB	100.00	99.44	99.80
4 MB	100.00	99.71	99.87
8 MB	100.00	99.86	99.94
16 MB	100.00	99.94	99.97
32 MB	100.00	99.97	99.99
64 MB	100.00	99.99	99.99
128 MB	100.00	99.99	100.00

See also: [Modern Operating Systems, Sec. 4.4.1, Disk Space Management].

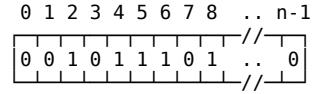
- Block size is chosen while creating the FS
  - Disk I/O performance is conflict with space utilization
    - smaller block size  $\Rightarrow$  better space utilization
    - larger block size  $\Rightarrow$  better disk I/O performance
- ```
$ dumpe2fs /dev/sda1 | grep "Block size"
```

## Keeping Track of Free Blocks

### 1. Linked List



### 2. Bit map (n blocks)



$$bit[i] = \begin{cases} 0 & \Rightarrow \text{block}[i] \text{ is free} \\ 1 & \Rightarrow \text{block}[i] \text{ is occupied} \end{cases}$$

## Journaling File Systems

### Operations required to remove a file in UNIX:

1. Remove the file from its directory
  - set inode number to 0
2. Release the i-node to the pool of free i-nodes
  - clear the bit in inode bitmap
3. Return all the disk blocks to the pool of free disk blocks
  - clear the bits in block bitmap

What if crash occurs between 1 and 2, or between 2 and 3?

Suppose that the first step is completed and then the system crashes. The i-node and file blocks will not be accessible from any file, but will also not be available for reassignment; they are just off in limbo somewhere, decreasing the available resources. If the crash occurs after the second step, only the blocks are lost.

If the order of operations is changed and the i-node is released first, then after rebooting, the i-node may be reassigned, but the old directory entry will continue to point to it, hence to the wrong file. If the blocks are released first, then a crash before the i-node is cleared will mean that a valid directory entry points to an i-node listing blocks now in the free storage pool and which are likely to be reused shortly, leading to two or more files randomly sharing the same blocks. None of these outcomes are good [Modern Operating Systems, Sec. 4.3.6, Journaling File Systems].

See also: [Operating System Concepts Essentials, Sec. 15.7.3, Journaling].

## Journaling File Systems

Keep a log of what the file system is going to do before it does it

- so that if the system crashes before it can do its planned work, upon rebooting the system can look in the log to see what was going on at the time of the crash and finish the job.
- NTFS, EXT3, and ReiserFS use journaling among others

# 19 Ext2 File System

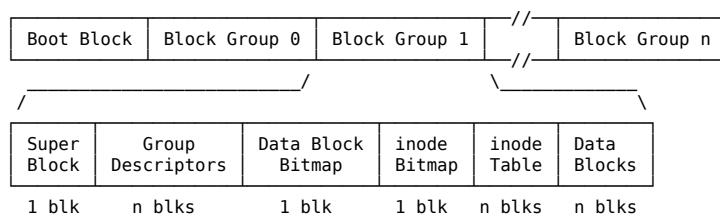
References:

- [The Second Extended File System Internal Layout, The Second Extended File System].
- [Operating System Concepts Essentials, Sec. 15.7, File Systems].
- [The Linux Kernel, Chap. 9, The File System].
- [“Design and Implementation of the Second Extended Filesystem”, Design and Implementation of the Second Extended Filesystem].
- [Analyzing a filesystem, Analyzing a filesystem].

## 19.1 Ext2 File System Layout

Ext2 File System

Physical Layout



See also: [Operating System Concepts Essentials, Sec. 15.7.2, The Linux ext2fs File System].

## 19.2 Ext2 Block groups

Ext2 Block groups

The partition is divided into **Block Groups**

- Block groups are same size — easy locating
- Kernel tries to keep a file's data blocks in the same block group — reduce fragmentation
- Backup critical info in each block group
- The Ext2 inodes for each block group are kept in the *inode table*
- The *inode-bitmap* keeps track of allocated and unallocated inodes

When allocating a file, ext2fs must first select the block group for that file[Operating System Concepts Essentials, Sec. 15.7.2, The Linux ext2fs File System, P. 625].

- For data blocks, it attempts to allocate the file to the block group to which the file's inode has been allocated.
- For inode allocations, it selects the block group in which the file's parent directory resides, for nondirectory files.
- Directory files are not kept together but rather are dispersed throughout the available block groups.

These policies are designed not only to keep related information within the same block group but also to spread out the disk load among the disk's block groups to reduce the fragmentation of any one area of the disk.

Group descriptor

- Each block group has a group descriptor
- All the group descriptors together make the *group descriptor table*
- The table is stored along with the superblock
- *Block Bitmap*: tracks free blocks
- *Inode Bitmap*: tracks free inodes

- **Inode Table:** all inodes in this block group

$\begin{array}{l} \text{Free blocks count} \\ \text{Free Inodes count} \\ \text{Used dir count} \end{array}$  } counters  
# dumpe2fs /dev/sda1

## Maths

**Given**  $\begin{array}{l} \text{block size} = 4K \\ \text{block bitmap} = 1 \text{ blk} \end{array}$ , then

$$\text{blocks per group} = 8 \text{ bits} \times 4K = 32K$$

### How large is a group?

$$\text{group size} = 32K \times 4K = 128M$$

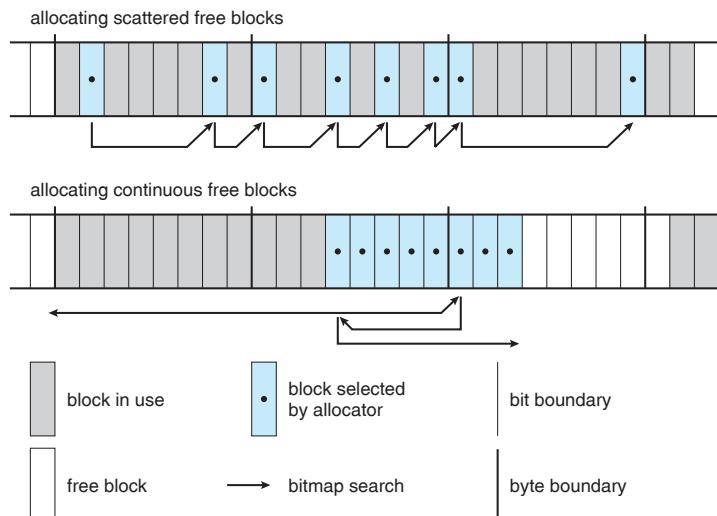
### How many block groups are there?

$$\approx \frac{\text{partition size}}{\text{group size}} = \frac{\text{partition size}}{128M}$$

### How many files can I have in max?

$$\approx \frac{\text{partition size}}{\text{block size}} = \frac{\text{partition size}}{4K}$$

## Ext2 Block Allocation Policies



**More about block bitmap** It maintains a bitmap of all free blocks in a block group. When allocating the first blocks for a new file, it starts searching for a free block from the beginning of the block group; when extending a file, it continues the search from the block most recently allocated to the file. The search is performed in two stages. First, ext2fs searches for an entire free byte in the bitmap; if it fails to find one, it looks for any free bit. The search for free bytes aims to allocate disk space in chunks of at least eight blocks where possible[*Operating System Concepts Essentials*, Sec. 15.7.2, *The Linux ext2fs File System*].

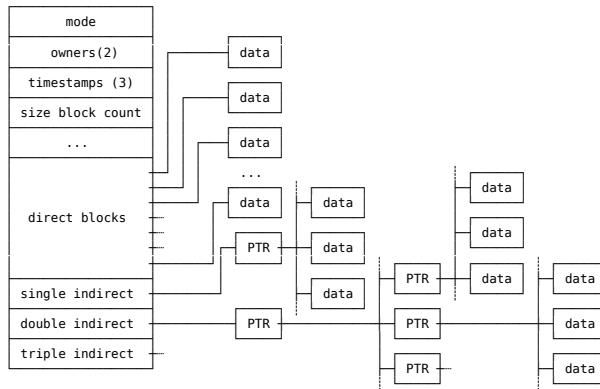
Once a free block has been identified, the search is extended backward until an allocated block is encountered. When a free byte is found in the bitmap, this backward extension prevents ext2fs from leaving a hole between the most recently allocated block in the previous nonzero byte and the zero byte found. Once the next block to be allocated has been found by either bit or byte search, ext2fs extends the allocation forward for up to eight blocks and *preallocates* these extra blocks to the file. This preallocation helps to reduce fragmentation during interleaved writes to separate files and also reduces the CPU cost of disk allocation by allocating multiple blocks simultaneously. *The preallocated blocks are returned to the free-space bitmap when the file is closed.*

Fig. 331 illustrates the allocation policies. Each row represents a sequence of set and unset bits in an allocation bitmap, indicating used and free blocks on disk. In the first case, if we can find any free blocks sufficiently near the start of the search, then we allocate them no matter how fragmented they may be. The fragmentation is partially compensated for by the fact that the blocks are close together and can probably all be read without any disk seeks, and allocating them all to one file is better in the long run than allocating isolated blocks to separate files once large

free areas become scarce on disk. In the second case, we have not immediately found a free block close by, so we search forward for an entire free byte in the bitmap. If we allocated that byte as a whole, we would end up creating a fragmented area of free space between it and the allocation preceding it, so before allocating we back up to make this allocation flush with the allocation preceding it, and then we allocate forward to satisfy the default allocation of eight blocks.

### 19.3 Ext2 Inode

#### Ext2 inode



#### Ext2 inode

**Mode:** holds two pieces of information

1. Is it a {file|dir|sym-link|blk-dev|char-dev|FIFO}?
2. Permissions

**Owner info:** Owners' ID of this file or directory

**Size:** The size of the file in bytes

**Timestamps:** Accessed, created, last modified time

**Datablocks:** 15 pointers to data blocks ( $12 + S + D + T$ )

#### Max File Size

**Given:**

$$\begin{cases} \text{block size} = 4K \\ \text{pointer size} = 4B \end{cases},$$

**We get:**

$$\text{Max File Size} = \text{number of pointers} \times \text{block size}$$

$$\begin{aligned} &= (\underbrace{12}_{\text{direct}} + \underbrace{1K}_{\text{1-indirect}} + \underbrace{1K \times 1K}_{\text{2-indirect}} + \underbrace{1K \times 1K \times 1K}_{\text{3-indirect}}) \times 4K \\ &= 48K + 4M + 4G + 4T \end{aligned}$$

### 19.4 Ext2 Superblock

#### Ext2 Superblock

**Magic Number:** 0xEF53

**Revision Level:** determines what new features are available

**Mount Count and Maximum Mount Count:** determines if the system should be fully checked

**Block Group Number:** indicates the block group holding this superblock

**Block Size:** usually  $4K$

**Blocks per Group:**  $8\text{bits} \times \text{block size}$

**Free Blocks:** System-wide free blocks

**Free Inodes:** System-wide free inodes

**First Inode:** First inode number in the file system

See more:

```
# dumpe2fs /dev/sda1
```

## Ext2 File Types

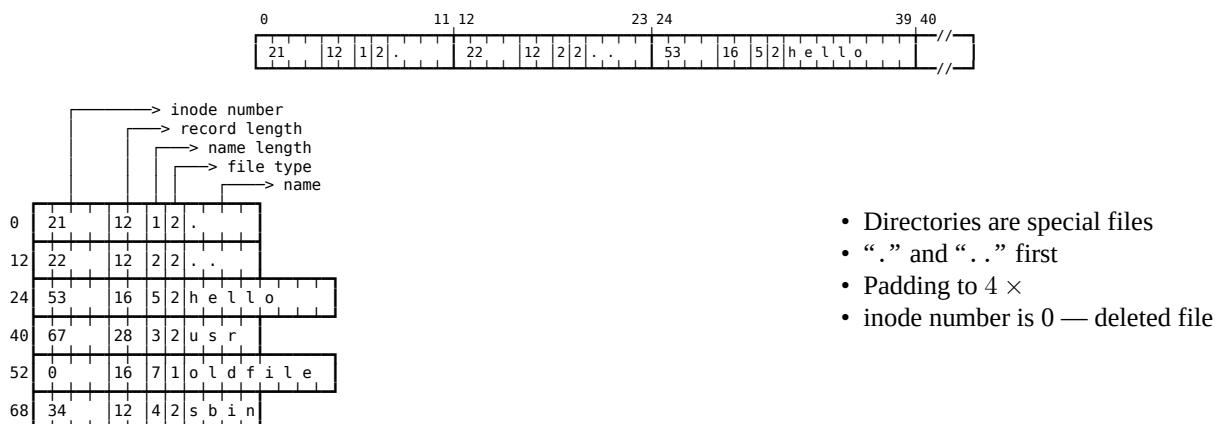
| File type | Description      |
|-----------|------------------|
| 0         | Unknown          |
| 1         | Regular file     |
| 2         | Directory        |
| 3         | Character device |
| 4         | Block device     |
| 5         | Named pipe       |
| 6         | Socket           |
| 7         | Symbolic link    |

**Device file, pipe, and socket:** No data blocks are required. All info is stored in the inode

**Fast symbolic link:** Short path name (< 60 chars) needs no data block. Can be stored in the 15 pointer fields

## 19.5 Ext2 Directory

### Ext2 Directories



Directory files are stored on disk just like normal files, although their contents are interpreted differently. Each block in a directory file consists of a linked list of entries; each entry contains the length of the entry, the name of a file, and the inode number of the inode to which that entry refers[*Operating System Concepts Essentials*, Sec. 15.7.2, *The Linux ext2fs File System*].

## 20 Vitural File Systems

### Many different FS are in use

#### Windows

uses drive letter (C:, D:, ...) to identify each FS

#### UNIX

integrates multiple FS into a single structure

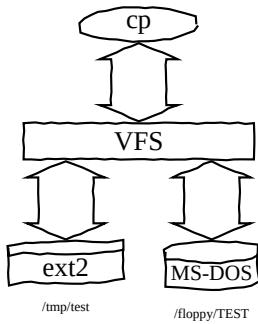
- From user's view, there is only one FS hierarchy

\$ man fs

Windows handles these disparate file systems by identifying each one with a different drive letter, as in C:, D:, etc. When a process opens a file, the drive letter is explicitly or implicitly present so Windows knows which file system to pass the request to. There is no attempt to integrate heterogeneous file systems into a unified whole[*Modern Operating Systems*, Sec. 4.3.7, *Virtual File Systems*].

A filesystem is a hierarchical storage of data adhering to a specific structure. Filesystems contain files, directories, and associated control information. Typical operations performed on filesystems are creation, deletion, and mounting. In Unix, filesystems are mounted at a specific mount point in a global hierarchy known as a *namespace*. This enables all mounted filesystems to appear as entries in a single tree. Contrast this single, unified tree with the behavior of DOS and Windows, which break the file namespace up into drive letters, such as C:. This breaks the namespace up among device and partition boundaries, “leaking” hardware details into the filesystem abstraction. As this delineation may be arbitrary and even confusing to the user, it is inferior to Linux’s unified namespace[*Linux Kernel Development*, Sec. 13.3, *Unix Filesystems*].

\$ cp /floppy/TEST /tmp/test



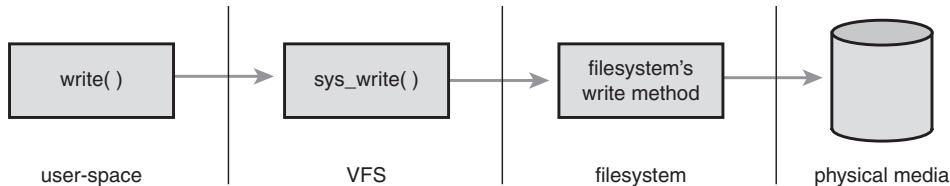
```

1 inf = open("/floppy/TEST", O_RDONLY, 0);
2 outf = open("/tmp/test",
3             O_WRONLY|O_CREAT|O_TRUNC, 0600);
4 do {
5     i = read(inf, buf, 4096);
6     write(outf, buf, i);
7 } while (i);
8 close(outf);
9 close(inf);

```

Where `/floppy` is the mount point of an MS-DOS diskette and `/tmp` is a normal Second Extended Filesystem (Ext2) directory. The VFS is an abstraction layer between the application program and the filesystem implementations. Therefore, the `cp` program is not required to know the filesystem types of `/floppy/TEST` and `/tmp/test`. Instead, `cp` interacts with the VFS by means of generic system calls known to anyone who has done Unix programming [*Understanding The Linux Kernel*, Sec. 12.1].

```
ret = write(fd, buf, len);
```

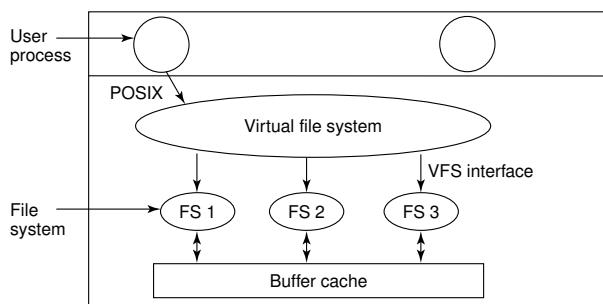


This system call writes the `len` bytes pointed to by `buf` into the current position in the file represented by the file descriptor `fd`. This system call is first handled by a generic `sys_write()` system call that determines the actual file writing method for the filesystem on which `fd` resides. The generic `write` system call then invokes this method, which is part of the filesystem implementation, to write the data to the media (or whatever this filesystem does on `write`). Fig. 340 shows the flow from user-space's `write()` call through the data arriving on the physical media. On one side of the system call is the generic VFS interface, providing the frontend to user-space; on the other side is the filesystem-specific backend, dealing with the implementation details [*Linux Kernel Development*, Sec. 13.2, P. 262].

## Virtual File Systems

### Put common parts of all FS in a separate layer

- It's a layer in the kernel
- It's a common interface to several kinds of file systems
- It calls the underlying concrete FS to actual manage the data



To the VFS layer and the rest of the kernel, however, each filesystem looks the same. They all support notions such as files and directories, and they all support operations such as creating and deleting files. ... In fact, nothing in the kernel needs to understand the underlying details of the filesystems, except the filesystems themselves [*Linux Kernel Development*, Sec. 13.2, *Filesystem Abstraction Layer*].

The key idea is to abstract out that part of the file system that is common to all file systems and put that code in a separate layer that calls the underlying concrete file systems to actual manage the data [*Modern Operating Systems*, Sec. 4.3.7, *Virtual File Systems*]. The overall structure is illustrated in Fig 341.

All system calls relating to files are directed to the virtual file system for initial processing. These calls, coming from user processes, are the standard POSIX calls, such as `open`, `read`, `write`, `lseek`, and so on. Thus the VFS has an “upper” interface to user processes and it is the well-known POSIX interface.

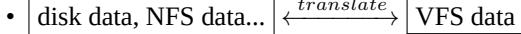
The VFS also has a “lower” interface to the concrete file systems, which is labeled VFS interface in fig. 341 . This interface consists of several dozen function calls that the VFS can make to each file system to get work done.

Thus to create a new file system that works with the VFS, the designers of the new file system must make sure that it supplies the function calls the VFS requires.

### Virtual File System

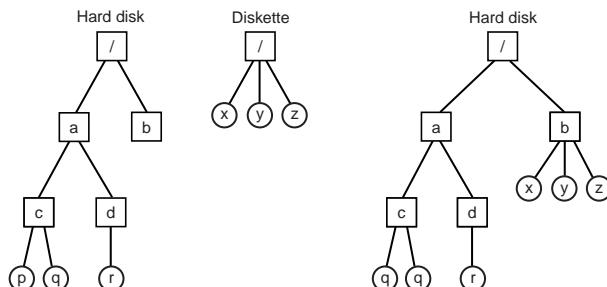
- Manages kernel level file abstractions in one format for all file systems
- Receives system call requests from user level (e.g. `write`, `open`, `stat`, `link`)
- Interacts with a specific file system based on mount point traversal
- Receives requests from other parts of the kernel, mostly from memory management

### Real File Systems

- managing file & directory data
- managing meta-data: timestamps, owners, protection, etc.
-  **disk data, NFS data...**  $\xleftrightarrow{\text{translate}}$  **VFS data**

Historically, Unix has provided four basic filesystem-related abstractions: files, directory entries, inodes, and mount points. ... Traditionally, Unix filesystems implement these notions as part of their physical on-disk layout. For example, file information is stored as an inode in a separate block on the disk; directories are files; control information is stored centrally in a superblock, and so on. The Unix file concepts are physically mapped on to the storage medium. The Linux VFS is designed to work with filesystems that understand and implement such concepts. Non-Unix filesystems, such as FAT or NTFS, still work in Linux, but their filesystem code must provide the appearance of these concepts. For example, *even if a filesystem does not support distinct inodes, it must assemble the inode data structure in memory as if it did*. Or if a filesystem treats directories as a special object, to the VFS they must represent directories as mere files. *Often, this involves some special processing done on-the-fly by the non-Unix filesystems* to cope with the Unix paradigm and the requirements of the VFS. Such filesystems still work, however, and the overhead is not unreasonable [*Linux Kernel Development*, Sec. 13.3, *Unix Filesystems*].

### File System Mounting



### A FS must be mounted before it can be used

#### Mount — The file system is registered with the VFS

- The superblock is read into the VFS superblock
- The table of addresses of functions the VFS requires is read into the VFS superblock
- The FS' topology info is mapped onto the VFS superblock data structure

#### The VFS keeps a list of the mounted file systems together with their superblocks

The VFS superblock contains:

- Device, blocksize
- Pointer to the *root inode*
- Pointer to a set of superblock routines
- Pointer to `file_system_type` data structure
- more...

See also: [*Linux Kernel Development*, Sec. 13.13, *Data Structures Associated with Filesystems*].

- `struct file_system_type`: There is only one `file_system_type` per filesystem, regardless of how many instances of the filesystem are mounted on the system, or whether the filesystem is even mounted at all.
- `struct vfsmount`: represents a specific instance of a filesystem — in other words, a mount point.

## V-node

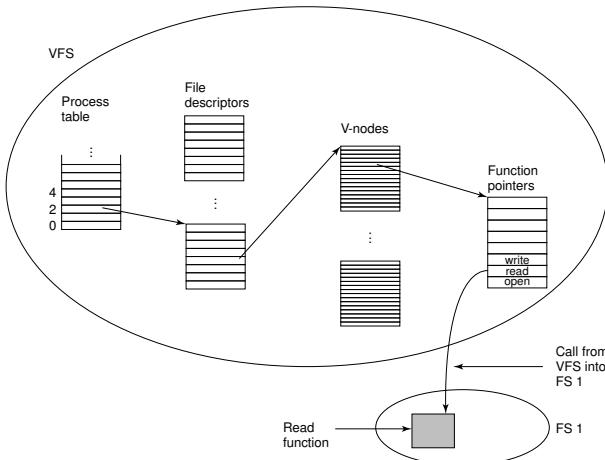
- Every file/directory in the VFS has a VFS inode, kept in the VFS *inode cache*
- The real FS builds the VFS inode from its own info

## Like the EXT2 inodes, the VFS inodes describe

- files and directories within the system
- the contents and topology of the Virtual File System

## VFS Operation

*read()*



To understand how the VFS works, let us run through an example chronologically. When the system is booted, the root file system is registered with the VFS. In addition, when other file systems are mounted, either at boot time or during operation, they, too must register with the VFS. *When a file system registers, what it basically does is provide a list of the addresses of the functions the VFS requires, either as one long call vector (table) or as several of them, one per VFS object, as the VFS demands.* Thus once a file system has registered with the VFS, the VFS knows how to, say, read a block from it — it simply calls the fourth (or whatever) function in the vector supplied by the file system. Similarly, the VFS then also knows how to carry out every other function the concrete file system must supply: it just calls the function whose address was supplied when the file system registered[*Modern Operating Systems*, Sec. 4.3.7, *Virtual File Systems*, P. 288].

After a file system has been mounted, it can be used. For example, if a file system has been mounted on `/usr` and a process makes the call

```
open( "/usr/include/unistd.h", O_RDONLY);
```

while parsing the path, the VFS sees that a new file system has been mounted on `/usr` and locates its superblock by searching the list of superblocks of mounted file systems. Having done this, it can find the root directory of the mounted file system and look up the path `include/unistd.h` there. The VFS then creates a v-node and makes a call to the concrete file system to return all the information in the file's i-node. This information is copied into the v-node (in RAM), along with other information, most importantly the pointer to the table of functions to call for operations on v-nodes, such as `read`, `write`, `close`, and so on.

After the v-node has been created, the VFS makes an entry in the file descriptor table for the calling process and sets it to point to the new v-node. (For the purists, the file descriptor actually points to another data structure that contains the current file position and a pointer to the v-node, but this detail is not important for our purposes here.) Finally, the VFS returns the file descriptor to the caller so it can use it to `read`, `write`, and `close` the file.

Later when the process does a `read` using the file descriptor, the VFS locates the v-node from the process and file descriptor tables and follows the pointer to the table of functions, all of which are addresses within the concrete file system on which the requested file resides. The function that handles `read` is now called and code within the concrete file system goes and gets the requested block. *The VFS has no idea whether the data are coming from the local disk, a remote file system over the network, a CD-ROM, a USB stick, or something different.* The data structures involved are shown in Fig 346. Starting with the caller's process number and the file descriptor, successively the v-node, read function pointer, and access function within the concrete file system are located.

In this manner, it becomes relatively straightforward to add new file systems. To make one, the designers first get a list of function calls the VFS expects and then write their file system to provide all of them. Alternatively, if the file system already exists, then they have to provide wrapper functions that do what the VFS needs, usually by making one or more native calls to the concrete file system.

See also: [*Linux Kernel Development*, Sec. 13.14, *Data Structures Associated with a Process*].

## Linux VFS

### The Common File Model

All other filesystems must map their own concepts into the common file model

For example, FAT filesystems do not have inodes.

- The main components of the common file model are
  - superblock** – information about mounted filesystem
  - inode** – information about a specific file
  - file** – information about an open file
  - dentry** – information about directory entry
- Geared toward Unix FS

**Dentry** Note that because the VFS treats directories as normal files, there is not a specific directory object. Recall from earlier in this chapter that a dentry represents a component in a path, which might include a regular file. In other words, a dentry is not the same as a directory, but a directory is just another kind of file. Got it?[*Linux Kernel Development*, Sec. 13.4, *VFS Objects and Their Data Structures*]

**The operations objects** An operations object is contained within each of these primary objects. These objects describe the methods that the kernel invokes against the primary objects[*Linux Kernel Development*, Sec. 13.4, *VFS Objects and Their Data Structures*]:

- The **super\_operations** object, which contains the methods that the kernel can invoke on a specific filesystem, such as `write_inode()` and `sync_fs()`
- The **inode\_operations** object, which contains the methods that the kernel can invoke on a specific file, such as `create()` and `link()`
- The **dentry\_operations** object, which contains the methods that the kernel can invoke on a specific directory entry, such as `d_compare()` and `d_delete()`
- The **file\_operations** object, which contains the methods that a process can invoke on an open file, such as `read()` and `write()`

The operations objects are implemented as a structure of pointers to functions that operate on the parent object. For many methods, the objects can inherit a generic function if basic functionality is sufficient. Otherwise, the specific instance of the particular filesystem fills in the pointers with its own filesystem-specific methods.

### The Superblock Object

- is implemented by each FS and is used to store information describing that specific FS
- usually corresponds to the *filesystem superblock* or the *filesystem control block*
- Filesystems that are not disk-based (such as `sysfs`, `proc`) generate the superblock on-the-fly and store it in memory
- `struct super_block` in `<linux/fs.h>`
- `s_op` in `struct super_block` — `struct super_operations` — the superblock operations table
  - Each item in this table is a pointer to a function that operates on a superblock object

The code for creating, managing, and destroying superblock objects lives in `fs/super.c`. A superblock object is created and initialized via the `alloc_super()` function. When mounted, a filesystem invokes this function, reads its superblock off of the disk, and fills in its superblock object[*Linux Kernel Development*, Sec. 13.5, *The Superblock Object*].

When a filesystem needs to perform an operation on its superblock, it follows the pointers from its superblock object to the desired method. For example, if a filesystem wanted to write to its superblock, it would invoke

```
sb->s_op->write_super(sb);
```

In this call, `sb` is a pointer to the filesystem's superblock. Following that pointer into `s_op` yields the superblock operations table and ultimately the desired `write_super()` function, which is then invoked. Note how the `write_super()` call must be passed a superblock, despite the method being associated with one. This is because of the lack of object-oriented support in C. In C++, a call such as the following would suffice:

```
sb.write_super();
```

In C, there is no way for the method to easily obtain its parent, so you have to pass it[*Linux Kernel Development*, Sec. 13.6, *Superblock Operations*].

### The Inode Object

- For Unix-style filesystems, this information is simply read from the on-disk inode
- For others, the inode object is constructed in memory in whatever manner is applicable to the filesystem

- struct inode in <linux/fs.h>
- An inode represents each file on a FS, but the inode object is constructed in memory only as files are accessed
  - includes special files, such as device files or pipes
- i\_op → struct inode\_operations

### The Dentry Object

- components in a path
- makes path name lookup easier
- struct dentry in <linux/dcache.h>
- created on-the-fly from a string representation of a path name

### Dentry State

- used
- unused
- negative

### Dentry Cache

consists of three parts:

1. Lists of “used” dentries
2. A doubly linked “least recently used” list of unused and negative dentry objects
3. A hash table and hashing function used to quickly resolve a given path into the associated dentry object

A **used** dentry corresponds to a valid inode (d\_inode points to an associated inode) and indicates that there are one or more users of the object (d\_count is positive). A used dentry is in use by the VFS and points to valid data and, thus, cannot be discarded [Linux Kernel Development, Sec. 13.9.1, *Dentry State*].

An *unused* dentry corresponds to a valid inode (d\_inode points to an inode), but the VFS is not currently using the dentry object (d\_count is zero). Because the dentry object still points to a valid object, the dentry is kept around — cached — in case it is needed again. Because the dentry has not been destroyed prematurely, the dentry need not be re-created if it is needed in the future, and path name lookups can complete quicker than if the dentry was not cached. If it is necessary to reclaim memory, however, the dentry can be discarded because it is not in active use.

A *negative* dentry is not associated with a valid inode (d\_inode is NULL) because either the inode was deleted or the path name was never correct to begin with. The dentry is kept around, however, so that future lookups are resolved quickly. For example, consider a daemon that continually tries to open and read a config file that is not present. The open() system calls continually returns ENOENT, but not until after the kernel constructs the path, walks the on-disk directory structure, and verifies the file’s inexistence. Because even this failed lookup is expensive, caching the “negative” results are worthwhile. Although a negative dentry is useful, it can be destroyed if memory is at a premium because nothing is actually using it.

A dentry object can also be freed, sitting in the slab object cache, as discussed in the previous chapter. In that case, there is no valid reference to the dentry object in any VFS or any filesystem code.

### The File Object

- is the in-memory representation of an open file
- open() ⇒ create; close() ⇒ destroy
- there can be multiple file objects in existence for the same file
  - Because multiple processes can open and manipulate a file at the same time
- struct file in <linux/fs.h>

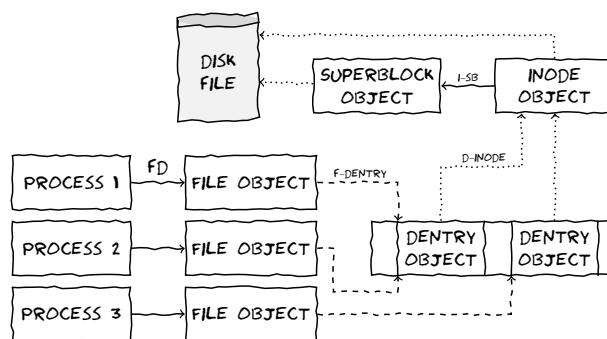


Fig. 353 illustrates with a simple example how processes interact with files. Three different processes have opened the same file, two of them using the same hard link. In this case, each of the three processes uses its own file

object, while only two dentry objects are required one for each hard link. Both dentry objects refer to the same inode object, which identifies the superblock object and, together with the latter, the common disk file[*Understanding The Linux Kernel*, Sec. 12.1.1].

## References

- [1] Wikipedia. File system — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=File\\_system&oldid=646603121](http://en.wikipedia.org/w/index.php?title=File_system&oldid=646603121).
- [2] Wikipedia. Computer file — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Computer\\_file&oldid=647724614](http://en.wikipedia.org/w/index.php?title=Computer_file&oldid=647724614).
- [3] Wikipedia. Inode — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Inode&oldid=647736522>.
- [4] Wikipedia. Ext2 — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Ext2&oldid=642312602>.
- [5] Wikipedia. Virtual file system — Wikipedia, The Free Encyclopedia. 2014. [http://en.wikipedia.org/w/index.php?title=Virtual\\_file\\_system&oldid=640354992](http://en.wikipedia.org/w/index.php?title=Virtual_file_system&oldid=640354992).

## Part V

# Input/Output

## 21 Principles of I/O Hardware

### I/O Hardware

#### Different people, different view

- *Electrical engineers*: chips, wires, power supplies, motors...
- *Programmers*: interface presented to the software
  - the commands the hardware accepts
  - the functions it carries out
  - the errors that can be reported back
  - ...

### I/O Devices

Roughly two Categories:

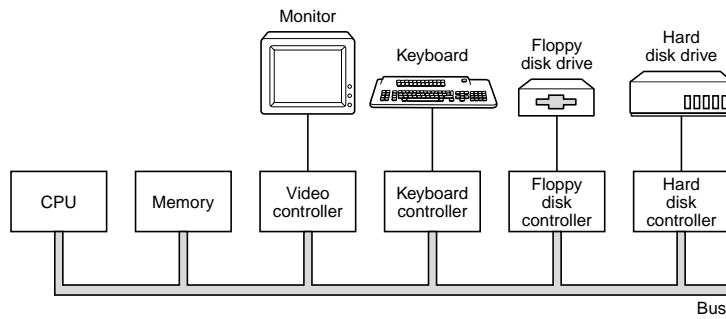
1. *Block devices*: store information in fix-size blocks
2. *Character devices*: deal with streams of characters

This classification scheme is not perfect. Some devices just do not fit in. Clocks, for example, are not block addressable. Nor do they generate or accept character streams. All they do is cause interrupts at well-defined intervals. Memory-mapped screens do not fit the model well either. Still, the model of block and character devices is general enough that it can be used as a basis for making some of the operating system software dealing with I/O device independent. The file system, for example, deals just with abstract block devices and leaves the device-dependent part to lower-level software[*Modern Operating Systems*, Sec. 5.1.1, *I/O Devices*].

### Device Controllers

I/O units usually consist of

1. a mechanical component
2. an electronic component is called the *device controller* or *adapter*
  - e.g. video adapter, network adapter...



**Port:** A connection point. A device communicates with the machine via a port — for example, a serial port.

**Bus:** A set of wires and a rigidly defined protocol that specifies a set of messages that can be sent on the wires.

**Controller:** A collection of electronics that can operate a port, a bus, or a device.

e.g. serial port controller, SCSI bus controller, disk controller...

## The Controller's Job

### Examples:

- *Disk controllers:* convert the serial bit stream into a block of bytes and perform any error correction necessary
- *Monitor controllers:* read bytes containing the characters to be displayed from memory and generates the signals used to modulate the CRT beam to cause it to write on the screen

The interface between the controller and the device is often a very low-level interface. A disk, for example, might be formatted with 10,000 sectors of 512 bytes per track. What actually comes off the drive, however, is a serial bit stream, starting with a *preamble*, then the 4096 bits in a sector, and finally a checksum, also called an *Error-Correcting Code (ECC)*. The preamble is written when the disk is formatted and contains the cylinder and sector number, the sector size, and similar data, as well as synchronization information [Modern Operating Systems, Sec. 5.1.2, *Device Controllers*].

The controller's job is to convert the serial bit stream into a block of bytes and perform any error correction necessary. The block of bytes is typically first assembled, bit by bit, in a buffer inside the controller. After its checksum has been verified and the block has been declared to be error free, it can then be copied to main memory.

The controller for a monitor also works as a bit serial device at an equally low level. It reads bytes containing the characters to be displayed from memory and generates the signals used to modulate the CRT beam to cause it to write on the screen. The controller also generates the signals for making the CRT beam do a horizontal retrace after it has finished a scan line, as well as the signals for making it do a vertical retrace after the entire screen has been scanned. If it were not for the CRT controller, the operating system programmer would have to explicitly program the analog scanning of the tube. With the controller, the operating system initializes the controller with a few parameters, such as the number of characters or pixels per line and number of lines per screen, and lets the controller take care of actually driving the beam. Flat-screen TFT displays are different, but just as complicated.

## Inside The Controllers

**Control registers:** for communicating with the CPU (R/W, On/Off...)

**Data buffer:** for example, a video RAM

Q: How the CPU communicates with the control registers and the device data buffers?

A: Usually two ways...

(a) Each control register is assigned an I/O port number, then the CPU can do, for example,

– read in control register PORT and store the result in CPU register REG.

IN REG, PORT

– write the contents of REG to a control register

OUT PORT, REG

*Note:* the address spaces for memory and I/O are different. For example,

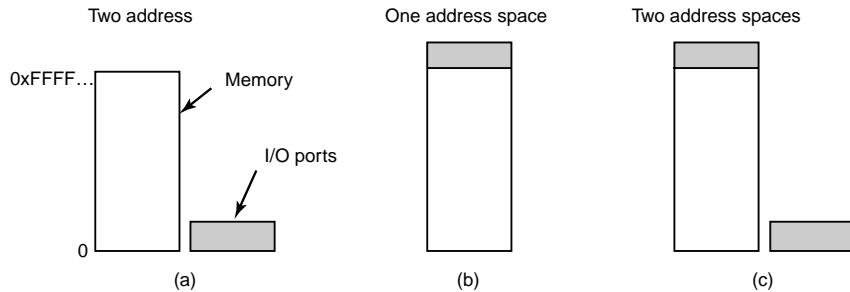
– read the contents of I/O port 4 and puts it in R0

IN R0, 4 ; 4 is a port number

– read the contents of memory word 4 and puts it in R0

MOV R0, 4 ; 4 is a memory address

## Address spaces



- (a) Separate I/O and memory space
- (b) *Memory-mapped I/O*: map all the control registers into the memory space. Each control register is assigned a unique memory address.
- (c) *Hybrid*: with memory-mapped I/O data buffers and separate I/O ports for the control registers. For example, Intel Pentium
  - Memory addresses 640K to 1M being reserved for device data buffers
  - I/O ports 0 through 64K.

How can the processor give commands and data to a controller to accomplish an I/O transfer? The short answer is that the controller has one or more registers for data and control signals. The processor communicates with the controller by reading and writing bit patterns in these registers. *One way in which this communication can occur is through the use of special I/O instructions that specify the transfer of a byte or word to an I/O port address.* The I/O instruction triggers bus lines to select the proper device and to move bits into or out of a device register. *Alternatively, the device controller can support memory-mapped I/O.* In this case, the device-control registers are mapped into the address space of the processor. The CPU executes I/O requests using the standard data-transfer instructions to read and write the device-control registers. Some systems use both techniques. For instance, PCs use I/O instructions [*Operating System Concepts Essentials*, Sec. 12.2, *I/O Hardware*].

How do these schemes work? In all cases, when the CPU wants to read a word, either from memory or from an I/O port, it puts the address it needs on the bus' address lines and then asserts a READ signal on a bus' control line. A second signal line is used to tell whether I/O space or memory space is needed. If it is memory space, the memory responds to the request. If it is I/O space, the I/O device responds to the request. If there is only memory space [as in (b)], every memory module and every I/O device compares the address lines to the range of addresses that it services. If the address falls in its range, it responds to the request. Since no address is ever assigned to both memory and an I/O device, there is no ambiguity and no conflict [*Modern Operating Systems*, Sec. 5.1.3, *Memory-mapped I/O*].

### Advantages of Memory-mapped I/O

#### No assembly code is needed (IN, OUT...)

With memory-mapped I/O, device control registers are just variables in memory and can be addressed in C the same way as any other variables. Thus with memory-mapped I/O, a I/O device driver can be written entirely in C.

#### No special protection mechanism is needed

The I/O address space is part of the kernel space, thus cannot be touched directly by any user space process.

The two schemes for addressing the controllers have different strengths and weaknesses. Let us start with the advantages of memory-mapped I/O. First, if special I/O instructions are needed to read and write the device control registers, access to them requires the use of assembly code since there is no way to execute an IN or OUT instruction in C or C++. Calling such a procedure adds overhead to controlling I/O. In contrast, with memory-mapped I/O, device control registers are just variables in memory and can be addressed in C the same way as any other variables. Thus with memory-mapped I/O, a I/O device driver can be written entirely in C. Without memory-mapped I/O, some assembly code is needed [*Modern Operating Systems*, Sec. 5.1.3, *Memory-mapped I/O*].

Second, with memory-mapped I/O, no special protection mechanism is needed to keep user processes from performing I/O. All the operating system has to do is refrain from putting that portion of the address space containing the control registers in any user's virtual address space. Better yet, if each device has its control registers on a different page of the address space, the operating system can give a user control over specific devices but not others by simply including the desired pages in its page table. Such a scheme can allow different device drivers to be placed in different address spaces, not only reducing kernel size but also keeping one driver from interfering with others.

Third, with memory-mapped I/O, every instruction that can reference memory can also reference control registers. For example, if there is an instruction, TEST, that tests a memory word for 0, it can also be used to test a control register for 0, which might be the signal that the device is idle and can accept a new command. The assembly language code might look like this:

```

LOOP: TEST PORT_4 ;check if port 4 is 0
      BEQ READY ;if it is 0, go to ready
      BRANCH LOOP ;otherwise, continue testing
READY:

```

If memory-mapped I/O is not present, the control register must first be read into the CPU, then tested, requiring two instructions instead of one. In the case of the loop given above, a fourth instruction has to be added, slightly slowing down the responsiveness of detecting an idle device.

## Disadvantages of Memory-mapped I/O

### Caching problem

- Caching a device control register would be disastrous
- Selectively disabling caching adds extra complexity to both hardware and the OS

### Problem with multiple buses

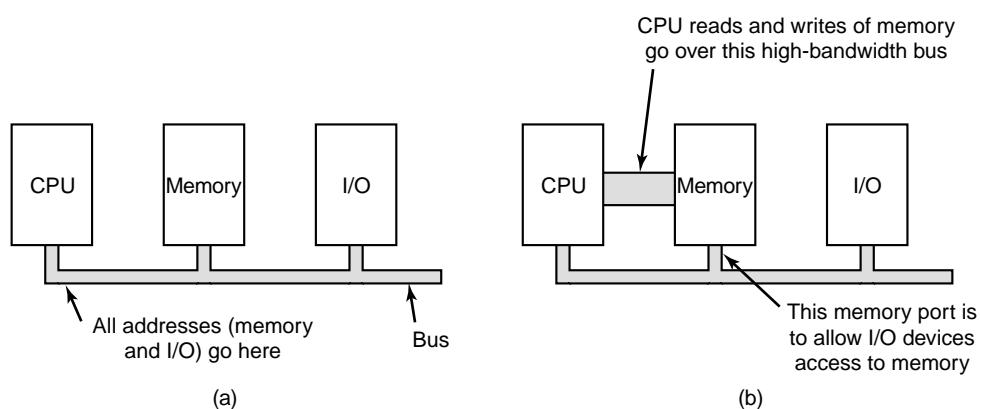


Fig. 2: (a) A single-bus architecture. (b) A dual-bus memory architecture.

In computer design, practically everything involves trade-offs, and that is the case here too. Memory-mapped I/O also has its disadvantages. First, most computers nowadays have some form of caching of memory words. Caching a device control register would be disastrous. Consider the assembly code loop given above in the presence of caching. The first reference to PORT\_4 would cause it to be cached. Subsequent references would just take the value from the cache and not even ask the device. Then when the device finally became ready, the software would have no way of finding out. Instead, the loop would go on forever[*Modern Operating Systems*, Sec. 5.1.3, *Memory-mapped I/O*].

To prevent this situation with memory-mapped I/O, the hardware has to be equipped with the ability to selectively disable caching, for example, on a per page basis. This feature adds extra complexity to both the hardware and the operating system, which has to manage the selective caching.

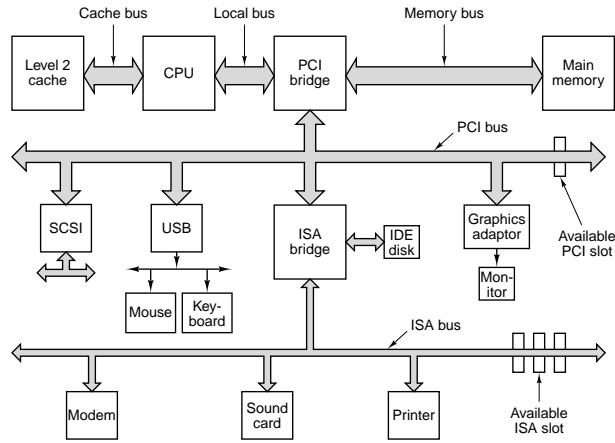
Second, if there is only one address space, then all memory modules and all I/O devices must examine all memory references to see which ones to respond to. If the computer has a single bus, as in Fig 2(a), having everyone look at every address is straightforward.

However, the trend in modern personal computers is to have a dedicated high-speed memory bus, as shown in Fig 2(b), a property also found in mainframes, incidentally. This bus is tailored to optimize memory performance, with no compromises for the sake of slow I/O devices. Pentium systems can have multiple buses (memory, PCI, SCSI, USB, ISA), as shown in Fig 3.

The trouble with having a separate memory bus on memory-mapped machines is that the I/O devices have no way of seeing memory addresses as they go by on the memory bus, so they have no way of responding to them. Again, special measures have to be taken to make memory-mapped I/O work on a system with multiple buses. One possibility is to first send all memory references to the memory. If the memory fails to respond, then the CPU tries the other buses. This design can be made to work but requires additional hardware complexity.

A second possible design is to put a snooping device on the memory bus to pass all addresses presented to potentially interested I/O devices. The problem here is that I/O devices may not be able to process requests at the speed the memory can.

A third possible design, which is the one used on the Pentium configuration of Fig 3, is to filter addresses in the PCI bridge chip. This chip contains range registers that are preloaded at boot time. For example, 640K to 1M could be marked as a nonmemory range. Addresses that fall within one of the ranges marked as nonmemory are forwarded onto the PCI bus instead of to memory. The disadvantage of this scheme is the need for figuring out



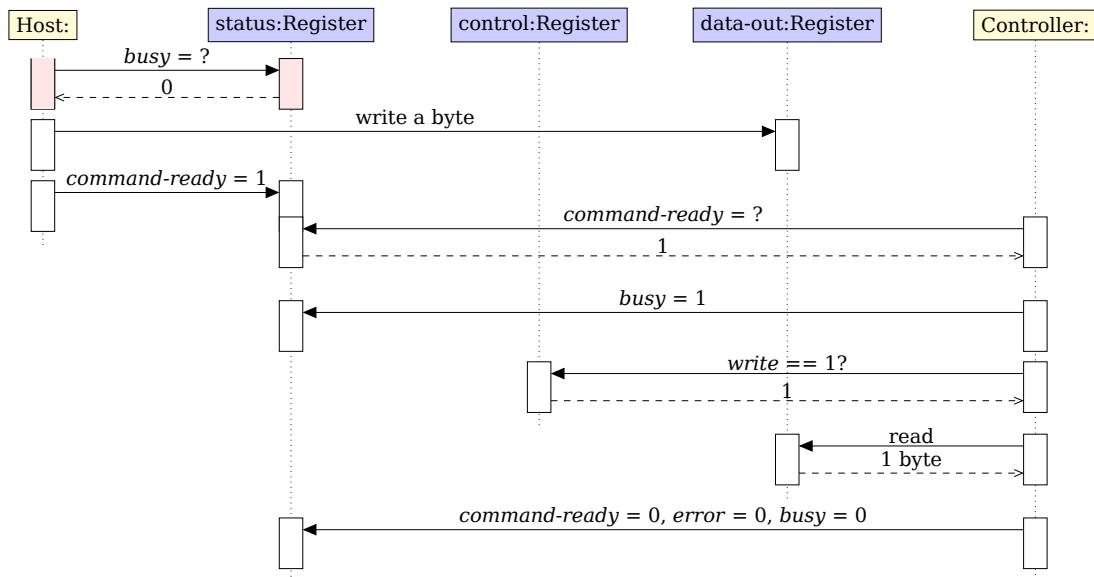
**Fig. 3:** The structure of a large Pentium system

at boot time which memory addresses are not really memory addresses. Thus each scheme has arguments for and against it, so compromises and trade-offs are inevitable.

## 21.1 Programmed I/O

### Programmed I/O

#### Handshaking



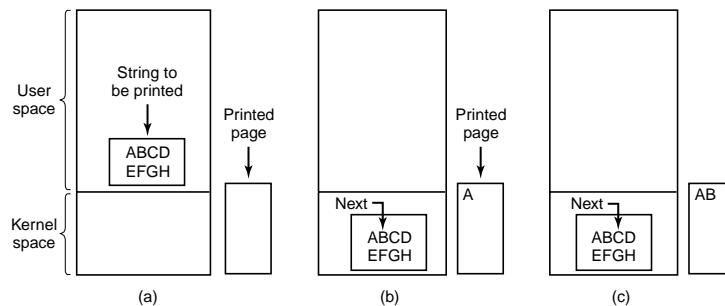
For this example, the host writes output through a port, coordinating with the controller by handshaking as follows[*Operating System Concepts Essentials*, Sec. 12.2.1, *Polling*].

1. The host repeatedly reads the *busy* bit until that bit becomes clear.
2. The host sets the *write* bit in the command register and writes a byte into the *data-out* register.
3. The host sets the *command-ready* bit.
4. When the controller notices that the *command-ready* bit is set, it sets the *busy* bit.
5. The controller reads the command register and sees the *write* command. It reads the *data-out* register to get the byte and does the I/O to the device.
6. The controller clears the *command-ready* bit, clears the *error* bit in the status register to indicate that the device I/O succeeded, and clears the *busy* bit to indicate that it is finished.

This loop is repeated for each byte.

### Example

#### Steps in printing a string



```

1 | copy_from_user(buffer, p, count);           /* p is the kernel buffer */
2 | for (i = 0; i < count; i++) {               /* loop on every character */
3 |     while(*printer_status_reg != READY);    /* loop until ready */
4 |     *printer_data_register = p[i];           /* output one character */
5 | }
6 | return_to_user();

```

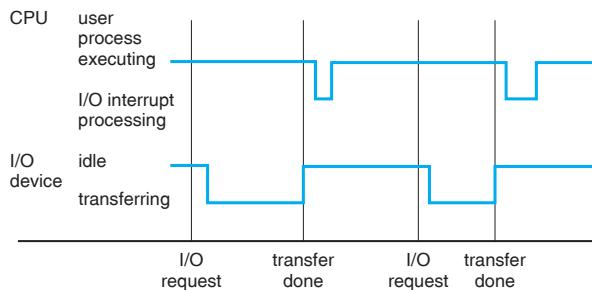
See also: [Modern Operating Systems, Sec. 5.2.2, Programmed I/O].

## 21.2 Interrupt-Driven I/O

### Pulling Is Inefficient

#### A better way

**Interrupt** The controller notifies the CPU when the device is ready for service.



See also: [Operating System Concepts Essentials, Sec. 12.2.2, Interrupts].

### Example

*Writing a string to the printer using interrupt-driven I/O*

#### When the print system call is made...

```

1 | copy_from_user(buffer, p, count);
2 | enable_interrupts();
3 | while(*printer_status_reg != READY);
4 | *printer_data_register = p[0];
5 | scheduler();

```

#### Interrupt service procedure for the printer

```

1 | if (count == 0) {
2 |     unblock_user();
3 | } else {
4 |     *printer_data_register = p[1];
5 |     count = count - 1;
6 |     i = i + 1;
7 | }
8 | acknowledge_interrupt();
9 | return_from_interrupt();

```

See also: [Modern Operating Systems, Sec. 5.2.3, Interrupt-Driven I/O].

## 21.3 Direct Memory Access (DMA)

### Direct Memory Access (DMA)

**Programmed I/O (PIO)** tying up the CPU full time until all the I/O is done (busy waiting)

**Interrupt-Driven I/O** an interrupt occurs on every character (wastes CPU time)

**Direct Memory Access (DMA)** Uses a special-purpose processor (DMA controller) to work with the I/O controller

### Example

#### Printing a string using DMA

In essence, DMA is programmed I/O, only with the DMA controller doing all the work, instead of the main CPU.

When the print system call is made...

```
1 copy_from_user(buffer, p, count);
2 set_up_DMA_controller();
3 scheduler();
```

### Interrupt service procedure

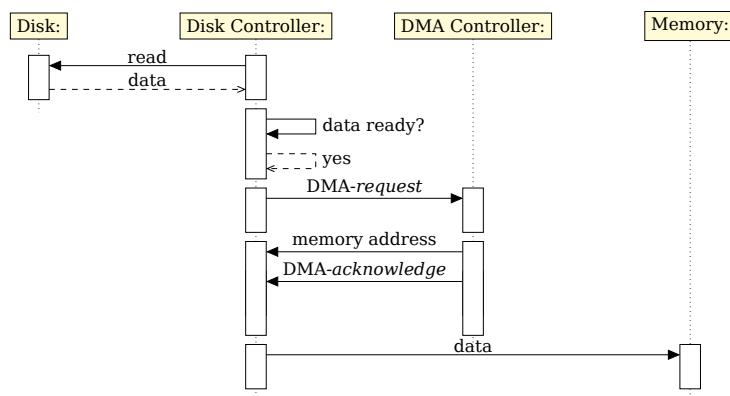
```
1 acknowledge_interrupt();
2 unblock_user();
3 return_from_interrupt();
```

The big win with DMA is reducing the number of interrupts from one per character to one per buffer printed.

The operating system can only use DMA if the hardware has a DMA controller, which most systems do. Sometimes this controller is integrated into disk controllers and other controllers, but such a design requires a separate DMA controller for each device. More commonly, a single DMA controller is available (e.g., on the parentboard) for regulating transfers to multiple devices, often concurrently [Modern Operating Systems, Sec. 5.1.4, *Direct Memory Access (DMA)*].

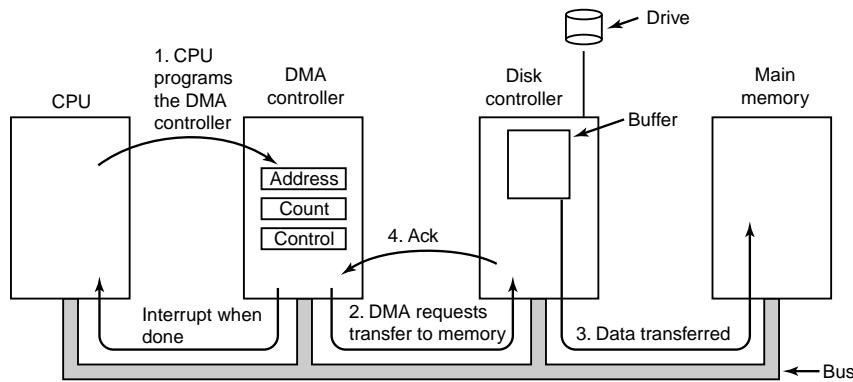
### DMA Handshaking Example

*Read from disk*



### The DMA controller

- has access to the system bus independent of the CPU
- contains several registers that can be written and read by the CPU. These includes
  - a memory address register
  - a byte count register
  - one or more control registers
    - \* specify the I/O port to use
    - \* the direction of the transfer (read/write)
    - \* the transfer unit (byte/word)
    - \* the number of bytes to transfer in one burst.

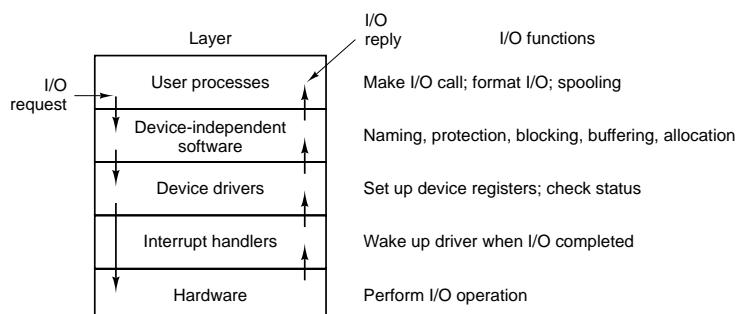


When DMA is used, ..., first the CPU programs the DMA controller by setting its registers so it knows what to transfer where (step 1 in Fig 373). It also issues a command to the disk controller telling it to read data from the disk into its internal buffer and verify the checksum. When valid data are in the disk controller's buffer, DMA can begin [Modern Operating Systems, Sec. 5.1.4, Direct Memory Access (DMA)].

The DMA controller initiates the transfer by issuing a read request over the bus to the disk controller (step 2). This read request looks like any other read request, and the disk controller does not know or care whether it came from the CPU or from a DMA controller. Typically, the memory address to write to is on the bus' address lines so when the disk controller fetches the next word from its internal buffer, it knows where to write it. The write to memory is another standard bus cycle (step 3). When the write is complete, the disk controller sends an acknowledgement signal to the DMA controller, also over the bus (step 4). The DMA controller then increments the memory address to use and decrements the byte count. If the byte count is still greater than 0, steps 2 through 4 are repeated until the count reaches 0. At that time, the DMA controller interrupts the CPU to let it know that the transfer is now complete. When the operating system starts up, it does not have to copy the disk block to memory; it is already there.

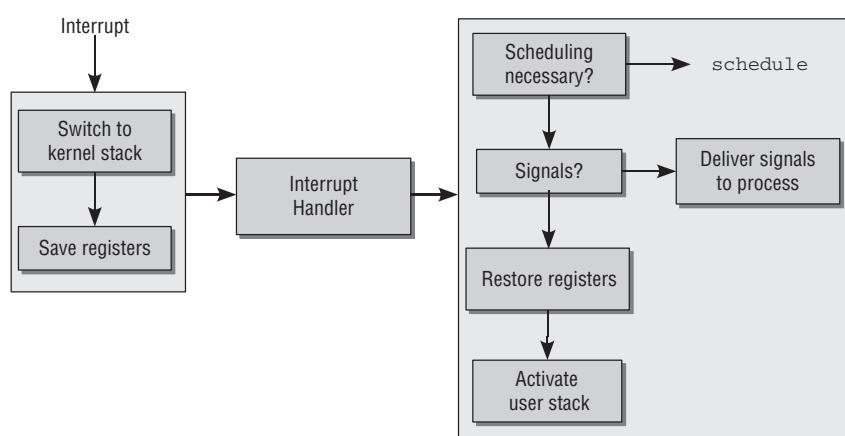
## 22 I/O Software Layers

### I/O Software Layers



### 22.1 Interrupt Handlers

#### Interrupt Handlers

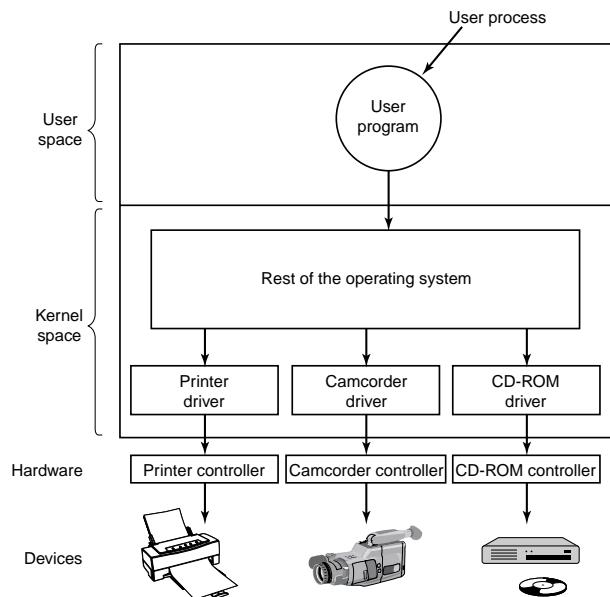


See also:

- [Professional Linux Kernel Architecture, Sec. 14.1.3, Processing Interrupts].
- [Modern Operating Systems, Sec. 5.3.1, Interrupt Handlers].

## 22.2 Device Drivers

### Device Drivers



See also: [Modern Operating Systems, Sec. 5.3.2, Device Drivers].

## 22.3 Device-Independent I/O Software

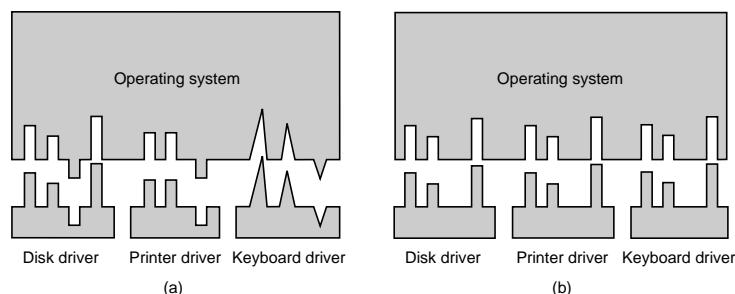
### Device-Independent I/O Software

#### Functions

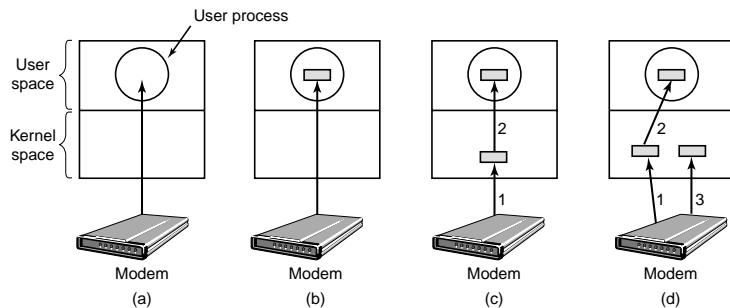
- Uniform interfacing for device drivers
- Buffering
- Error reporting
- Allocating and releasing dedicated devices
- Providing a device-independent block size

See also: [Modern Operating Systems, Sec. 5.3.3, Device-Independent I/O Software].

### Uniform Interfacing for Device Drivers



### Buffering



## 22.4 User-Space I/O Software

### User-Space I/O Software

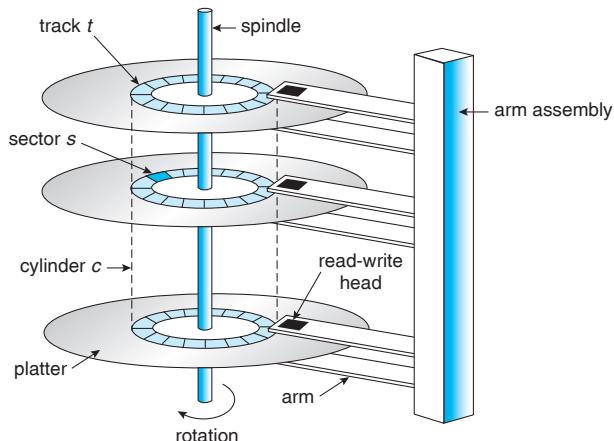
#### Examples:

- stdio.h, printf(), scanf()...
- spooling(printing, USENET news...)

See also: [Modern Operating Systems, Sec. 5.3.4, User-Space I/O Software].

## 23 Disks

### Disk



A read-write head “flies” just above each surface of every platter. The heads are attached to a disk arm that moves all the heads as a unit. The surface of a platter is logically divided into circular tracks, which are subdivided into sectors. The set of tracks that are at one arm position makes up a cylinder. There may be thousands of concentric cylinders in a disk drive, and each track may contain hundreds of sectors. The storage capacity of common disk drives is measured in gigabytes[Operating System Concepts Essentials, Sec. 11.1.1, Magnetic Disks].

When the disk is in use, a drive motor spins it at high speed. Most drives rotate 60 to 200 times per second. Disk speed has two parts. The transfer rate is the rate at which data flow between the drive and the computer. The positioning time, sometimes called the random-access time, consists of the time necessary to move the disk arm to the desired cylinder, called the seek time, and the time necessary for the desired sector to rotate to the disk head, called the rotational latency. Typical disks can transfer several megabytes of data per second, and they have seek times and rotational latencies of several milliseconds.

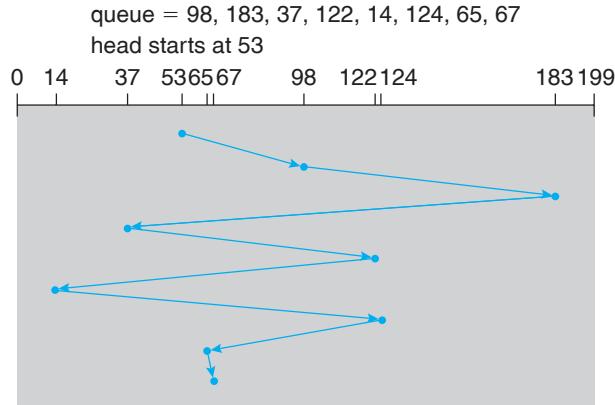
...

A disk drive is attached to a computer by a set of wires called an I/O bus. Several kinds of buses are available, including enhanced integrated drive electronics (EIDE), advanced technology attachment (ATA), serial ATA (SATA), universal serial bus (USB), fiber channel (FC), and small computer-systems interface (SCSI) buses. The data transfers on a bus are carried out by special electronic processors called controllers. The host controller is the controller at the computer end of the bus. A disk controller is built into each disk drive. To perform a disk I/O operation, the computer places a command into the host controller, typically using memory-mapped I/O ports, as described in Section 8.7.3. The host controller then sends the command via messages to the disk controller, and the disk controller operates the disk-drive hardware to carry out the command. Disk controllers usually have a built-in cache. Data transfer at the disk drive happens between the cache and the disk surface, and data transfer to the host, at fast electronic speeds, occurs between the cache and the host controller.

## 23.1 Disk Scheduling

### Disk Scheduling

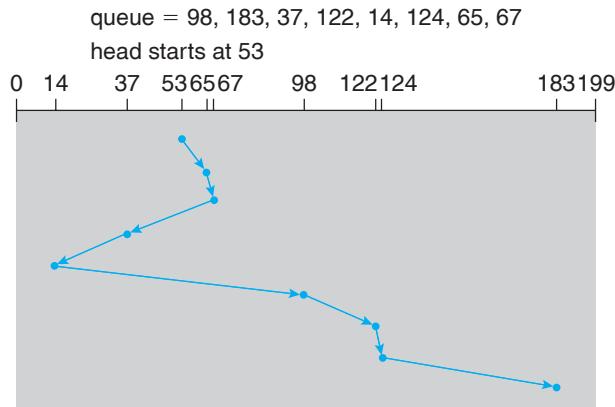
*First Come First Serve (FCFS)*



If the disk head is initially at cylinder 53, it will first move from 53 to 98, then to 183, 37, 122, 14, 124, 65, and finally to 67, for a total head movement of 640 cylinders. The wild swing from 122 to 14 and then back to 124 illustrates the problem with this schedule. If the requests for cylinders 37 and 14 could be serviced together, before or after the requests for 122 and 124, the total head movement could be decreased substantially, and performance could be thereby improved[*Operating System Concepts Essentials*, Sec. 11.4.1, *FCFS Scheduling*].

### Disk Scheduling

*Shortest Seek Time First (SSTF)*



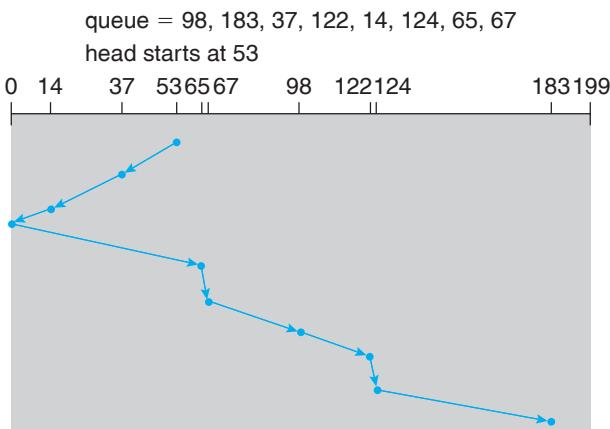
For our example request queue, the closest request to the initial head position (53) is at cylinder 65. Once we are at cylinder 65, the next closest request is at cylinder 67. From there, the request at cylinder 37 is closer than the one at 98, so 37 is served next. Continuing, we service the request at cylinder 14, then 98, 122, 124, and finally 183. This scheduling method results in a total head movement of only 236 cylinders—little more than one-third of the distance needed for FCFS scheduling of this request queue. Clearly, this algorithm gives a substantial improvement in performance[*Operating System Concepts Essentials*, Sec. 11.4.2, *SSTF Scheduling*].

SSTF scheduling is essentially a form of shortest-job-first (SJF) scheduling; and like SJF scheduling, *it may cause starvation of some requests*. Remember that requests may arrive at any time. Suppose that we have two requests in the queue, for cylinders 14 and 186, and while the request from 14 is being serviced, a new request near 14 arrives. This new request will be serviced next, making the request at 186 wait. While this request is being serviced, another request close to 14 could arrive. In theory, a continual stream of requests near one another could cause the request for cylinder 186 to wait indefinitely. This scenario becomes increasingly likely as the pending-request queue grows longer.

Although the SSTF algorithm is a substantial improvement over the FCFS algorithm, *it is not optimal*. In the example, we can do better by moving the head from 53 to 37, even though the latter is not closest, and then to 14, before turning around to service 65, 67, 98, 122, 124, and 183. This strategy reduces the total head movement to 208 cylinders.

## Disk Scheduling

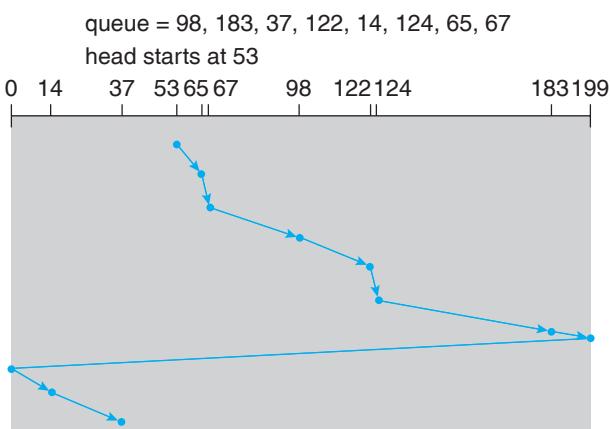
### SCAN Scheduling



See also: [Operating System Concepts Essentials, Sec. 11.4.3, SCAN Scheduling].

## Disk Scheduling

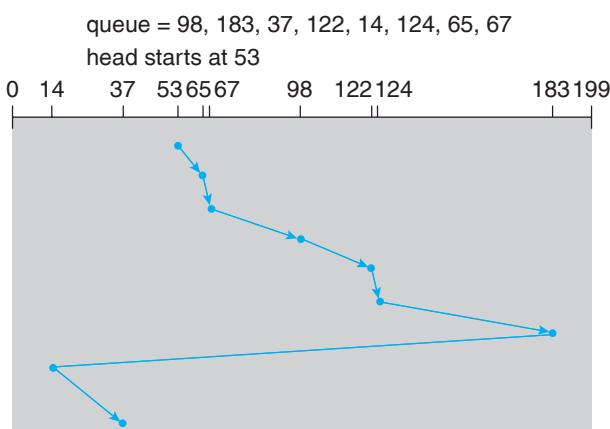
### Circular SCAN (C-SCAN) Scheduling



See also: [Operating System Concepts Essentials, Sec. 11.4.4, C-SCAN Scheduling].

## Disk Scheduling

### C-LOOK Scheduling



## The Best Scheduling Algorithm?

**Performance depends heavily on the number and types of requests**

File system design can be influential

- File-allocation method (contiguous? indexed?)

- Location of directories and index blocks

In Linux, the *deadline I/O scheduler* used in Version 2.6 works similarly to the elevator algorithm (C-SCAN) except that it also associates a deadline with each request, thus addressing the starvation issue. By default, the deadline for read requests is 0.5 second and that for write requests is 5 seconds. The deadline scheduler maintains a *sorted queue* of pending I/O operations ordered by sector number. However, it also maintains two other queues — a *read queue* for read operations and a *write queue* for write operations. These two queues are ordered according to deadline. Every I/O request is placed in both the sorted queue and either the read or the write queue, as appropriate. Ordinarily, I/O operations occur from the sorted queue. However, if a deadline expires for a request in either the read or the write queue, I/O operations are scheduled from the queue containing the expired request. This policy ensures that an I/O operation will wait no longer than its expiration time [*Operating System Concepts Essentials*, Sec. 15.8.1, *Block Devices*].

See also: [*Operating System Concepts Essentials*, Sec. 11.4.6, *Selection of a Disk-Scheduling Algorithm*].

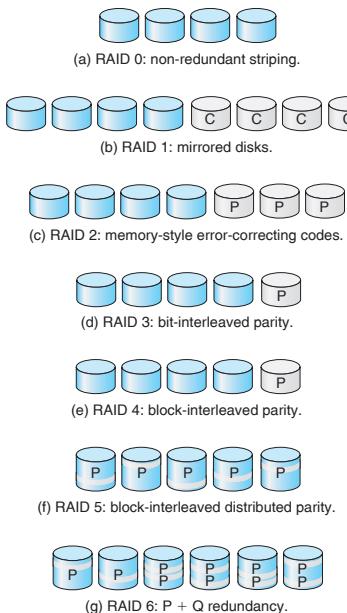
## 23.2 RAID Structure

### RAID Structure

#### Redundant Arrays of Independent/Inexpensive Disks

- ⌚ Replace SLED with RAID
- ⌚ Replace the disk controller card with a RAID controller
- ⌚ Better performance and better reliability
- ⌚ No software changes are required to use the RAID

### RAID Levels



### References

- [1] Wikipedia. Hard disk drive — Wikipedia, The Free Encyclopedia. 2015. [http://en.wikipedia.org/w/index.php?title=Hard\\_disk\\_drive&oldid=648161868](http://en.wikipedia.org/w/index.php?title=Hard_disk_drive&oldid=648161868).
- [2] Wikipedia. Input/output — Wikipedia, The Free Encyclopedia. 2015. <http://en.wikipedia.org/w/index.php?title=Input/output&oldid=646503450>.