User-Level Threads

- ☹ Lack of coordination between threads and OS kernel
  - Process as a whole gets one time slice
  - Same time slice, whether process has 1 thread or 1000 threads
  - Also – up to each thread to relinquish control to other threads in that process
- ☹ Requires non-blocking system calls (i.e. a multithreaded kernel)
  - Otherwise, entire process will blocked in the kernel, even if there are runnable threads left in the process
  - part of motivation for user-level threads was not to have to modify the OS
- ☹ If one thread causes a page fault(interrupt!), the entire process blocks