

Completely Fair Scheduling (CFS)

Linux's Process Scheduler

up to 2.4: simple, scaled poorly

- ▶ $O(n)$
- ▶ non-preemptive in kernel mode
- ▶ single run queue (cache? SMP?)

from 2.5 on: $O(1)$ scheduler

- 😊 140 priority lists — scaled well
- 😊 one run queue per CPU — true SMP support
- 😊 preemptive
- 😊 ideal for large server workloads
- 😞 showed latency on desktop systems

from 2.6.23 on: Completely Fair Scheduler (CFS)

- 😊 improved interactive performance