

# CFS

In practice it works like this:

- ▶ While a task is using the CPU, its `wait_runtime` decreases

$$\text{wait\_runtime} = \text{wait\_runtime} - \text{time\_running}$$

if: its `wait_runtime`  $\neq$   $\text{MAX}_{\text{wait\_runtime}}$  (among all processes)

then: it gets preempted

- ▶ Newly woken tasks (`wait_runtime = 0`) are put into the tree more and more to the right
- ▶ slowly but surely giving a chance for every task to become the “leftmost task” and thus get on the CPU within a deterministic amount of time